



LEARN ENOUGH

CUSTOM DOMAINS

TO BE DANGEROUS

DEVELOPER KNOWLEDGE

TUTORIAL BY

LEE DONAHOE AND MICHAEL HARTL

Learn Enough Custom Domains to Be Dangerous

Lee Donahoe and Michael Hartl

Contents

1 A name of our own	1
1.1 Introduction	2
1.2 DNS overview	4
1.2.1 Visiting a site via DNS lookup	10
1.3 Custom domain registration	18
1.3.1 What to register?	20
1.3.2 You've got a domain, now what?	24
2 DNS management	25
2.1 Cloudflare setup	25
2.1.1 Cloudflare features	26
2.1.2 Cloudflare signup	31
2.1.3 Connecting registrar nameservers	31
2.2 Custom domains at GitHub Pages	36
2.2.1 Configuring Cloudflare for GitHub Pages	36
2.2.2 Configuring GitHub Pages	41
2.3 Custom domains at Heroku	43
2.4 Canonical URLs	53
2.4.1 Cloudflare page rules	54
2.5 Profit!!	59
3 Custom email and analytics	61
3.1 Custom email	61
3.1.1 Introduction to email	61

3.1.2	Google mail	65
3.1.3	Google Workspace signup	68
3.1.4	Verify domain and configure MX records	69
3.2	Site Analytics	74
3.2.1	Add Snippet	75
3.3	Conclusion	79

About the authors

Learn Enough cofounder Lee Donahoe is an entrepreneur, designer, and front-end developer. In addition to doing the design for [Learn Enough](#), [Softcover](#), and the [Ruby on Rails Tutorial](#), he is also a cofounder and front-end developer for [Coveralls](#), a leading test coverage analysis service, and is tech cofounder of [Buck Mason](#), a men's clothing company once featured on ABC's [Shark Tank](#). Lee is a graduate of [USC](#), where he majored in Economics and studied Interactive Multimedia & Technologies.

[Michael Hartl](#) is the creator of the [Ruby on Rails Tutorial](#), one of the leading introductions to web development, and is cofounder and principal author at [Learn Enough](#). Previously, he was a physics instructor at the [California Institute of Technology](#) (Caltech), where he received a [Lifetime Achievement Award for Excellence in Teaching](#). He is a graduate of [Harvard College](#), has a [Ph.D. in Physics](#) from [Caltech](#), and is an alumnus of the [Y Combinator](#) entrepreneur program.

Chapter 1

A name of our own

Welcome to *Learn Enough Custom Domains to Be Dangerous!*

So, you want to make a custom website? Or perhaps you already have one? This tutorial is designed for anyone who'd rather have their website live at [example.com](#) than at [example.someoneelsesdomain.com](#)—in other words, at a domain you control and that no one can ever take away. As a bonus: your email address can be yourname@example.com instead of yourname1995@-someoneelsesdomain.com.

If you're not sure how to build a real website, we've got you covered there, too: an especially good time to read this tutorial is after you've put together and deployed your first live site on the Internet while following the [Learn Enough tutorial series](#)—especially the Learn Enough tutorials on [CSS & Layout](#), [JavaScript](#), and [Ruby](#), as well as the [Ruby on Rails Tutorial](#).

Finally, this tutorial is designed as a general reference for custom domains, useful even for experienced web professionals. Configuring a website to use a custom domain is generally quite complicated—and every time you need to buy and configure a new domain you'll find yourself following these same steps... which ends up happening for basically every project. This is why we decided it would probably be a good idea to put all that information into a living single document (which we'll update and expand over time), instead of having the content spread out in small sections over a number of tutorials.

1.1 Introduction

At first glance, this tutorial might look fairly long—the print-equivalent of over 80 pages—but most of the length is due to including a lot of screenshots. We’ve also included supplementary material on things like the Domain Name System ([Section 1.2](#)) because we thought it was worth explaining some of the underlying structure that makes these technologies work. (If you’re doing all this a second time, feel free to skip such sections.) Once you know how this process works, you can do all the steps one after another in quick succession, which should let you go quickly from having a test site on [GitHub](#) with a generic email account to a site that is on its own domain with a professional-looking email address. Getting everything done really only takes an afternoon.

The main focus in this tutorial is configuring a custom domain for a statically built site hosted on [GitHub Pages](#), as covered in the Learn Enough tutorials on [Git](#), [HTML](#), [CSS & Layout](#), and [JavaScript](#). There is also a bonus section ([Section 2.3](#)) covering custom domains for dynamic web applications, such as those created in [Learn Enough Ruby to Be Dangerous](#) and the [Ruby on Rails Tutorial](#). The process is similar for both. Additionally, this tutorial shows you how to set up a couple of useful Google services: Gmail, for custom email addresses ([Section 3.1](#)), and Google Analytics, so that you can collect information about who is on your site ([Section 3.2](#)).

Because of how quickly things change on the Web, it is entirely possible that some of the information in this guide might at times be out of date due to services changing interfaces or features ([Figure 1.1](#)).¹ We at Learn Enough are going to try to make sure that this tutorial is always kept up-to-date, but please understand that there is always a degree of variability in the interface for all these services—even if in the end they still mostly work the same. That means you should feel free to contact Learn Enough if you notice that some part of the tutorial is out of date (support@learnenough.com).

Even with the steps detailed here, in all likelihood you’ll still have to apply some of your [technical sophistication](#) in order to get everything to work. That’s just how it goes on the World Wide Web!

¹Image retrieved from <https://flic.kr/p/2a8hC> on 2019-11-20. Copyright © 2005 jimthompson and used under the terms of the [Creative Commons Attribution 2.0 Generic](#) license.



Figure 1.1: Sorry we broke your tutorial by changing our interface (not sorry)!

One last note: in the course of doing this tutorial, we are going to use a bunch of different services that you might have to pay for. We aren't making any money from endorsing any of these services; we just find them to be easy to use, and so we are suggesting that you use them, too (Figure 1.2).²

There are ways of doing much of what we discuss entirely for free, but it might come with the cost of dealing with the hassle of constant configuration and monitoring to make sure that everything is working as expected. When it comes to tech services, often it's better to pay someone to make your life a little easier.

1.2 DNS overview

This section gives an overview of how Internet domains work. If you're not interested in these details, or if you've seen them before, at this point you can skip right to custom domain registration in Section 1.3.

You may or may not know this, but when you open up your browser and type in the address of a site like **google.com**, you aren't *actually* entering a real address where the Google server can be found. Instead **google.com** can be thought of more as a screen name that is an easier way to get to Google's web services than if you had to remember the real machine address... which for us at the time of writing this, in the geographic location we're connecting from, is **172.217.4.174**.

You can find a site's IP address by using the **ping** command in your **command line** terminal, as shown in Listing 1.1 (which also has a little bit of the resulting output). Big sites like Google generally have many servers, so which one you connect to at any given time will vary based on the details of Google's system.

Listing 1.1: Pinging google.com from your command line reveals the server IP address.

²Image retrieved from <https://flic.kr/p/brd1K2> on 2019-11-20. Copyright © 2012 and used under the terms of the [Creative Commons Attribution 2.0 Generic license](#).



Figure 1.2: Totally not getting paid to say that.

```
$ ping google.com  
PING google.com (172.217.4.174): 56 data bytes  
64 bytes from 172.217.4.174: icmp_seq=0 ttl=57 time=305.001 ms  
. . .
```

The number in Listing 1.1 is called an [Internet Protocol \(IP\) address](#), and it is made up of four 8-bit numbers (also called *octets*) separated by periods ([Box 1.1](#)). These IP addresses are what define the location of machines on the network and allow for data to be transferred via the underlying system, which is known as [TCP/IP](#) (Transmission Control Protocol/Internet Protocol). TCP/IP is the suite of protocols that governs the flow of information around the network of networks that we call the Internet. The association of the domain name with an IP address is known as an *A record* (where “A” stands for “address”).

Box 1.1. Bits and bytes

A *bit*, or *binary digit*, is the basic unit of information in a computer. It represents one of two values (thus [binary](#)), typically written as either 0 or 1. A set of 8 bits is called a *byte*.

Because bits can take either of two values, the number 2 is especially important in computing, as are various powers of 2. For example, as noted in the text, IP numbers consist of four 8-bit numbers (octets) separated by periods. But what is an 8-bit number?

Well, each bit represents two possible values, so an 8-bit number represents $2^8 = 256$ possibilities. If we started counting at 1, the largest number allowable in an octet would be 256 itself, but because 0 is an allowable number the maximum is actually $2^8 - 1 = 255$.

Thanks to the explosion of Internet-connected devices like phones, computers, cars, and even refrigerators, the world has actually run out of IP addresses. The system described above using four 8 bit numbers is called [IPv4](#), and to add

more addresses we're in the process to a new system called **IPv6** (Box 1.2).³ For now, both systems are still being used during the transition.

Box 1.2. IPv4 to IPv6

Originally, IP addresses were divided up into blocks that were then assigned to corporations, universities, countries, etc. On top of that, some addresses were reserved for specific uses, making them unavailable to be assigned to machines; for instance, 127.0.0.1 is reserved as the internal *home address* for any computer. (Thus the computer joke “There’s no place like 127.0.0.1”, a reference to the famous quote from *The Wizard of Oz*.)



³Image retrieved from <https://flic.kr/p/2aGJ7q5> on 2019-11-20. Copyright © 2012 and used under the terms of the [Creative Commons Attribution 2.0 Generic license](#).

Even though mathematically there are $256^4 = 4,294,967,296$ (or just over four billion) addresses in the IPv4 system, a lot weren't able to be used, or they are parts of already-assigned blocks and so can't be used even if there is no machine at the address now. For a long time, we were able to get away with not having enough addresses by having entire networks with their own sets of internal IP addresses represented on the Internet by a single IP address, but that was just a strategy to buy time.

We've gradually been making the switch to a new addressing system, IPv6, that uses an address that can consist of 8 groups of 4 [hexadecimal digits](#). The format of the new IP addresses looks like this:

0000:0000:0000:0000:0000:0000:0000:0000

So Google's IPv6 address is much more complicated-looking:

2a00:1450:4007:0813:0000:0000:0000:200e

There are rules for shortening some of the sections if they start with a zero, or shortening the address if it contains multiple consecutive sections that are all zeros, but we aren't going to get into that here. We just wanted you to understand that there was an old system, and that now there is a new system we'll all eventually be moving to.

IP addresses are needed because all information that is sent or received on the Internet is chopped up into little pieces called *packets*, and each packet has an origination address and a destination address. So when you request a web page, the actual content is chopped up by the server into a bunch of chunks, flung out onto the Internet, and each chunk has the address for your computer written on it. Those chunks pass from server to server to get to your computer, and when they reach you, your browser reassembles everything into the web pages you are familiar with.

Imagine how much harder it would be to go to different sites on the Internet



Figure 1.3: Domain names are a lot easier than this would be!

if you had to remember something like the string of numbers in an IP address every time you wanted to go to Google’s site—or any other site (or even something much harder, like [Egyptian hieroglyphs](#) (Figure 1.3)). The people who created the modern Internet knew that navigating around the network using strings of numbers was not going to be a great way to encourage the growth of the network, and so they came up with methods to make things easier.

One of the first methods, and one that actually still exists even on modern computers, was to create a file locally called the *HOSTS* file, which would have the raw IP address and then next to it a more human-readable name. Every user would keep their own little directory of IP addresses and assign an easier-to-remember name for the resource. [Listing 1.2](#) shows what a modern *HOSTS* file looks like with an entry for Google.

Listing 1.2: Adding a host entry for Google.

/etc/hosts

```

## 
# Host Database
#
# localhost is used to configure the loopback interface
# when the system is booting. Do not change this entry.
## 
127.0.0.1 localhost
255.255.255.255 broadcasthost
::1           localhost

216.58.217.206 www.google.com

```

Having to keep a HOSTS file up to date with a bunch of different sites was kind of a pain, though. So, sometime in the ’80s, all the different groups and individuals started coming up with a new global system, called the Domain Name System (DNS), that would work like a phone book for the Internet, and would allow for users to navigate around the network using human-understandable words and phrases which would get **automagically** converted to machine addresses behind the scenes—no more editing local files to save the IP addresses of sites that you want to visit.

1.2.1 Visiting a site via DNS lookup

To get an idea of how DNS works, let’s take a look at what happens when we enter a web address—often called a *URL* (for “Uniform Resource Locator”)—into a browser.

First of all, you’ve seen addresses like **<https://www.google.com>** a million times, but why do we use addresses with all those different parts (Figure 1.4) instead of just entering **google** into the browser? What do the different parts mean?

- *protocol*: The **https://** part of the address tells the browser what sort of connection it should try making to a server. In the case of **http(s)**, we are telling the browser that it is going to make a hypertext transfer protocol request; the optional but common addition of the **s** at the end means that we are requiring the connection to the server to be *secure* and encrypted.

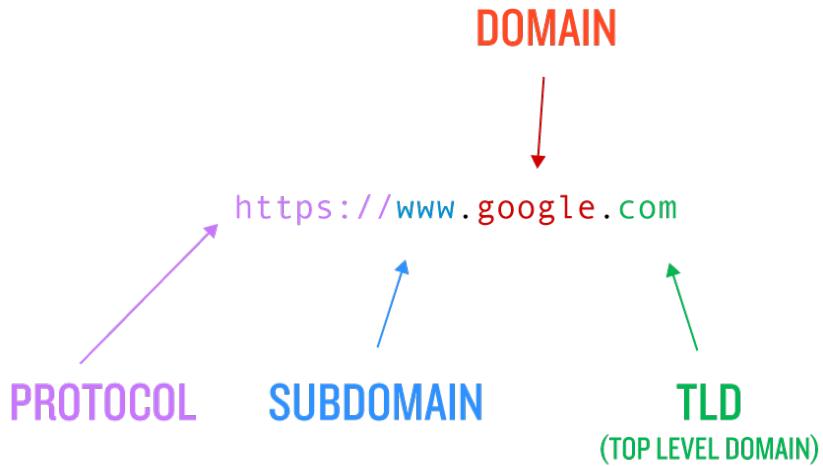


Figure 1.4: The anatomy of a URL.

HTTP requests are the kind of request that ask a server to send back a web page. There are a number of other protocols, like **ftp** (file transfer protocol) for sending and receiving large files and **smtp** (simple mail transfer protocol) for sending email.

- *top-level domain (TLD)*: The top-level domains are the highest level of the domain name system, and you can think of them as completely separate address books that keep a registry of unique names that people have paid for. Each TLD is given the authority to allow people to register domain names for only that TLD—so if you own **example.com** that doesn't mean that you also automatically own **example.org**.

If you want your site available at both addresses, the names would have to be registered with both TLD organizations. There are generic TLDs, like **.com** and **.org**, and then there are TLDs that have been created for individual countries, like **.us** for the United States, **.uk** for the United Kingdom, etc.

Additionally, the body that governs the Internet naming system has al-

lowed organizations to pay to create their own new TLDs, and so there are wild new options like `.ninja` and `.limo`. Ultimately, you can think of TLDs as being like area codes of phone numbers (only at the end instead of at the beginning).

- *domain name*: This is the name of the site or resource that has been registered with a TLD. Domain names have to be unique within a TLD in the same way that there can't be multiple identical phone numbers within a single area code.
- *subdomain*: The subdomain is an extra domain that is fully controlled by the owner of a domain, so there is no need to register or pay for new subdomains. This means that web traffic can go to `www.google.com`, email traffic can go to `mail.google.com`, and customer support traffic can go to `support.google.com`. Each one of those subdomains can be configured to show users a completely different site.

The `www` subdomain—for “World Wide Web”—was created to tell servers that the browser requesting information from that subdomain was looking to receive web pages that conform to the Web standard—in other words, HyperText Markup Language, or HTML, the language of the World Wide Web (see *Learn Enough HTML to Be Dangerous*). While it is possible to configure a site to serve web content to a browser request without a subdomain, like `https://google.com`—which is called a “naked” or “root” domain—there are *complicated reasons* why you probably don’t want to run your site on a root domain, and indeed most big sites permanently *redirect* the root domain to the corresponding address using `www`. Examples of such redirected URLs include `google.com`, `amazon.com`, `apple.com`, `microsoft.com`, and `facebook.com`. (We’ll learn how to make such redirects ourselves in [Section 2.4.1.](#).)

After you finish typing the address and hit enter, your computer first checks to see if you’ve visited that site before (or if it’s in the HOSTS file). If so, your computer already has the IP address and can thus connect immediately ([Figure 1.5](#)).

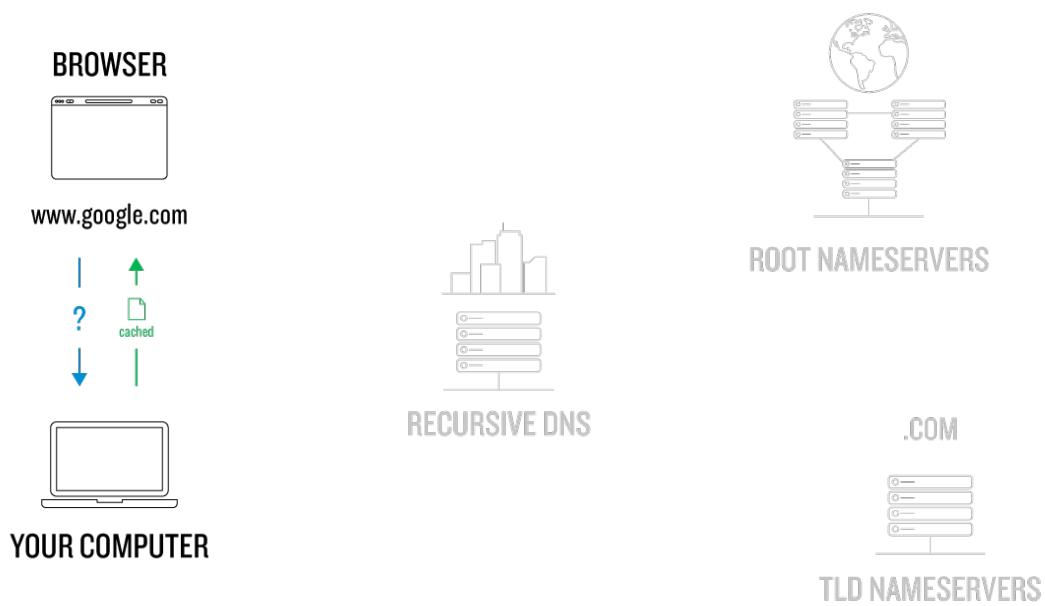


Figure 1.5: Connecting with a locally stored IP address.

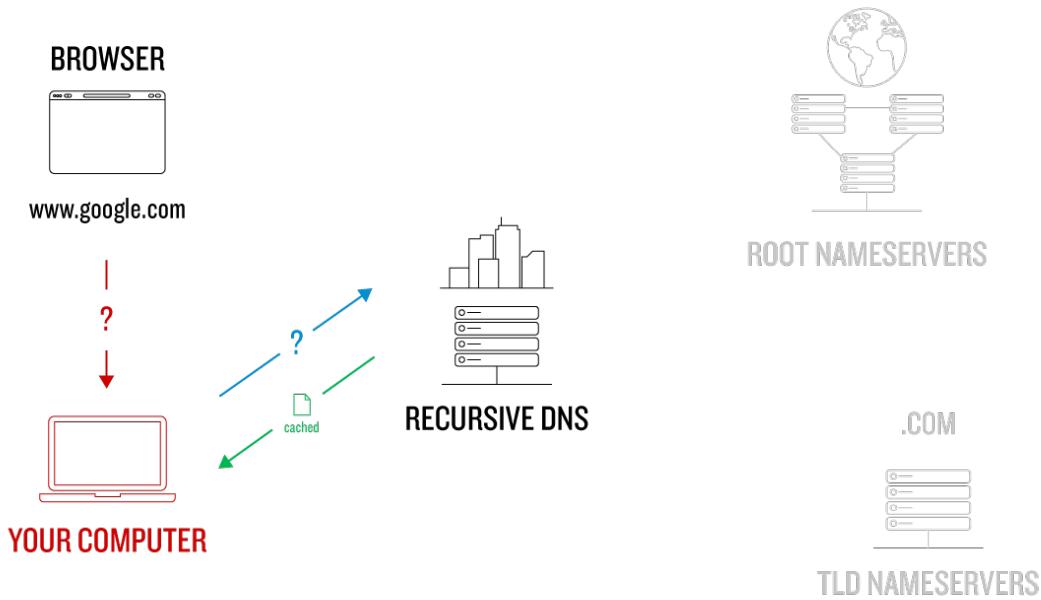


Figure 1.6: Obtaining an address from the nearest recursive DNS.

If the address isn't stored locally, the browser next reaches out to a server (run either by your Internet service provider or by a nearby publicly accessible server) called a *recursive DNS server*. It asks that server, "Do you know the machine address that corresponds to www.google.com?" These recursive servers have their own **cache** of site names and addresses, so most of the time your request is quickly found (Figure 1.6). As a side note, you can configure your computer or router to use specific recursive DNS servers (if you are concerned about security), and Google actually runs their own public servers that are a good option (Box 1.3).

Box 1.3. Using the Google DNS servers

You might not realize it, but every time that you connect to Wi-Fi or use the Internet, by default your device is given the address for a recursive DNS server. This address might be set on a modem that is provided by your Internet service

provider (ISP), it could be from the settings for a Wi-Fi router in a café that set their own custom DNS address, or it could be the DNS that your wireless phone provider gives you. While this system often works well in practice, one potential problem is that you have to trust that the server you are connecting to is going to give you the right addresses.

What do we mean by that?

Well, if you are connecting to sites through an unscrupulous service, like let's say your Internet service provider (because some actually got caught doing this), they could have their DNS server point you to their own web servers instead of to the site you wanted to connect to. Then they request the page you were looking to load, add their own content (like ads) onto the page (or otherwise mess with the content), and then pass everything back to you without your realizing that something bad happened. So you think you are getting Google search results, and you do, but your ISP also injected a bunch of their own ads alongside.

This is called DNS hijacking and it is clearly not good. One way to avoid this is to go into the settings for your router, your computer, and your phone, and set a custom trusted DNS server as the default.

Google runs their own DNS servers that you are free to use if you trust them (and we generally do), but that's up to you. The address for the servers are **8.8.8.8**, and **8.8.4.4**. You can always search on Google to find the addresses if you forget.

The method to add these servers as your default DNS differs by device and operating system, but if you look online you'll be able to find guides on how to do it. Once you set a preferred DNS server on your device, that setting will override any potentially malicious DNS server address set by a router, modem, or ISP.

If a recursive DNS server doesn't have a machine address listing that corresponds to the name you entered in your browser, then your request will be kicked up to the global domain name system. When that happens, the request is first sent to the global *root nameservers*. This is a network of thirteen servers that act kind of like a **switchboard**, and all they care about is directing your request to the correct top-level domain network. They don't know the answer to

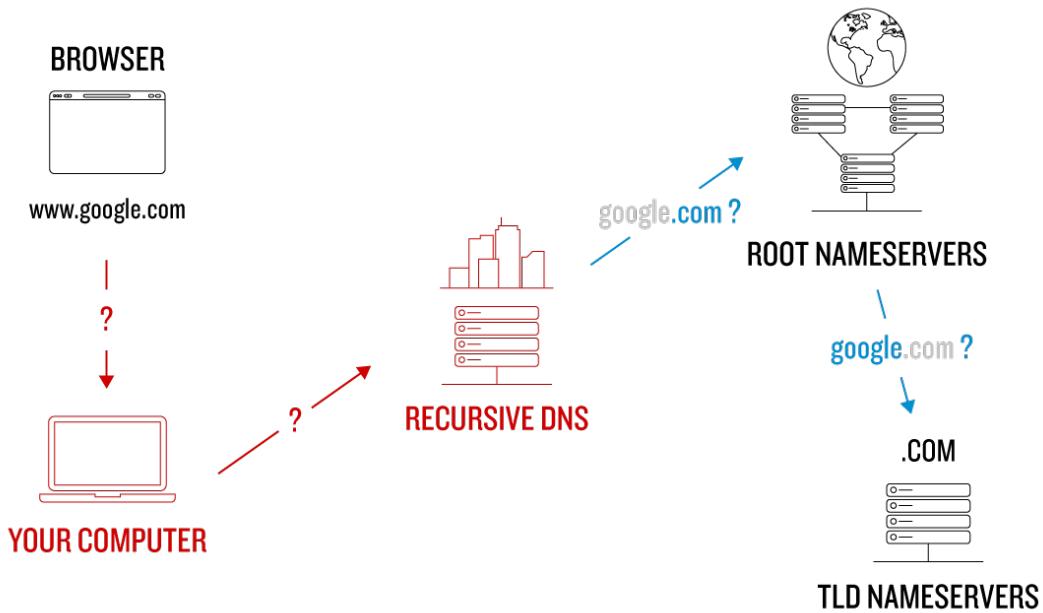


Figure 1.7: Using the root nameservers find the TLD’s nameservers.

“What is the IP address for `www.google.com`?”, but what they do know is the addresses for the network of servers that manage the `.com` TLD (Figure 1.7).

Your request then travels from the root nameservers down to the network of servers run by the top-level domain of the site that you are trying to reach. Your request for `www.google.com` is routed around inside the `.com` network until the address record for the site you are trying to reach is found. This will contain an entry that ultimately has the IP address for the server you are trying to reach.

A reply is sent back to your computer through the recursive DNS server that you originally contacted, so that the record can be cached on that server. As a result, the next time someone wants to go to the site you requested the whole lookup process won’t be necessary (Figure 1.8). At the end of the process, your computer receives the content of the DNS record, and your browser opens up a connection to the server.

All this happens in a fraction of a second, and the requests can bounce all

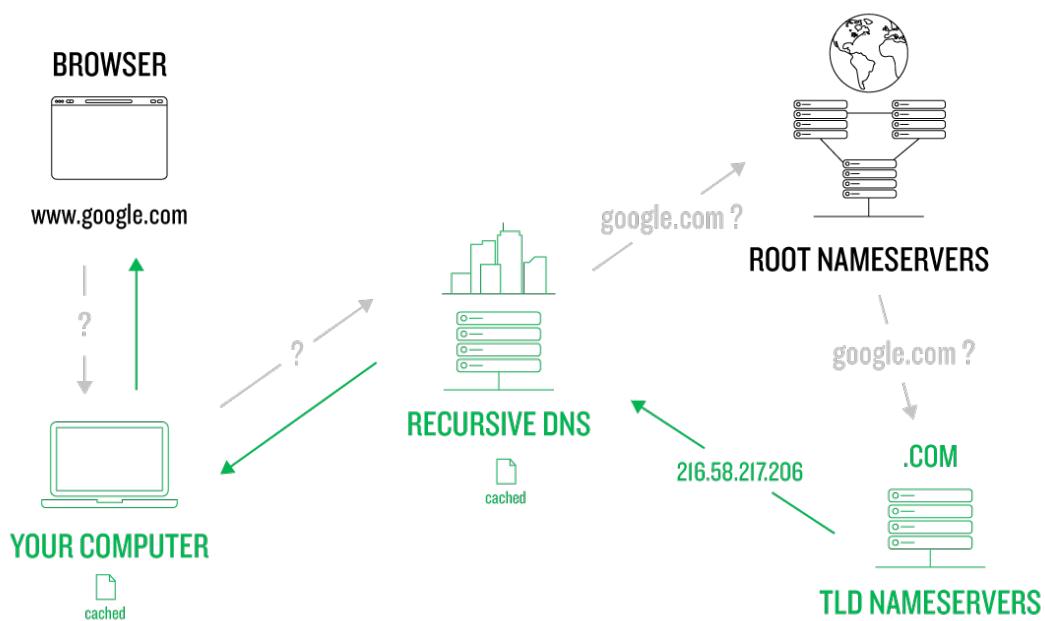


Figure 1.8: Caching the address for future use.

the way around the world without your noticing... Isn't technology amazing?

1.3 Custom domain registration

It's pretty hard to configure a custom domain unless you actually have one, so our first step is to *register* one.

Domain names have to be registered through an authorized registrar, and there are a bunch of them. You may even have seen some of their **bad commercials** on TV. Most of the big names—GoDaddy, Network Solutions, Hover, etc.—allow you to register domains with the majority of TLDs, including those TLDs managed by countries.

In general, domains aren't property that you own outright, like real estate. Instead, domains are more like an office you lease from a property owner, but with the right to renew the lease in perpetuity. The rights to the domain can be transferred and sold, but ultimately someone has to regularly pay the registrar for the right to the domain; otherwise, it will be released back to the pool of possible domains, where someone else can buy it.

The cost of the domain depends on the domain's TLD, and most registrars offer discounts on the yearly cost if you pay for a number of years at once. You'll also have to provide a bunch of required information when you register the domain, as the registrars are required by law to know who owns domains, and they have to make this public information. Most registrars offer a privacy service that will keep your personal information off of the Internet, though.

The domain registrar we recommend is [Hover.com](#), which has good customer service, reasonable prices, an intuitive interface, and a free domain privacy service (Figure 1.9).⁴ It's also nice that, unlike some [other registrars](#), Hover doesn't try to [nickel and dime](#) its customers with unnecessary services.

⁴Image retrieved from <https://flic.kr/p/89UQJ> on 2019-11-20. Copyright © 2006 and used under the terms of the [Creative Commons Attribution 2.0 Generic](#) license.



Figure 1.9: Totally not getting paid to say that (seriously).

1.3.1 What to register?

You've probably noticed that most global services and corporations use **.com**, which is generally considered to be the most desirable TLD. This TLD was created for general commercial activity by **ICANN** ("EYE-can"), the organization responsible for managing Internet names.⁵

So, which TLD should you pick? Well, a lot of that depends on what you want to do online, who your audience is, and the availability of the domain name that you want. If you are a big enough company or have enough name recognition in other media, you could just pick any TLD that has a domain name available that incorporates your established name. Unless you have a good reason to choose otherwise, a **.com** domain is probably the way to go, though there are a lot of different possibilities:

- There are *generic* TLDs, like **.com**, **.info**, **.net**, and **.org**, which can be used for general purposes.
- There are *generic restricted* TLDs, like **.biz**, **.name**, **.pro**, where there are rules on how domains registered to the TLD can be used.
- Each country has its own TLD, such as **.us** (United States), **.uk** (United Kingdom),⁶ **.it** (Italy), **.ly** (Libya), **.co** (Colombia), and **.io** (Indian Ocean territories). The **.io** domain has become popular in the tech industry, especially for a **quality domain name** whose **.com** equivalent is already in use or (especially) when it is being held hostage by a **domain squatter**.⁷
- There are the current *sponsored generic* TLDs, like **.aero**, **.asia**, **.cat** (not about cats but rather the **Catalan** linguistic and cultural community), **.coop**, **.edu**, **.gov**, **.in**, **.jobs**, **.mil**, **.mobi**, **.tel**, **.travel**, and **.xxx** (which can be

⁵ICANN was originally under the control of the US Government (via the Department of Commerce), but since 2016 it has been free from US Government oversight.

⁶The UK actually has its own sub-TLD, so many UK sites end in something like **.co.uk**.

⁷Technically, **domain squatting** refers only to the bad-faith registration of domains, such as those corresponding to known trademarks. In practice, though, the supposedly benign practice of "**domaining**" (buying up domains for future resale but doing nothing useful with them in the interim) is so annoying that many people use the term "squatter" for those who engage in either practice.

used only for purposes within a particular industry that you can probably guess).

- Finally, there is the “new generic TLD (gTLD) program”, which seeks to add an unlimited number of gTLDs (i.e., non-country TLDs). Any person or organization that wishes to have their own TLD can pay a hefty fee and apply to have their own TLD—some examples are things like .xyz, .ninja, .limo, as well as a large assortment of TLDs that use writing systems other than the Latin alphabet. Most of these new TLDs will be for entirely generic use, and some will be sponsored generic TLDs that can be used only for a specific purpose.

You should know, though, that many people consider these new domains, and any business that uses one, to be less reliable than the old ones, probably because they aren’t as familiar. Some of the country TLDs have become pretty well-accepted, such as [.io](#) mentioned above, though they often still cause confusion for less tech-savvy people. The problem is that if you tell a non-tech-savvy person to go to a site like <http://bit.ly>, you might find that they try to go to <http://bit.ly.com>. The key takeaway here is that you need to know your audience: you don’t want a domain that scares off potential users, you don’t want a TLD with lots of bad sites, and you don’t want a name that confuses people.

With those considerations in mind, let’s get back to registering a domain. The first step is to visit the [Hover.com](#) home page. At this point, you’ll have to choose a domain. For this tutorial, we’ll use [codedangerously.com](#), but of course you’ll need your own (since we already registered that one!).

If you don’t already have an idea for a domain name, there are actually a lot of free tools that are specifically designed to help you come up with one. They generally allow you to enter a bunch of different words or phrases to generate suggested domain names that are available to be registered. A couple of popular sites that we’ve used are [instantdomainsearch.com](#) and [domainr.com](#). Or you can just use the domain registrar itself, which is what we’ll do here.

On Hover, finding out if the domain name is available is as easy as entering the name in the giant input field on the home page. At this point, Hover

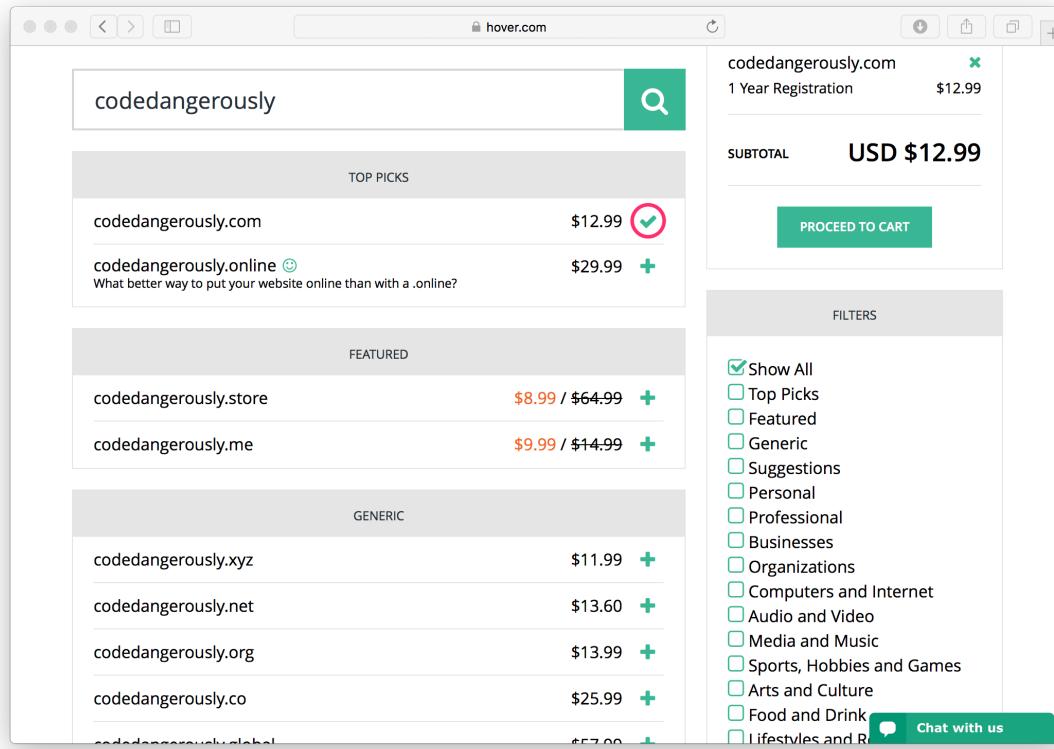


Figure 1.10: The different TLD options that are available to register.

will give you a long list of TLDs for which your name of choice is still available (Figure 1.10), as well as suggested domains containing word substitutions (farther down the page).

If you see that the domain name you want is available, and on the TLD you are looking for, add that domain to your cart!

You might have noticed that even when you find a domain name that you like, there are often a lot of very similar variations available on different TLDs, as well as domains that are close in spelling to your domain. If you have the extra money to spend, you can buy as many misspelled domains as you want and then permanently redirect users to the appropriate domain. For example, we own the domain [railstutorial.org](#)—which contains a common misspelling of “tutorial”—and automatically redirect it to [railstutorial.org](#). (As noted in Sec-



Figure 1.11: Your very own fantasy castle floating in the Internet sky!

tion 1.2, we'll look at how to do domain forwarding ourselves in Section 2.4.1.)

When you've selected all the domains that you want to buy (just the .com is probably fine), go ahead and go through all the checkout steps, including providing the necessary contact info for your domain registration. Thanks to Hover's free anonymization service, you don't have to worry about this information being made public.

If you've made it through all these steps, congratulations! You now own a domain name on the Internet. You are the proud owner (or at least long-term leaser) of virtual digital property (Figure 1.11).⁸

⁸Image retrieved from <https://flic.kr/p/AkF7w> on 2019-11-20. Copyright © 2007 jimthompson and used under the terms of the [Creative Commons Attribution 2.0 Generic](#) license.

1.3.2 You've got a domain, now what?

Once you've registered a domain, the next step is to set up the *DNS records* for your domain and edit the so-called *nameservers*. The nameserver property for your domain tells the DNS system exactly which server should be queried to find the domain details, such as which subdomains exist, which URLs should be forwarded, etc.

Right after you purchase a domain, the nameservers are usually set to the registrar's servers, and in the case of Hover that would be something like Listing 1.3.

Listing 1.3: Typical Hover.com nameservers

```
ns1.hover.com  
ns2.hover.com
```

Although you would be perfectly fine using the default DNS controls to complete this tutorial, we are going to set up a third-party service that provides a bunch of additional useful features. Setting up this service will involve changing the nameservers on the registrar to point to the new nameservers so that they can be the first point of contact for domain names in DNS lookups. Making these changes, and configuring the result, is the subject of Chapter 2.

Chapter 2

DNS management

Chapter 1 ended by noting that custom domains come with default DNS name-servers (Section 1.3.2) but indicated that we’re going to use a third-party service instead. In this chapter, we’ll set up this service, called *Cloudflare* (Section 2.1), and then configure it to associate custom domains with static websites located at GitHub Pages (Section 2.2). As a bonus, we’ll show how to make the same association for dynamic web applications deployed at Heroku (Section 2.3).

While GitHub Pages and Heroku are only two of an enormous number of options for hosting websites, the principles we cover here are generally applicable to web deployment. After completing this chapter, you’ll have the knowledge needed to configure pretty much any other custom domain setup as well.

Important note: GitHub changed how GitHub Pages works during the summer of 2022, so if you have read this guide in the past there will be some changes.

2.1 Cloudflare setup

So what is Cloudflare, and why are we going to be using it?

[Cloudflare](#) is a service that sits between the Internet and a website to provide the following benefits:

- Fast DNS management using a nice interface

- SSL support for setting up a secure connection to our site
- *Edge caching* to make the site’s content easy to access across the world
- Protecting the site from anyone who might try to take it offline using a *Distributed Denial of Service* (DDoS) attack
- ...and the basic service is **free to use** (with inexpensive pay plans as your website’s needs grow)

It would probably be useful to quickly explain what these features are in case you aren’t familiar with them.

2.1.1 Cloudflare features

First, most registrars allow you to do DNS management, but the interfaces can be less than optimal (to put it generously). Additionally, when you make DNS record changes on a registrar (like pointing to a new server IP address), the changes often aren’t immediate. Rather, they have to propagate throughout the DNS system—a process which can take 24–48 hours to complete as your update bounces around the DNS network ([Figure 2.1](#)).

Because Cloudflare sits between your servers and the world while passing web requests through, any server changes stay internal to the Cloudflare service. To the outside world, all requests still go first to the Cloudflare IP address. This means that the propagation of DNS changes (like subdomains and redirects) is nearly instantaneous—a feature that can be critical when administering a working site.

The next feature on the list, SSL (or [Secure Sockets Layer](#)), is a method for encrypting web traffic from a user’s computer to the server ([Figure 2.2](#)).

You’ve actually been using SSL for years, possibly without even knowing it: every time you go to a site with a little lock next to the site’s domain name, you’re using SSL ([Figure 2.3](#)).¹

¹The Internet is a dangerous place, and you can’t ever be certain that you really are safe, but SSL increased the chances considerably.

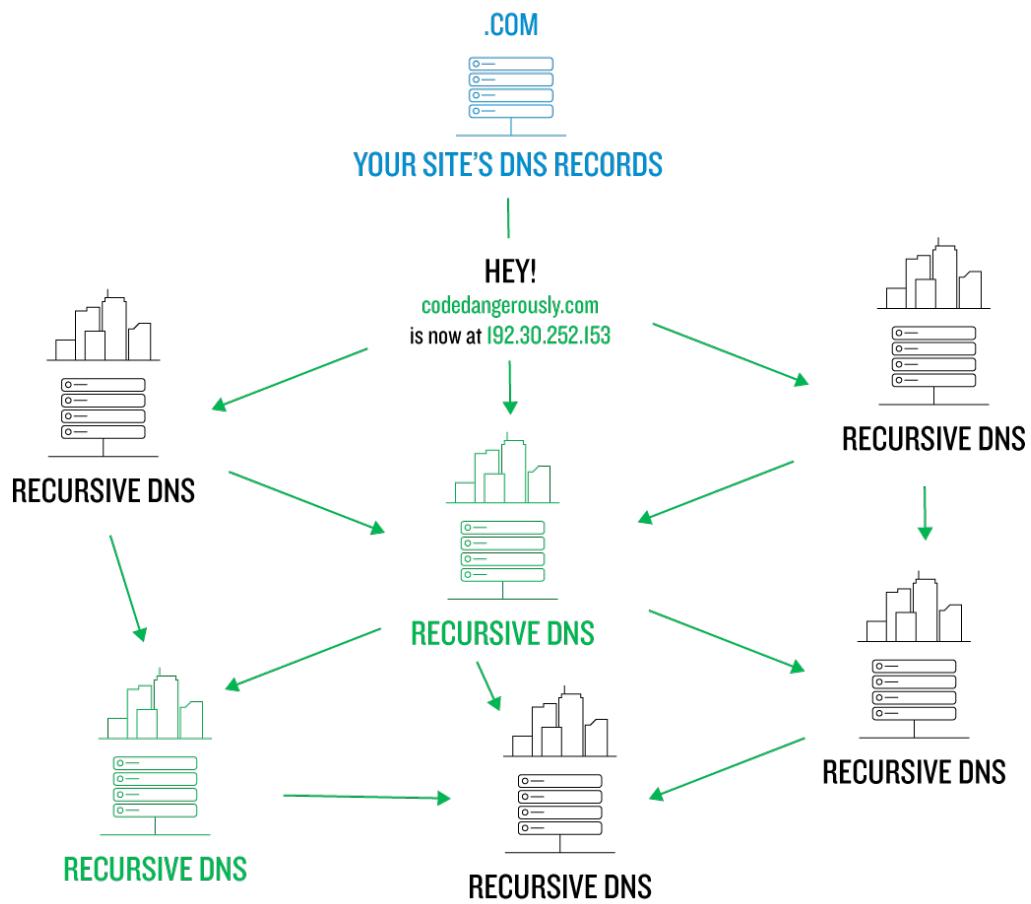


Figure 2.1: Potentially slow DNS propagation.

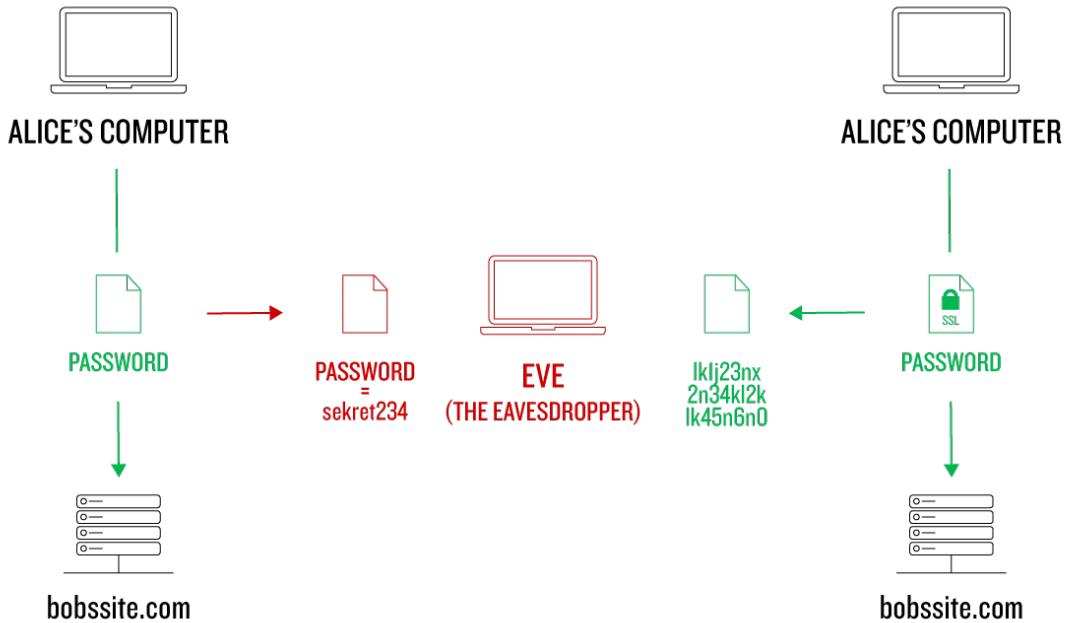


Figure 2.2: Encrypting data using SSL.

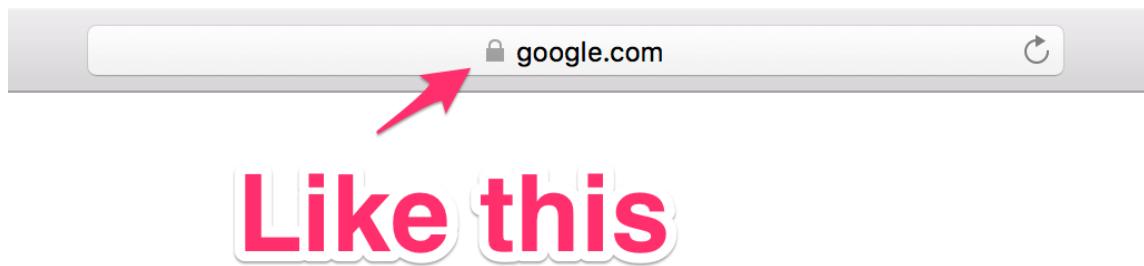


Figure 2.3: An SSL lock icon.

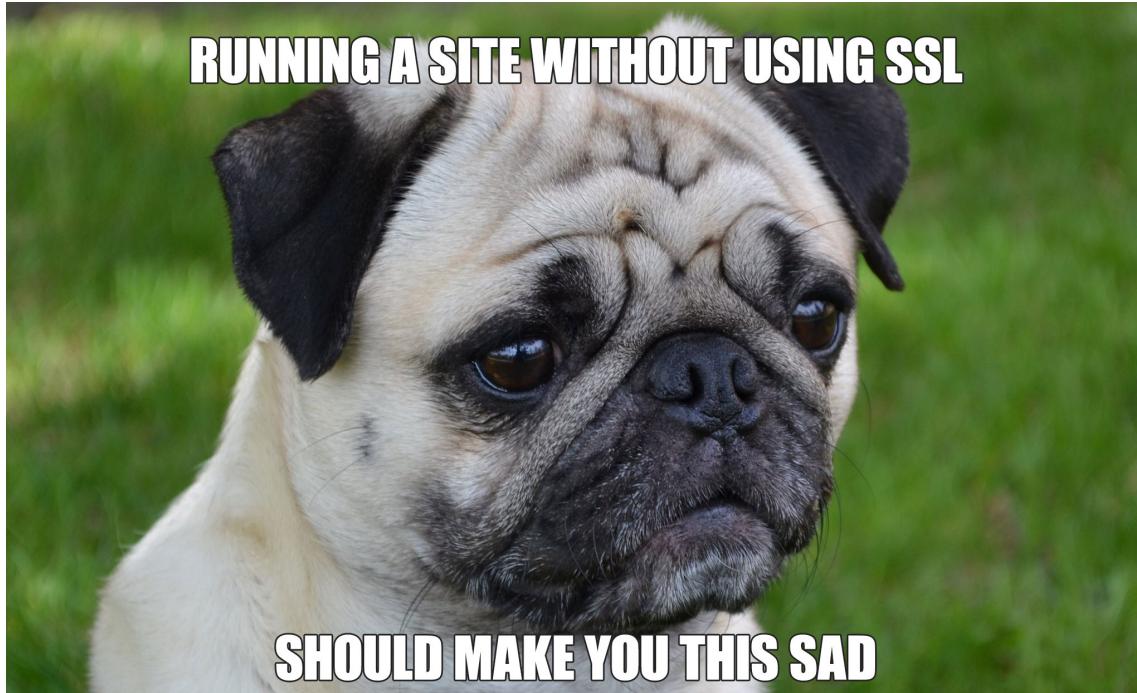


Figure 2.4: Don't do it! Bad dog!

SSL works by turning the packets of information sent between your browser and a server from plain old text into a jumble of letters and other characters. This makes it virtually impossible for someone to intercept the traffic and see what is being sent back and forth. In many cases, there is a lot of ceremony and pain associated with setting up SSL for a site, but when you use Cloudflare everything is already set up for you for free. Because non-SSL sites are so insecure, you shouldn't run a site these days without SSL—not even a static site like a blog (Figure 2.4).²

Next up on the Cloudflare feature list is *edge caching* (Figure 2.5). Edge caching involves automatically saving your site's content on servers around the world, allowing users to access the content from the server nearest to them (which speeds load times).

²Image retrieved from <https://flic.kr/p/2g1Yi5g> on 2020-02-18. Copyright © 2019 DaPuglet and used under the terms of the [Creative Commons Attribution 2.0 Generic](#) license.

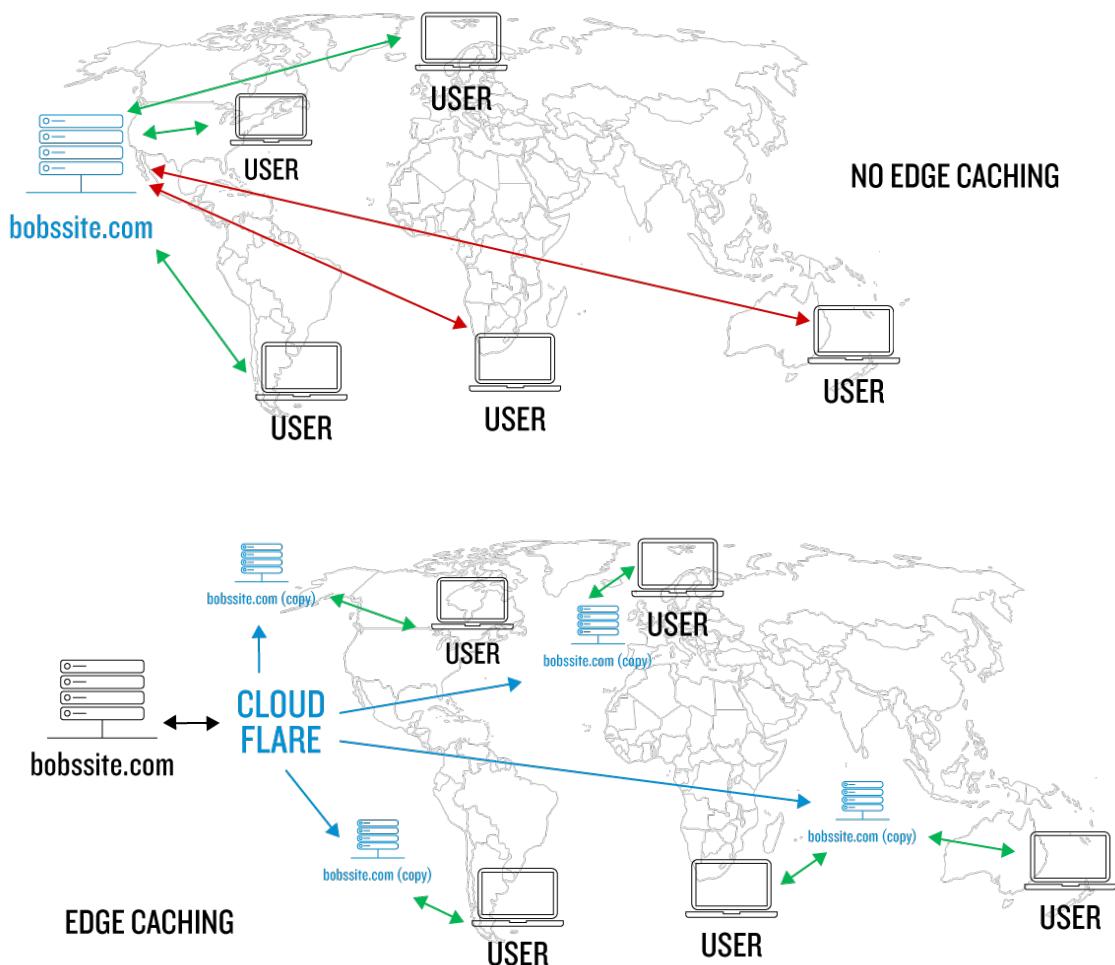


Figure 2.5: Optimizing load times with edge caching.

Lastly, Cloudflare protects the site from **DDoS** attacks. You might have heard of these on the news in reference to some group being accused of “hacking” another organization. DDoSes are online attacks that try to make entire sites inaccessible by flooding their servers with fake requests—imagine millions of fake users that all ask a server for a site at the same time.

There are a limited number of requests that a single server can handle at once, so if there are enough fake users making requests, real users won’t be able to get to the site. Cloudflare has the ability to notice when this type of attack is happening and can filter out the fake users to keep your site online and available. Again, you don’t have to configure anything to get this benefit.

All this sounds pretty good, right? No one wants to have their site overrun by computer zombies ([Figure 2.7](#)).³

These days, our first step after registering a domain is immediately setting up Cloudflare. The benefits are great, and unless you have a fairly complicated application the service is free.⁴ What’s not to love? Let’s get started.

2.1.2 Cloudflare signup

The first step in using Cloudflare (just like pretty much any service on the Internet) is to create an account. One added benefit, though, is that the signup process also walks you through adding a domain to Cloudflare. So, head on over to cloudflare.com and click the Sign Up button to get started with the process.

2.1.3 Connecting registrar nameservers

After you’ve created your account, Cloudflare automatically kicks you into the first step to add a domain to their service. Enter your newly purchased domain in the box and press the Add Site button, and then on the next screen select the

³Image retrieved from <https://flic.kr/p/9QeRNF> on 2020-02-18. Copyright © 2011 Marilyn Sherman and used under the terms of the [Creative Commons Attribution 2.0 Generic](#) license.

⁴We use the professional, paid tier for more complicated sites like learnenough.com and railstutorial.org, which is but one of many indications that Cloudflare’s [freemium](#) business model is working well.

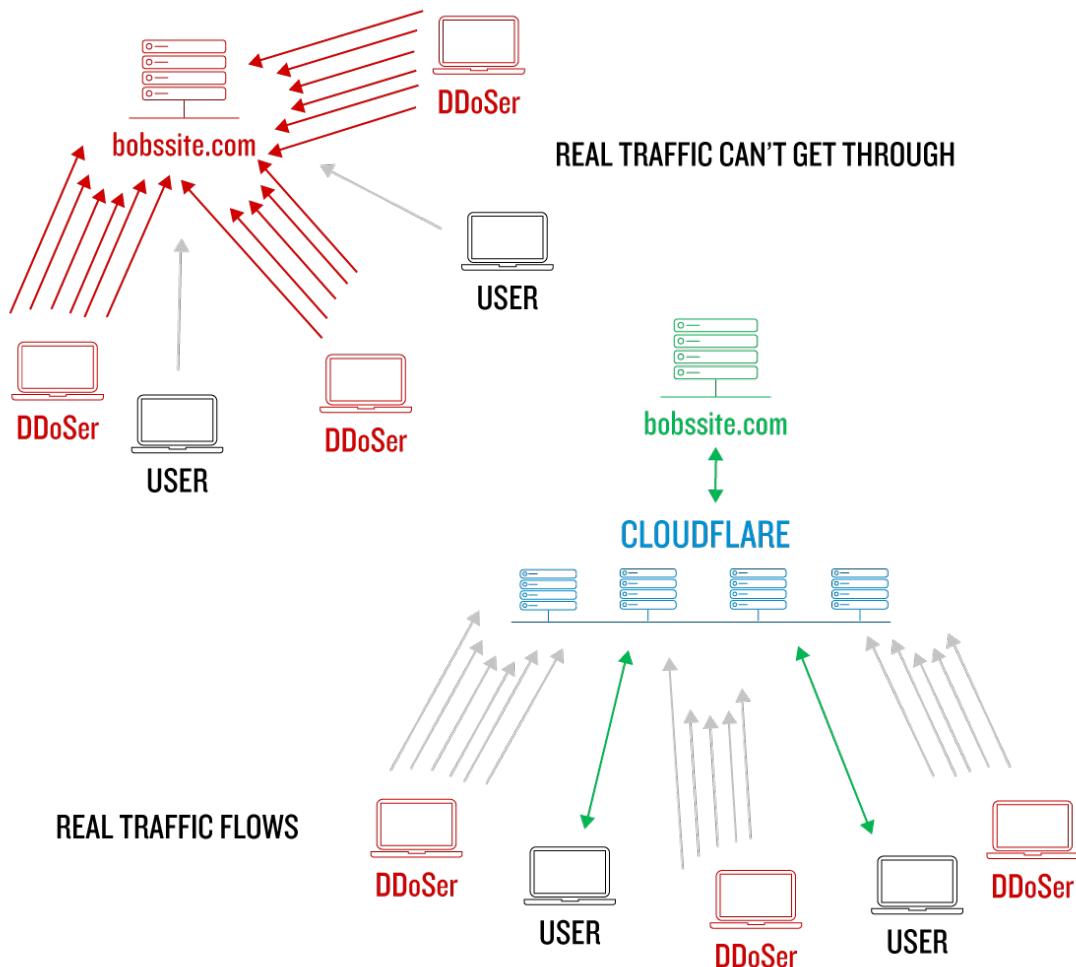


Figure 2.6: Protecting against DDoS attacks.



Figure 2.7: Or real zombies using computers—really any zombies are bad.

Type	Name	Value	TTL	Proxy status
A	*	points to 64.98.145.30	Automatic	
A	codedangerously.com	points to 64.98.145.30	Automatic	
A	www	points to 64.98.145.30	Automatic	
CNAME	mail	is an alias of mail.hover.com.cust.h...	Automatic	
MX	codedangerously.com	mail handled by mx.hover.com...	Automatic	

Figure 2.8: Our domain's DNS settings report.

free plan for this tutorial. When you confirm the plan selection Cloudflare will start adding your domain to their system.

It should take less than a minute for Cloudflare to pull in the domain records, and when it's done the next page will show you the result of the scan from the last step: a report with all of the DNS records associated with your domain.

As you can see in Figure 2.8, there are a few A records (mentioned briefly in Section 1.2), which are set to flow through Cloudflare (the orange cloud icon). There is also a so-called *CNAME record*, which is set up to skip the Cloudflare service (the gray cloud icon). (We'll discuss more about these record types in Section 2.2.1.)

Since we just purchased this domain, the current records aren't very com-

plicated, and there's nothing to configure on this step. However, if you were importing a domain that had a bunch of custom settings, this would be a good place to carefully look over the records before hitting "Continue".

At this point, you'll be taken to the last step of the sign up and domain-addition process: the nameserver configuration. On this page, you'll see a listing of the current nameservers for your domain (the ones from Hover mentioned in [Section 1.3.2](#)), and the new Cloudflare servers that you need to set for your domain. For the account we created, the nameservers appear as follows:

```
vern.ns.cloudflare.com  
zara.ns.cloudflare.com
```

Cloudflare has many different nameservers, though, so your results may vary.

At this point, you'll have to switch back to Hover (either in a different window or browser tab) and then click the Edit link in the header of the nameservers box. Switch back to the Cloudflare tab, and copy over the new nameservers supplied by Cloudflare to replace the old ones back back on the Hover tab. Save your changes and you should see your new nameservers on the box on the overview page for your domain.

Now, switch back to the Cloudflare window or tab and press the "Done, check nameservers button", and then on the next page that pops up go ahead and click the button to do the quick start guide.

Make sure that the Automatic HTTPS Rewrites and Always Use HTTPS options are enabled (you can ignore the other options). Click Finish to save the changes and you'll be taken to the overview page for the domain. If you're using a Rails app at Heroku ([Section 2.3](#)), be sure to set the `force_ssl` parameter to `false` as shown in [Listing 2.1](#).

It shouldn't take too long for your domain's nameserver change to resolve; you can hit the button to recheck the nameservers page every so if you'd like to force Cloudflare to check again. When the updates are public on your registrar, your domain will be active on Cloudflare...

So, congratulations!

But what on Earth did we just do? Well, now all page or DNS requests that are directed at your domain will first go through the Cloudflare service, which

means that we can change the way those requests are handled using Cloudflare’s user interface. For example, we’ll be able to set up subdomains (Section 2.2.1), and URL redirects (Section 2.4.1).

2.2 Custom domains at GitHub Pages

In this section, we’ll explain how to connect a custom domain to GitHub Pages. This is a great (free!) option for hosting static websites on custom domains.⁵ (Indeed, the [Learn Enough blog](#) runs on a custom domain at GitHub Pages.) It’s also a good exercise even if your primary interest is hosting dynamic websites (Section 2.3), since many of the principles are the same, so we suggest completing this section in any case.

If you don’t have a GitHub account, you should follow the steps from Chapter 2 of [Learn Enough Git to Be Dangerous](#) at this time. If you don’t already have an example website at GitHub Pages, we suggest following the steps in Chapter 1 of [Learn Enough HTML to Be Dangerous](#) to create and deploy a simple “hello, world!” app. ([Learn Enough CSS & Layout to Be Dangerous](#) develops a much more sophisticated website using the same basic technique.) In what follows, we’ll assume you’ve set up such a site (or equivalent).

2.2.1 Configuring Cloudflare for GitHub Pages

Our first step is to tell Cloudflare that our site is located at GitHub Pages. To get started, click on the DNS menu item to go to the DNS settings for your domain. The result should be a list of DNS records that looks something like Figure 2.9.

So, what are these different types of records? We’ll go in order and make the changes needed to get our site running off of GitHub Pages.

The *A records* that you see are also known as *address records* (Section 2.1.3); the top A record is known as an *apex record*. These settings are usually used to

⁵Here *static* refers to static server-side assets like HTML, CSS, JavaScript, etc. The site’s behavior itself still might be dynamic, typically due to JavaScript executing on the user’s browser. An example of a statically generated site yielding dynamic client-side behavior appears in the image gallery developed in the final chapter of [Learn Enough JavaScript to Be Dangerous](#).

The screenshot shows the Cloudflare dashboard for the domain `codedangerously.com`. The sidebar on the left lists various services: Overview, Analytics, DNS (selected), Email (Beta), SSL/TLS, Firewall, Access, Speed, Caching, and Workers. The main content area is titled "DNS" and displays a message: "A few more steps are required to complete your setup." Below this is a section titled "DNS management for `codedangerously.com`". It features a search bar, an "Advanced" button, and a "Add record" button. A table lists the current DNS records:

Type	Name	Content	Proxy status	TTL	Actions
A	*	64.98.145.30	DNS only	Auto	Edit
A	<code>codedangerously.c...</code>	64.98.145.30	Proxied	Auto	Edit
A	www	64.98.145.30	Proxied	Auto	Edit
CNAME	mail	mail.hover.com.cust.h...	Proxied	Auto	Edit
MX	<code>codedangerously.c...</code>	mx.hover.com.cu...	DNS only	Auto	Edit

At the bottom, there is a section titled "Cloudflare Nameservers" with the instruction: "To use Cloudflare, [change your nameservers](#), or authoritative DNS servers. These are your assigned".

Figure 2.9: How your DNS settings should look before we start making changes.

define the way that the *root domain* (e.g., `codedangerously.com`, without the `www.`) is handled, and they should always point to a valid IP address.

In the default settings transferred over from Hover, the `www` subdomain is set up with an A record that also points to a server's IP address. As was briefly mentioned in [Section 1.2.1](#), there are a lot of technical reasons why it's a good idea to use a `www` subdomain, but the easiest way to describe why you should direct traffic there, and not to the root domain, is that down the road it gives you much more flexibility in how you deal with the traffic coming to your site. For our example site, we are going to get rid of this A record for the subdomain, and instead add a CNAME record for the `www` subdomain.

Let's start the cleanup. Go ahead and click the Edit on the right-hand side of the screen for the `www` DNS record, and delete the record. (If you are using a registrar that didn't automatically add an A record for `www`, you don't need to do anything at this point.)

Depending on which registrar you used, you may or may not also have an A record with an asterisk `*` in the name field (the first field in [Figure 2.9](#)). That's a wildcard record that handles requests to any random subdomain of your URL, which allows random requests like `blarglebargle.codedangerously.com` to pass through to the server. But notice that there is no little orange cloud icon by the `*` record in [Figure 2.9](#)—that means the traffic to those random addresses is not being handled by Cloudflare, and will instead flow straight to the server, which we definitely don't want.

To prevent this unwanted behavior, remove the wildcard A record, as well as the CNAME record pointing to `mail.hover.com.cust.hostedmail` (we'll set up a different mail system in [Chapter 3](#) that won't need a CNAME record).

GitHub Pages IP addresses

Now, to have our site be served from GitHub Pages, we need to point the A record to the IP addresses of GitHub's servers instead of to the current address (yes, you could have just edited the A records, but we wanted to step you through removing and adding things from scratch). You can find the addresses in the [GitHub Pages documentation](#); at the time of this writing, the IP numbers for the GitHub Pages servers are as follows:

```
185.199.108.153  
185.199.109.153  
185.199.110.153  
185.199.111.153
```

As we noted at the beginning of this chapter, there have been some changes recently to the GitHub Pages service and to accomodate the changes we will have to do DNS setup a little differently. GitHub now issues an SSL/TLS certificate for sites hosted by GitHub Pages and that process can be blocked if the DNS records that point our new domain to GitHub are first proxied through Cloudflare. Once all the setup has been completed, and GitHub has issued encryption certifications, we'll be able to turn the proxying back on to get the benefits of Cloudflare's features.

Where there was a single A record before, we are going to add a second backup server address (you don't want your site to be unavailable because of a server issue that you can't control, do you?). To change the IP address for the existing record, just click edit and then paste the new IP address into the field. If you are currently setting up a GitHub Pages site, make sure to disable the proxy. Once everything has been updated click "Save" to update the record.

Now let's add the second record. Above the list of records, you'll see a button that is labeled "+ Add record". Clicking that causes Cloudflare to reveal a row of inputs. Make sure that the dropdown menu is set to **A**, and then add **yourdomain.com** in the Name field and the second GitHub IP number (e.g., **185.199.109.153**) into the address field (Figure 2.10), and once again disable the proxying if this is for GitHub Pages. Then click "Save" to finish.

If you see two A records with GitHub IP addresses in the list, good job! You did it!

CNAME records

Let's now take a look at CNAME records, which are *canonical name* records. They are a type of DNS record that allows you to point traffic to any domain you want (such as a URL like **codedangerously.com**). This is in contrast to A records, which point only at IP addresses (like **185.199.108.153**).

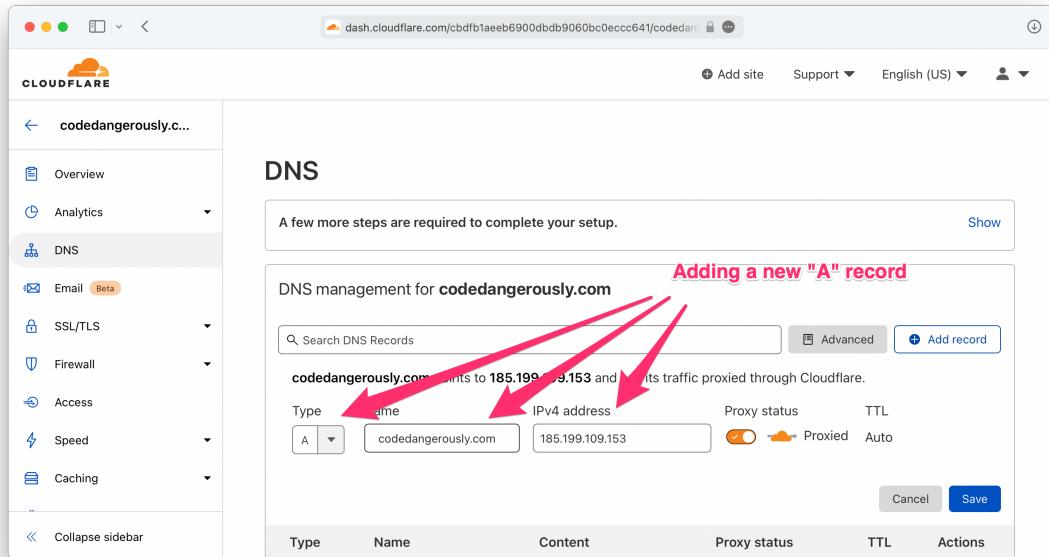


Figure 2.10: Adding the second alternate GitHub Pages server IP.

CNAMEs are often used to create [aliases](#) on sites so that visitors who go to a subdomain like `help.codedangerously.com` are instead permanently redirected to, say, `docs.codedangerously.com`. In our case, we want traffic that comes to our site using a `www` subdomain—i.e., `www.codedangerously.com`—to see the site’s homepage, which we can do by adding a CNAME record called `www` that points to the root domain.

To start, click the “+Add record” once again if the add record input fields aren’t showing (if you never left the Cloudflare page, they should still be visible). To add the CNAME record, click the dropdown at the beginning of the new record input form and select CNAME from the list. Then add `www` in the “Name” field and add your site domain in the “Target” field.

Hit the “Save” button and you’re done... or you would have been if GitHub’s service hadn’t changed. We have left this here because this is how you would configure the `www` subdomain if you are using any other service.

For GitHub pages though, we need to change the target for our subdomain CNAME record. Head back over to the [GitHub Pages documentation](#) page,

and if you scroll down a little past the IP addresses that we used earlier, you’ll see a new section regarding the `www` subdomain. One aspect of the changes that GitHub has made is that they are now handling the redirect from the root domain to the `www` subdomain for sites, and as a part of that they now require that the target for the `www` subdomain be `yourusername.github.io` if your repo is owned by a personal account. Or if the repo is owned by an organization then the target will need to be `organizationprofilename.github.io`. So head back to Cloudflare, add the correct target in the field, be sure to disable the proxy once again, and hit save.

There is one last step in the process required to make GitHub happy. On the Cloudflare sidebar menu, click on the “SSL/TLS” link to go to that settings page. In the first section, be sure to click the “Off (not secure)” option to disable any Cloudflare security. GitHub Pages seems to be happiest when there is nothing standing in the middle during the next step where Pages is enabled for your repo.

That’s all the DNS setup needed to get our new domain to work with GitHub Pages. A note: there is one last DNS record at the bottom that we haven’t altered, the *MX record*. That is a *mail exchanger* record that directs email traffic to your email server. We’ll be editing that record in [Section 3.1](#).

2.2.2 Configuring GitHub Pages

Now that we’ve configured Cloudflare for GitHub Pages, all we need to do is tell Pages about our custom domain. To get started, point your browser at your repository’s GitHub page, click on the Settings tab and then click the Pages option in the menu on the left ([Figure 2.11](#)).

Click on the dropdown menu that is likely showing the text None, and select the name of your project’s default branch. When we initially wrote this tutorial, the standard default branch name was `master`, but the current GitHub default is now `main` (although you can easily [change the default branch name](#) to `master` or any other name of your choice).⁶

⁶See the Learn Enough blog post “[Default Git Branch Name with Learn Enough and the Rails Tutorial](#)” for more information.

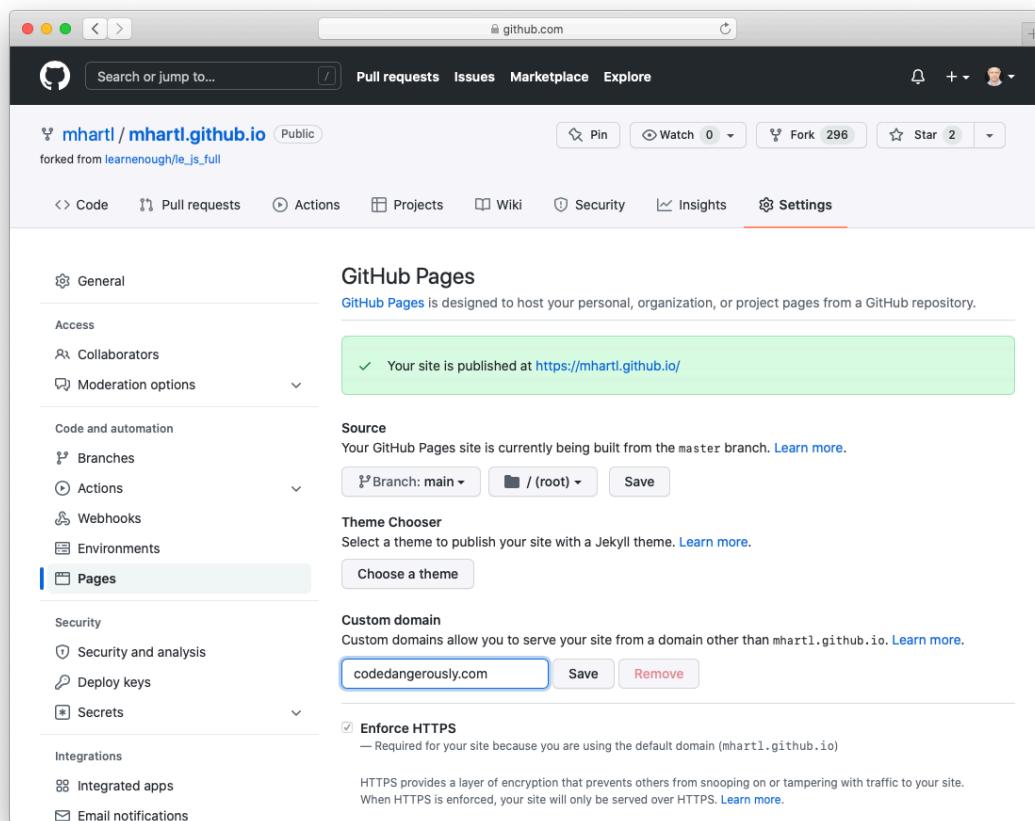


Figure 2.11: Settings for GitHub Pages.

After selecting the branch, add your site's domain name into the Custom Domain input field, including the `www`, and that's it! Depending on the internet weather it could take some time for GitHub to recognize the DNS changes that we've made, and in the meantime you might see a DNS error message or two, but eventually there will be a message that says that the DNS checks were successful and there should be a box letting you know that a SSL certificate has been provisioned. This whole thing might take 15–20 minutes, so have patience!

Now that GitHub issues their own certificate for your site, the whole setup is actually more secure than before and the web traffic from users' browsers, through Cloudflare, and to GitHub pages will be fully encrypted. If you open up a new tab or window in your browser and go to your custom domain, you should now see your GitHub Pages site ([Figure 2.12](#)).

At this point, if you are using GitHub Pages you can both this next section as well as [Section 2.4](#), but for any other service you will still need to create a page rule to redirect traffic from the root domain to the `www` subdomain. If you were to add this to a GitHub Pages site you end up creating a situation where there would be endless redirects as the Cloudflare page rule fights tooth and nail with the new GitHub page redirect that is trying to the same thing. Eventually your browser would throw up an error and there would be much sadness.

If however you have become a truly dangerous programmer and are referencing this section while setting up a site hosted by some service other than GitHub Pages or Heroku, first congratulations on being awesome, and second, there's only one significant task left, which is to arrange for our site to use only *canonical URLs*, a step that is covered in [Section 2.4](#).

Unless you want to deploy a web application to Heroku in [Section 2.3](#), you can skip right to [Section 2.4](#) at this time.

2.3 Custom domains at Heroku

In this section, we'll explain how to use a custom domain for a dynamic web application deployed to [Heroku](#). As one of many examples of such a site, consider... [learnenough.com](#) itself ([Figure 2.13](#))! That's right: the Learn Enough

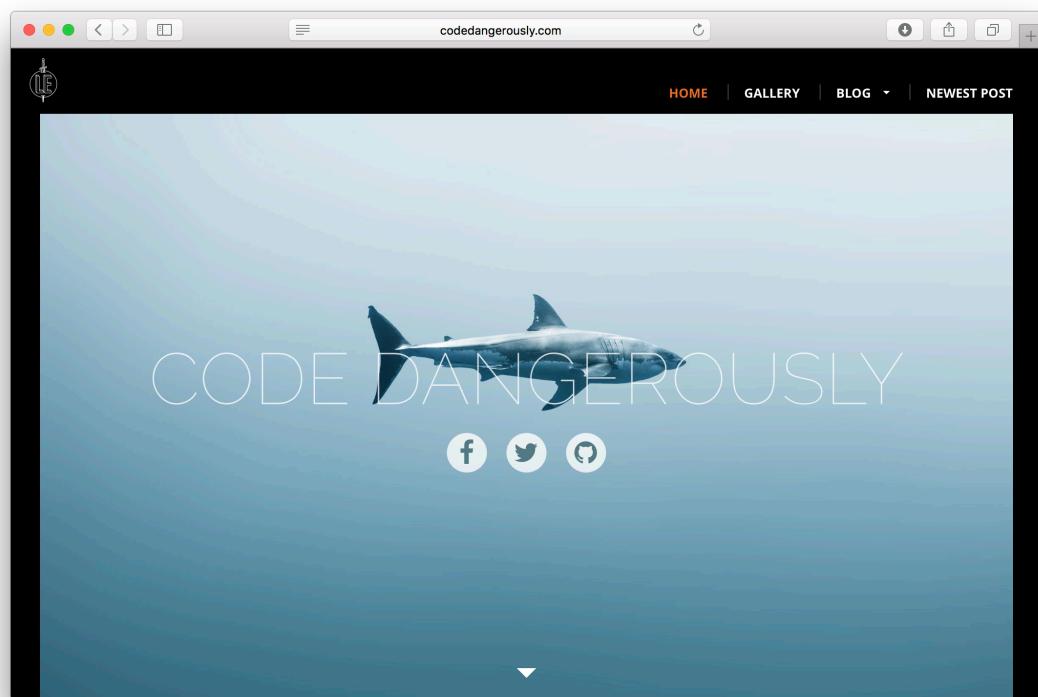


Figure 2.12: Can you believe all the steps it took to get to this place?

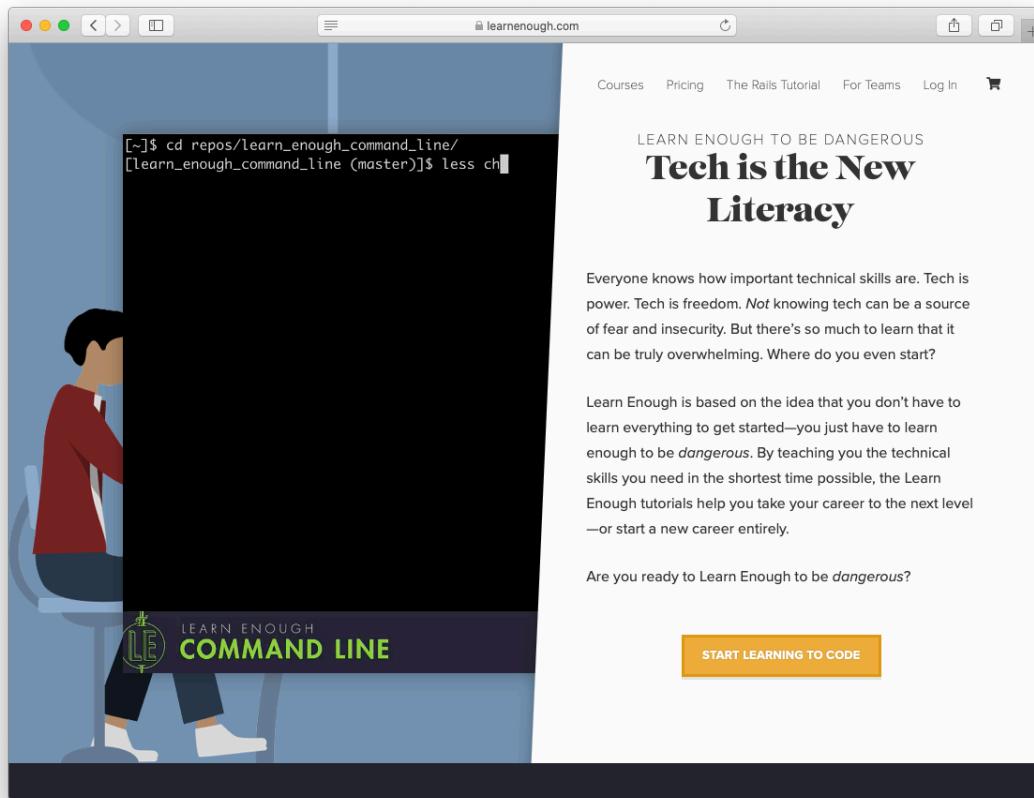


Figure 2.13: An example of a Heroku-deployed app running on a custom domain.

website is a Heroku app configured at Cloudflare and running on a custom domain.⁷ If you don't currently have an app running at Heroku, you can create one by following Chapter 1 of the *Ruby on Rails Tutorial*. If you don't want to follow these steps right now, you can skip to Section 2.4 without loss of continuity.

To get started configuring Heroku for a custom domain, log in to your Heroku account and open up the page for your application. Then click the “Set-

⁷The learnenough.com domain is actually registered at [Namecheap](#) rather than at [Hover](#), but the principles are the same.

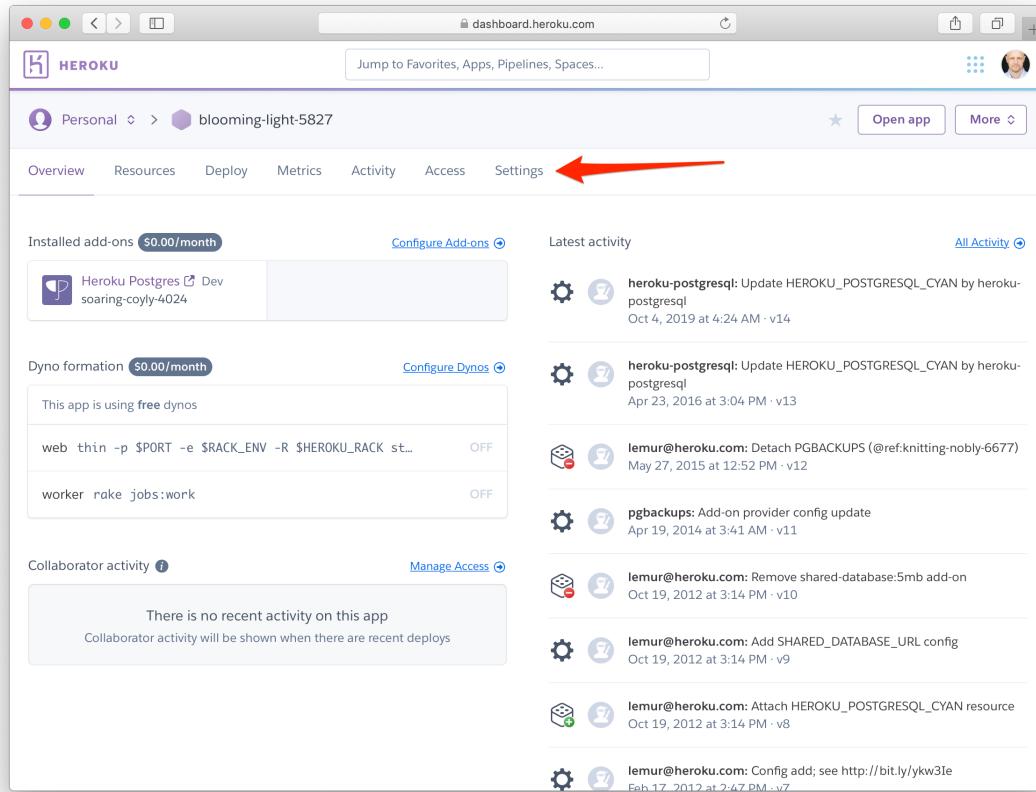


Figure 2.14: The Settings link for a Heroku app

tings” link in the navigation ([Figure 2.14](#)).

Next, scroll down to the “Domains and Certificates” section and click the “Add domain” button ([Figure 2.15](#)) to bring up the sidebar interface to add a new custom domain. (When you do this, note that it means you won’t be serving that URL from GitHub Pages any longer. To deploy both GitHub Pages and Heroku apps with custom domains, you can either use two different custom domains or use different subdomains.) You should enter the **www** subdomain version of your custom domain, which for us is **www.codedangerously.com** ([Figure 2.16](#)).

After clicking “Next”, Heroku will show you a unique “DNS target” URL,

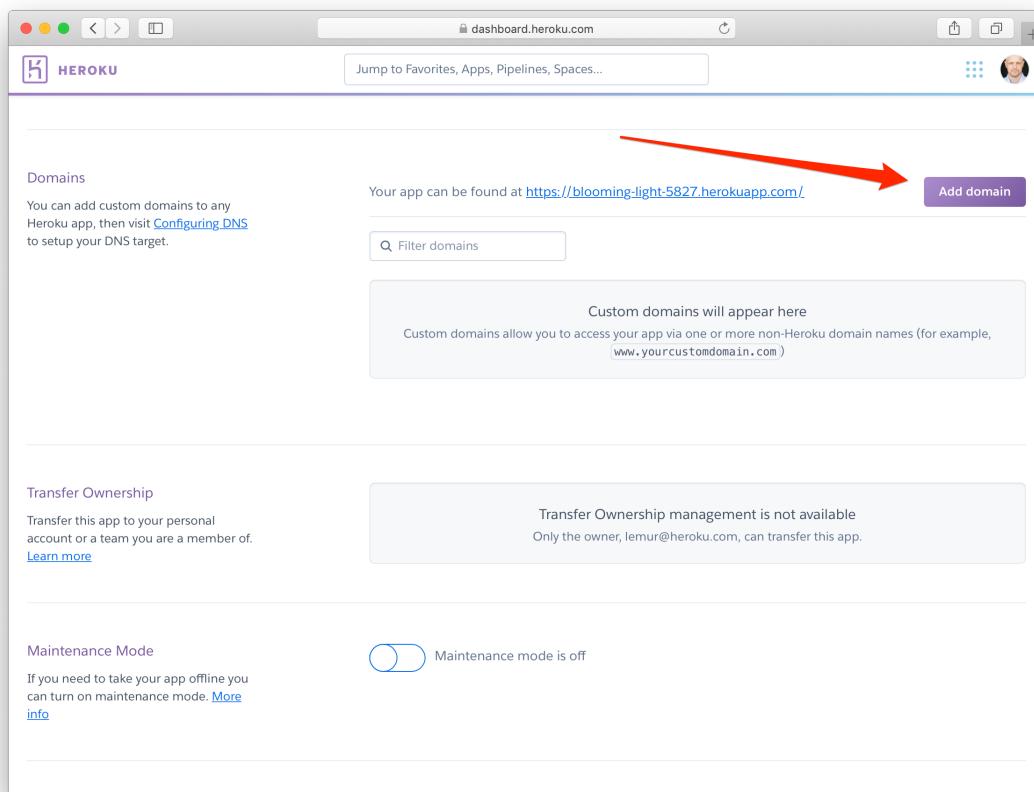


Figure 2.15: The link to add a Heroku domain.

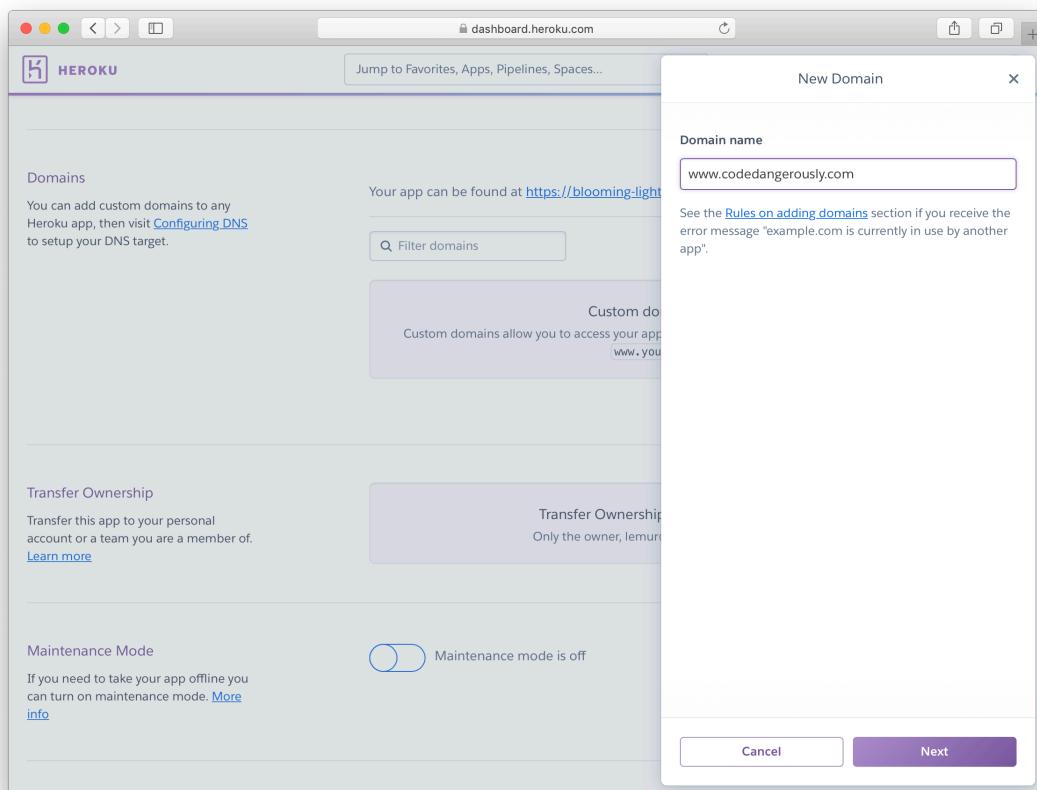


Figure 2.16: The interface for adding a Heroku custom domain.

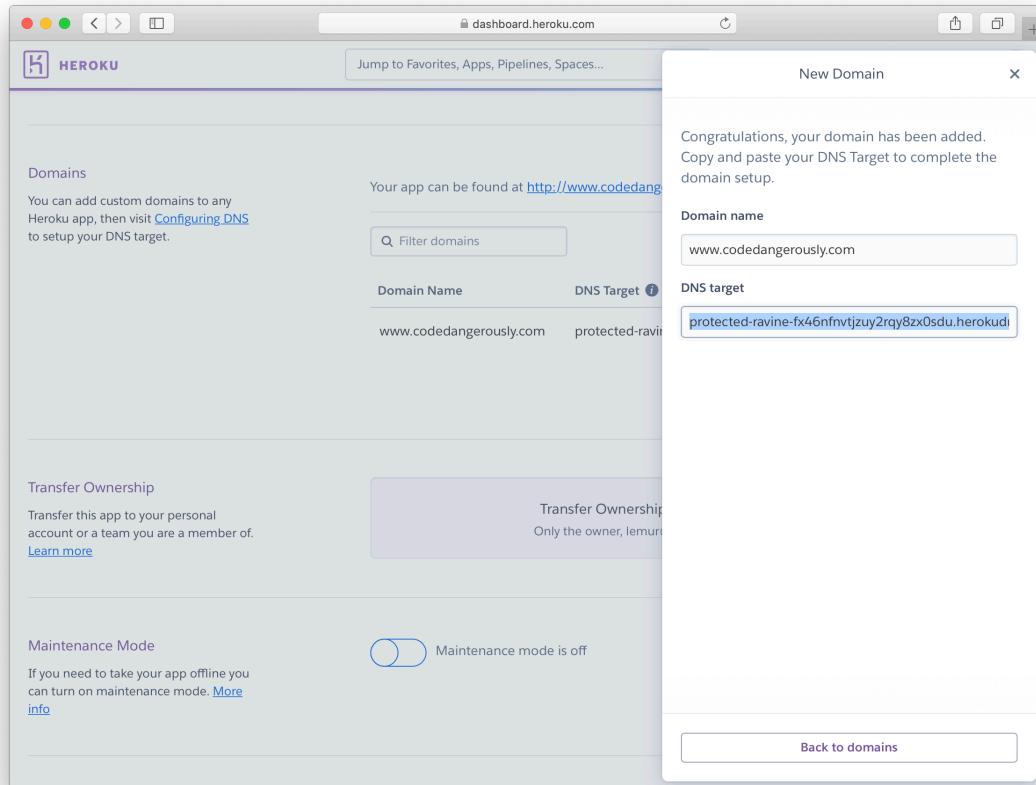


Figure 2.17: Copying the DNS target on the New Domain page

which you should copy for future use (Figure 2.17). A fragment of the URL appears on the domain dashboard (Figure 2.18); to see the whole string again, click on the pencil icon (circled in Figure 2.18) to bring up the Edit Domain page (Figure 2.19).

At this point, there's only one step left: tell Cloudflare about the Heroku configuration. When configuring Cloudflare for a Heroku-deployed site, you should *remove* the A records entirely and create CNAME records for both the root domain and for the `www` subdomain. The content of each record should be the DNS target from your version of Figure 2.17, which in our case is this:

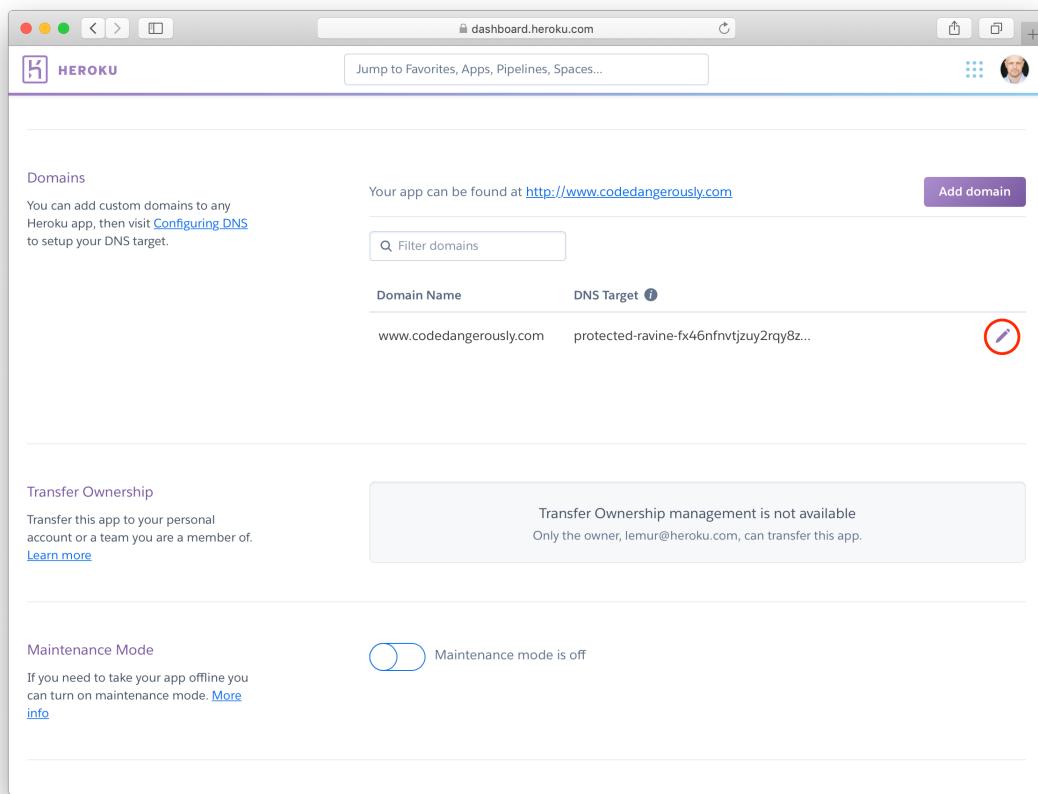


Figure 2.18: The Heroku dashboard with a custom domain.

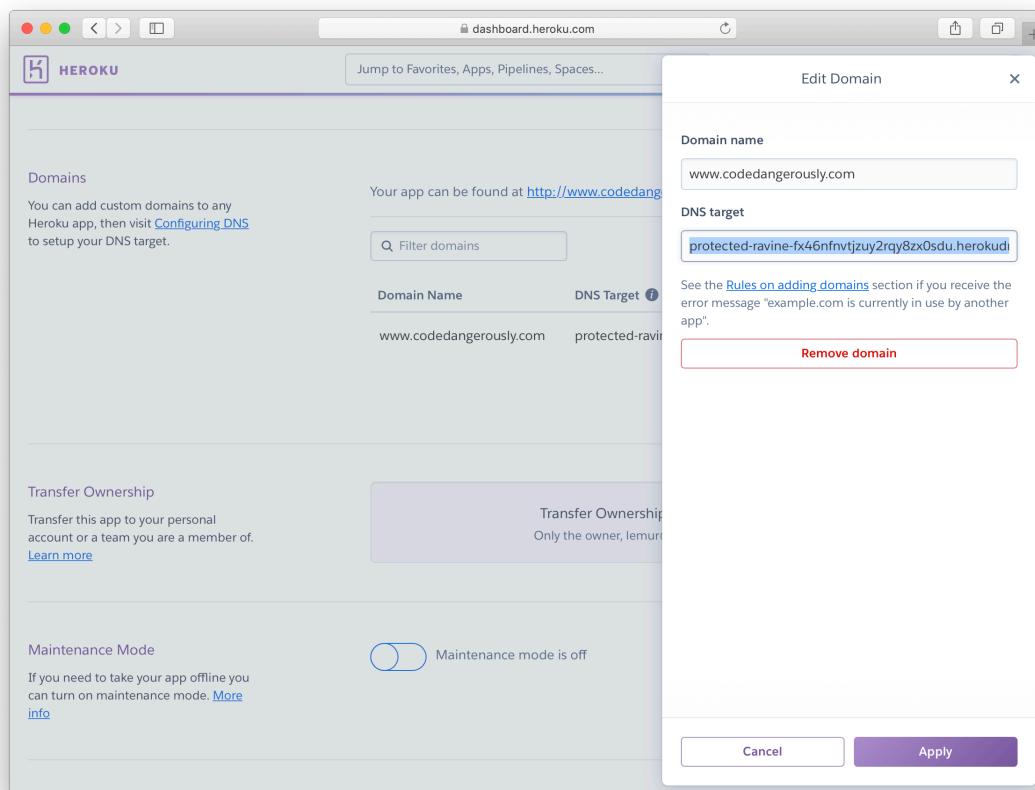


Figure 2.19: Viewing the DNS target on the Edit Domain page.

Type	Name	Content	TTL	Proxy status
CNAME	codedangerously.com	protected-ravine-fx46nfntjzuy2rqy8zx0...	Auto	Proxied
CNAME	www	protected-ravine-fx46nfntjzuy2rqy8zx0...	Auto	Proxied
MX	codedangerously.com	mx.hover.com.cust.hostedemail...	10	DNS only

Figure 2.20: Adding the Heroku domain target at Cloudflare.

`protected-ravine-fx46nfntjzuy2rqy8zx0sdu.herokuapp.com`

The Cloudflare DNS configuration for `codedangerously.com` is shown in Figure 2.20.

Because we added only the `www` version of our domain at Heroku, the root domain configuration won't actually resolve to the site. This is OK, though, because we're going to set up canonical URLs (Section 2.4) so that the root domain automatically redirects to the `www` version using a Cloudflare page rule (Section 2.4.1). The root DNS setting at Cloudflare is necessary for this redirect to work properly, but the bare domain `codedangerously.com` won't ever get accessed directly.

Note: If you’re using a custom domain to host a website created using [Ruby on Rails](#), you’ll need to configure your app *not* to force SSL, as shown in Listing 2.1, and then redeploy your app to Heroku ([git push heroku](#)). See Chapter 7 of the [Ruby on Rails Tutorial](#) for more information.

Listing 2.1: Configuring a Rails application to *not* to use SSL in production.

config/environments/production.rb

```
Rails.application.configure do
  .
  .
  .
  # Force all access to the app over SSL, use Strict-Transport-Security,
  # and use secure cookies.
  config.force_ssl = false
  .
  .
  .
end
```

2.4 Canonical URLs

Whether you’re using a static site as developed in [Section 2.2](#) or a dynamic web application as developed in [Section 2.3](#), there’s one detail left to finish: configure the site to use *canonical URLs* to ensure a consistent and secure experience for all users. (The word “canonical” has its origins in religious [canon](#), and is now used in a variety of [technical contexts](#) ([Figure 2.21](#))).⁸

Important note: GitHub changed how GitHub Pages works during the summer of 2022, so in case you missed this notice before, you can skip this section!

In our case, standardizing our site’s URLs to use a canonical form involves two principal tasks:

1. Ensure that connections to the site are always secure.
2. Ensure that all pages use the standard [www](#) form of the site rather than the root domain.

⁸Image retrieved from <https://flic.kr/p/dmDByU> on 2020-02-18. Copyright © 2012 Ozzy Delaney and used under the terms of the [Creative Commons Attribution 2.0 Generic](#) license.



Figure 2.21: This is the wrong type of cannon to be thinking of.

The first step is already complete since secure connections are enforced automatically by the Always Use HTTPS setting enabled in [Section 2.1.3](#). Completing the second step is the goal of [Section 2.4.1](#).

2.4.1 Cloudflare page rules

At this point, our website is properly configured to serve the home page from the address `www.codedangerously.com`. In order to avoid serving content from two different domains (which can complicate things like session cookies), we'd like to arrange for the `www` version to be the *only* address for the home page for our website. In particular, we want any traffic pointed at the root domain `codedangerously.com` to be automatically redirected to `www.codedangerously.com`. We can accomplish these kinds of redirects using Cloudflare *page rules*.

Page rules are a powerful and flexible tool, and we'll barely be scratching the surface here, but performing a redirect is one of their most common and important applications. To get started, click on Rules > Page Rules on the main

Rules

Page Rules

[← Back](#)

Create a Page Rule for codedangerously.com

If the URL matches: By using the asterisk (*) character, you can create dynamic patterns that can match many URLs, rather than just one. All URLs are case insensitive. [Learn more](#)

codedangerously.com/*

Then the settings are:

Forwarding URL

301 - Permanent Redirect

https://www.codedangerously.com/\$1

You cannot add any additional settings with "Forwarding URL" selected.

[Cancel](#)

[Save as Draft](#)

[Save and Deploy](#)

Figure 2.22: Creating a new page rule.

Cloudflare menu, and then click the Create Page Rule button to open the Page Rule interface ([Figure 2.22](#)).

One thing we could do is to forward *everything* at the root domain to the **www** equivalent:

codedangerously.com -> www.codedangerously.com

There's one problem, though: what if someone hits a subpage of our site? Rather than do this:

```
codedangerously.com/blarg -> www.codedangerously.com
```

We'd rather do this:

```
codedangerously.com/blarg -> www.codedangerously.com/blarg
```

This way the root and `www` domains will truly be equivalent, and *any* attempt to access the site via the root URL will end up on the corresponding `www` version.

We can arrange for this behavior using a *wildcard* (a concept mentioned briefly in [Section 2.2.1](#)), which will dynamically match whatever the user enters after the domain name, and keep it intact after the redirect. So if they enter `codedangerously.com/blarg`, the page rule will make sure they are forwarded automatically to `www.codedangerously.com/blarg`. We don't want to lose that `/blarg` at the end of the address!

To configure Cloudflare to use a wildcard and make a redirect, we'll start with the interface shown in [Figure 2.22](#). In the big input field, add your domain name, a forward slash, and then an asterisk to represent the wildcard—in our case, `codedangerously.com/*`. Then click the “+ Add a Setting” link, and in the dropdown choose “Forwarding URL”. You'll now see more options, including two kinds of *redirects*: 301 and 302 ([Box 2.1](#)).

Box 2.1. 301 versus 302 redirects

As we saw in the page rule dropdown, there are two types of redirects: *301 - permanent* and *302 - temporary*. You can probably guess the intended purpose of each redirect from the names, but what you might not realize is that your choice has an effect on search engine results.

If you choose a 302 redirect, then search engines like Google will assume that at some point the redirect will be removed and the page in question will be a destination that users can reach. That means that the search engines don't combine the traffic—both URLs exist as separate objects. That isn't good for analytics or search engine optimization (SEO) if you really intended for the redirect to be permanent.

If you set the redirect to 301, then search engines consider the URLs to be exactly the same. So even if a user came to your site through a redirected URL, your main URL would still get credit for purposes of traffic estimates and SEO.

In practice, the vast majority of page rule redirects will be 301s. Indeed, such permanent redirects are so common that “301” is often used as a verb, as in “[Please 301 your old links instead of breaking them.](#)”

For canonical URLs, we want the redirects to be permanent, so set “Status Code” to “301 - Permanent Redirect”, and in the “Destination URL” field add in the full URL for your site (including the protocol string `https://`) followed by `/$1`:

```
https://www.example.com/$1
```

Here `$1` represents the text matched by the first wildcard. In other words,

```
example.com/*
```

matches

```
example.com/blarg
```

and makes the string `blarg` available as `$1`. As a result, the code

```
https://www.example.com/$1
```

produces the canonical URL

```
https://www.example.com/blarg
```

In the present case, there is only one wildcard, but it's possible to match on multiple wildcards as well (Box 2.2).

Box 2.2. Matching URLs

As you might be able to guess, page rule matches are numbered sequentially, so if we had two wildcards in a rule they would be available as \$1 and \$2. For example, `*.codedangerously.com/*` is a pattern to match any subdomain; in this case, `foo.codedangerously.com/bar` would put `foo` in \$1 and `bar` in \$2. (Note that the free Cloudflare tier supports subdomain wildcards only for CNAME records that are defined explicitly. Redirecting *all* subdomains requires upgrading to the [Business plan](#).)

The same permanent forwarding can be used to direct users who requested a completely different domain over to your main site. So if we purchased `dangerouscoding.com`, we could forward any `dangerouscoding.com` URL to `codedangerously.com` using the above wildcard method to preserve the specific page they requested. The matching rule in this case would be

```
*dangerouscoding.com/*
```

where the leading wildcard matches all subdomains as well as the root domain (i.e., both `subdomain.dangerouscoding.com` and `dangerouscoding.com`). Assuming we want to ignore the subdomain, the 301 redirect would then be

```
https://www.codedangerously.com/$2
```

where we use \$2 because the subpage is matched by the second wildcard.

In our case, the forwarding URL is



Figure 2.23: It works, it really works!

```
https://www.codedangerously.com/$1
```

Click the “Save and Deploy” button to save your new rule. To see if the redirect worked, you can try visiting the [root domain](#) in your browser. For example, if we enter codedangerously.com into a browser we get redirected to www.codedangerously.com, as shown in [Figure 2.23](#).

Another convenient way of checking URL forwarding is to use the `curl` command at the command line (as covered in [Learn Enough Command Line to Be Dangerous](#)). In particular, we can use the `--head` option to return just the HTTP header (rather than the whole page):

```
1 $ curl --head codedangerously.com
2 HTTP/1.1 301 Moved Permanently
3 Date: Thu, 13 Feb 2020 16:47:00 GMT
4 Connection: keep-alive
5 Cache-Control: max-age=3600
6 Expires: Thu, 13 Feb 2020 17:47:00 GMT
7 Location: https://www.codedangerously.com/
8 Server: cloudflare
9 CF-RAY: 5648481ded567896-LAX
```

We see in line 2 that the HTTP status code is **301** as required, with the forwarding URL (line 7) including the **www** subdomain and the required secure **https://** protocol indicator.

2.5 Profit!!

Whether you got to this point after setting up a GitHub Pages site or a Heroku app, you now have an edge-cached site that protects you from DDoS attacks while making DNS changes easy and fast. Congratulations!

The only steps left, should you wish to follow them, are setting up email at your domain and adding analytics to collect data on your site's users. These are the subjects of [Chapter 3](#).

Chapter 3

Custom email and analytics

By the end of [Chapter 2](#), we've accomplished the main task we set out to do: we have our own website on a custom domain. In this chapter, we include two refinements: custom email ([Section 3.1](#)) and site analytics ([Section 3.2](#)). They're in the same chapter because the solutions we recommend are both provided by the same company (Google) and therefore fit nicely together.

3.1 Custom email

In this section, we'll first give an overview of how email works and then go into the details of using [Google Workspace](#) to support custom email addresses at your domain. The end result will be the ability to receive email at yourname@yourdomain.com.

3.1.1 Introduction to email

Everyone uses email, but few know how it works. Ultimately, the system is surprisingly simple—email is really just plain text documents being sent back and forth with a little extra data needed to make things work.

The first real email message was sent by Ray Tomlinson in 1971 after he came up with an innovative way to format email addresses using the `@` sign. This



Figure 3.1: The two parts of an email.

allowed people to specify both the name and the domain for the destination for any given email (Figure 3.1).

Before Tomlinson came up with his addressing idea, there wasn't really any standardized way to send messages to people who were on a different network or computer (this was the time of **mainframes**, so often many people would share one giant room-sized computer). You could send messages to people on the same system as you, but not outside it (Figure 2.21).¹

The convention of using the @ sign to denote the computer or network made email capable of being a universal messaging protocol. Eventually, groups of engineers came together to create a set of properties that every properly formatted email needed to have in order to be delivered—things like “to:” and “from:” fields. All that extra information is contained in what is called the *email header*. The message, and any other files like attachments, are a part of the *email body*.

When you send an email message, your properly formatted email is transferred as a text file to an outgoing email SMTP (Simple Mail Transfer Protocol) server. This is a software application that collects emails being sent by authorized users of that server and places them in a **queue**. The server then looks at each email's header to see what the destination is. Next, the SMTP server queries the DNS system to find the destination domain's *mail exchanger (MX) records*, which specify the mail server responsible for accepting email messages on that domain. When the location of the destination server is returned,

¹Image retrieved from <https://flic.kr/p/bu2gfG> on 2020-02-18. Copyright © 2012 Robert Nix and used under the terms of the [Creative Commons Attribution 2.0 Generic](#) license.



Figure 3.2: Well Bob, if I'm reading this right, you'll be able to send messages alllllll the way over to Mary there. Let's call it 10ft.

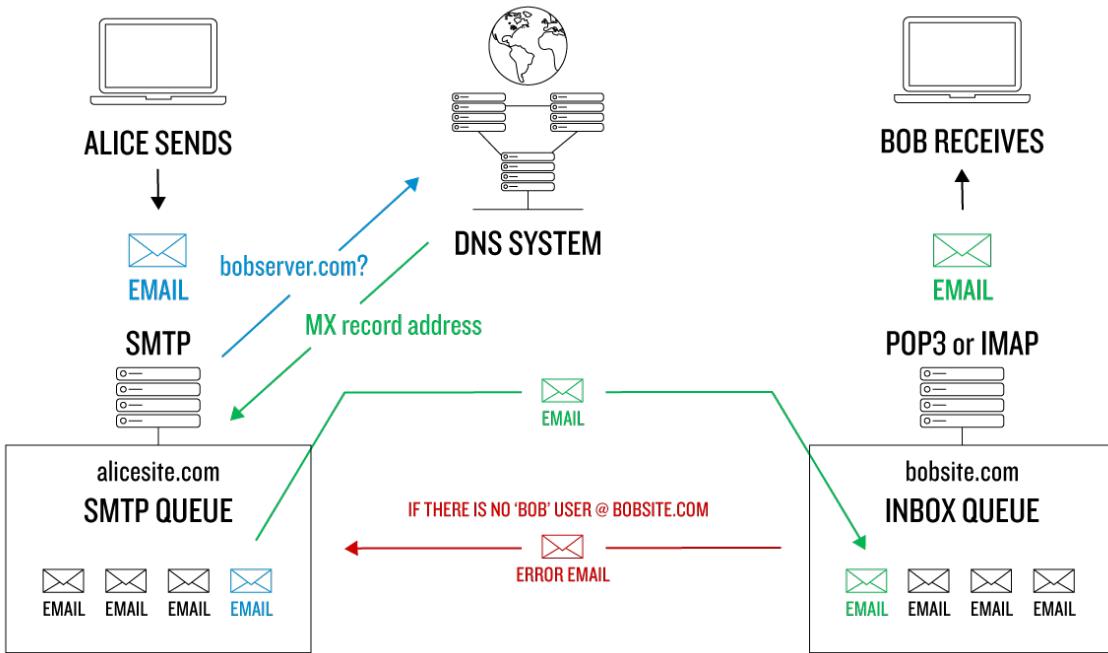


Figure 3.3: Sending an email between outgoing and incoming email servers.

the SMTP server opens up a connection to the receiving server and sends the file over the Internet (Figure 3.3).

It's certainly possible to run an email server on your own hardware, and then set your DNS for your domain to point messages to that server... but why would you want to give yourself the headache of having to deal with managing and updating hardware and software?

If you are running your own server, then more than likely you are going to be tied to a single machine that might go down or lose access to the Internet. That's to say nothing of the risk you run of having your servers hacked (there's an [American politician](#) you might be able to ask [how poorly that can go](#) (Figure 3.4)).²

It's a lot easier to use a global email service that can take care of all the security and uptime requirements, while also providing additional features like

²Image retrieved from <https://flic.kr/p/zQ8tnC> on 2019-11-20. Copyright © 2015 jimthompson and used under the terms of the [Creative Commons Attribution 2.0 Generic](#)) license.



Figure 3.4: No comment.

advanced spam filtering. This is the option we recommend.

3.1.2 Google mail

The specific global email service we recommend is Google mail (Gmail). We use Gmail because, as one of Google's core consumer-facing businesses, there is support for Gmail on a large variety of devices and other services. It is also available and accessible basically anywhere at any time since it is running on Google's core infrastructure—that means you don't have to worry about some rinky-dink server outage taking down your access to email. Things have to go very wrong on the Internet to prevent you from being able to access Gmail.

Gmail's spam filtering is also top-notch, as they leverage both Google's expertise in machine learning and the millions of users that are on Gmail to figure out how to distinguish legitimate email from [spam](#). If you are one of the global few who never has had a Gmail account, and has had to deal with inadequate spam filtering, we can promise you that with Gmail you rarely end up with spam in your inbox, or with real messages in your spam folder. It feels



Figure 3.5: Spam is always full of lies.

like a solved problem (Figure 3.5).³

Lastly, we actually like the Gmail interface... We know that some people don't, but hey, you know what they say about opinions. We find that searching email on Google is always fast and accurate (something that we can't say about other systems that we've used in the past), and there are nice little power-user features that make life easier (like the ability to unsend a message within 30 seconds or the dots-and-pluses trick (Box 3.1)).

Box 3.1. Gmail dots and pluses trick

Many people don't know this, but when you have a Gmail email address, you actually have a basically infinite set of sub-email addresses that can tie into your account without needing to do any setup. This is a nice way to pre-filter your email by giving different people, or websites that you don't fully trust, a slightly different address. Then you can easily have Gmail move any mail sent to an email variant into a folder.

What do we mean? Let's say you have the email address `yourname@gmail.com`.

The first trick is that Gmail doesn't care about dots in your email address. So, if you tell someone to email you at `your.name@gmail.com`, that will still end up in the inbox for `yourname@gmail.com`—and you'll be able to see that it was sent to `your.name@gmail.com`. The lone caveat here is that you can't have

³Image retrieved from <https://flic.kr/p/8jBWuu> on 2020-02-18. Copyright © 2010 GuillermoJM and used under the terms of the [Creative Commons Attribution 2.0 Generic](#) license.

two periods in a row. So `your..name@gmail.com` wouldn't work, but there would be nothing wrong with `y.o.u.r.n.a.m.e@gmail.com`. Eventually, though, you run out of variations...

The second trick is that you have the ability to create a limitless number of *ad hoc* emails by adding a plus sign, +, to the end of your regular email and then whatever additional text you desire.

In other words

`yourname+travelsite@gmail.com`

and

`yourname+shadysite@gmail.com`

both end up going to `yourname@gmail.com`. This lets you use distinct email addresses for different sites.

So what are the downsides of using Gmail for your custom domain? Well, unlike the free service (the one where your email is `something@gmail.com`), you have to pay for it. Google offers custom Gmail via its general suite of cloud productivity tools called [Google Workspace](#) (formerly G Suite), which (as of this writing) costs \$6/user/mo. for the Business Starter plan. (Note that each user can have up to 30 *aliases*, or other email addresses that all go to the same inbox, so there's no need to pay for multiple users in this case.) For more details, see the [Workspace pricing page](#).

Some people have also expressed privacy concerns about Google's services. While this may apply to Gmail's consumer service, which serves ads based on email content, Workspace is different, as noted in Google's [FAQ](#):

Does Google use my organization's data in Google Workspace services or Cloud Platform for advertising purposes?

No. There are no ads in Google Workspace [Services](#) or Google Cloud Platform, and we have no plans to change this in the future.

We do not scan for advertising purposes in Gmail or other Google Workspace services. Google does not collect or use data in Google Workspace services for advertising purposes.

The process is different for our free offerings and the consumer space. For information on our free consumer products, be sure to check Google's [Privacy and Terms](#) page for more consumer tools and information relating to consumer privacy.

There's an old axiom that applies here: "If you aren't paying for a product or service, then you are the product." (For example, Facebook is "free", but their product is your personal information, which they use for ad targeting.) With Gmail for work, you *are* paying, so you are *not* the product.

If you'd really rather use an alternate service (such as privacy-focused [Proton](#)), the setup is going to end up being similar to what we cover here. This means that the steps discussed in this section will be useful in any case.

3.1.3 Google Workspace signup

Before we start, a little caveat: Google changes their services all the time so you might need to use a heaping spoonful of technical sophistication ([technical sophistication](#)). We visited the site one day to go through all the steps and create an account, and then the next day went back to double-check a step in the process only to find that the signup flow was entirely changed. They still wanted the same information—everything was just arranged totally differently! Then, shortly after, Google changed the name of its service from G Suite to Google Workspace... it's just a mess.

To sign up for Google Workspace, head over to the [Workspace home page](#) and get started by clicking the Get Started button... or whatever the equivalent is today. Again, Google might be testing different language for buttons—who knows? Proceed through whatever random information they need from you to finish the account creation—the only important step is that when you are asked about your domain name, make sure to put in the domain you registered (without www, so like codedangerously.com).

Now that the initial setup is done, Google will want you to make changes to your DNS records ([Section 2.1](#)) so that Google can verify that you actually own the domain and have added the correct records for setting up email ([Section 3.1.4](#)). That sounds more complicated than it is, but thankfully by now you are an expert at editing DNS records!

3.1.4 Verify domain and configure MX records

As mentioned above, Google wants to make sure that you are in fact the owner of the domain that you are trying to create an email account for, and luckily the DNS system provides a lot of opportunities to add information to a domain's records. For Gmail, Google just wants you to add an extra TXT record (which is what it sounds like: a text record on your site's DNS configuration) consisting of a unique string of letters and numbers).

Because of the public nature of the DNS records, Google can just query your domain's information and easily see that the verification strings match, confirming that you are in fact the person who controls the domain. If you haven't already, click past the verification-instructions welcome page to get to the details on how to make Google happy. Copy the verification text on the page ([Figure 3.6](#)).

In another tab or window, open up your site's DNS records in Cloudflare, click the “+Add record” button and select **TEXT** from the new record dropdown. In the Name field, add in an **@** sign, which is a shorthand way of saying use whatever the root domain is for this record. (In other words, for us **@** is the same as **codedangerously.com**.) Click into the Content field and paste in your verification text, as shown in [Figure 3.7](#).

After hitting the Save button, you should see the new record in the list. As a last step, delete any old MX records that are not pointing to the Google servers (a step mentioned briefly when discussing A records in [Section 2.2.1](#)).

We are almost done, but we still need to add the actual DNS records that point to Google's servers to handle email. Hop back over to your Google Workspace admin console tab and click the Verify My Domain button at the

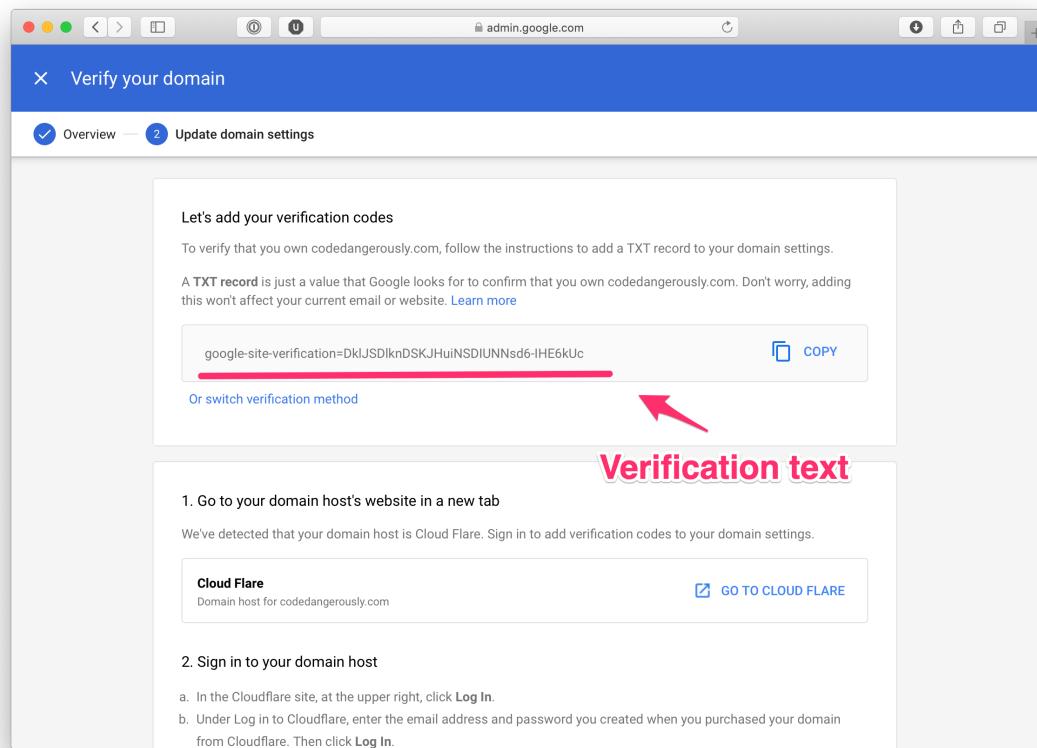


Figure 3.6: The TXT record verification string (not our real one!).

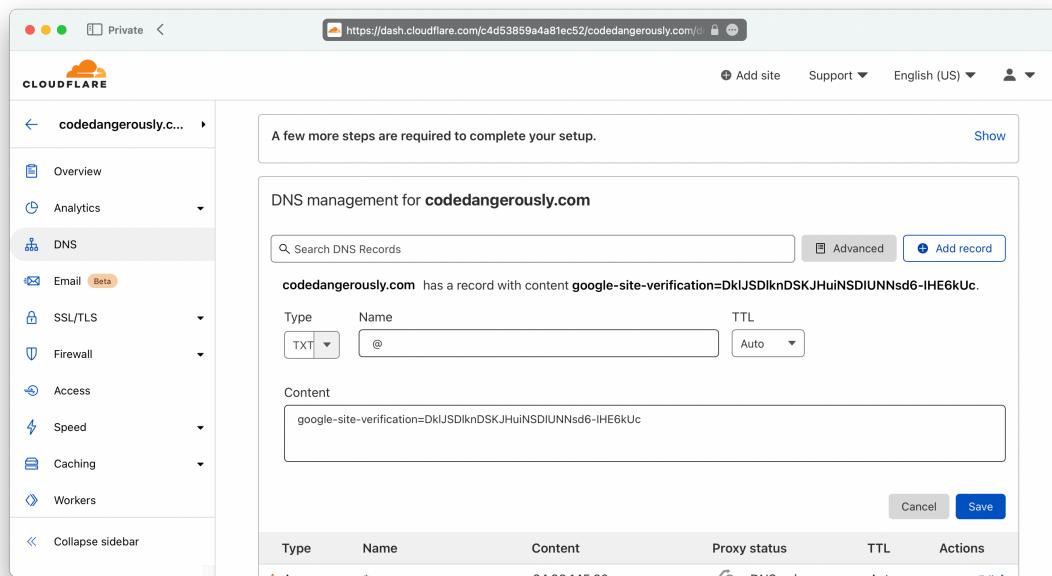


Figure 3.7: Adding the new TXT record (still not our real verification key!).

bottom of the page.⁴

If you are working with any other people on your project, you can use the admin console setup process to add other users to your Google Workspace account. It is a pretty self-explanatory process, and if you don't want to deal with it now, you don't need to.

To get your email up and running, click the Activate link to bring up a page with instructions that show how to set up your DNS records so that email gets routed to the correct mail servers. In particular, we'll be adding MX records to the DNS settings at Cloudflare. To add the first one, copy the first value from the Google Workspace setup, which for us is **ASPMX.L.GOOGLE.COM**. Then, on the Cloudflare DNS management page, perform the following steps:

- Add a new record, making sure that **MX** is selected in the Type dropdown.
- Add **@** in the Name field so that the record points to the domain root.
- Add the server name that you copied from Google Workspace into the Mail server field.
- For this first record, set a priority of 1.

Now head back over to your Google Workspace setup tab, and then bounce back and forth between that and Cloudflare to finish entering the other MX records listed in the table, making sure that you are setting the priority on each one to the specified level (1, 5, or 10 from Google's MX info).

When you are done, click the Activate Gmail button and let Google check to see if the DNS records have been correctly added. If they have been, that's it—you are all done!

At this point, Google will (could... may... really who knows?!) take you to a confirmation page to let you know that your Google Workspace account is set up. If you want to check out your new Gmail account, click the Google Apps menu icon at the top right, and then select Gmail (Figure 3.8).

⁴If we weren't using Cloudflare, the changes could theoretically take up to 72 hours to go into effect (Section 2.1.1)—which is all the more reason to use Cloudflare.

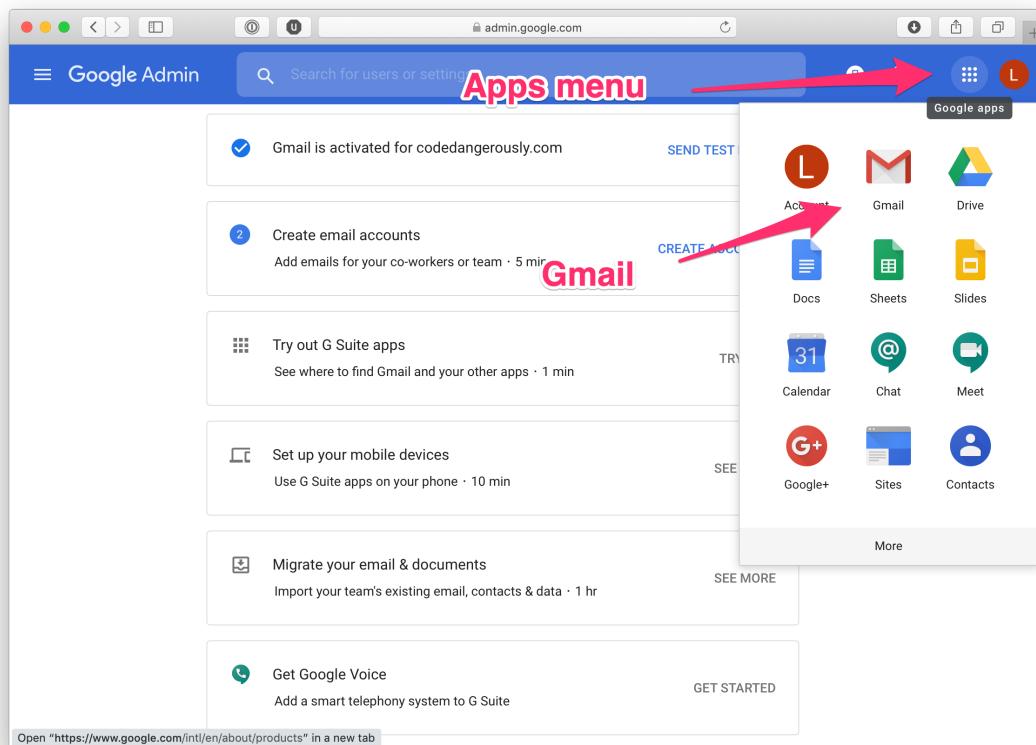


Figure 3.8: It's all done, done.

In the future, you can access your new email inbox by going to mail.google.com and signing in with the new account that you created during the mail setup process. If you need to get back to the console and can't remember what the address is, it's super-easy: admin.google.com. Once you've logged in there, you can add additional users to your account, update billing information, or add additional services.

Speaking of additional services... since you have a Google account now, why not sign up for Google Analytics so that you can get information about who is visiting your site?

3.2 Site Analytics

Google Analytics is probably the most useful free service that Google has for the owners of websites and web applications. By adding just a tiny little snippet of code to your site, you can have Google track not just how many people are visiting your site, but also the browser and operating system they are using, whether they are on desktop or mobile, how they found your site, how long they stayed on your site, which pages they viewed, which page they left your site from... and a lot more.

To get started, visit the [Google Analytics home page](https://www.google.com/analytics) and click the big Sign Up for Free button. (As always, use your technical sophistication if Google has changed their interface since we wrote this.) From there, you'll be asked to sign in using the Google account that you just created.

Once you are in, you'll be on the Analytics setup track. Click the Get Started or Sign Up button to go to the site information page. Fill out the fields on the page with the requested information about your site, and then at the bottom of the page press the Get Tracking ID button.

After you agree to the Terms of Service, the site will take you to a page in your new analytics account that will have the little bit of code you need to add to your site for the tracking to work. What you are looking for is the code snippet that looks something like Listing 3.1.

Listing 3.1: This is the code that makes Google Analytics work.

```
<script>
(function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function(){
(i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new Date();a=s.createElement(o),
m=s.getElementsByTagName(o)[0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)
})(window,document,'script','https://www.google-analytics.com/analytics.js','ga');
ga('create', 'UA-XXXXXXX-1', 'auto');
ga('send', 'pageview');
</script>
```

(We've removed one level of indentation from the script to get it to fit on the page, but you can just copy-and-paste from Google.) Note that you should always use the exact code supplied by Google; Listing 3.1 is included only as an example of the kind of code you can expect.

3.2.1 Add Snippet

One option is to paste the snippet from Listing 3.1 on every page of your site, but in general this is rather inconvenient. Instead, we're going to assume that you are using some sort of templating framework like Jekyll (as covered in *Learn Enough CSS & Layout to Be Dangerous*) or Ruby on Rails (as covered in the *Ruby on Rails Tutorial*).

Whichever system you are using, you should paste the analytics snippet into the `<head>` section of your site. In *Learn Enough CSS & Layout to Be Dangerous*, this section is located in `_includes/head.html`, as shown in Listing 3.2.

Listing 3.2: Where to put the code snippet in your site code.

```
_includes/head.html

<head>
  {% if page.title %}
    <title>{{ page.title }} | Test Page</title>
  {% else %}
    <title>Test Page: Don't Panic</title>
  {% endif %}
  {% if page.description %}
    <meta name="description" content="{{ page.description }}">
```

```

{%- else %}
    <meta name="description" content="This is a dangerous site.">
{%- endif %}
<link href="/favicon.png" rel="icon">
<meta charset="utf-8">
<meta name=viewport content="width=device-width, initial-scale=1">
<link rel="stylesheet" href="/fonts/font-awesome-4.7.0/css/font-awesome.min.css">
<link rel="stylesheet" href="/css/main.css">

<script>
(function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function(){
(i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new Date();a=s.createElement(o),
m=s.getElementsByTagName(o)[0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)
})(window,document,'script','https://www.google-analytics.com/analytics.js','ga');
ga('create', 'UA-XXXXXXXXX-1', 'auto');
ga('send', 'pageview');
</script>
</head>

```

Save your changes, deploy your site... and that's it! If you want to force Google Analytics to push some test visitors to the site and make sure everything is working, you can click the “Send test traffic” button ([Figure 3.9](#)).

If you've looked for information on how to set up Analytics in the past, you might have seen someone suggest that you should put the Google Analytics code snippet at the very end of your site right above the closing `</body>` tag. You can absolutely do that, and the analytics will still work, but the main reason for doing that has changed.

It used to be the case that when your site tried to load Analytics it would actually hold up the rest of your site from loading until the Analytics code finished downloading—not a good thing if the connection to Google is slow. Since then though, the Google Analytics code has switched to using an *asynchronous loading* method that makes it so that the rest of your site load isn't held up.

That said... of course there is a caveat. With Google Analytics, you can add in a bunch of custom tagging and tracking functionality—things that require parts of your site to load first to set up tracking code. If you are using a lot of custom Analytics tracking, then it might be a good idea to move the Analytics initialization to the bottom of your page template above the closing `</body>` tag.

Once you are all set up, you can sign into the Analytics dashboard by going to analytics.google.com. On the dashboard, you will have a ton of options for

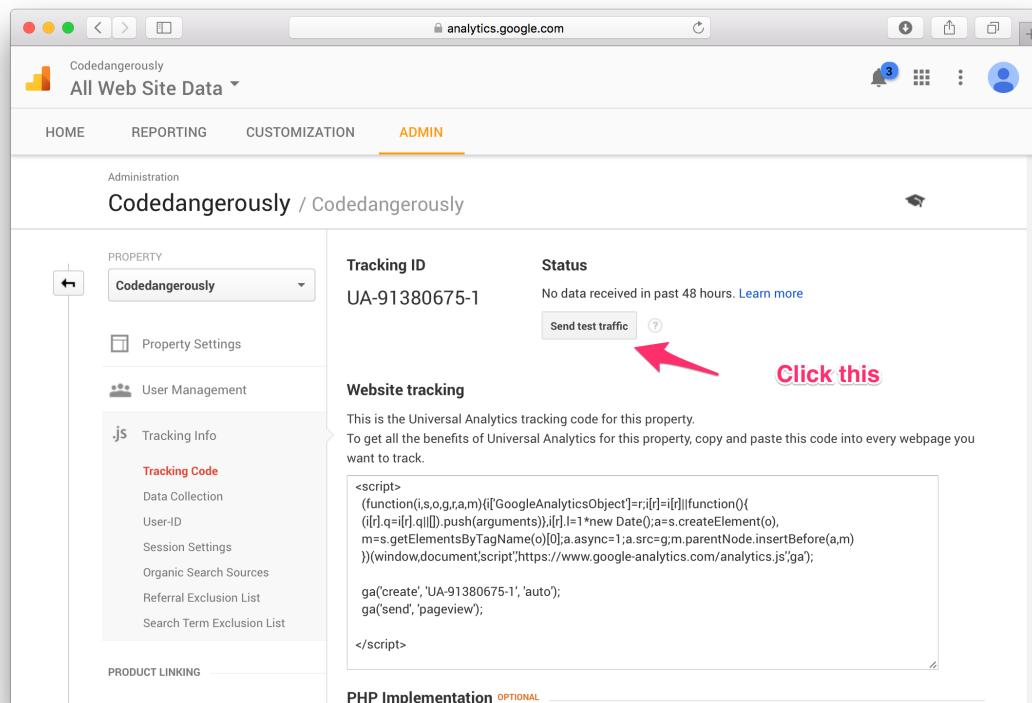


Figure 3.9: The tracking code page for your site.

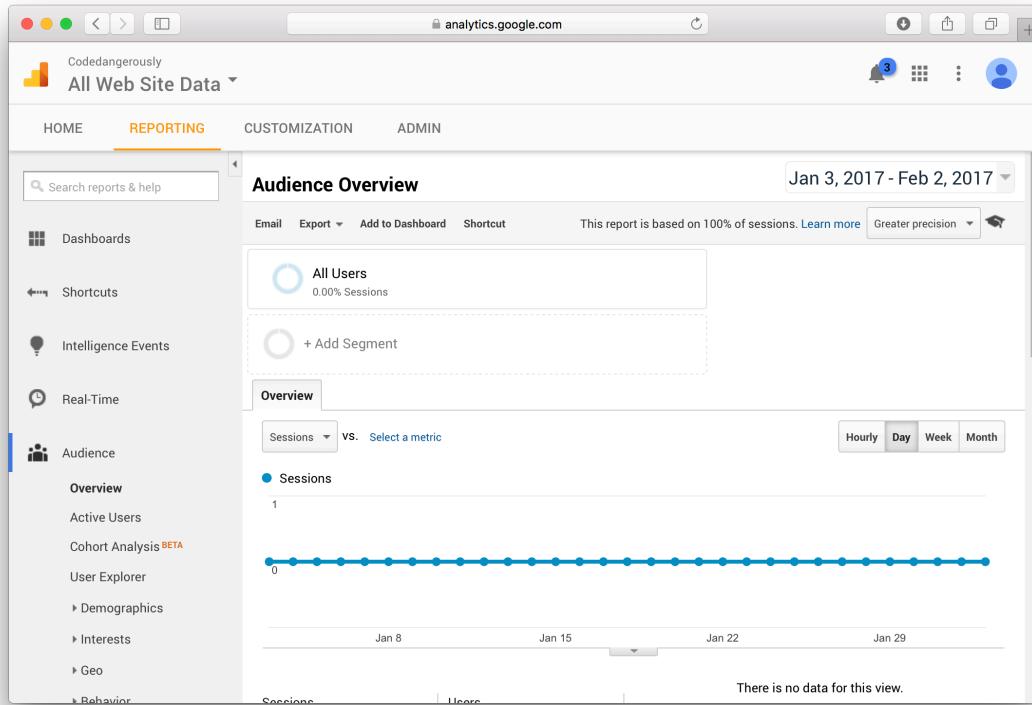


Figure 3.10: Hooray analytics, boo no visitors. So lonely.

all the different ways to slice and dice your data, or even see the users who are on the site in real time (Figure 3.10).

Google Analytics is a *really* deep service, and we aren't going to dive into how to use it—that would require a book of its own (and there are many out there). This goal here was just to get you up and running with the service, and in the future we might write up an Analytics tutorial of our own. Until then, you'll just need to search around the [interwebs](#) and find guides and tutorials to teach you the dark arts of Google Analytics.

3.3 Conclusion

So there you go! If you made it all the way through the tutorial, you now have your site running on a custom domain, you are using Cloudflare for DNS management, you have a custom email address through Google Workspace, and you are using Google Analytics to see who is visiting your site. You've gone from just having a rinky-dink page on the Internet to actually looking like a legitimate site or business. Great work!