



THE IMPOSTER'S HANDBOOK

A PRIMER FOR SELF-TAUGHT DEVELOPERS

ROB CONERY | BIG MACHINE

FOREWORDS BY SCOTT HANSELMAN AND CHAD FOWLER

by Rob Conery

© 2016 Big Machine, Inc. All rights reserved.

All rights reserved. No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law. For permission requests, write to the publisher, addressed “Attention: Rob,” at the address below.

Published by: **Big Machine, Inc** 111 Jackson, Seattle WA 96714

Publication date: November, 2016

CHANGELOG

This book was written in the open, with quite a lot of feedback from many people. The original release was in August of 2016, with version 0.0.1. The book has changed quite a lot since then.

2.0 - Current Version

What else can I say? This is a *gigantic* rethink of this book. The changes are quite extensive:

- Split the book into 4 parts
- Rearranged the chapters to flow better from history to concepts and finally practice
- Added Stack, Queue, Trees, Dictionary and Trie to Data Structures chapter
- Added Depth-first and Breadth-first traversal to Simple Algorithms
- Split BDD out from Testing chapter
- Split Unix out into its own Part
- Moved everything out of Softcover so I could have more control over rendering, formatting and writing
- Removed *all* text-based code samples in favor of images. I know this might make some people cranky but the code is available on Github and copying from a book is fail.
-

1.3

- Complete rewrite of the Functional Programming chapter
- Reorganization and separation of a number of chapters, including Testing, Functional Programming, Design Patterns, and Design Principles.
- Numerous typo fixes
- Fix to a Lambda Calculus bug where a wrong reduction was being applied.

1.2

- Completely rewrote the Lambda Calculus chapter. I was never happy with the way the chapter ended up, and now I'm quite excited. I fixed a few inaccuracies and also added a number of code samples (using ES2016). Of all the updates, this one is the biggest.
- Complete rewrite of Complexity Theory chapter (P/NP). A number of people have told me that this chapter needed some work, so I simplified it as much as possible while also making it a bit more accessible.
- Better formatting. I changed to using Softcover for the editing, which allowed me to have much cleaner code samples along with better math equations. Overall the layout is more "book-like" and I think you'll find it much improved.
- Slight restructuring. Some of the chapters went rather long (Computational Theory is one), so I broke a few

things apart to keep focus.

1.1

- Many typos and minor code issues fixed. Fixed a flaw with formatting and kindles.
- Removed the chapter on learning a language

1.0 (Initial Release)

- Clean up, formatting and final checks.

0.0.5

- I added back the PDF version and it's now nicely formatted.
- I'm using a different formatting tool for the mobi file (more on this below)
- I updated/smoothed out/redid the entire "Hacking Linux" section and put it back in. It's now called "Essential Unix Tools"
- As always: quite a number of grammatical/typo fixes

0.0.4

This one is a massive overhaul! I moved the entire book out of Gitbook due to formatting errors. The format I wanted to use (epub 3) was not supported, so much of the HTML underlying the book caused validation errors.

I moved into iBooks Author at that point and reformatted everything. A better font, better layout, better... everything.

I also removed the Hacking Linux section from the book entirely. It didn't add anything to the overall theme. In addition:

- Broke apart the Comp Sci, Baby chapter into many smaller chapters and reordered a few things.
- I added graphs to the Data Structures chapter
- I added code and shortest paths (Bellman-Ford and Dijkstra) to Algorithms
- I added dynamic programming, Fibonacci and Sieve of Eratosthenes to Algorithms as well
- I expanded on Big-O to make it a bit more clear and I also added complexity analysis to the algorithms chapter
- I completely rewrote, again, P=NP?. This time, I think I have it. At least my brother thinks so, which is a small victory!
- Finally, there were numerous spelling errors and typos that were fixed thanks to your feedback. I was also able to use the Proofreader functionality of iBooks Author to tighten up the grammatical stuff (before it goes to final editing).

0.0.3

- Numerous typos and grammatical flow issues fixed.
- Added a chapter on Impostor's Syndrome
- Stubbed out chapter markers for probability, math and more algorithms (to be completed for the next rev)
- Rewrote (completely) P=NP? - The P=NP? chapter worked but there were some errors and also some holes. I investigated a whole lot more and rewrote the entire thing. I quite like it – please have a read.
- Fixed Omega reduction error where I omitted the substitution step.
- Fixed O($n + 1$) errors where I confused things a bit
- Tightened up Ruby language description with a blurb about YARV
- Added chapter on Learning a New Language (task)
- I removed most of the Linux essentials section so I can rework it into a series of walkthroughs. It will come back, but in a different form.
-

0.0.2

- Fixed a mess of spelling errors and goofy errors (like spelling von Neumann's name wrong in the chapter header).
- Expanded on Big-O Notation, quite a lot.
- Expanded definition of a kernel in the Linux chapters.
- Tweaked a number of epub validation and styling problems.
- Fixed a layout issue caused by Gitbook when rendering. H2 headers would cause a page break, which doesn't make any sense but... OK. Moved those to H3s and H3s to H4s.
- Expanded on Linux permissions and group ownership.
- Expanded on loading shell scripts into the shell and executing them.

0.0.1

Initial preview release.

FOREWORDS

Scott Hanselman

I never did a formal Computer Science course. My background was in Software Engineering. For a long time I didn't understand the difference, but later I realized that the practice of software engineering is vastly different from the science of computers.

Software Engineering is about project management and testing and profiling and iterating and SHIPPING. Computer Science is about the theory of data structures and $O(n)$ notation and mathy things and oh I don't know about Computer Science.

Fast forward 25 years and I often wonder if it's too late for me to go back to school and "learn computer science." I'm a decent programmer and a competent engineer but there are...gaps. Gaps I need to fill. Gaps that tickle my imposter syndrome.

I've written extensively on Imposter Syndrome. I'm even teased about it now, which kind of makes it worse. "How can YOU have imposter syndrome?" Well, because I'm always learning and the more I learn the more I realize I don't know. There's just SO MUCH out there, it's overwhelming.

Even more overwhelming are the missing fundamentals. Like when you're in a meeting and someone throws out "Oh, so like a Markov Chain?" and you're like "Ah, um, ya, it's...ah...totally like that!"

If only there were other people who felt the same way. If only there was a book to help me fill these gaps.

Ah! Turns out there is. You're holding it.

Scott Hanselman @shanselman August 12, 2016 Portland, Oregon

Scott Hanselman has been computering, blogging, and teaching for many years. He works on Open Source .NET at Microsoft and has absolutely no idea what he's doing.

Chad Fowler

I've been honored to be asked to write the foreword for several books over the course of my career. Each time, my first reaction is something like "Oh wow, how exciting! What an honor! HEY LOOK SOMEONE WANTS ME TO WRITE THEIR FOREWORD!!!" My second reaction is almost always, "Oh no! Why me? Why would they ask *me* of all people? I'm just a saxophone player who stumbled into computer programming. I have no idea what I'm doing!"

No offense to Rob intended, but this may be the first foreword I feel qualified to write. Finally, a book whose very title defines my qualifications not just to write the foreword but to participate in this great industry of ours. A handbook for impostors. It's about time.

You know that friend, classmate, or family member who seems to waste too many precious hours of his or her life sitting in front of a computer screen or television, mouth gaping, eyes dilated, repetitively playing the same video game? In my late teens and early twenties, that was me. I made my living as a musician and had studied jazz saxophone performance and classical composition in school. I was an extremely dedicated student and serious musician. Then I got *seriously* addicted to id Software's Doom. I played Doom *all the time*. I was either at a gig making money or at home playing the game. My fellow musicians thought I was really starting to flake out. I didn't practice during the day or try to write music. Just Doom.

I kind of knew how personal computers worked and was pretty good at debugging problems with them, especially when those problems got in the way of me playing a successful game of Doom death-match. I got so fascinated by the virtual 3D rendered world and gameplay of Doom that my curiosity led me to start learning about programming at around age 20. I remember the day I asked a friend how programs go from text I type into a word processor to something that can actually execute and do something. He gave me an ad hoc five minute explanation of how compilers work, which in different company I'd be ashamed to admit served my understanding of that topic for several years of professional work.

Playing Doom and reading about C programming on the still-small, budding internet taught me all I knew at the beginning about networking, programming, binary file formats (We hacked those executables with hex editors for fun. Don't ask me why!), and generally gave me a mental model for how computer systems hung together. With this hard-won knowledge, I accidentally scored my first job in the industry. The same friend who had explained compilers to me (thanks, Walter!) *literally* applied for a computer support job on my behalf. At the "interview", the hiring manager said "Walter says you're good. When can you start?"

So, ya, I really stumbled into this industry by accident. From that point on, though, I did a lot of stuff on purpose. I identified the areas of computer technology that were most interesting to me and systematically learned everything I could about them. I treated the process like a game. Like a World of Warcraft skill tree, I worked my way up various paths until I got bored, intimidated, or distracted by something else. I covered a lot of territory over many hours of reading, asking questions of co-workers, and experimentation.

This is how I moved rather quickly from computer support to network administration to network automation. It was at this time that the DotCom bubble was really starting to inflate. I went from simple scripting to object oriented programming to co-creating a model/view/controller framework in Java for a major corporation and playing the role of "Senior Software Architect" only a few short years after packing the saxophone away and going full time into software development.

Things have gone pretty well since then, but I've never gotten over that nagging feeling that I just don't *belong* here. You know what I mean? You know what I mean. You're talking about something you know well like some database-backed Web application and a co-worker whips out Big-O notation and shuts you down. Or you're talking about which language to use on the next project, and in a heated discussion the word "monad" is invoked.

Oh no. I don't know what to say about this. How do I respond? How can I stop this conversation in a way that doesn't expose me for the fraud I am? WHAT THE HELL IS A MONAD?

Haha, ya.

I hope that non-response made sense, you think as you walk toward the restroom, pretending that's why you had to suddenly leave the discussion.

In daily work, I find myself to be at least as effective as the average programmer. I see problems and I come up with solutions. I implement them pretty quickly. They tend to work. When performance is bad, I fix it. When code is hard to understand I find a way to make it easier to understand. I'm pretty good at it I guess.

But I didn't go to college for this stuff. I went to college and studied my ass off, but all I have to show for it is an extremely vast array of esoteric music knowledge that would bore the hell out of anyone who isn't a musician. In college you learn about algorithms. That sounds hard. When I write code, I don't use algorithms I think. Or do I? I'm not sure. I don't invoke them by name most of the time. I just write code. These college programmers must be writing stuff that's unimaginably more sophisticated since their code has **algorithms!**

And how can my code perform well if I didn't use Big-O notation to describe its performance and complexity characteristics? What the hell does "complexity" even mean in this context? I must be wasting **so** many processor cycles and **so** much memory. It's a miracle my code performs OK, but it does.

I think most of my success in the field of computer software development comes from my belief that:

1. A computer is a machine. In some cases it's a machine I own. I could break it into tiny pieces if I wanted.

2. These machines aren't made of magic. They're made of parts that are pretty simple. They're made in ways that tens of thousands of people understand. And they're made to conform to standards in many cases.
3. It's possible to learn about what these components are, how they work, and how they fit together. They're just little bits of metal and plastic with electricity flowing through them.
4. Everything starts with this simple foundation and grows as simple blocks on top.
5. All of the hard sounding stuff that college programmers say is just chunks of knowledge with names I don't know yet.
6. Most of this stuff can be learned by a young adult in four years while they're also trying to figure out who they are, what they want to do with their lives, and how to do as little work as possible while enjoying youth.
7. If someone can learn *all this stuff* in just a four year degree, it's probably pretty easy to hunt down what they learn and learn it myself one concept at a time.
8. Finally, and most important, somehow I get good work done and it doesn't fall apart. All this stuff I don't know must be just a *bonus* on top of what I've already learned.

All this is just stuff you can learn! Wow. In fact, the entirety of human knowledge is just a collection of stuff that you can learn if you want to. That's a pretty amazing realization when you fully absorb it. A university doesn't magically anoint you with ability when you graduate. In fact, most people seem to leave university with very little actual ability and a whole lot of knowledge. Ability comes from the combination of knowledge, practice, and aptitude.

So, what separates us impostors from the Real Deal? Knowledge, practice, and aptitude. That's it. Knowledge is attainable, practice is do-able, and we just have to live with aptitude. Oh well.

Here's a big secret I've discovered: I'm not the only impostor in the industry. The first time I met Rob, he interviewed me for a podcast. We ended up talking about Impostor Syndrome. On hearing my interview, several people told me "Wow, I feel like that too!" The more I dig, the more I think we *all* feel like that at some points in our careers.

So, welcome to Impostor Club! I'm kinda bummed now that I know it's not as exclusive as I had thought, but it's nice to have company I guess.

Anyway, now we have a handbook.

Ironically, reading and using this handbook *might* cause you to graduate from the club. If that happens, I wish you luck as you enter full scale computer software programmer status. Let me know how it feels. I'll miss you when you're gone.

Chad Fowler August 12, 2016 Memphis, Tennessee

*Chad Fowler is the CTO of [Wunderlist](#), which is now a part of Microsoft. He is also a Venture Capitalist with [BlueYard Capital](#) where he **really** feels like an imposter.*

PREFACE

A Matter of Principle

Back in November of 2008, Jeff Atwood published a post that proved somewhat embarrassing. It was entitled [Your Favorite NP-Complete Cheat](#) and it sparked a bit of a controversy at the time, specifically this quote (which Jeff later redacted):

... Nobody can define what makes a problem NP-complete, exactly, but you'll know it when you see it.

It turns out that [many people can define what makes a problem NP-complete](#). When I read this I had no idea what any of it meant anyway, but I felt badly for Jeff. It's something that keeps me awake at night: *the fear of being utterly, publicly, horribly wrong*. But there's something that scares me a bit more.

Spreading ignorance is my true nightmare. There's a subtle (yet incredibly important) difference between trying and failing at something vs asserting a complete falsehood as truth. Not only do you look foolish, but you make life harder for people *who truly know the subject* you've failed to understand properly. These people now have to spend their time fixing the mess you've created.

You might have the best of intentions; you might believe, in the moment, that your experience is sufficient enough to speak out on a given subject. *None of that matters if you're wrong*. If you cross the line between "I think that..." and "This is the way it is", *you better be right*.

An anonymous commenter left a wonderfully eloquent reply to Jeff's post that perfectly captures the essence of why it is so important to *triple check* your subject when asserting something publicly. I love this comment. I go back and read it once a year or so to help keep me grounded, to make sure I go deeper, take an extra day or week ... to care that much more whenever I write a post on a topic or give a presentation.

I want to share this comment with you. It has haunted me with every chapter written in this book. I hope that when you're done here, this comment will haunt you as well. I am using a screen shot because I don't know who the author is and I wanted to capture his/her/their exact words closely.

M

M1116

Nov '08

For me, spreading ignorance (or bad information due to ignorance) is an issue.

If you are gonna talk about subject X, make sure you know subject X, well enough to talk about it. At least, make sure that you are not making gross errors about subject X. Is it really too much to ask?

It bothers me on both the practical level and as a matter of principle.

I've been working as a professional programmer for years, and I've encountered many without basic scientific background. And that's FINE! Not everyone needs it, or wants it. But then these people read a blog post, and it sounds right, so they believe it. After all, they lack the knowledge to figure out which parts are true and which are don't. That's why they are reading the blog!

And they learn stuff that is WRONG. And then they are going to implement this stuff, and argue about it, and generally BELIEVE that anything having to do with CS is either unpractical or is easily enough learned in 30 minutes of reading.

And then I have to work with these people, and manage them! They have no grasp of how much they simply don't KNOW. And at some point, their knowledge simply ain't gonna cut it. And they are going to argue with me, or write me off as a user of fancy computer science jargon. I mean, it's just register allocation, right? How hard could it be? It's only BSP tree optimization, let's just check all the options!

So I am trying to combat this ignorance for practical, selfish reasons. Programmers need to understand the problems they work on. They need to understand when they are out of their depths, and its time to hit the books, or call someone who knows. Or reject that project, or bump up the cost and time estimate. At least map out the areas of your understanding, so you'll know when you're on treacherous ground.

The other reason is a matter of principle. Ignorance is pretty bad, and I reject mediocrity for its own sake. In less fancy terms, if you are writing technical posts, get the technical details RIGHT! Isn't that a given? But apparently I am an asshole for pointing it out 😂

On Being Wrong

I will be wrong in this book. It might seem odd to start off this way, but I think it's important. Without failure, we do not learn and grow. As much as anything: *this book represents my failure in equal measure with any success I might have*. If we're going to get anywhere in our careers, we must seek out challenge and, along with it, failure. Failure is growth, growth is learning.

Embrace it.

Now, to this book: it is a compendium in summary form. I cannot possibly dive into things in as much detail as they deserve. For that, I do apologize ... but it's not the point of my writing this book.

I wrote this for people like me: self-taught developers who have never taken the time to learn foundational concepts. If you're brand new to programming, this is not the book for you. I expect you to have been programming for a few years, at least.

I also wrote this book for fellow explorers. If you're the entitled lazy type who expect full demos with code you can copy and paste – you won't find that here. I'm going to give you enough detail so that you can have an intelligent conversation about a topic. Enough, hopefully, to get you interested. If you want to know more: that's up to you. Explore away my explorer friend!



Your Feedback

A book is like any program you might write: *full of bugs*. I will be releasing new editions of this over time and I'm sure they will be largely influenced by good people like you. I plan on failing and, with that, *I plan on learning*. This book is a living thing! I've researched as much as I possibly could in order to surface every detail on every subject – but I will fail at this.

Our industry has been around long enough that many of these details have eroded, changed meaning over time, and have resurfaced anew. Many concepts have stayed intact... others have not.

Much of what you're about to read **is heavily arguable** and I've tried to point out these little arguable points, and also where I share my opinion. As Obi-wan once said:

...you're going to find that many of the truths we cling to depend greatly on our own point of view.

For instance: I wrote a section on software patterns and practices (SOLID, BDD, TDD, etc). **These subjects make people argue a lot.** I wrote another on databases (relational and non) and another on programming languages. It's simply not possible to write about these subjects and 1) keep them objective and 2) avoid kicking a hornet's nest.

So why did I include them? I believe that *some* discussion is better than none at all. I've been as careful as I can to highlight contentious subjects and places where I have an opinion - but you may find that I could finesse a bit more. Or you might find that I'm flat out wrong.

Which I think is **outstanding**. You have my email address (I emailed this book to you) - please feel free to drop

me a line! I'll happily amend things if I need to or, at the very least, you and I can have a discussion.

This is a living book. *It will change* and I want you to be a part of that!

Cover Photo Credit

The image used for the cover of this book (and the inspiration for the splash image on the website) comes from NASA/JPL:



The image is entitled HD 40307g, which is a “super earth”:

Twice as big in volume as the Earth, HD 40307g straddle the line between “Super Earth” and “mini-Neptune” and scientists aren’t sure if it has a rocky surface or one that’s buried beneath thick layers of gas and ice. One thing is certain, though: at eight times the Earth’s mass, its gravitational pull is much, much stronger.

David Delgado, the creative director for this series of images describes the inspiration for HD 40307g’s groovy image:

As we discussed ideas for a poster about super Earths – bigger planets, more massive, with more

gravity – we asked, “Why would that be a cool place to visit? We saw an ad for people jumping off mountains in the Alps wearing squirrel suits, and it hit us that this could be a planet for thrill-seekers.

When I saw this image for the first time I thought this is what it felt like when I learned to program. Just a wild rush of freakishly fun new stuff that was intellectually challenging while at the same time feeling relevant and meaningful. I get to build stuff? And have fun!?!? Sign me up!

This is also what this last year has been like writing this book. I can't imagine a better image for the cover of this book. Many thanks to NASA/JPL for allowing reuse.

A Word On Formatting

This ebook is intended to be read on a white background with flow enabled. There are quite a few choices you could make with your reader in terms of font and background and most of them should translate fine. If things look weird, however, it might be due to how your reader is scaling this book.

Here are the details. I wrote this edition using Scrivener and have 3 specific compilation settings:

1. **Epub.** The headers follow a progressive 14/18/24/36 point setting and use Helvetica as the font. The body text is 12pt. Iowan Old Style which should render fine on Apple devices.
2. **Mobi.** I scaled everything down for Kindles. The body text is 11pt Iowan Old Style which will render on some devices; otherwise it's your device default. Headers are 12/14/18/24 Helvetica, but your device will probably override that too.
3. **PDF.** This is where I had the most control, so I used 18/24/36/48 Bebas Neue as the header font and Iowan Old Style for the body at 13pt. It lays out quite nicely.

Of course what I see and what you see are probably going to be different. Here are the devices I tested this thing with:

- iBooks on my computer
- iBooks on my iPhone 7+, iPad and iPad pro.
- Kindle app on my iPhone 7+, iPad and iPad pro
- Kindle app on my computer

Each one of these apps/devices gloriously does its own thing with the layout depending on your settings. All of the line height/font settings I have in place are tossed in favor of your device's settings, which is as it should be.

But, once again, if it looks weird to you, try scaling the font size up or down. When things line up, that's probably my compiler target.

Code Samples

All code samples are *images*. I know this runs counter to what people do with blogs etc - the problem is that ebook readers wipe out formatting as I mention above. I want the code to stand out and to be legible; it wasn't that before. If you want to copy/paste the code you can do so from my [Github repo](#).

If you need to see the code close up, try zooming the image. If you're on iBooks, double tap the image to have it come up on the screen. These are high resolution PNGs so you should be able to zoom in really well and have the image remain clear.

TOC?

I purposely removed the inline table of contents because it took up a ton of useless space. *Every single PDF and ebook reader* has TOC functionality built in; this book is big enough already. I know that other books do it a different way... this is just my preference.

Finally: It's my hope that you can **set your reader to a white background**, 12pt. font (Iowan or Georgia, ideally), sit back and enjoy.

Rob Conery, August 2017, Seattle WA.

PART 1: HISTORY

CHAPTER ONE: COMPUTATION

As programmers, we instruct machines to solve problems for us. How did this even become possible? Algorithms have existed for millennia, but to use them you had to do things by hand, which is error prone. People began writing the solutions to complex problems in books, called “Mathematical Tables”. These people were called “computers” and in 19th and 20th centuries their jobs were given to machines which were also called “computers”. This is the story of how it happened.



Let's get deep for a minute. Moving beyond computers and programming, going behind everything we take for granted today. Let's dim the lights, turn on some Pink Floyd and get completely radical for a minute. Let's consider the very nature of our universe and what it means to compute things...

This might sound absurd, but stay with me. In April of 2016, scientists and philosophers gathered in New York to discuss that very question:

The idea that the universe is a simulation sounds more like the plot of “The Matrix,” but it is also a legitimate scientific hypothesis. Researchers pondered the controversial notion Tuesday at the annual Isaac Asimov Memorial Debate here at the American Museum of Natural History.

Moderator Neil deGrasse Tyson, director of the museum’s Hayden Planetarium, put the odds at 50-50 that our entire existence is a program on someone else’s hard drive. “I think the likelihood may be very high,” he said.

The problem with well-meaning articles like this one is that they tend to go for the low-hanging fruit, dumbing-down the core of what should otherwise be a compelling idea – simply to appeal to a mass readership.

Which is unfortunate, because the question has a solid foundation: the physical universe is indeed computed. The very nature of the physical universe describes a progressive system based on rules:

- **Cause and effect:** what you and I think of as “conditional branching”
- **Consistent, repeated patterns** and structure: the magical numbers pi, phi, and e (among others) hint at a design we can only perceive a part of. Indeed, Plato suggested that the world we see is just an abstraction of its true form.
- **Loops:** the two ideas above (cause and effect, pattern repetition) interact repeatedly over time. A day is a repeated cycle of light and dark, life is the same way – as is a planet orbiting the sun, the sun orbiting the center of the galaxy.

This has the appearance of what we think of today as a program – either one very big one or many small ones interacting. I suppose it depends on whether your Deity of Choice is into microservices or not...

Nature's Strange Programs

At this point you might think I’m probably taking this whole idea a bit too far – but I’d like to share a story with you about [Cicadas](#).

I was watching a fascinating show a few weeks back called The Code, which is all about math and some of the interesting patterns in the universe. At one point it discussed a small insect that lives in North America, the

[magicicada](#), which has the strange trait of living under ground for long periods during its young life.

These periods range from 7 years to 17 years although the most common periods are from 13 to 17 years. That's a very long time, but that's not the most surprising thing about them. The strangest aspect of these creatures is that they emerge from the ground on schedule: 7 years, 13 years, or 17 years.

Those are prime numbers. This is not an accident; it's a programmed response to a set of variables. Trip out, man...

The emergence, as it's called, is a short time (4-6 weeks) that the cicadas go above ground to mate. They sprout wings, climb trees and fly around making a buzzing noise (the males) to attract a mate. Why the prime number thing? There are a few explanations – and this is where things get a bit strange.

Predatory Waves

We know that predators eat prey, and each is constantly trying to evolve to maximize their chances of staying alive. Cats have amazing hearing, eyesight and stealth whereas mice counter that with speed, paranoia and a rapid breeding rate. You would think this kind of thing balances, but it doesn't.

It comes and goes in waves as described in [this article from Scientific American](#):

As far back as the seventeen-hundreds, fur trappers for the Hudson's Bay Company noted that while in some years they would collect an enormous number of Canadian lynx pelts, in the following years hardly any of the wild snow cats could be found ... Later research revealed that the rise and fall ... of the lynx population correlated with the rise and fall of the lynx's favorite food: the snowshoe hare. A bountiful year for the hares meant a plentiful year for lynxes, while dismal hare years were often followed by bad lynx years. The hare booms and busts followed, on average, a ten-year cycle...

A recent hypothesis is that the population of hares rises and falls due to a mixture of population pressure and predation: when hares overpopulate their environment, the population becomes stressed ... which can lead to decreased reproduction, resulting in a drop in next year's hare count.

This much makes sense and isn't overwhelmingly strange ... until this theory is applied to the cicada:

Now, imagine an animal that emerges every twelve years, like a cicada. According to the paleontologist Stephen J. Gould, in his essay "Of Bamboo, Cicadas, and the Economy of Adam Smith," these kind of boom-and-bust population cycles can be devastating to creatures with a long development phase. Since most predators have a two-to-ten-year population cycle, the twelve-year cicadas would be a feast for any predator with a two-, three-, four-, or six-year cycle. By this reasoning, any cicada with a development span that is easily divisible by the smaller numbers of a predator's population cycle is vulnerable.

This is where the prime number thing comes in (from the same article):

Prime numbers, however, can only be divided by themselves and one... Cicadas that emerge at prime-numbered year intervals ... would find themselves relatively immune to predator population cycles, since it is mathematically unlikely for a short-cycled predator to exist on the same cycle. In Gould's example, a cicada that emerges every seventeen years and has a predator with a five-year life cycle will only face a peak predator population once every eighty-five (5×17) years, giving it an enormous advantage over less well-adapted cicadas.

Who would have thought a tiny insect could master math in this way?

Overlapping Emergences

Another fascinating theory behind the prime-numbered emergence of these cicadas is the need to avoid overlapping with other cicada species. We're dealing with prime numbers here, and it just so happens that the multiples of 13 and 17 overlap the least of any numbers below and immediately after them:

CICADA EMERGENCE

1	2	3	4	5	6	7	8	9	10	11	12		14	15	16	
18	19	20	21	22	23	24	25		27	28	29	30	31	32	33	
35	36	37	38		40	41	42	43	44	45	46	47	48	49	50	
	53	54	55	56	57	58	59	60	61	62	63	64		66	67	
69	70	71	72	73	74	75	76	77		79	80	81	82	83	84	
86	87	88	89	90		92	93	94	95	96	97	98	99	100	101	
103		105	106	107	108	109	110	111	112	113	114	115	116		118	
120	121	122	123	124	125	126	127	128	129		131	132	133	134	135	
137	138	139	140	141	142	143	144	145	146	147	148	149	150	151	152	
154	155		157	158	159	160	161	162	163	164	165	166	167	168		
171	172	173	174	175	176	177	178	179	180	181		183	184	185	186	
188	189	190	191	192	193	194		196	197	198	199	200	201	202	203	
205	206	207		209	210	211	212	213	214	215	216	217	218	219	220	

Natural Computation

This is natural computation, there is simply no getting around that. There's nothing magical or mystical about what these cicadas do – they're adhering to the patterns and structure of physical world.

Bernoulli's [Weak Law of Large Numbers](#) states that the more you observe the results of a set of seemingly random events, the more the results will converge on some type of relationship or truth. Flip a coin 100 times, you'll have some random results. The more you flip it, the more your results will converge on a fifty-fifty distribution of heads vs. tails.

We're seeing the same thing with cicada emergences. After millions and millions of years of evolution, a natural emergence pattern is resolving that is based on math and the need to avoid overlap. We're seeing the computational machinery behind evolution itself, revealed by the Weak Law of Large Numbers. Prime number distribution in *two completely different population controls* (predators and emergence).

Fascinating stuff. Let's get back on track...

What Is Computation?

Human beings have understood that there is some process at work in the natural world, and they've struggled to express it in a way other than mysticism and spirituality. This is a tall order: how do you explain the mechanics behind the physical world?

Philosophers, mathematicians and logicians throughout history have been able to explain highly complex processes with equations and numbers. [The Sieve of Eratosthenes](#), for example, is a simple algorithm for finding all the prime numbers in a bound set. It was described at around 250 B.C. We'll be playing around with this algorithm later in the book.

Algorithms have been around for millennia, and if you ever wanted to use one you needed to break out your rocks and sticks, calculating things by hand.

In the early 17th century people who were really good at math were hired to sit down and calculate things. From encoding and decoding ciphers to tide tables and ballistic trajectories – if you needed to know the answer to a mathematical question you could hire a human computer to figure it out for you, or you could just go buy a book of Mathematical Tables in which they wrote their calculations for general use.

In the early 19th century, when the Industrial Revolution was in full swing and steam machines were powering humans forward, a number of mathematicians wondered if it was possible to have these amazing machines execute calculations along with pulling rail cars and lifting heavy loads. Human computers were error-prone, and their mistakes were costly, so the need was there.

These machine-minded mathematicians, however, had to find the answer to a very basic question: **What does it mean to compute something?** How do you tell that to a machine?

For philosophers this goes even deeper. If our very existence is a computed process, do we become gods when writing programs and creating applications? In a sense, yes; we're the creator and omniscient controller of every aspect of the programs we write.

Taking this further: if we truly live in a “lazy” universe that repeats old patterns instead of creating new, novel ones then it makes sense that all of the patterns, machinery and computation in our universe is repeated in the digital one of our creation.

Going even further: this suggests that our existence could very well be a repetition of this set of patterns, machinery and computation itself! A derivation, if you will. This would suggest that god (in whatever sense you observe the word) is probably a machine.

Deep stuff – and that's where I'll leave it.

Computation In The Steampunk Age

If I were to ask you what is the square root of 94478389393 – would you know the answer? It's unlikely – but if you did know the answer somehow, it would be due to some type of algorithm in your head that allows you to step through a process of deduction. Or maybe you have eidetic memory?

If you did have a photographic memory, you would be in luck if you lived a few hundred years ago. Before we had calculators, going all the way back to 200BC, people used to write down mathematical calculations in books called mathematical tables. These were a gigantic pile of numbers corresponding to various calculations that are relevant in some way. Statistics, navigation, trajectory calculations for your trebuchet – when you needed to run some numbers you typically grabbed one of these books to help you.

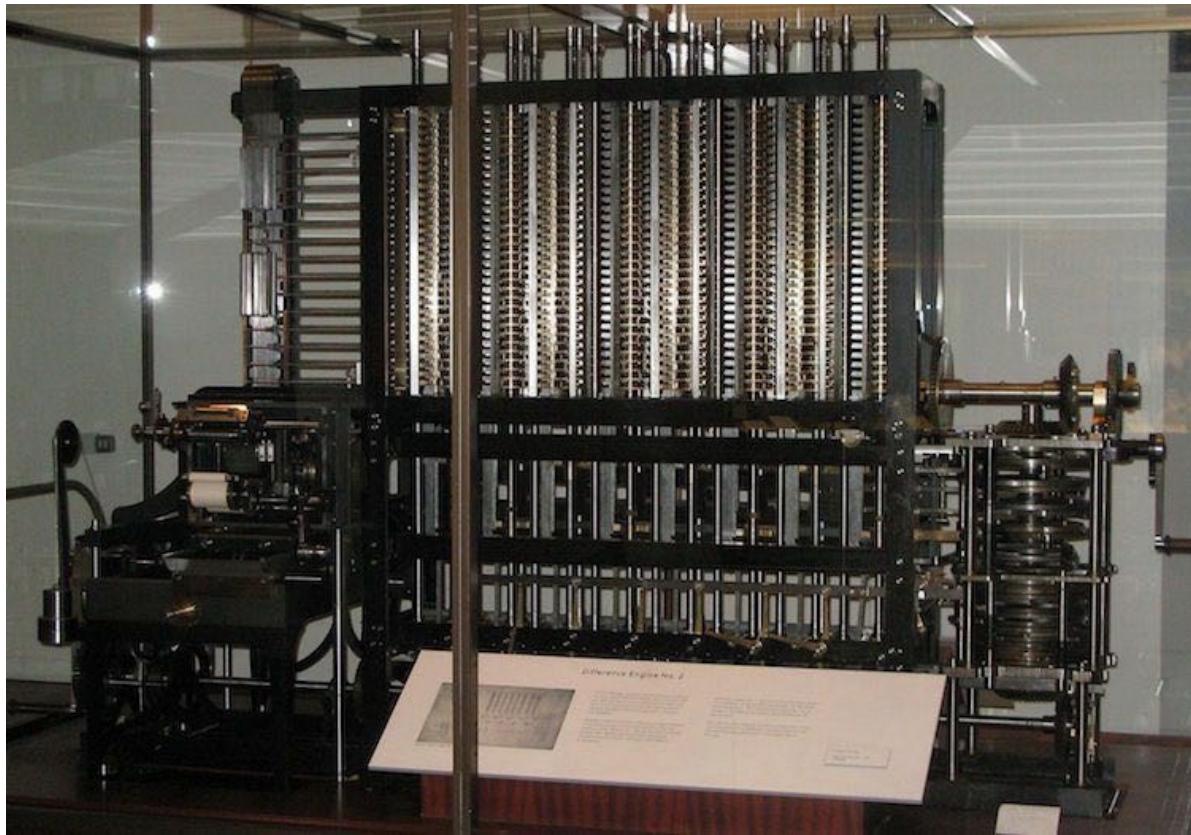
The problem was that these books were prone to error. They were created by human computers, people that sat around all day long for months and years on end, and just figured out calculations. Of course, this means that errors can sneak in, and when they did, it took a while to find the problem.

Errors like this are annoying and, in some cases, deadly. Ships were reported to have run aground because of errors in navigation tables – which were traced to errors in the mathematical tables used to generate them.

Charles Babbage, a mathematician, philosopher, and engineer decided that it was time to fix this problem. The industrial revolution was in full swing and machines were changing humanity at a rapid pace. Babbage believed they could also change mathematics by removing the need for a human being to be involved in routine calculations.

The Difference Engine

In the early 1820s Babbage designed The Difference Engine, a mechanical computer run by a steam engine. His idea was that, through a series of pulleys and gears, you could compute the values of simple functions.



Difference Engine #2. Babbage conceived a number of machines, none of them were completely built, however. The machine you see here was built by the London Museum of Science in the 1990s based on Babbage's plans. Image credit: Computer History Museum.

Babbage's machine could be "programmed" by setting up the gears in a particular way. So, if you wanted to create a set of squares, you would align the gears and start cranking the handle. The number tables would be printed for you and, when you were done, a little bell would ring.

Babbage had found a way to rid mathematical tables of errors, and the English government was all over it. They gave him some money to produce his machine, but he only got to part of it before the project imploded. After 10 years of designing, redesigning and arguing with other scientists – the government pulled the plug.

Which was OK with Babbage, he had another idea.

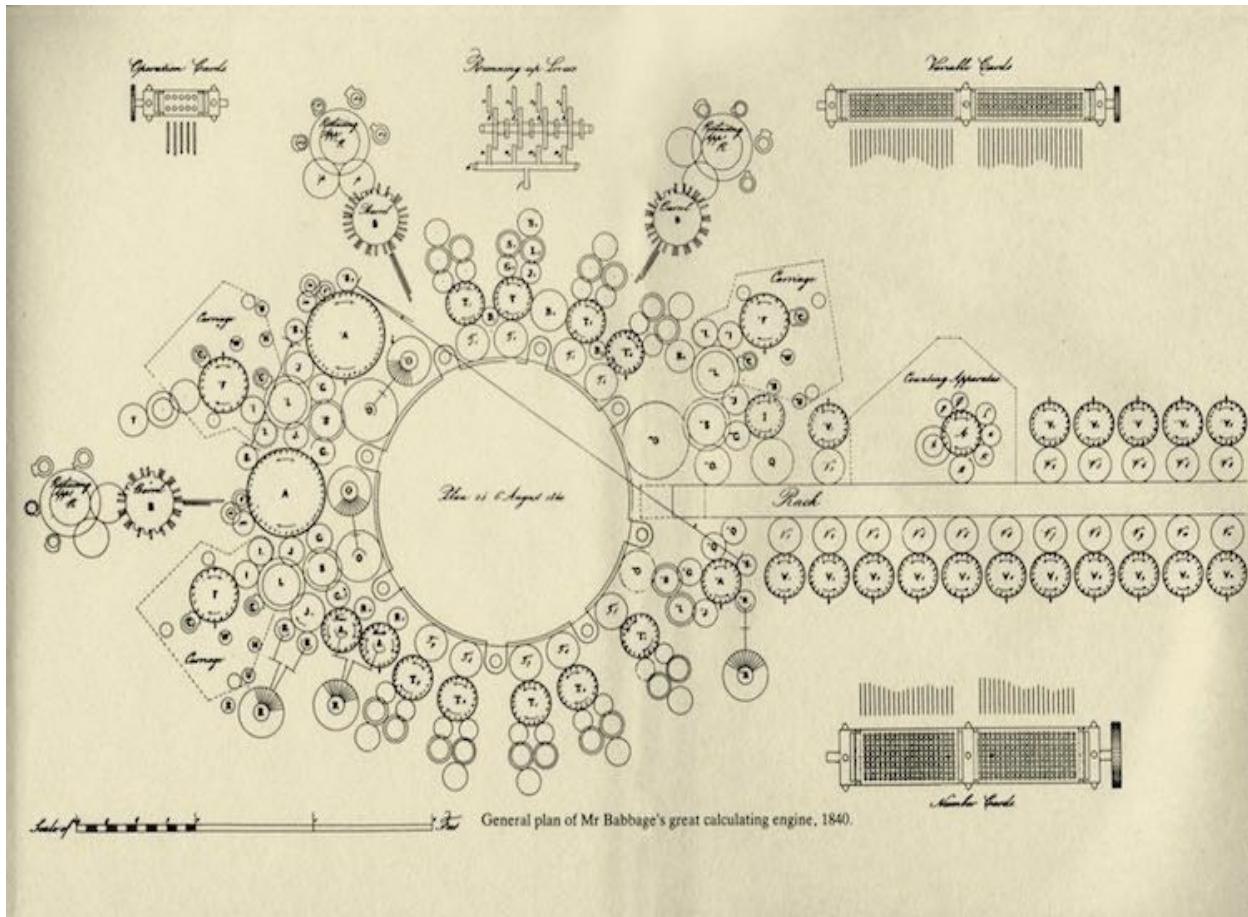
The Analytical Engine

Babbage believed his difference engine could do more. He was inspired by the Jacquard Loom, which was a programmable loom that could create complex patterns in textiles. The loom used a series of pins moving up and down to direct the pattern – and Babbage thought he could use the same idea for his machine using punch cards.

The idea was a simple one: tell the machine what to do by punching a series of holes in a card, which would, in turn, affect gear rotation. An instruction set and the data to act on – a blueprint for the modern computer.

According to the plans, this machine had memory, could do conditional branching, loops, and work with [variables and constants](#):

The programming language to be employed by users was akin to modern day assembly languages. Loops and conditional branching were possible, and so the language as conceived would have been Turing-complete as later defined by Alan Turing.



Plans for the Analytical Engine. Image credit: Computer History Museum.

Babbage knew he was onto something:

As soon as an Analytical Engine exists, it will necessarily guide the future course of the science. Whenever any result is sought by its aid, the question will then arise: By what course of calculation can these results be arrived at by the machine in the shortest time?

The interesting thing is that Babbage was focused on mathematical calculations. Someone else realized his machine could do so much more.

Ada Lovelace

Ada Lovelace is widely regarded as the very first programmer, although some contend that this statement is not only arguable, it also undermines her true importance: understanding the true affect of Babbage's Analytical Engine.

She was a brilliant mathematician and, interestingly, was the daughter of Lord Byron. In 1833 she met Babbage and instantly recognized the utility of what he wanted to build. They worked together often, and she helped him expand his notions of what the machine could do.

In 1840 Lovelace was asked to help with the translation of a talk Babbage had given at the University of Turin.

She did, and in the process added extensive notes and examples to clarify certain points. One of these notes, Note G, [stood out above the others](#):

Ada Lovelace's notes were labeled alphabetically from A to G. In note G, she describes an algorithm for the Analytical Engine to compute Bernoulli numbers. It is considered the first published algorithm ever specifically tailored for implementation on a computer, and Ada Lovelace has often been cited as the first computer programmer for this reason.

Diagram for the computation by the Engine of the Numbers of Bernoulli. See Note G. (page 722 et seq.)

Number of Operation. z	Nature of Operation.	Variables acted upon.	Variables receiving results.	Indication of change in the value on any Variable.	Data										Working Variables.										Result Variables.							
					Statement of Results.					v_{V_1}	v_{V_2}	v_{V_3}	v_{V_4}	v_{V_5}	v_{V_6}	v_{V_7}	v_{V_8}	v_{V_9}	$v_{V_{10}}$	$v_{V_{11}}$	$v_{V_{12}}$	$v_{V_{13}}$	$v_{V_{21}}$	$v_{V_{22}}$	$v_{V_{23}}$	$v_{V_{24}}$						
					1	2	n	1	2	1	2	n	2n	2n	2n	1	2	n	2n	2n	2n	1	2	n	2n	2n	2n	1	2	n		
1	X	$v_{V_2} \times v_{V_3}$	$v_{V_4}, v_{V_5}, v_{V_6}$	$\begin{cases} v_{V_2} = v_{V_2} \\ v_{V_3} = v_{V_3} \end{cases} = 2n$	2	n	2n	2n	2n	1	2	n	2n	2n	2n	1	2	n	2n	2n	2n	1	2	n	2n	2n	2n	1	2	n	
2	-	$v_{V_4} - v_{V_1}$	v_{V_4}	$\begin{cases} v_{V_4} = v_{V_4} \\ v_{V_1} = v_{V_1} \end{cases} = 2n - 1$	1	2n - 1	1	
3	+	$v_{V_3} + v_{V_1}$	v_{V_3}	$\begin{cases} v_{V_3} = v_{V_3} \\ v_{V_1} = v_{V_1} \end{cases} = 2n + 1$	1	2n + 1	1		
4	+	$2v_{V_5} + 2v_{V_4}$	$v_{V_{11}}$	$\begin{cases} 2v_{V_5} = v_{V_5} \\ 2v_{V_4} = v_{V_4} \end{cases} = \frac{2n - 1}{2n + 1}$	0	0		
5	+	$v_{V_{11}} + v_{V_2}$	$v_{V_{11}}$	$\begin{cases} v_{V_{11}} = v_{V_{11}} \\ v_{V_2} = v_{V_2} \end{cases} = \frac{2n - 1}{2 \cdot 2n + 1}$	2		
6	-	$v_{V_{13}} - 2v_{V_{11}}$	$v_{V_{13}}$	$\begin{cases} 2v_{V_{11}} = v_{V_{11}} \\ v_{V_{13}} = v_{V_{13}} \end{cases} = \frac{1}{2} \cdot \frac{2n - 1}{2n + 1} = A_0$	0			
7	-	$v_{V_3} - v_{V_1}$	$v_{V_{10}}$	$\begin{cases} v_{V_3} = v_{V_3} \\ v_{V_1} = v_{V_1} \end{cases} = n - 1 (= 3)$	1	...	n	n - 1		
8	+	$v_{V_2} + v_{V_7}$	v_{V_7}	$\begin{cases} v_{V_2} = v_{V_2} \\ v_{V_7} = v_{V_7} \end{cases} = 2 + 0 = 2$	2	2		
9	+	$v_{V_6} + v_{V_7}$	$v_{V_{11}}$	$\begin{cases} v_{V_6} = v_{V_6} \\ v_{V_{11}} = v_{V_{11}} \end{cases} = \frac{2n}{2} = A_1$	2n	2	$\frac{2n}{2} = A_1$		
10	X	$v_{V_2} \times v_{V_{11}}$	$v_{V_{12}}$	$\begin{cases} v_{V_2} = v_{V_2} \\ v_{V_{11}} = v_{V_{11}} \end{cases} = B_1 \cdot \frac{2n}{2} = B_1 A_1$	$\frac{2n}{2} = A_1$	$B_1 \cdot \frac{2n}{2} = B_1 A_1$		
11	+	$v_{V_{12}} + v_{V_{13}}$	$v_{V_{13}}$	$\begin{cases} v_{V_{12}} = v_{V_{12}} \\ v_{V_{13}} = v_{V_{13}} \end{cases} = -\frac{1}{2} \cdot \frac{2n - 1}{2n + 1} + B_1 \cdot \frac{2n}{2}$	0	0	$\left\{ -\frac{1}{2} \cdot \frac{2n - 1}{2n + 1} + B_1 \cdot \frac{2n}{2} \right\}$		
12	-	$v_{V_{10}} - v_{V_1}$	$v_{V_{10}}$	$\begin{cases} v_{V_{10}} = v_{V_{10}} \\ v_{V_1} = v_{V_1} \end{cases} = n - 2 (= 2)$	1	n - 2	
13	-	$v_{V_6} - v_{V_1}$	v_{V_6}	$\begin{cases} v_{V_6} = v_{V_6} \\ v_{V_1} = v_{V_1} \end{cases} = 2n - 1$	1	2n - 1	1	2n - 1	1	
14	+	$v_{V_1} + v_{V_7}$	v_{V_7}	$\begin{cases} v_{V_1} = v_{V_1} \\ v_{V_7} = v_{V_7} \end{cases} = 2 + 1 = 3$	1	3	3		
15	+	$2v_{V_6} + 2v_{V_7}$	v_{V_8}	$\begin{cases} 2v_{V_6} = v_{V_6} \\ 2v_{V_7} = v_{V_7} \end{cases} = \frac{2n - 1}{3}$	2n - 1	3	2n - 1	3	2n - 1	3	2n - 1	3	2n - 1	3	2n - 1	3	2n - 1	3	2n - 1	3	2n - 1	3	2n - 1	3	2n - 1	3	
16	X	$v_{V_8} \times v_{V_{11}}$	$v_{V_{11}}$	$\begin{cases} v_{V_8} = v_{V_8} \\ v_{V_{11}} = v_{V_{11}} \end{cases} = \frac{2n}{3} \cdot \frac{2n - 1}{3}$	0	$\frac{2n}{3} \cdot \frac{2n - 1}{3}$		
17	-	$v_{V_6} - v_{V_1}$	v_{V_6}	$\begin{cases} v_{V_6} = v_{V_6} \\ v_{V_1} = v_{V_1} \end{cases} = 2n - 2$	1	2n - 2	1	2n - 2	1	2n - 2	1	2n - 2	1	2n - 2	1	2n - 2	1	
18	+	$v_{V_1} + v_{V_7}$	v_{V_7}	$\begin{cases} v_{V_1} = v_{V_1} \\ v_{V_7} = v_{V_7} \end{cases} = 3 + 1 = 4$	1	4	1	4	1	4	1	4	1	4	1	4	1	
19	+	$3v_{V_6} + 3v_{V_7}$	v_{V_9}	$\begin{cases} 3v_{V_6} = v_{V_6} \\ 3v_{V_7} = v_{V_7} \end{cases} = \frac{2n - 2}{4}$	2n - 2	4	2n - 2	4	2n - 2	4	$\left\{ \frac{2n}{4} \cdot \frac{2n - 1}{3} \cdot \frac{2n - 2}{3} \right\} = A_2$		
20	X	$v_{V_9} \times v_{V_{11}}$	$v_{V_{11}}$	$\begin{cases} v_{V_9} = v_{V_9} \\ v_{V_{11}} = v_{V_{11}} \end{cases} = \frac{2n}{3} \cdot \frac{2n - 1}{3} \cdot \frac{2n - 2}{3} = A_3$	0	A_3		
21	X	$v_{V_{22}} \times v_{V_{12}}$	$v_{V_{12}}$	$\begin{cases} v_{V_{22}} = v_{V_{22}} \\ v_{V_{12}} = v_{V_{12}} \end{cases} = B_3 \cdot \frac{2n}{3} \cdot \frac{2n - 1}{3} \cdot \frac{2n - 2}{3} = B_3 A_3$	0	$B_3 A_3$	B_3		
22	+	$v_{V_{12}} + v_{V_{13}}$	$v_{V_{13}}$	$\begin{cases} v_{V_{12}} = v_{V_{12}} \\ v_{V_{13}} = v_{V_{13}} \end{cases} = A_0 + B_1 A_1 + B_3 A_3$	0	0	$\left\{ A_0 + B_1 A_1 + B_3 A_3 \right\}$	
23	-	$v_{V_{10}} - v_{V_1}$	$v_{V_{10}}$	$\begin{cases} v_{V_{10}} = v_{V_{10}} \\ v_{V_1} = v_{V_1} \end{cases} = n - 3 (= 1)$	1	n - 3	
24	+	$4v_{V_{13}} + 4v_{V_{14}}$	$v_{V_{24}}$	$\begin{cases} 4v_{V_{13}} = v_{V_{13}} \\ 4v_{V_{14}} = v_{V_{14}} \end{cases} = B_7$	B_7
25	+	$v_{V_1} + v_{V_3}$	v_{V_3}	$\begin{cases} v_{V_1} = v_{V_1} \\ v_{V_3} = v_{V_3} \end{cases} = n + 1 = 4 + 1 = 5$	1	...	n + 1	0	0	B_7

Lovelace's Note G. Image Credit: sophiararebooks.com.

Besides Bernoulli's numbers, she added this gem:

[The Analytical Engine] might act upon other things besides number, were objects found whose

mutual fundamental relations could be expressed by those of the abstract science of operations, and which should be also susceptible of adaptations to the action of the operating notation and mechanism of the engine...Supposing, for instance, that the fundamental relations of pitched sounds in the science of harmony and of musical composition were susceptible of such expression and adaptations, the engine might compose elaborate and scientific pieces of music of any degree of complexity or extent.

Some historians question how much Lovelace contributed to the first programs written for Babbage's Analytical Engine. We'll never really know the true answer to this question, but we do know that Lovelace completely understood the potential power of this machine.

Lost For a Century

The Analytical Engine was a bit of an oddity for the time. Babbage would muse on his plans until his death in 1871. Others picked up interest in his work, some even recreating parts of what Babbage specified in his plans – but none of these subsequent machines worked properly, and no one really seemed to care.

Babbage's Analytical Engine faded into scientific obscurity.

Researchers began to design electronic computers in the 1930s and 40s and although they were aware of Charles Babbage and the Analytical Engine, they hadn't taken the time to analyze Babbage's ideas:

The Mark I showed no influence from the Analytical Engine and lacked the Analytical Engine's most prescient architectural feature, conditional branching. J. Presper Eckert and John W. Mauchly similarly were not aware of the details of Babbage's Analytical Engine work prior to the completion of their design for the first electronic general-purpose computer, the ENIAC.

The need for faster, better computation ramped up in the 1930s and 40s as war broke out across Europe and the rest of the world. This caused much "parallel thinking", if you will, about models of computation and the idea of a computer.

War and Computers

The World Wars of the early 20th century pushed mathematicians and engineers to create the modern computer, as we know it. ENIAC, the world's first electronic computer (which we'll discuss in the next chapter) was designed to calculate artillery tables and ballistic trajectories for use on the battle field.

In 1939 Alan Turing designed and built the [Bombe](#), which helped the allies compute the daily settings for the German encryption device called Enigma. The machine depicted in the movie *The Imitation Game* named "Christopher", was the Bombe.

A few years later Tommy Flowers designed [Colossus](#), based on Turing's idea of a computational machine which, again, we'll discuss in the next chapter. This led to the Colossus Mark 1 which was the forerunner to ENIAC which led to EDVAC, which is the modern digital computer as we know it.

In terms of computation, all the hard work was done by Babbage, Lovelace, Turing and Church. They figured out how a machine could solve problems for us, coming up with blueprints, rules, and abstractions. All they needed was the engineering.

The rest of the 20th century handled that part. Growing up I clearly remember the advent of the home computer. My parents bought me a [TRS-80](#) when I was 11 and my brother taught me how to use a for loop (using BASIC) to print a tree on the screen.

Good friends of mine got an Apple II in the early 80s and I still, very fondly, remember playing King's Quest and Castle Wolfenstein on days when we should have been outside riding our bikes or playing baseball or football.

Today we have phones in our pockets that are exponentially more powerful than these machines – *but they use the exact same methods of computation* laid out by Babbage, Lovelace, Turing and Church. I bought my first iPhone right when it came out, back in June of 2007. That was almost 10 years ago!

What will happen in the next 10 years?

The Future of Computation

There are many smart people wondering where computational theory is taking us next, as you might imagine. Computer scientists are hard at work exploring holographic memory and quantum computing, and many others are wondering if there are things we might have missed in the last century.

My brother helped me out quite a bit with this chapter. He was good enough to listen to my ideas and push me in interesting directions – and didn't laugh at me too much.

So, to wrap this part of the book up, I wanted to do a blurb on where computational theory and computer science is going, in general. My brother asked me if I had ever heard of the [Ubiquity Symposia](#):

A Ubiquity symposium is an organized debate around a proposition or point of view. It is a means to explore a complex issue from multiple perspectives. An early example of a symposium on teaching computer science appeared in Communications of the ACM (December 1989).

Sounds very academic, doesn't it? But that's where you find the thinkers doing their thinking.

Anyway, he pointed out to me that there was a symposium on exactly my question: [What Is Computation](#):

What is computation? This has always been the most fundamental question of our field. In the 1930s, as the field was starting, the answer was that computation was the action of people who operated calculator machines. By the late 1940s, the answer was that computation was steps carried out by automated computers to produce definite outputs. That definition did very well: it remained the standard for nearly fifty years. But it is now being challenged. People in many fields have accepted that computational thinking is a way of approaching science and engineering. The Internet is full of servers that provide nonstop computation endlessly. Researchers in biology and physics have claimed the discovery of natural computational processes that have nothing to do with computers. How must our definition evolve to answer the challenges of brains computing, algorithms never terminating by design, computation as a natural occurrence, and computation without computers?

All these definitions frame computation as the actions of an agent carrying out computational steps. New definitions will focus on new agents: their matches to real systems, their explanatory and predictive powers, and their ability to support new designs. There have been some real surprises about what can be a computational agent and more lie ahead.

To get some answers, we invited leading thinkers in computing to tell us what they see. This symposium is their forum. We will release one of their essays every week for the next fifteen

weeks.

Jackpot. I asked my brother which of the papers he thought I should read first, and his answer was very typically my brother:

Well mine, of course!

Of course he wrote a paper on this subject ... why wouldn't that happen? Sigh ... big brothers...

So this is where I will leave this subject: staring off into the future while thinking about the past. I invite you to read through the papers in this symposia, starting with [my brother's](#) of course. They're rather short and quite interesting if you like academics.

CHAPTER TWO: COMPLEXITY

Complexity Theory deals with how difficult a given problem is. There are multiple classifications (P, NP, NP-Complete, NP-Hard, Exp, etc) that identify classes of problems based on the idea of time complexity. Over the years complexity theorists have found that many problems are related in some abstract way, and can be reduced, quite interestingly, from one to another.

As a programmer, you solve problems. You should know if the problem is considered easy, hard, or impossible. In addition, you should be able to identify the type or classification of a problem that you're trying to solve.



Every summer I get together with friends from college for the weekend. We'll rent a cabin in the mountains, go on hikes, fish and try not to hurt ourselves playing American football. I look forward to it every year; it's the only chance I get to see most of these people.

I'm sure you have a group of friends just like this one. People you've known for a very long time. Seeing them is always a treat... deciding what you want to do when you see them... well that's when things get more than a bit complex.

Making decisions as individuals is hard enough. Adding more people to the decision process takes something that was hard and makes it feel almost impossible.

Consider this scenario: my friends and I have just returned from a day fishing for trout in a high mountain lake. We're tired, hungry, and more than a little thirsty. We've cleaned ourselves up and someone asks:

So, anyone up for going somewhere for a pint?

There are 6 of us standing in the same room. I'm sure the scene isn't hard to imagine – we look from one to the other, trying to gauge what our answer will be. Someone says "sure", another says "yeah where?" and a third says "I don't drink but I'm happy to tag along if I can find something to eat".

Uh oh. Group inertia!

You've had these discussions, I'm sure. They're not very fun because it's difficult to figure out what's going to work out the best for everyone; there are just too many parameters to consider. Joe doesn't like bitter beers, Kim wants a cocktail, Kevin doesn't drink but is craving a pretzel and Olga wants to go her favorite pub 20 miles away because she knows the bartender and can get us a discount.

These situations are usually resolved by one person (typically the loudest) suggesting that everyone get in the damn car, we'll start out at place X and move on from there. This approach works remarkably well for a small group of people. If the group were to grow, however, that's when things get really complex.

What do I mean by really complex, however? In normal conversation we could just assume that this means "it's really hard to figure out". As programmers, however, we should have a more precise way of thinking about complexity. We deal with it every day – it's what we do. When your boss (or client) asks you the Big Question: how hard is it to add feature X?, how do you respond?

If you're like me, you might typically say "I'll have a look and let you know". Maybe you've done feature X before and you might say "it's doable".

Shouldn't we have better words for this conversation? Perhaps a way to more clearly define exactly how complex a given problem might be? Yes. We, as programmers need to be able to do this. The ability to clearly understand the complexity of a problem can save your boss or client tons of time and money. You might even avoid getting fired... like me... which is something I'll explain in more detail in just a few minutes.

Let's head back into the hills and see if we can figure out where my friends and I can go grab a pint and a bite to eat. Our situation is getting complicated, let's see if we can figure out just how complicated it's getting.

Simple Solutions And Polynomial Time (P)

My 6 friends and I are standing in the middle of the room, staring at each other, trying to figure out where we're going to go. People are stating their criteria and asking questions at the same time, and it's turning into a bit of a muddle.

Aaron, whose cabin we're staying at, speaks up and says:

Here's a list of the two pubs in town, and Olga's favorite pub 20 miles away. Put a check next to the one you want to go to. We'll go to the one with the most checks.

This type of solution is simple to implement and, best of all for the 6 of us, happens in a reasonable time frame. What's a reasonable time frame you ask? Good question! The answer is a bit abstract, but important nonetheless: it's less than the amount of time that would prevent us from doing it in the first place.

If deciding on a pub took us 15 minutes, we might get a bit frazzled, but we'd still do it. 30 minutes and we're grumpy for sure – anything more than that would be anarchy.

Thinking about this in terms of programming, we can think of "a reasonable time" as the time we might consider it acceptable and relevant to execute a routine for a given set of inputs. Sorting 100 things should come back in less than a millisecond, for instance. Sorting 1,000,000 things we might be willing to wait for a second or two.

Crunching analytics on billions of records, we might wait a few hours. Indexing a huge corpus of free text might take an entire night, or possibly a few days!

Using Aaron's democratic sorting process for selecting a pub is, best of all, fast. It also has another advantage: it scales really well. If our group of 6 turned into a group of 12, the process would take twice as long. A group of 15 would take 2.5 times as long.

Notice that I'm talking about the complexity of our democratic sorting process using time? This is the key to understanding Complexity Theory: you think about complexity in terms of time as you scale the inputs that go into the algorithm that you're using to solve the problem.

We can describe the way our democratic sorting process will scale using a bit of math:

$$T = 2x$$

Where T is the overall process time in minutes and x represents the number of friends. The equation that describes this scaling is a very simple one, and if you're a mathematician, you would call this type of a equation a polynomial equation:

In mathematics, a polynomial is an expression consisting of variables and coefficients which only employs the operations of addition, subtraction, multiplication, and non-negative integer exponents. An example of a polynomial of a single variable x is $x^2 - 4x + 7$. An example in three variables is $x^3 + 2xyz^2 - yz + 1$.

When the complexity of an algorithm scales according to a polynomial equation, mathematicians say that it scales in “P time”, where P stands for *polynomial*. These algorithms are the simplest and often the most boring to work with. List operations, for example, are P time algorithms. Things like sorting, searching, zipping and enumerating all happen in P time.

Complexity Theory, however, isn’t exactly concerned with the algorithms that are used to solve a problem. It focuses on the problems themselves. We wouldn’t say that our democratic sorting process executes in P time, that only matters to the programmers. A mathematician is much more concerned with the problem itself, and how difficult it is to solve. If one were sitting next to me, she would say that our pub selection problem is “in P”.

In more concrete terms, P is a complexity class that describes the set of all problems solvable in P time. Things like sorting and searching arrays, arranging your sock drawer and finding your cat who’s been wandering the neighborhood.

In other words: *simple problems*. These tend to be less interesting and not typically what programmers want to spend their time solving. Have you ever had to create a sorting algorithm by hand? I have! I had to do it for this book – and I’m going to make you do it with me in a few chapters. It was an interesting exercise but I don’t think I’d want to be a programmer if writing sorting algorithms was what I did all day.

99% of the time you and I are working with a class of problems that are not in P. They are significantly more complex (and therefore a bit more interesting) and have a classification all their own.

Hard Problems (Exp)

Aaron’s democratic pub selection process will indeed help us find a pub in short order, but it’s a bit too simplistic. It would be nice if we had a chance to consider everyone else’s opinion before making our own choice.

Olga speaks up with a suggestion:

How about each of us takes 30 seconds to explain where we want to go and why. Then we can decide after that.

This is a much nicer way to do things, and also a bit more complicated. Let’s think about Olga’s idea in terms of math.

Each of the 6 of us will be making a decision. With Aaron’s democratic method, there were only 6 decisions + 1: each of us decided where we wanted to go – that’s 6 decisions, plus one more in deciding the winner.

With Olga’s way of doing things, each of us needs to make 6 decisions apiece:

- I need to decide what I want to do
- I need to listen to the other 5 people in the room and then decide if I’m going to change my mind

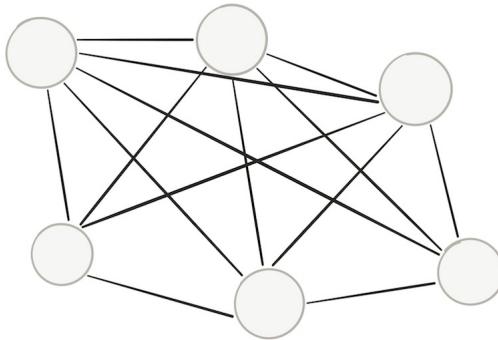
That’s 6 people * 6 decisions, which is 36 total decisions. If this is all it took to make a group decision like ours, we would still be in P time because the equation for calculating the time required is still a polynomial equation:

$$T = n^2$$

This is not the extent of what we need to do, unfortunately. We’re trying to *optimize* the selection of where we go – which means that I need to make a decision based on Aaron’s, and Aaron’s is based on Olga’s, and Olga’s is based on Todd’s which is also based on Kim’s…

If your brain is starting to hurt while thinking this through, I don’t blame you. It’s hard for any human mind to grasp the complexity of this type of process! In mathematical terms, this is a *factorial process*, or $n!$ where n is the number of people involved.

You can visualize this with a graph:



Each node on our graph represents one of my friends and each line represents a “weighting” in terms of a decision. If all we cared about was a measurement for each of the nodes, we would have a simple 36-part decision process. We want more than this, however. We want to optimize the combination of decisions, which is a factorial process.

This long-winded explanation is a way of saying that this problem is extremely complex. Rather than use the word “extremely”, however, a mathematician would say it’s “exponentially complex”, or that it is solvable in Exp time.

If we were to add a 7th person to our group – say my friend Kaveh finally shows up – the time to solve this problem would increase exponentially by an order of magnitude.

In pure mathematical terms, exponential time problems are solvable in:

$$T = 2^n$$

Where **T** is time, **n** are the inputs.

Exponential time is a very, very, very long time. It encompasses problems that could take millenia to complete... even millions, or billions of years. The sun can blow up 10 times over and our Exp problem could still be executing... chugging along... trying to figure out the optimal solution for us.

With small cases like ours, it would take us a long time to decide where to go, but it's likely that the sun would still be around when we're finished. Just because a problem is in Exp doesn't mean that it will always take a long time to solve. This is just a classification of complexity. Keep this in the back of your mind – I'll be coming back to it in a bit.

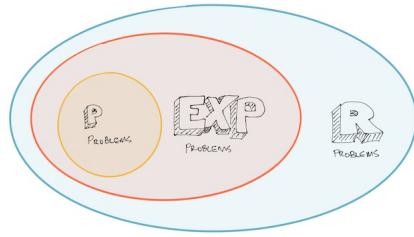
Almost every problem we can think of is in Exp. In fact, all of the problems we care about solving are in Exp. I'll get to explaining why that's true in just a bit – first, let's see what lies beyond exponential time.

All Solvable Problems (R)

Simply put, R represents all of the problems that are solvable in finite time. In other words: *any problem that we can solve is classified in R*.

You might be confused by some of this as these definitions overlap. Problems in P are “solvable”, aren't they? As are problems in Exp? Doesn't this mean they're also in R?

Yes. These classifications are sets of problems within given timescales. A problem that takes 10 seconds to solve is still solvable within finite time. It's also solvable in exponential time. These timescales overlap! It might be easier to see this in an illustration:



Infinitely Complex Problems Beyond R

What problems could possibly take an infinite amount of time to solve? Believe it or not, some deceptively simple ones. There are people who make a living researching these things, if you can believe it.

There's a very famous problem that lies beyond R, and it's known as *The Halting Problem*. You might have heard of it – it was introduced by Alan Turing in the very same paper he used to introduce the Turing Machine. We'll get to Turing, his machine, and the Halting Problem in detail later on – but I'll summarize it very quickly here.

Imagine you've written a program and you want to know if it will end at some point (aka "halt"). Will it return an answer? Will it crash? As it turns out: *it's not possible to write a generalized program that can decide if any other program will halt*. This type of problem is called *undecidable*.

The proof of this is rather elegant and I'll get into it in a later chapter. It has to do with feeding our halt-checking program back to itself, tweaking the internals before we do, which will throw the machine into an infinite loop. The "infinite loop" part is the thing that makes this problem undecidable.

There are quite a few other undecidable problems as well. In fact, if you dive into Complexity Theory a bit more you'll find that most of the problems that exist are undecidable, the proof of which is quite fascinating and a bit beyond my ability to explain. If you want to know more, have a look at MIT Open Courseware. There are hours of material up there.

Let's get back to the Halting Problem, however. I mention that it's solvable beyond R, but the strange thing is that the classification for this problem is something else entirely. It's classified as NP-Hard.

What does that mean? Let's dive into NP.

Determinism, Nondeterminism, and Magical Guesses

We have our three main complexity classes defined:

- P, which is the set of problems solvable in polynomial time. In other words: easy
- Exp, which is the set of problems solvable in exponential time. In other words: hard
- R, which is the set of all solvable problems

What do I mean, however, by the term solvable? As programmers we might think of writing a program to solve a given problem, or maybe constructing an algorithm of some kind. For the sake of moving things along, let's assume that these are the same things, and I'll talk about solutions in the context of a computer program.

When you sit down to solve a problem (let's say it's sorting an array), you lay out a set of instructions in code for the processor to execute. Conditional branches, loops, variables, etc. Hopefully you test your code because you need to make sure that the computer can determine the correct answer.

This is how we write computer programs: deterministically. We execute a routine and get a result. Based on that result, we'll execute another one, and then another. Each result links together until we arrive at an answer.

There is another way to arrive at this result, which is nondeterministically. In a nondeterministic system we execute a routine and get a result, from that point we might execute one of a series of next steps. We don't know which next step from that series will be executed – it will be decided at that time.

Once that next step is executed (whatever it is), we might execute one of another series of steps – once again having no idea which step in that series it will be until it's chosen. We keep on doing this until we arrive at the answer. When we do have our answer, there's simply no way of determining how we got there, at least from our deterministic program's point of view.

You can think of nondeterminism as using a series of lucky guesses or random chance (each of which is always correct) to figure out the answer to a given problem. If we were able to program a computer this way, things would get very interesting indeed.

Let's revisit the pub selection problem above, where my friends and I are trying to decide where to go have a pint and, possibly, a bit to eat. If I could use a nondeterministic algorithm to help us make this decision, we could leave in very short order! I could phrase the problem in a very particular way, and then off we would go:

- "If we go to the Leaking Moose, would everyone be optimally happy?" Answer: no
- "If we go to the Iron Tongue, would everyone be optimally happy?" Answer: no
- "If we go to the Cougar's Kitchen, would everyone be optimally happy?" Answer: yes

There it is! To arrive at this answer all I needed to do was to iterate over the possible choices and ask a simple yes or no question. My nondeterministic algorithm was able to arrive at the answer directly, because it's capable of correct lucky guesses every single time.

I'll explain more in a second (there is an explanation to this, trust me) – but it's important to understand what just happened. We took an exponentially complex problem and solved it in P time using a lucky nondeterministic process.

This might sound a bit magical, but there are quite a few people studying the notion of nondeterminism, trying to figure out if it's possible to have a nondeterministic algorithm that can solve exponentially complex problems like our pub selection problem. They like this notion so much that they've identified a whole group of exponentially complex problems that can be solved in P time if only we had this amazing nondeterministic algorithm.

These problems are classified as NP: problems solvable in P time given a nondeterministic algorithm. NP can be thought of as a bit of a mythical time frame, because we simply don't know if a nondeterministic algorithm is possible. If it is possible, then we can say that a given problem (like our pub selection problem) is solvable in nondeterministic polynomial time; which is what NP stands for.

More formally: a problem is classifiable in nondeterministic polynomial time (NP) if it is:

- Solvable in exponential time (Exp)
- Verifiable in polynomial time (P)
- Also solvable in polynomial time by nondeterministic methods

NP is where the action is. From the programs we write on a daily basis to the games we play at night with our family and friends: NP problems are what our brains enjoy the most.

Does P=NP?

We know, so far, that problems in P are pretty simple and can be solved in polynomial time using deterministic methods (aka "computers"). Problems in NP are quite a bit more complex, but become very simple if we have nondeterministic methods.

But we don't have those just yet. The question is: **will we ever have the ability to solve problems nondeterministically?** Some people believe it's just a matter of time before we have a computer chip or algorithm capable of nondeterministic processing. Others think it's a bunch of silliness. If a nondeterministic algorithm is ever developed that can solve these problems, then all of the problems in NP suddenly become solvable in P time. Put in mathematical terms: $P=NP$. If such an algorithm could never exist, then $P \neq NP$. We can't make either claim right now because we simply don't know the answer.

I explained, above, that you can think of nondeterminism as a “lucky guess”, which is sort of true. You can also think of it as a computer’s ability to make the right choice given a set of possibilities. There are languages that are capable of carrying these ideas out. Prolog, for example, allows you to define two functions with the same definition. The programmer does not tell the runtime which to choose (which would be deterministic), the runtime decides. There are various ways that this is carried out, including “backtracking”, where one function is tried, and if it fails, the runtime will back up and try the second function with the same definition. This is brute force, and won’t solve our pub selection in P time... yet. Someday? Maybe?

Who knows...

Reductions and NP

I was able to use nondeterminism to solve the pub selection problem in P time by reducing the problem into a series of yes/no questions. The original problem was something a bit different:

What are the optimal pubs for optimal group happiness?

This type of problem is exceedingly difficult in that we’re trying to optimize a combination of things: people and pubs. For most people (myself included), this isn’t something your brain is equipped to handle – it’s just not possible to keep all of the permutations together in your mind! These types of problems are called combinatorial optimizations and, as I mention, are very complex. But how complex are they?

Consider the definition of NP as it relates to this problem:

- Solvable in exponential time (check)
- Verifiable in polynomial time (no)
- Solvable in polynomial time by nondeterministic methods (no)

This means that this problem is not classifiable within NP and is, instead, in Exp. Sort of what you might expect as this problem is exceptionally hard.

Our reduction to a decision problem, however, is a different matter:

If we go to the Leaking Moose, would everyone be optimally happy?

This problem is solvable in exponential time, so we know it’s in Exp. It’s verifiable in polynomial time – all I have to do is ask if people are happy once we’re at the Leaking Moose. Finally, we already know that our nondeterministic algorithm can solve this problem in P time because we just did it.

That means our decision problem is classifiable in NP. This is where things get weird because, in essence, it’s the same problem! I just reduced it from an optimization problem to a decision problem!

This type of thing is very common - where one version of a problem is classified differently than its decision problem reduction. The good news (I think) is that there are two classifications for just such an issue:

- NP-Hard: problems that can be reduced to other problems in NP, but are not within NP themselves
- NP-Complete: decision problems classifiable in NP

I don’t blame you if your head is swimming right now. I know mine is... and I’m the one writing this book! If it helps: decision problems are almost always NP-Complete because of the ability to verify the answer to the problem (make sure it’s yes). Combinatorial problems are NP-Hard.

Before we move on to some examples, I want to revisit a statement I made above, about The Halting Problem

being classified as NP-Hard. If you've stayed with me through this whole explanation you might be able to reason through the answer.

As we know: solving The Halting Problem is beyond R; solvable only with an infinite amount of time. It is, however, a simple decision problem: will this program halt? Because of that, other problems in NP can be reduced to it, and that's all you need for a classification of NP-Hard.

NP-Complete and Decisions

One of the main goals of turning a complex problem into a decision problem is the idea of verification. With our initial, combinatorial pub selection problem, we can only verify that we have indeed gone to the optimal pub for our group by going through every iteration of the decision process for each person and each pub. A pub crawl might not sound so bad, but visiting each friend (and each combination of friends) to ask if they like the place would get more than a little tedious.

The decision problem variation, however, is very easy to verify, as we did above.

These types of problems were formalized by Leonid Levin and Stephen Cook in the early 1970s, when they formulated the Cook-Levin Theorem. This theorem states that any problem in NP can be reduced to what's known as The Boolean Satisfiability Problem (or SAT):

In computer science, the Boolean Satisfiability Problem ... is the problem of determining if there exists an interpretation that satisfies a given Boolean formula. In other words, it asks whether the variables of a given Boolean formula can be consistently replaced by the values TRUE or FALSE in such a way that the formula evaluates to TRUE.

This problem sounds intimidating, but it's really not. The best way to think about it is to consider a very long **if** statement (pseudo code):

```
if ((x and y) and (x and not y)) or ((x or y) and (x or not y)) then  
/.../
```

Given this, SAT wants to know what values for **x** and **y** will return true?

Imagine a flow chart, with lots of branches, end points and ultimately a satisfiable terminus which returns true. This is a variation of SAT.

Now, imagine the last program you wrote (or the one you're writing now). There are business rules, database queries, and probably some type of UI that a user can use. At its core, this program is a bunch of decisions and can be described by a large decision tree. This is also a variation of SAT.

I suppose if you look hard enough, you could find a software project that exists to solve a problem in P. Most of them don't, however, if they expect to deliver any kind of value. Software exists to make complicated things easier. Given this, most of what we do on a daily basis is to deal with problems in NP.

This is an important thing to recognize, mostly because problems can usually be reduced from one to the other, as we've seen. Put another way: it's highly likely that a mathematician somewhere has tried to solve the very problem that you're working on right now. To understand what I mean, we need to take a small trip back in time.

Classic NP-Complete Problems

In the early 1970s, mathematician Richard Karp expanded Cook and Levin's work with his paper Reducibility

Among Combinatorial Problems. In this paper, Karp showed that you could reduce a number of NP problems to SAT in polynomial time, and he came up with a list called “Karp’s 21 NP-Complete Problems”:

The problems [edit]

Karp’s 21 problems are shown below, many with their original names. The nesting indicates the direction of the reductions used. For example, **Knapsack** was shown to be NP-complete by reducing **Exact cover** to **Knapsack**.

- **Satisfiability**: the boolean satisfiability problem for formulas in **conjunctive normal form** (often referred to as SAT)
 - **0–1 integer programming** (A variation in which only the restrictions must be satisfied, with no optimization)
 - **Clique** (see also **independent set problem**)
 - **Set packing**
 - **Vertex cover**
 - **Set covering**
 - **Feedback node set**
 - **Feedback arc set**
 - **Directed Hamilton circuit** (Karp’s name, now usually called **Directed Hamiltonian cycle**)
 - **Undirected Hamilton circuit** (Karp’s name, now usually called **Undirected Hamiltonian cycle**)
 - **Satisfiability with at most 3 literals per clause** (equivalent to 3-SAT)
 - **Chromatic number** (also called the **Graph Coloring Problem**)
 - **Clique cover**
 - **Exact cover**
 - **Hitting set**
 - **Steiner tree**
 - **3-dimensional matching**
 - **Knapsack** (Karp’s definition of Knapsack is closer to **Subset sum**)
 - **Job sequencing**
 - **Partition**
 - **Max cut**

Figure 1.2: Karp’s 21 NP-Complete problems.

One of the best things you can do for your career is to get to know these and other NP-Complete problems, at least at a high level. They’re fascinating to understand! Who knows? You might even save your job someday...

Let’s take a look at a few.

Knapsack

This problem has been around for over 100 years, and is a combinatorial optimization problem that centers on packing a bag for the weekend:

Given a set of items, each with a weight and a value, determine the number of each item to include in a collection so that the total weight is less than or equal to a given limit and the total value is as large as possible.

Once again, we’re trying to optimize combinations of values: price and weight for a limited space.

Clique

This problem was first formulated in 1935, and has to do with graphs and graph theory:

Consider a social network, where the graph's vertices represent people, and the graph's edges represent mutual acquaintance. Then a clique represents a subset of people who all know each other, and algorithms for finding cliques can be used to discover these groups of mutual friends.

You can see how the solution to this problem could apply to any social aspect of an application, or grouping of like “things” that are self-assembling.

Bin Packing

This problem is a variation of Knapsack, and is once again a combinatorial optimization problem that you encounter quite often in the programming world:

... objects of different volumes must be packed into a finite number of bins or containers each of volume V in a way that minimizes the number of bins used.

If you've ever had to pack up your house or apartment and move, you've had to deal with the Bin Packing problem. This is a classic NP-Hard problem because of its combinatorial nature, and the fact that verifying that you've optimally packed things up means that you have to carry out every possible iteration to prove yours is the best.

We can reduce this to a decision problem, however, by iterating over bin configurations and asking if the current configuration is optimal. This reduction turns the combinatorial problem into a decision problem, which would classify it as NP-Complete.

Traveling Salesman

The classic NP-Hard problem of trying to figure out the cheapest way to send a salesman on a trip:

Given a list of cities and the distances between each pair of cities, what is the shortest possible route that visits each city exactly once and returns to the origin city?

Once again, a combinatorial optimization problem, which we can recognize now as NP-Hard. You can reduce this in P time to an NP-Complete decision problem by simply enumerating through every valid path and asking the question “is there a path that’s shorter”?

Approximations And Laziness

If you follow any of the links for these problems, you'll see that there are algorithms that exist which “solve” them in some cases in P time. These are called approximations and are very useful if you can tolerate their inexact nature.

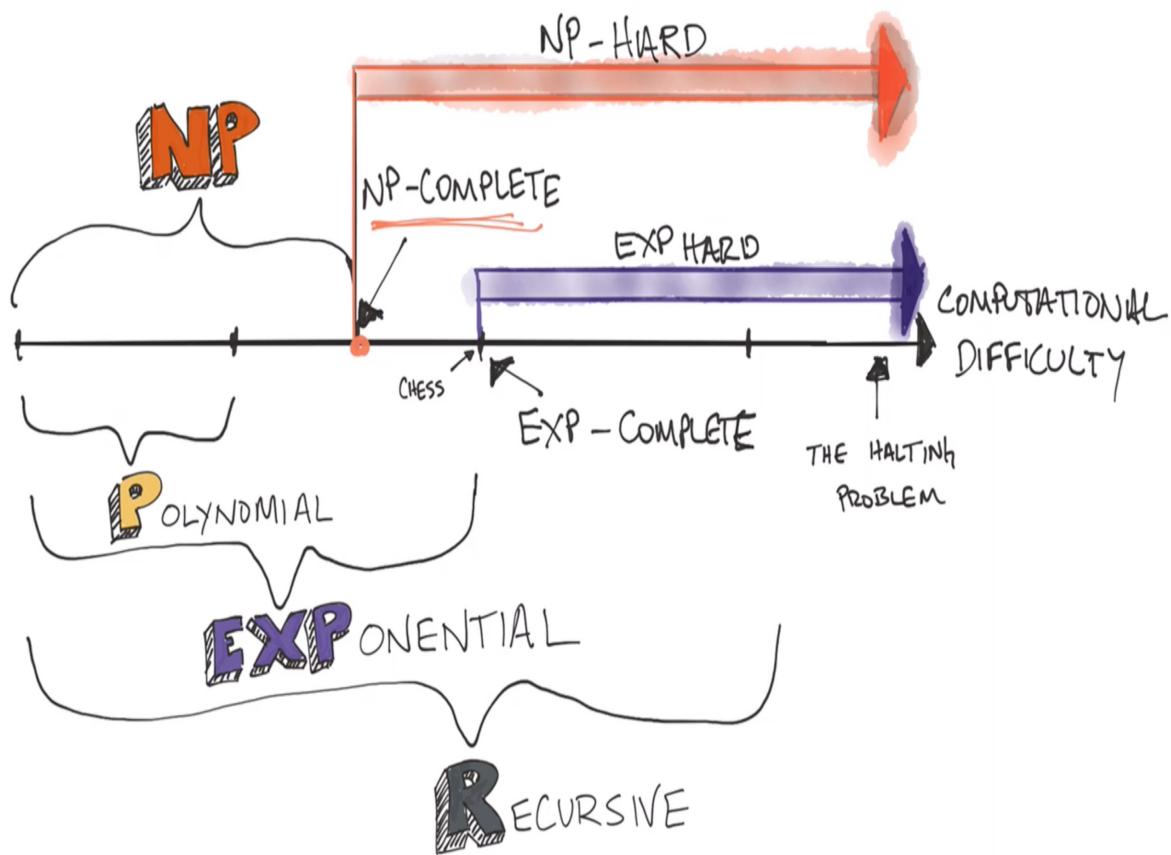
For Traveling Salesman, you could start from Los Angeles and head to the next nearest city, which (for our example) might be San Francisco. When you get there, you see that of all the destinations you could get to, Reno is the next nearest, so you go there. This approach is called “nearest neighbor” and is classified as a “greedy algorithm”, which means you do what suits your current position and value on the graph. Nearest neighbor usually returns a path within 25% of the shortest one, on average. Would that work for you? For many companies it just might.

In The Real World

I'd like to share a story with you about getting fired, and how I could have avoided it. Before I do, let's go over what we know so far:

- Simple, rather boring problems can be solved in a short amount of time; what a mathematician would call "P" time.
- More difficult problems, such as a group of people trying to decide on the optimal location for a pint and some food, are more complex and solvable in exponential time, or "Exp".
- Some problems are so complex that they can't be solved in all the time we can possibly have (R, or "finite time"), and are simply undecidable. For instance: Turing's Halting Problem.
- There are special subclasses of problems that are of great interest to us, specifically: exponentially complex problems that we can solve in P time with a nondeterministic lucky guess. These problems are in a classification called "NP". We can further divide this complexity class into decision problems that we can quickly verify (NP-Complete) and more complicated problems that are in Exp, but that other problems in NP can be reduced to in P time (NP-Hard).

We can visualize this on a timeline:



It might take a few reads to have this all sink in, but you might be wondering something: why in the world

should I care. For that, I'd like to share a story.

Where's The Party?

In 2010 I was hired to build a Ruby on Rails site for a client that wanted to enrich the lives of college freshman. The idea was simple: let's help these new students make friends and find interesting things to go do.

Simple enough! I had been building web sites for the previous 12 years, this should be straightforward. Indeed it was! I built the core of the site, integrated the designs, and all was well until a status meeting we had one Tuesday afternoon:

So, Rob how's the progress on the algorithm?

Not knowing what algorithm they were talking about, I asked. They responded:

You know, the one that is going to match students with other students and places to go, etc. We were told you were good with this kind of thing...

Oh no. It's true: one of the reasons I got the job was because I had a background in analytics, but I didn't think I'd be doing any of it for this site! They hired me to build a Rails app!

This is where my ego took over. I hate saying "no" to clients, and somewhere in my mind was a cowboy looking for adventure... so I said what turned out to be some fatal words: tell me more...

The "algorithm" turned out to be this:

- Gather a basic interest profile from new signups. Things like favorite music, extrovert/introvert, favorite actor, etc.
- Create a tagging system for submitted events. Anyone could submit an event to the system - but they had to tag it so we could match on it.
- Optimally match new students to other students and also to events that they would find interesting.

At this point in the chapter you should know what we're dealing with here. That's right: a combinatorial optimization. I could probably reduce this problem to any of Karp's NP-Complete problems, but (here, now, in the future) I don't need to do that to know that this problem is going to require some lengthy discussions.

Back then, I simply said "let me take a look".

You probably know where this is going to end up. **I got fired.** I looked into various ways to solve the problem but quickly realized that this isn't something that:

1. A Ruby on Rails site could handle
2. Anything that could happen in real time
3. I would even know how to start implementing

I pushed back as best I could, the client insisted this is why I was hired, I asked them to show me where I agreed to this in the contract... but it didn't matter. I got fired and the project blew up.

If I knew then what I know now about exponential time algorithms, approximations, decision problems and all the other lovely things discussed in this chapter – I might have been able to walk the client through the problem and why, no matter how many developers they hired then fired, they would never get what they wanted.

Maybe.

Have you worked on a project like this? I'm sure you have. Or, if not, it's likely you will. People try to tackle NP-Complete and NP-Hard problems all the time – after all we have machine learning and big data don't we?

The good news for you? Hopefully you will now be able to spot these efforts from the start and, hopefully, help your team think approach the problem with care.

CHAPTER THREE: LAMBDA CALCULUS

Before there were computers or programming languages, Alonzo Church came up with a set of rules for working with functions, what he termed lambdas. These rules allow you to compute anything that can be computed. You use Lambda Calculus every day when you write code. Do you know how it works? As a programmer, understanding Lambda Calculus can enhance your skills.



In the early 20th century, mathematicians began to ponder what it means to have a machine solve problems for you. The question was a simple one: how do you compute something?

The steps to solving certain problems (aka: algorithms) had been known for millennia; the trick was to be able to give these algorithms to a machine. But how? More than that: is there a limit to what a machine can calculate? Are there solutions that machines simply cannot compute?

This led to some interesting discoveries in the early 20th century, most notably by two men: Alan Turing and Alonzo Church. We'll talk about Alan Turing in a later chapter.

This section is about Alonzo Church's contribution: the **Lambda Calculus**. I should note here, as I've done in so many chapters, that I could spend volumes diving into the details of Lambda Calculus. What you will read here is a simple summary of the very basics. I do, however, think the missing details are quite important and if you care, I would urge you to have a look online – there are quite a few resources.

So, consider this a gentle introduction, the results of my recent headlong dive into the subject. Hopefully you will read enough to ignite your curiosity – which it should! What you're about to read is the foundation of computer programming.

Credit Where Due

I've read quite a few articles and text books on Lambda Calculus and I wanted to list them here, as I would never have been able to understand the basics otherwise.

At the top of the list is [this detailed explanation](#) of Y Combinator and Ω Combinator from Ayaka Nonaka. It is outstanding. I wanted to add some details about combinators and almost gave up, until I found this post.

Next is [Jim Weirich's amazing keynote](#) on the Y Combinator. I remember watching it years ago, having my mind blown. I watched it three times over when writing this chapter and most of it still goes over my head. I link to it again below.

Ben Hall has an [amazing Github repository](#) that has all kinds of Church encoding magic. If you find yourself lost in any of this, go study the code in his repo. The Church encoding you see below is based directly on his work.

Finally, I'd like to thank James G (last name omitted... but you know who you are) for his patience helping me understand a small but crucial aspect of reduction and substitution. In the original version of this chapter I managed to get a few things wrong; I wouldn't consider them critical, but they were wrong nonetheless. After a number of lengthy, wonderful emails James convinced me of these errors, which I was able to cross-check and verify easily. Thanks James!

The Basics

Alonzo Church introduced Lambda Calculus in the 1930s as he was studying the foundations of mathematics. As a programmer, you might recognize the description:

The λ -calculus is, at heart, a simple notation for functions and application. The main ideas are applying a function to an argument and forming functions by abstraction. The syntax of basic λ -calculus is quite sparse, making it an elegant, focused notation for representing functions. Functions and arguments are on a par with one another. The result is an intensional theory of functions as rules of computation, contrasting with an extensional theory of functions as sets of ordered pairs. Despite its sparse syntax, the expressiveness and flexibility of the λ -calculus make it a cornucopia of logic and mathematics.

A lambda is simply an anonymous function that can be thought of as a value. Most modern programming languages have some notion of an anonymous function, expression, or lambda. Lambda Calculus is where this idea arose. In fact Lambda Calculus is the foundation of what we consider programming today.

Lambda Calculus is an abstract notation that describes formal mathematical logic. There are no numbers or types; only functions. Like modern functional languages, a function in Lambda Calculus can be treated as a value. By arranging these functions carefully, you can build out some very interesting structures.

In this chapter we'll use Lambda Calculus directly, but we'll also jump into ES2016 (JavaScript) to see how some of the ideas might work with modern code.

Some Rules

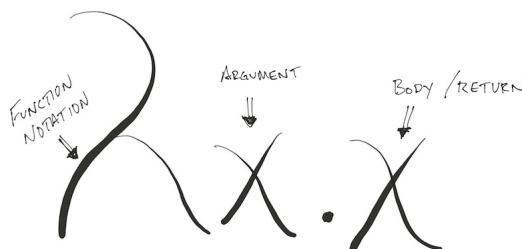
The rules of Lambda Calculus are rather simple:

- There are only functions, nothing else. No data types (strings, numbers, etc) of any kind
- You can substitute functions using variables
- You can reduce one function to another through application
- You can combine multiple terms in Lambda Calculus to create a higher-order function called a combinator, which is where the fun begins

There are rules which you can apply to these functions and combinators, and we'll get into a bunch of them using Lambda Calculus and also a bit of JavaScript.

Anatomy

Let's see our first function:



This is a Lambda Calculus function, sometimes called a “term” or “expression”. The first thing to notice is the Greek lambda (λ) on the left, which denotes a lambda function. The next bit is the x , which is the argument to the function. The final part of the expression is the body, which is the third x and is segregated from the argument by a ..

Using JavaScript, you could think of this function with this code:

```
function(x){return x};
```

The **function** keyword is equivalent to the λ symbol. Our function takes an argument **x** and returns **x**. While this is, indeed, an anonymous function in JavaScript, there is a more applicable syntax that can be used with ES6:

```
(x => x);
```

This is a pure lambda expression in the following ways:

- It takes in an argument, **x**, and since we're using a single line the value **x** is also returned
- It is a functional closure, which means we can set this expression to a variable and invoke it anywhere without worrying about scoping issues
- lambdas are the same as values

In this way, JavaScript follows Lambda Calculus conventions rather closely.

Function Application

Let's revisit our first function:

```
λx.x
```

This function takes an argument **x** and returns it. We can apply a value to this function like so:

```
λx.x (y)
```

This notation means “substitute all occurrences of **x** with **y**” and is called *function application*. It’s just the same as if we did this in JavaScript:

```
(x => x)(y)  
//or  
function(x){return x}(y) //returns y;
```

This function: $\lambda x.x$ in Lambda Calculus has a special name: the *identity function*. Whatever you pass to it is returned. It’s also called the *I Combinator* – which we’ll discuss in a bit.

Consider this function and application:

```
λx.y (z)
```

What do you think will happen here? It would be the same as doing this in JavaScript:

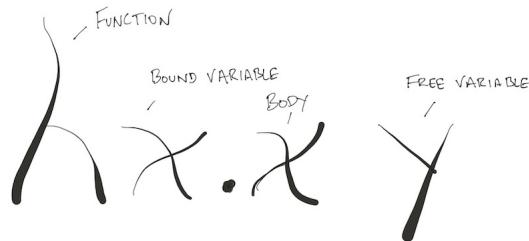
```
let z = 3;
function(x){return y}(z);
```

Here I'm using the value 3 just for clarity, though it's worth restating that there are no numbers in Lambda Calculus. There are *representations* of numbers, which we'll get to in a minute, but no strict numbers.

With this function: $\lambda x . y$, it doesn't matter what you pass into it, some value y will be returned. This is called the constant function as it returns a constant value of y no matter what you pass in for x .

Bound and Free Variables

When you apply a value to a function, there are some rules. The first thing you have to know is what *types of variables you're working with*. Note: I'm not referring to data types as there are none in Lambda Calculus; I'm referring to one of two different designations for variables:



In this expression we're dealing with two variables: x and y . The variable x is in the function head as the argument and thus is bound to this expression. The variable y on the other hand is not bound to any lambda function, so it's known as "free".

The difference between the two is very important. I can substitute x in this function if (and only if) I make sure to substitute every occurrence of the bound variable x . So if I wanted to change x to z , for instance, I could:

```
λz.z y
```

This wouldn't change the meaning of the function. If, however, I decided to change y we could run into trouble. To understand why, we need to dive into substitutions and reductions.

Substitution and Reduction

When you apply a value to a function, you substitute that value in the function itself. Substitution is *left-applicative*, which means you start on the left and move to the right as required.

For instance, let's substitute and then reduce this function:

```
λx.x 3
```

The first thing to do is to state that all bound occurrences (that's important, they have to be bound) of x will be replaced with 3 in the body, x :

```
λx.x 3  
x[x:=3]  
3
```

This notation: **x[x:=3]** can be read as “replace all bound occurrences of **x** in body **x** with 3”. Now, I know I said that there are no numbers in Lambda Calculus - and this is true! I’m using 3 here simply to show how substitution and reduction work.

Let’s apply this to JavaScript. Consider this expression:

```
let y = 3;  
var fn = (x => y)(1);
```

What do you think the value of **fn** will be? The same principles apply here: we’re substituting all instances of **x** with 1 in our equation, which doesn’t matter as we’re returning **y**, which is set to 3. This is, once again, the constant function so it’s no surprise that whatever we set **x** to doesn’t matter.

Let’s try this again with the identity function:

```
let y = 3;  
var fn = (x => x)(1);
```

In this case **fn** would “reduce” to 1.

Applying Multiple Values

You can create lambda functions that handle multiple values as well. It would be tempting to write them like this:

```
λxy.t
```

But this would be incorrect. Functions in Lambda Calculus only have an arity of 1 - which means they can only have one argument or one bound variable. To write this expression properly we would need to reorder things:

```
λx.λy.t
```

This is called currying: breaking a function with multiple arguments into chained functions with an arity of 1. We’ll talk about currying more in a later chapter.

In this expression, the lambda function **x** has a body of **λy . t**. This might look confusing, but if you remember that everything in Lambda Calculus is a function, it makes a bit more sense.

This corresponds to this lambda in JavaScript:

```
(x => y => x + y);
```

Reductions get interesting at this point. Consider this function:

```
λx.λy.y x
```

We need to apply the substitution from left to right, so our first substitution will be “replace all bound values of **x** with **x**”:

```
λx.λy.y x  
λy.y[x:=x]  
λy.y
```

There is no substitution in the body, **λy.y**, which leaves us with **λx.(λy.y)**, which is the constant function, which means we can reduce further to **λy.y** and we’re set.

We can see this reduction in action with JavaScript as well:

```
let first = (y => y);  
var second = (x => first(x));  
console.log(second(first));  
//[Function: first]
```

In more concrete terms, let’s revisit the JavaScript expression above (**x => y => x + y**). This time we’ll reduce it by running a substitution for **x** and **y**... in other words “invoking” the lambda:

```
(x => y => x + y)(2)(3);
```

In this code, **x** is set to 2 and then passed to **y**. The resulting function would then be **y => 2 + y**, which we apply 3 to, resulting in our answer, 5. This is straight up Lambda Calculus.

But what about this?

```
(x => x)(y => y)(2)(3);
```

This won’t work and we’ll get an error. Can you reason as to why? Let’s make it work and then step through it:

```
(x => x)(y => y)(2); //2
```

This substitution and reduction happens in the exact same way, but the result of one lambda is passed to the next. So, **y => y** is passed to **x => x**, which returns whatever was passed to it as it’s the identity function.

This leaves the value 2, which is then passed to **y => y**, which is once again the identity function that returns 2.

To see this in more detail, set **y => y** to **y => 5**. In this case **y => 5** is passed to **x => x** which, once again, just returns **y => 5**. If we pass 2 to that it’s ignored completely, so we get an answer of 5.

Just like any mathematical operation, the use of parentheses can affect the order of operations in lambda expressions. Consider this function:

```
λx.(λy.y x)
```

How do you think this would reduce? The parentheses dictate the order of operations so we would first reduce the body of the function like so:

```
λx.(λy.y x)
λx.(y[y:=x])
λx.x
```

If we omit the parentheses, something different happens:

```
λx.λy.y x
(λy.y)[x:=x]
λy.y
```

A completely different result.

These are the basics of Lambda Calculus, now let's see what we can do with it!

Church Encoding

Lambda Calculus looks a lot like programming doesn't it? Unfortunately this can cause confusion for a number of people trying to grasp it. Programming languages have numbers, strings, control structures and conditional branching built in; Lambda Calculus just has functions.

It's important to muse on this for a bit before we push forward. Put yourself in a classroom seat at Princeton, back in the early 1930s. Alonzo Church is trying to represent constructs that we take for granted today: conditional branches, loops, and higher order functions that can be used to compute things.

In order to do this, he set about creating *representations* for various values and operations. This is called *Church encoding*, and allows us to use booleans, numbers, conditional statements, and loops to construct things called *combinators* which are, unsurprisingly, combinations of functions that do a thing. We'll take a look at those later on; for now let's do the simplest of computations with Church encoding.

Booleans

How would you convey a simple conditional construct, like an if statement, using only symbolic functions? We need to have something like this:

```
if true then x
else y
```

The first step is to have a notion of what "true" means in Lambda Calculus, and we have that with the following representation:

```
λx.λy.x
```

The first bound variable is returned for a true statement in Lambda Calculus. Conversely, a false statement returns the second bound variable:

```
λx.λy.y
```

Formalizing this to JavaScript we might have:

```
let True = (x => y => x);
let False = (x => y => y);
True(true)(false) //true
False(true)(false) //false
```

That's a great start! Now we need to stretch this further in order to have a conditional statement. We need to evaluate 3 things; returning the first value if true, second if false. Using Lambda Calculus and leveraging Church encoding, we can use this expression:

```
λx.λy.λz.x y z
```

Translating this to JavaScript:

```
let If = (x => y => z => x(y)(z));
```

Our first argument, x, will be set to the function under evaluation. The arguments y and z will then be given to x to evaluate. The result of that will be the result of our If function.

Now, let's apply things. We can reuse our definitions of True and False to see if our If works:

```
let True = (x => y => x);
let False = (x => y => y);
let If = (x => y => z => x(y)(z));

If(True)("TRUE")("oops..");//TRUE
If(False)("oops")("FALSE");//FALSE
```

I'll be honest, this confused the heck out of me when I first tried to get through it. One of the only ways I could make it make sense was to draw it out, and when I did it all came together:

The True FUNCTION
↓
If = $(x \Rightarrow y \Rightarrow z \Rightarrow x(y)(z))$
↑
"TRUE"
argument
"oops"
↓
INVOKED True(), passing
"TRUE" and "oops"

We have booleans and conditional statements, now we just need to work with some values.

Numbers

You can represent numbers in Lambda Calculus by arranging a function to, basically, encapsulate and call itself.

It might be easier to see this rather than to explain:

```
λf.λx.x = 0  
λf.λx.f (x) = 1  
λf.λx.f (f(x)) = 2  
λf.λx.f (f(f(x))) = 3
```

This might look rather arbitrary at first, but there is some logic at work here. Consider this statement:

```
f(x) = x
```

A function of **x** is equal to **x**. This means the function essentially has no value at all. This, therefore, represents 0:

```
λf.λx.x
```

There is no reduction to be done here as nothing is applied. Thus **f** has a representative value of 0.

Now, consider this statement:

```
λf.λx.f x  
(λx.f) x[f:=x]  
λx.x
```

The application of this function results in the identity function, which is a bit like multiplying by 1.

Let's try this with JavaScript. The first thing we need to do is to set up a **calculate** function that will figure out how many times a given function is called. This function will accept a Church number and invoke it with a return that is itself a function. Every time that result function is invoked, it will increment itself:

```
let calculate = f => f(x => x + 1)(0);
```

Now let's define some Church numbers:

```

let calculate = f => f(x => x + 1)(0);

let zero = f => x => x;
let one = f => x => f(x);
let two = f => x => f(f(x));
let three = f => x => f(f(f(x)));

calculate(zero) // 0
calculate(one) // 1
calculate(two) // 2
calculate(three) // 3

```

Very nice! We now have a way of representing numerical values with nothing but a set of functions.

There is a whole lot more we can do here; things like addition, subtraction, multiplication, simple list operations and so on. In fact, we can encode our way to a Turing complete programming language. Which is really no surprise as the exercise we've just gone through is the foundation of what we think of as *programming itself*.

Combinators

We have booleans, numbers, and conditional branching; let's see what else we can do with Lambda Calculus. Before we get there, I should note that we've already been working with *combinators*, which are simply combinations of smaller functions that do a thing.

More precisely, a [combinator is defined as](#):

... a higher-order function that uses only function application and earlier defined combinators to define a result from its arguments.

In other words, we have lambda expressions and bound variables, nothing else. Let's see what we can do.

Loops and The Omega Combinator

Every good programming language needs to have the ability to loop, so let's see if we can implement that. We've seen that the identity function returns whatever is passed to it, so you might be tempted to do something like this:

```

λx.x (λx.x)
x[x:=λx.x]
λx.x

```

This works once, but not multiple times. We can see this with JavaScript and Church numbers:

```

let zero = f => x => x;
let one = f => x => f(x);
let two = f => x => f(f(x));

```

A previously-defined lambda can call itself, but there's no way to call a lambda from within its own definition. We've been working with JavaScript a lot and we're able to work with variables and labels so it might seem like we

can do this, but with pure Lambda Calculus, there are no such conventions as functions don't have a name.

But what if you could create recursion as a function itself? We already have a bit of a start with the identity function:

```
λx.x (λx.x)  
x[x:=λx.x]  
λx.x
```

In this application, however, it only returns what's on the right. How can we get it to repeat? What if we added one more application of x?

```
λx.xx (λx.xx)  
x[x:=λx.xx]  
λx.xx λx.xx
```

The Omega Combinator

Well now, look at that. A function that not only returns what it's given, but also its replica. This is called the Ω combinator, or “looping” combinator as it allows for recursively looping over a given function.

What does this look like in JavaScript? Let's see:

```
let Omega = x => x(x);  
console.log(Omega(Omega));  
//RangeError: Maximum call stack size exceeded
```

A recursive loop. This is interesting, but how can we apply some other function to this recursion? In other words, we have a basic looping structure, but it doesn't really do anything.

What would be more fun is to have a recursive function that would execute a given function once per iteration. In other words, what you and I might consider a for/each construct.

The Y Combinator

The Y combinator is a recursive, fixed-point function. It's sometimes called the “fixed-point combinator” as well. If you're like me you might not know what “fixed-point” means.

It is, essentially, when the result of a given function is the same as the input:

$$yf = f(yf) \text{ for all } f$$

Consider this function:

$$f(x) = x^2 - 3x + 4$$

If I set $x=2$, then the input will equal the output. This is the fixed point of this function. We want to expand on this idea, but instead of passing in a value to the Y combinator, we want to pass in a function, invoke it, and then get that function back so we can invoke it again later on.

This is what the Y combinator does: you give it a function as an argument which it invokes. It then returns that function back out to you.

More formally, the Y combinator is defined as:

$$Y = \lambda f. (\lambda x. f(xx))(\lambda x. f(xx))$$

It looks just like the Ω combinator doesn't it? The main difference, however, is that instead of just replicating xx , we're replicating the application $f(xx)$, where f is the initial function itself. In addition, the Y combinator allows you to pass in a second expression which represents the number of times you want the loop to execute. We'll see this below.

If you're stuck on this (which I was, for a very, very long time), go through the building up of the Ω combinator to see how the recursion/replication works. Ideally you should see the exact same pattern at work here; all we're doing is adding an additional function definition (f) and then invoking it.

In JavaScript it looks a bit different. You can't reference f in the second lambda definition as you'll get an error, so you need to wrap it with an additional lambda:

```
let Y = f => (x => x(x))(x => f(y => x(x)(y)));
```

The Y-Combinator in JavaScript

This is interesting, but what can you actually do with it? Let's see.

We need a function to iterate over, which you can think of as the code block you might pass to a **foreach** statement. As with all computer science books, Fibonacci must appear at some point, so let's use that:

```
let fib => f => n => n <= 1 ? n : f(n-1) + f(n-2);
```

If you don't recall, Fibonacci defines a set where a given number is the sum of the two previous numbers in the set. The first few numbers are:

```
1, 1, 2, 3, 5, 8, 13...
```

The expression above simply takes in a number, n , and checks to see if it's less than or equal to 1. If it is, then it's returned. If it's not, then the function calls itself to figure out previous values.

There is a better, faster way to write this which I'll go into in the chapter on algorithms. It's important to note, however, that it wouldn't be a pure lambda expression if I did that. We want a self-contained function expression. No variables, no conditionals, etc. Just functions - Lambda Calculus style.

Let's use Y Combinator to implement our function:

```
let Y = f => (x => x(x))(x => f(y => x(x)(y)));
let fib = f => n => n <= 1 ? n : f(n-1) + f(n-2);
let yFib = Y(fib);
yFib(10);
//55
```

Yes! It works! We've just used Y combinator to execute our **fib** function 10 times, which yields the result 55.

Want to play around with some other combinators and see what you can create? Head over to [Ben's Github repo](#) and have a play!

Summary

My head is starting to hurt a bit; it might be time to move on. Hopefully you have a grasp on the basics of what Lambda Calculus is and why it's important. Simple functions that do a simple thing, which you can arrange carefully to do more complicated things in a symbolic way.

A very powerful way to compute things, which allowed Alonzo Church to make this claim:

All total functions are computable.

Put another way: *if it can be computed, Lambda Calculus can compute it.* Turing agreed with this, and in 1937 he proved that his Turing machine provided the same computational abilities as Lambda Calculus, which led to the Church-Turing conjecture. We'll discuss that in a later chapter.

Today, this type of statement is kind shrug-worthy. We have powerful computers that can compute almost anything – so what? Back then, however, people with pencils and papers were the ones doing the calculations and they were limited in their ability.

This led mathematicians to ponder *just what could be computed* in terms of complexity – the human brain is only capable of so much. What was missing was a method of computation that could guarantee a result if the thing to do the computing had unlimited resources. Like a mechanical computer.

Those came along in just a few short years after Church and Turing's work and guess what we use today for programming these things?

Lambda Calculus.

CHAPTER FOUR: THEORY TO MACHINERY

Babbage, Lovelace, Church and Turing laid the foundation for how a machine could be used to compute things, but how did this actually come about? The abstract notions of machinery and computation led engineers like Jon von Neumann, JP Eckert and John Mauchly to design and build the very first electronic computer. Every computer that exists today is based on these abstractions and designs.

Over the years, the notion of an abstract ability to compute things has taken on many forms. Let's visit that now, starting off with a little history, once again. We'll visit Plato and ponder the true nature of things, drop in on Bernoulli in the 1500s and wind our way to Russia in the early 1900s. We'll visit Blechley Park and Alan Turing in the early 20th century, eventually ending up back in the United States with John von Neumann, the creator of the modern computer.



Probability and The Theory of Forms

At sometime around 400 BC, Plato mused that the world we see and experience is only a partial representation of its true form. In other words: an abstraction. The real world is either hidden from us, or we're unable to perceive it.

He based this notion on his observations of nature: that there is a symmetry to the world, to our very existence, that lies just beyond our ability to fully perceive it. Phi, the Golden Ratio, pi, and e are examples of some cosmic machinery that is just outside of our grasp.

To many, the natural world appears to be a collection of random events, colliding and separating with no guiding purpose. To a mathematician, however, these random events will converge on an apparent truth if we simply study them for a long enough time.

The Italian mathematician Gerolamo Cardano suggested that statistical calculations become more accurate the longer you run them. Jacob Bernoulli proved this in 1713 when he announced [The Law of Large Numbers](#) in his publication [Ars Conjectandi](#) which was published after he died.

This law has two variations (weak and strong) – but you can think of it this way: if you flip a coin long enough, the statistical average will come closer and closer to 50% heads, 50% tails. This might seem obvious – after all there are only two sides to a coin and why wouldn't it be a fifty-fifty distribution?

The simple answer is that reality tends to do its own thing most of the time, with apparent disregard for our mathematical postulations. The fact that we can completely rely on statistical models over a long enough period is astounding.

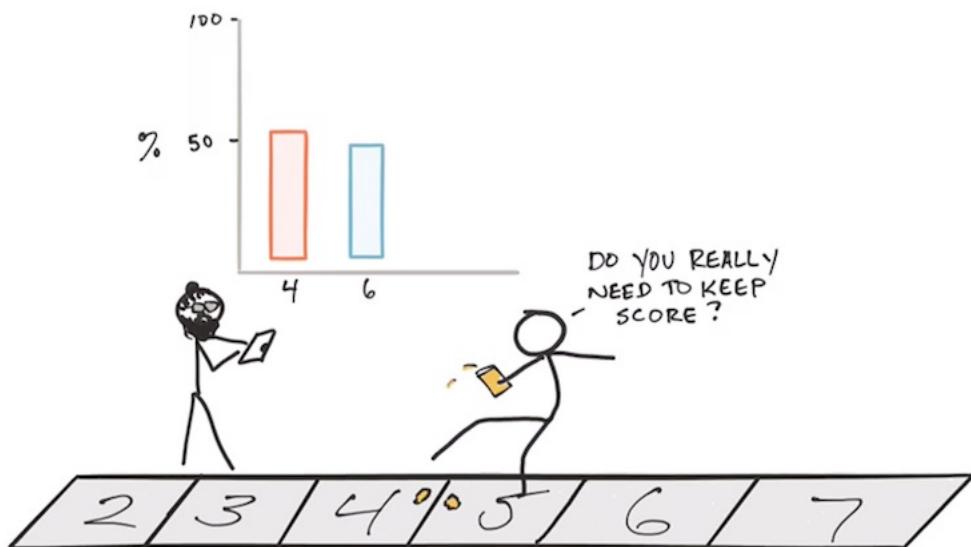
This is what keeps casinos in business. All they have to do is to make sure the mathematical odds of each of their games is in their favor, and over time the money they make on those games will reflect those odds increasingly closely. It doesn't matter if you walk in tomorrow and win all their money – statistics says they will get it back if they wait long enough.

Flipping a coin, however, is just a single event. Playing a game of craps or a hand of black jack consists of a series of events, one dependent on the next. Does the law of large numbers still hold?

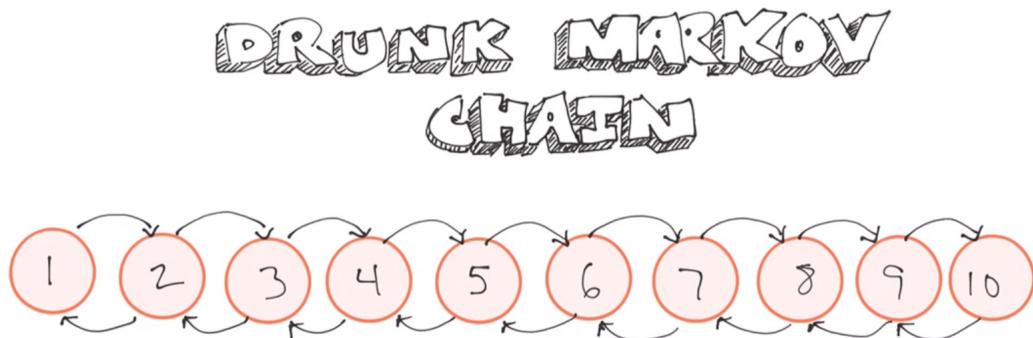
The Russian mathematician Andrey Markov said yes, and set about to prove it in the early 1900s with what has become known as the Markov chain.

If you've ever used a flow chart, then you'll recognize a Markov chain. It is usually described graphically as a set of states with rules that allow transition between those states:

A famous Markov chain is the so-called “drunkard’s walk”, a random walk on the number line where, at each step, the position may change by +1 or -1 with equal probability. From any position there are two possible transitions, to the next or previous integer. The transition probabilities depend only on the current position, not on the manner in which the position was reached. For example, the transition probabilities from 5 to 4 and 5 to 6 are both 0.5, and all other transition probabilities from 5 are 0. These probabilities are independent of whether the system was previously in 4 or 6.



A Markov chain looks a bit like a flow chart, and in many ways that's exactly what it is: a probability graph of related events. This fits nicely with the notion of an abstract computing machine.



In the above diagram, each orange circle is a “state” and the arrows dictate transitions possible between each state. Except for 1 and 10, the only transitions possible at any state along the chain is the next direct state or the one

prior. Notice also that none of the states allows for a “loop back”, or a transition back to itself.

This diagram represents a very simple abstract computation, or a “machine”.

Finite State Machine

If we focus only on the notion of state and transitions, we can turn a Markov chain into a computational device called a Finite State Machine (or Finite Automata).

A simple machine does simple work for us. A hammer swung in the hand and striking a nail translates various forces into striking power, which drives a nail into wood. It doesn’t get much simpler than that.

The nail, as it is hit, moves through various states, from outside the wood to partially inside the wood to fully inside the wood. The nail transitions from state to state based on actions imparted to it by the hammer.

The action of the hammer transitioning the nail through various states constitute a state machine in the simplest sense.

Deterministic

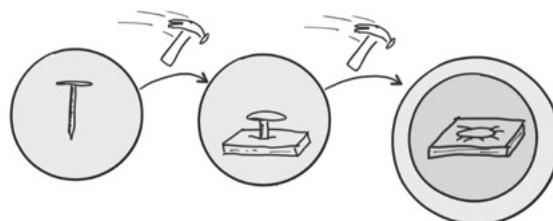
We discussed determinism a few chapters back, and we also discussed Finite State Machines a little bit. Let’s go a bit deeper.

A Finite State Machine (FSM) is an abstract, mathematical model of computation that either accepts or rejects something. We can apply this abstraction to our hammer and nail process above, so we can describe things mathematically. We start by diagramming the state of the nail:

- Outside the wood is the starting state
- Partially driven into the wood is another state
- Fully driven into the wood is a final state, also called acceptance
- A bent nail is also a final state, but not the state we want so it is a rejection state



We can relate these states together through various transitions or actions, which is the striking of the hammer:



This is a deterministic finite state machine, which means that every state has a single transition for every action until we finally reach acceptance or rejection. In other words: when we hit the nail, it will go into the wood until its all the way in, which is our accepted state (denoted by a double circle) – there is no other course of action.

Being a good programmer (which I'm sure you are), you're probably starting to poke holes in this diagram. That's exactly what you should be doing! It's why these things exist!

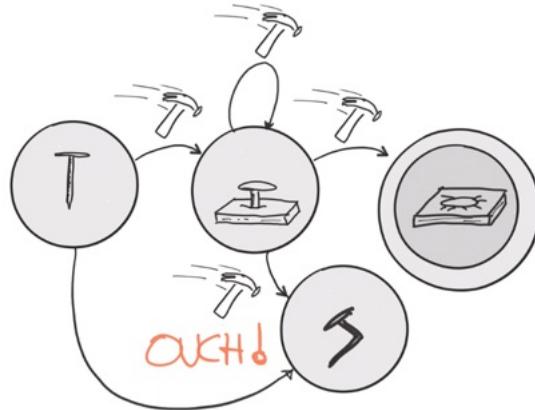
Our FSM only describes moving the nail through three states with only three actions. Furthermore, there is no rejection state for when I partially hit the nail and bend it. How do we account for that?

Nondeterministic

A non-deterministic finite state machine has one or more transitions for a set of actions, and these transitions are essentially “random”. A nondeterministic machine is capable of transitioning from one state to the next as a result of some random choice, completely independent of a prior state.

As a programmer, you can think of a nondeterministic machine as a program that will produce different results when run multiple times – even with the same input given.

Let's update our FSM to be nondeterministic by introducing partial hits which can sometimes bend the nail:



Here, I have two possible transitions for the initial state: I hit the nail, or I miss and bend the nail (yelling loudly when I do). I've added an additional action on the second state as well – the one that loops back to itself. This is the action taken when the nail has not been fully driven into the wood.

The conditional step of hitting the nail multiple times is deterministic, however the random step of bending the nail is not as we don't know if it will happen before we start hammering. Sort of. There are all kinds of variables and probabilities at play here from muscle twitches to fatigue, hammer integrity and wind direction. It will all come together at some point and the probabilities are so out there that we might as well consider this to be a random event.

Alphabets and Language

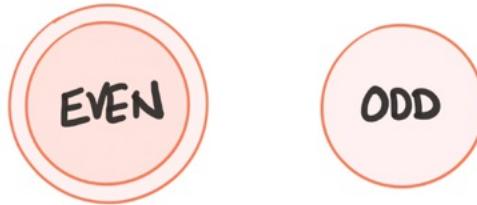
Let's move away from hammers and into the world of code and machine processes. You'll often hear people talking about FSMs working over a particular alphabet, string, or language. This is particularly relevant in computer science.

At the lowest level, our alphabet is a bunch of 1s and 0s – bits and bytes on a disk or in memory that we need to understand. This is how computer science people thought about the world decades ago: as holes on a piece of paper fed into a machine.

Let's do another FSM, but this time we'll do some real computering. Let's take a string of bits and write a

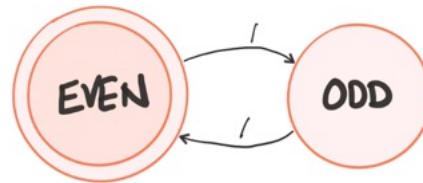
routine that will tell us if there is an even number of 1s in the supplied string.

We'll start out with the two possible states:

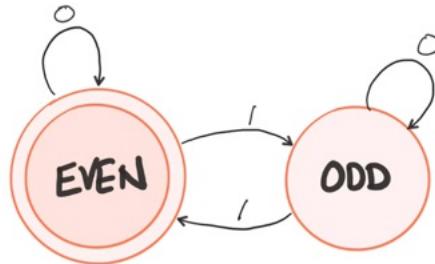


The double-circle here is our accept state, and if that's also our final state we have acceptance. Now we need to account for the action of reading each character in our string, one at a time.

If we read a 1, we'll have an odd number of ones. If we read a 1 while in the odd state we'll transition to the even state. Another 1 and we're back to the odd state:



What happens if we read in a 0? We do nothing and just keep on reading:



Limitations

Here is the part where we come to the naming of things. An FSM relies on a simple, known or finite set of conditions. The inputs are 1s or 0s, we have 1Mb of RAM to work with, or maybe 10Mb of hard drive space – these are finite conditions and tend to describe simpler constructs.

Street lights, alarm clocks, vending machines – these are perfectly good Finite State Machines. An iPhone is not.

To understand this, think about our 1s and 0s example above. What if we wanted to sum all of the 1s and then divide that by the sum of all the 0s to get a randomization factor? In short: we can't describe this with an FSM. There is no notion of a summing operation.

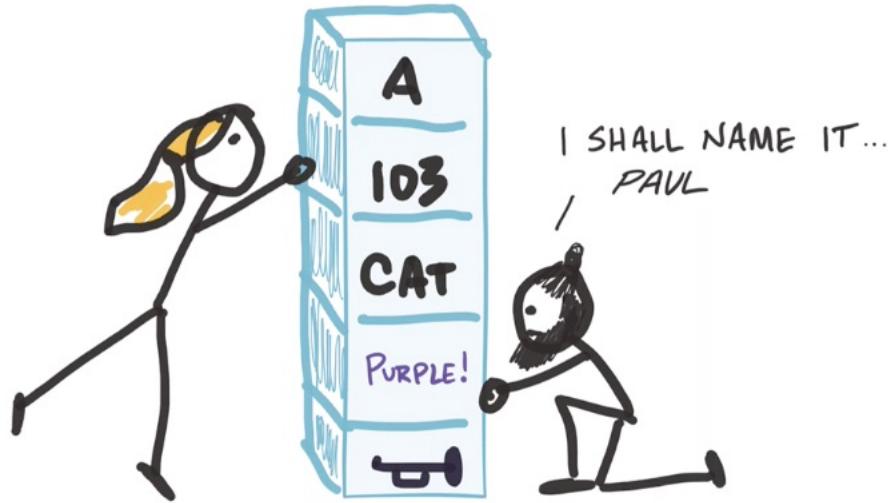
The basic problem is this: *you can't store state outside of the current state*. In other words: there is no persistence, or RAM.

We can't calculate the digits of pi, or perform map/reduce operations. We can, however, parse Regular Expressions (ugh).

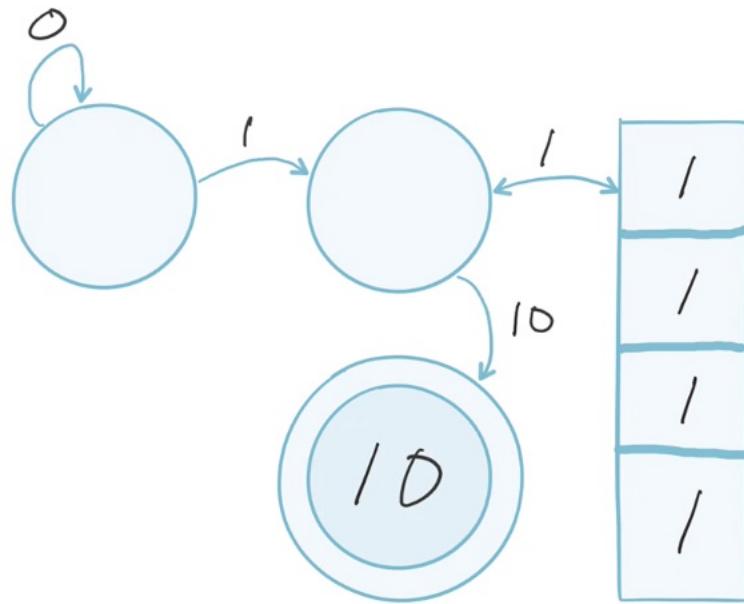
Much of the world runs on FSMs, but for what you and I do to make a living – we need a computational model that allows for a little more complexity.

Pushdown Machine

If we want to sum things with our computation, we need to move away from the concept of an FSM and add computational power to our machine. To sum things, our machine will at the very least need to remember the last state. We can add a facility for this called a stack.



A stack is quite literally that: a stack of data that you can add to (push) or remove from (pop). If we alter our FSM to have a stack, we can now compute a whole new set of problems:



This, however, is no longer a Finite State Machine – it's called a Pushdown Machine (also called Pushdown Automata or PDA). This machine has a lot more computational power because the state transitions are decided by three things working together:

- The input symbol (a 1 or a 0)
- The current state
- The stack symbol

With this new power we can compute running totals of 1s, creating a summing operation, and a whole lot more.

The notion of a stack is very powerful, but it's also a bit limited. You can only compute something that the space on the stack can handle. In other words, if we had 10 possible states and only 5 slots in our stack and we wanted to run a summing operation as before, we could easily run out of memory and our application would crash.

The Pushdown Machine is bounded by its stack and is, therefore, limited in terms of what it can compute. But what if the stack didn't have a limit?

This problem was solved by Alan Turing.

Turing's Machine

In the mid 1930s, Alan Turing published what has become one of the cornerstones of computational theory as well as computer science in general: his paper, entitled *On Computable Numbers*.

Written when he was just 24 years old, *On Computable Numbers* described a computational process in which you stripped away every “convenience” and introduced a machine with a read/write head and some tape. The tape has a set of cells, each of which holds a simple symbol of some kind, and you can move that tape under the head so the machine could read from it, or write to it:

We may compare a man in the process of computing a real number to a machine which is only capable of a finite number of conditions $q_1: q_2: \dots: q_I$; which will be called “m-configurations”. The machine is supplied with a “tape” (the analogue of paper) running through it, and divided into sections (called “squares”) each capable of bearing a “symbol”. At any moment there is just one square, say the r -th, bearing the symbol $\Sigma(r)$ which is “in the machine”. We may call this square the “scanned square”. The symbol on the scanned square may be called the “scanned symbol”. The “scanned symbol” is the only one of which the machine is, so to speak, “directly aware”. However, by altering its m-configuration the machine can effectively remember some of the symbols which it has “seen” (scanned) previously. The possible behaviour of the machine at any moment is determined by the m-configuration q_n and the scanned symbol $\Sigma(r)$. This pair $q_n, \Sigma(r)$ will be called the “configuration”: thus the configuration determines the possible behaviour of the machine. In some of the configurations in which the scanned square is blank (i.e. bears no symbol) the machine writes down a new symbol on the scanned square: in other configurations it erases the scanned symbol. The machine may also change the square which is being scanned, but only by shifting it one place to right or left. In addition to any of these operations the m-configuration may be changed... It is my contention that these operations include all those which are used in the computation of a number.

Each “machine” that Turing describes is designed to take another machine as its input, which allows for massive computational power. Alonzo Church had the same notion with Lambda Calculus: functions that take other functions as arguments. This flexible structure is the foundation of modern computer science.

We can write a Turing Machine (using mathematical notation) to describe the number 4 (M4). We can write another machine to describe the number 6 (M6) – and yet another to perform multiplication on a set of numbers (MX).

We can then write a final Turing Machine that accepts M4 and M6, and uses MX to run the multiplication. This is central to Turing’s idea: small, concise machines orchestrated to derive a final result . 1s and 0s are all we need to describe and compute any problem we can describe to the machine.

This led Turing to claim something rather extraordinary:

If an algorithm is computable, a Turing Machine can compute it

This is another way of stating what Alonzo Church asserted two years prior:

All total functions are computable

This became known as the **Church-Turing Conjecture**, and was an extraordinary claim for the time. Prior to the idea that machines could calculate something for us, mathematicians would declare that a given problem was effectively calculable if someone could sit down and figure it out with a pencil and some paper.

Church and Turing, however, gave us a different way to run calculations and, therefore, a different way to think about what we can compute, in general. Put in a simpler way: **if you can describe it with a Turing machine or with a set of functions, it can be computed.**

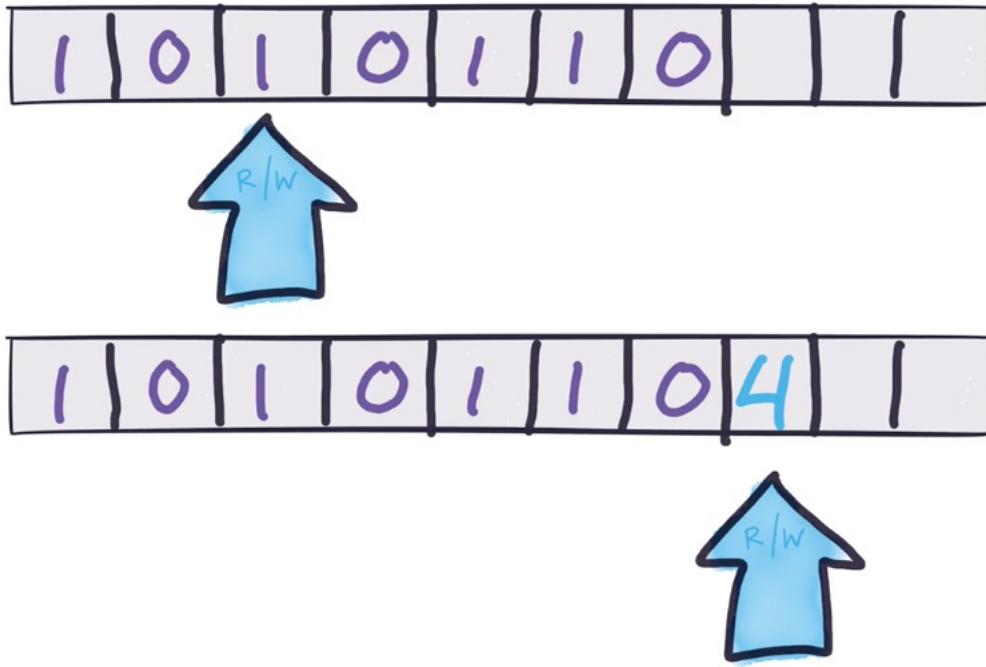
Basics

A Turing machine is an abstract computational device that can compute anything that you can program. It’s important to note that this is not supposed to be a real machine although some have actually made one.

It’s an incredibly simple construct that can simulate any computational algorithm, no matter how complex.

A Turing Machine has four main parts:

- A set of symbols, defined in an alphabet or language the machine can understand (usually just 0s and 1s)
- An infinitely long “tape” which contains a single row of symbols that can be anything
- A read/write head that reads from the tape and writes back to it
- Rules for reading and writing from the tape



If you create an instruction set that is capable of running on a Turing machine, it is said to be Turing Complete. All you need for Turing-completeness is:

- Conditional branching
- Loops
- Variables and memory

Stripping To The Essence

Turing was trying to create a model of computation that could, essentially, scale to any problem thrown at it. A Turing machine has more computational power than the machines previously discussed because a Turing machine has an infinite amount of storage (in the form of a “tape”) to read from and to store state to. Now, being good computer scientists we might get stuck on the notion of an “infinite” bit of tape. Such a thing might be a bit hard to come by in the real world.

Believe it or not, it doesn’t matter; this is a conceptual machine after all. The amount of tape use, the speed of the tape through the machine and/or the symbolic alphabet that you choose to use with it – none of these variables effect the overall computational power of a Turing machine. As we learned last chapter: *if you can describe it with a Turing machine, it’s computable.*

It’s a rather wonderful computational model, and it’s derived from the idea of stripping the notion of computation itself to the barest minimum, as Turing further described in his paper:

Computing is normally done by writing certain symbols on paper. We may suppose this paper is divided into squares like a child’s arithmetic book. In elementary arithmetic the two-dimensional character of the paper is sometimes used. But such a use is always avoidable, and I think that it will be agreed that the two-dimensional character of paper is no essential of computation. I assume then that the computation is carried out on one-dimensional paper, i.e. on a tape divided into squares.



1	2	3	4	5	6	7	8	9
2	4	6	8	10	12	14	16	18
3	6	9	12	15	18	21	24	27
4	8	12	16	20	24	28	32	36
5	10	15	20	25	30	35	40	45
6	12	18	24	30	36	42	48	54
7	14	21	28	35	42	49	56	63
8	16	24	32	40	48	56	64	72
9	18	27	36	45	54	63	72	81

Reading and Writing

Given that a Turing Machine can write to, as well as read from, a tape full of symbols, we now can store an “infinite” amount of information (assuming we have an infinite tape).

A quick tangent here. The term *infinite* is used a lot when discussing Turing Machines – don’t get hung up on it. If you need to, translate it as “just enough”. In other words, your computer can store massive amounts of information – but someday it will fill up. At that point you can go get another hard drive – and someday you’ll fill that up. But then you’ll get another drive … the space your machine has on disk has nothing to do with its computing power – same with the size of RAM. These things do have a lot to do with how long you’ll have to wait, however...

State, in a Turing Machine, is stored on the tape that it is given to read from:

The behaviour of the computer at any moment is determined by the symbols which he is observing, and his “state of mind” at that moment. We may suppose that there is a bound B to the number of symbols or squares which the computer can observe at one moment. If he wishes to observe more, he must use successive observations.

The “tape” is being observed by moving under what Turing called a “head”, something that can both read and write. As the tape moves, the machine observes the symbols and the “state of mind” of the machine changes because the information within changes.

By the way: it’s quite fun to see Turing’s treatment of the machine in human terms.

We may now construct a machine to do the work of this computer. To each state of mind of the computer corresponds an “m-configuration” of the machine.

Multiple machines, multiple configurations – all usable by other machines with different configurations: this is the essence of the Turing machine.

Universal Turing Machine

One very interesting feature of a Turing Machine is that it can read in and simulate another Turing machine, even a copy of itself. In the 1930s, when *On Computable Numbers* was written, machines were built for a specific purpose. A drill made holes, an elevator lifted things, etc so it made sense that computing machines would be purpose-built in the same way.

Turing proposed something altogether different. His abstract machine was able to simulate any other abstract machine – making it a universal computational device.

This idea had quite a profound effect, as you can imagine. The question what can we solve? was no longer interesting – with the Church-Turing Conjecture the question evolved to something a bit more profound: what can't we solve?.

In a wonderful bit of cheek, Turing provided the answer in the very same paper he used to introduce his machine.

The Von Neumann Machine

We started this book by thinking about complexity and computation in general. We then got a little steampunk, and explored the early 1800s and Babbage's Analytical Engine – only to find out that he and Ada Lovelace had designed the first Turing-complete language out of punch cards.

We explored Turing machines in this chapter and got to know both the Church-Turing Conjecture as well as its counterpart: The Halting Problem.

We have a couple of ways of computing anything that is computable – but we're still stuck in the land of theory. This changed in the early 1940s with John von Neumann, JP Eckert and John Mauchly.

Oh, and Turing too, as it turns out.

That Tape Can Hold a Whole Lot More

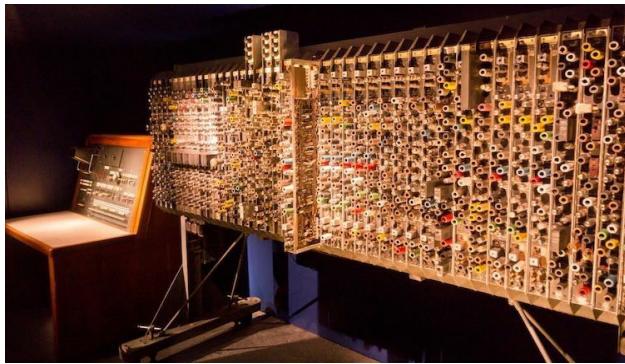
We know that a Turing machine works on a very simple idea: a read/write head and an infinite supply of tape. The machine has some type of instruction set that influences the current state of the machine. It's simple enough to conceive of a tape and a read/write head – but what about that instruction set? Where does it live?

Initially, the idea of producing a working computer was that you would produce it for a specific purpose – maybe a calculator or prime number generator – a dedicated machine hard-wired to do a dedicated process. This seemed to defeat the notion of a “universal computing device”, but the hardware simply wasn't there yet.

At almost the same time, Turing and von Neumann were plotting to fix this. They both came up with the same answer: make the instructions themselves part of the tape. This is what we call a program today.

The Automatic Computing Engine

In 1946, just after the war, Turing designed and submitted plans to the National Physics Laboratory (NPL) for an electronic “general purpose calculator” called the Automatic Computing Engine – using the term Engine in homage to Charles Babbage. A number of these machines were created, and they stored their processing instructions as part of the machine's data.



Turing's ACE Machine. Image credit: Antoine Taveneaux. You can see this on display at the London Science Museum.

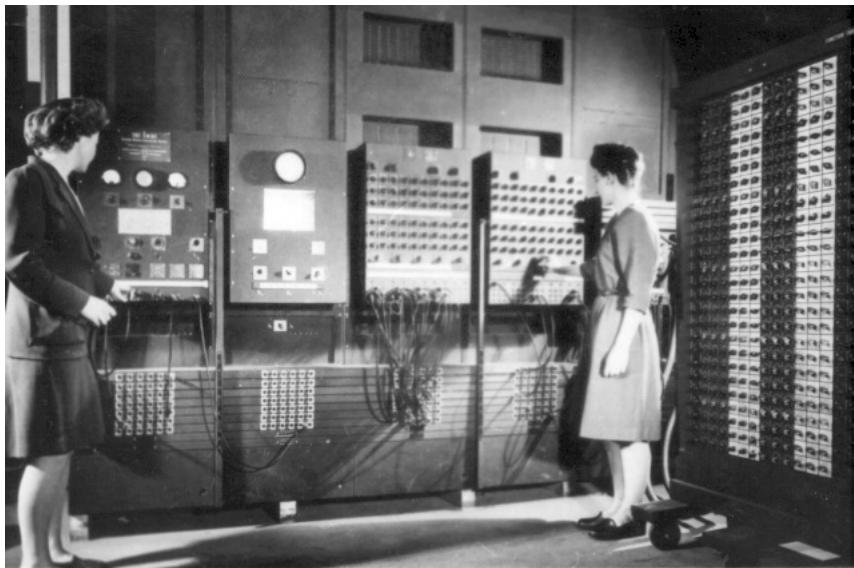
In an interesting historical side note: the NPL wasn't sure that creating such a machine would work. They were unaware of Turing's work during the war, where he worked with an electronic computer every day (Colossus) to help break Nazi ciphers. He couldn't tell them this, however, as he was sworn to secrecy.

ENIAC and EDVAC

In the early 1940s JP Eckert and John Mauchly were designing ENIAC, the world's first electronic computer. Governments were realizing quickly that computers would give them the edge during wartime, and the race was on.

The existence of Colossus, the machine that helped crack German codes during World War II – is known today. Back then, however, this was a big secret that no one outside of Blechley Park knew about. I bring this up because the initial designs for computers, underway with various science teams around the globe, were done in secret and without much sharing.

ENIAC, for instance, was developed to calculate ballistic trajectories and other top-secret calculations. It was Turing-complete and programmable, but to program it you had to physically move wires around.



ENIAC Operators

In 1943 Eckert and Mauchly decided to improve this design by storing the instruction set as a program that was stored next to the data itself. This design was called the EDVAC.

By this time the Manhattan Project (the US Atomic Bomb project) was rolling, and von Neumann, being a member of the project, needed computing power.

He took interest in ENIAC and later EDVAC which led him to write a report about the project, detailing the idea of a stored-program machine. This is where things get controversial. For one reason or another – maybe because the paper was an initial draft – von Neumann’s name was the only name on the report, even though it was mainly Eckert and Mauchly’s idea.

As well as Turing’s – but he couldn’t tell them that because it was still a secret.

This gets further complicated by a colleague of von Neumann’s, who decided to circulate the first draft to other scientists. Soft publishing it, if you will. These scientists got very interested in the paper and ... well ... we now have the generally incorrect attribution of stored programs to von Neumann.

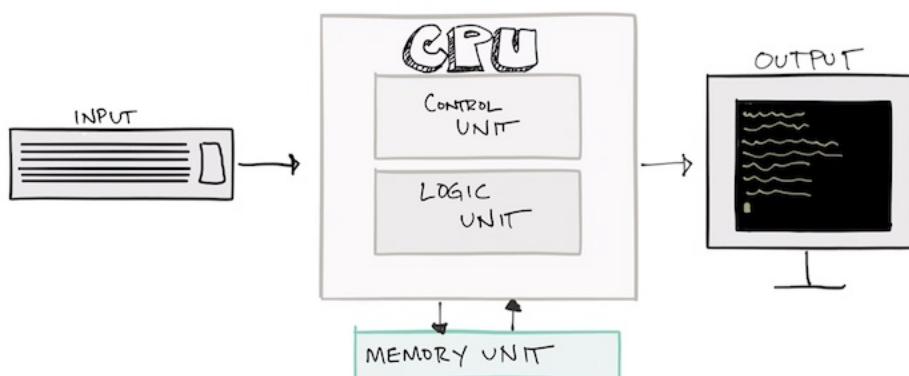
One of von Neumann’s colleagues, Stan Frankel is quoted as saying:

Many people have acclaimed von Neumann as the “father of the computer” (in a modern sense of the term) but I am sure that he would never have made that mistake himself. He might well be called the midwife, perhaps, but he firmly emphasized to me, and to others I am sure, that the fundamental conception is owing to Turing—in so far as not anticipated by Babbage... Both Turing and von Neumann, of course, also made substantial contributions to the “reduction to practice” of these concepts but I would not regard these as comparable in importance with the introduction and explication of the concept of a computer able to store in its memory its program of activities and of modifying that program in the course of these activities.

As you might imagine, Eckert and Mauchly were not happy about this. Either way, this new idea now has a name.

Von Neumann Architecture

Historical silliness aside, we have now arrived at a modern day computer and the birth of modern computer science. von Neumann’s machine should look reasonably familiar:



von Neumann Architecture

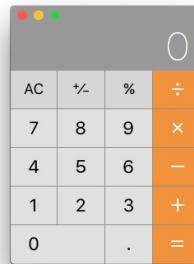
This machine architecture allows us to create, store, and deliver functions as a set of symbols. Thinking about this in depth can twist your brain a little bit.

A Turing machine is an abstract construct and the initial attempts to create one ended up with rather large bits of machinery that needed to be fiddled with by hand in order to run a given calculation. von Neumann was able to abstract that fiddling, which meant that Turing Machines could be built with nothing more than code and Turing's original vision of a "Universal Turing Machine" was complete.

To grasp this in a more real sense – think of a pocket calculator.



It does only one thing: math. Now open the calculator on your computer.



My Mac's Calculator.

It does exactly the same thing as the calculator above, but there's no hardware involved. I can also change it to do scientific and programming calculations on the fly, which is kind of magical. **The Mac calculator is as much of a machine as the one you hold in your hand**, but it's made of nothing more than pixels.

You can thank Mauchly, Eckert and von Neumann for this. Machines running within machines. When you consider various code execution environments (like the JVM, the CLR, or Google's V8 Engine) that are running on virtualized machines in the cloud this whole idea tends to start spiraling.

How many abstract machines are involved to execute the code you're writing today? When you run a VM or Docker (or some container functionality of your choice), these are machines within machines executing machines made up of smaller machines within other machines...

It's machines all the way down.

PART 2: CONCEPTS

Resources

You can find the code used in parts of this section up at [my github repo](#). In addition, you [can buy 17 video walkthroughs](#) of the algorithms you see in this chapter and others [from here](#).

The videos cover some new ground as well, and use languages other than C#.

CHAPTER FIVE: THE BIG-O

Have you ever written some code that you were rather proud of? It's a pretty good feeling to see a test pass so you can move on to solving another problem or, better yet, writing more tests.

Any coder can solve a problem given enough time, *solving it well*, however – that's what we want to do! But what does that even mean?

Simply put: it means that your code does what the spec requires, it can scale, and it's written in a way that other developers will understand *your intentions* in the future. This chapter focuses on the second part of that sentence: the “it can scale” part.

How can you demonstrate that your code can scale using something more than just waving your arms? **You can use Big-O.**



The specific way that you can think about Big-O notation is that it *mathematically describes the complexity of an algorithm in terms of time and space*. It is intimidating and quite a few developers I've encountered (coworkers, at conferences and so on) will instantly withdraw from a conversation at the first mention of “Order N” or the like.

It can come off as a bit elitist, to be sure, if you examine someone else's code and casually drop Big-O into a conversation. It can also mean that you'll pass your next interview when you interviewer asks you about the complexity of some code you've written and whether you can improve it to be $\log n$ vs. n^2 ...



A Word About Data Structures

Before we get going I think it's worth addressing the examples you're about to see. I'll be using a basic array for every single one of them, and if you're not a JavaScript developer you might think to yourself *that's ridiculous! I would never write code like that!* And I would understand.

.NET, for example, has a remarkable amount of list types, both generic and non, that allow you to write rather powerful, exact code for the task at hand. Elixir and other functional libraries allow you to choose freely between an enumerable, a list, a dictionary or a map.

Your choice of data structure depends squarely on the type of data you're working with and then what you're trying to do with that data. *But do you know why you're making these choices?* Do you know why these structures exist in the first place?

Each data structure in .NET (or Java or Elixir) was created for use in a particular type of algorithm. If you weren't formally trained in data structures (as I wasn't), you would just pattern your choice from what you've read in blogs or been told by a more senior developer. This works fine for many, but you're reading this book because you want to go a bit deeper, to learn the concepts that underly so many decisions you've made in the past.

To ask "why this data structure?" is to also ask "how are we using this data", which then naturally leads to "which algorithms are we going to implement?". The answers to these questions are somewhat interdependent, which presents a problem for me in terms of writing this book in that *I have to start somewhere*.

So I'm going to level the playing field. All we're going to use in this chapter is an array and some integers - much like any coding interview. I'm doing this because I want to be able to focus on *complexity* as a means for making educated choices about algorithms and, correspondingly, the data structures you choose to work with.

We'll get into data structures and algorithms later in the book, for now I ask you to suspend what you know about various list types and to just go with the flow. Yes, the examples are contrived, but only because I didn't have a choice!

Onward...

A Super Simple First Step: O(1)

Let's just jump right in and see what kind of mess we can make here, shall we? Let's say I have an array of 5 numbers:

```
const nums = [1,2,3,4,5];
```

Now, let's say I ask you to get the first number from this array. Obviously if you look at it can you say "oh - sure that's a 1". A program doesn't have eyes, however, so it needs a way to pull that number out.

Being the smart person you are, you decide that it's a simple matter of using an index – the very first index as a matter of fact. We're using JavaScript here so we can get the value thus:

```
const nums = [1,2,3,4,5];
const firstNumber = nums[0];
```

Now comes the question: *how complex was that operation?* If you were like me just a few years ago I would have said "not complicated at all – I just took the first element of the array". This is a correct answer, but we can be more specific by thinking about complexity in terms of *operation per input*.

We have 5 inputs here because there are 5 elements in the array. How many operations did we need to perform on these inputs to derive a result for our algorithm? *Only one* as it turns out. How many operations would we need to perform if there were 100 elements in the array? Or 1000? Maybe 1,000,000,000? Still: *only one*. We just take the very first element at index 0.

We can capture that inelegantly long paragraph in a more scientific way by saying that our algorithm was “on the order of 1 operation for all possible inputs”, or better yet: $O(1)$. This notation is pronounced “order 1” or, more casually, *constant time*. For all inputs to our algorithm there is and will always be **only one operation required**.

As you’ve probably guessed, $O(1)$ algorithms are pretty efficient and also quite desirable!

Iterations and Order(n)

Now that you were able to figure out how to pull the first item from our array, let’s try something more complicated: *let’s sum the items in the array*.

Again, let’s use some code:

```
const nums = [1, 2, 3, 4, 5];
let sum = 0;
for(let num of nums){
    sum += num;
}
```

Now we get to ask ourselves the same question: *how many operations do we have per input to our algorithm?* This time the answer is different: we have to add each number to a running sum, so we have to **operate on each one**. This means *one operation per input*.

Using Big-O time notation, we would say this is $O(n)$, or “Order n” where n is the number of inputs. This type of algorithm is also referred to as “linear”, or that it has “linear scaling” (think of describing a line on a graph: $y = 2x$ or something of the sort).

Analysis

This type of scaling is common when you calculate a result by iterating over a collection of values like we’re doing above. It’s a simple way of doing things, but it has implications in terms of complexity which we can see if we ask ourselves a simple question: *how does this scale?*

As opposed to our $O(1)$ algorithm, our CPU is doing a bit more work in our summing operation above – in fact it’s doing n times the amount of work of our $O(1)$ algorithm. If we have an array of 10 items it won’t matter; but what happens if our array has 1,000,000 elements? Now we need to worry a bit as we have 1,000,000 operations to perform.

It might seem academic, but it’s a good question to ponder whenever you write a loop (or worse: a *nested loop* which we’ll address in the next section): “is there a way I can make this algorithm a bit more efficient?”. For our summing operation - *no*, there isn’t. We have to consider every element. For other things, however, sometimes a bit of math might do the trick.

Consider our very same array:

```
const nums = [1, 2, 3, 4, 5];
```

What if I told you to write a summing function for a sorted, contiguous array of integers that starts with the number 1? Ahhhh this time things are different as we know a bit more about our input.

We can use a bit of [interesting math](#) here, specifically something that Carl Friedrich Gauss figured out while in grade school. If you want the full explanation, follow the link above; otherwise I’ll just get to it.

We can use this equation to figure out the sum of the series $[1...n]$:

$$S = n(n + 1)/2$$

Plugging this into our example array, we would have:

$$5(5 + 1)/2 = 15$$

How do we know what n is? Simple! It's the very last element of the array. Now we can change our algorithm a bit:

```
const sumContiguousArray = function(ary){  
    //get the last item  
    const lastItem = ary[ary.length - 1];  
    //Gauss's trick  
    return lastItem * (lastItem + 1) / 2;  
}  
const ary = [1, 2, 3, 4, 5];  
const sumOfArray = sumContiguousArray(ary);
```

The answer here will be 15!

Notice that we're not running an iteration? That makes our $O(n)$ algorithm *a whole lot faster* as we're now in constant time thanks to Mr. Gauss. You might be thinking “but wait, if we have to figure out the length of the array *and* pull off the last element *and* run the calculation – isn't that $O(3)$ or something?”.

This is something we should get straight about Big-O right up front: yes that would be the literal complexity, but we're not interested in that. All we care about is that it's *constant time*, meaning that the time complexity will not change based on the number of inputs.

Constant time algorithms are always referred to as $O(1)$. Same with linear time. An algorithm might literally be $O(n + 5)$ but in Big-O that's just $O(n)$.

The Not-so Good Approach: $O(n^2)$

Let's up the complexity a bit. It was suggested in today's standup that I don't make very good integer arrays, and that it's possible that I might have duplicated one of the elements. I insisted that I did not! However we decided it might be a good idea if you were to create a routine to verify this.

There are some simple solutions to this, and some that are quite a bit more efficient. Let's start with the simple, brute force solution which require a nested **for** loop:

```

const nums = [1,2,3,4,5];
const hasDuplicates = function(nums){
  for(let num of nums){
    //let's go through the list again and have a look
    //at all the other numbers so we can compare
    for(let otherNum of nums){
      //is there another element equal to this one?
      if(otherNum === num){
        //if so...
        return true;
      }
    }
    //only arrive here if there are no dups
    return false;
  }
}

```

It works, but it's not ideal. Here we're iterating over our array, which we already know is $O(n)$, and another iteration inside it, which is another $O(n)$. For every item in our list, we have to loop our list again to calculate what we need. This type of complexity is $O(n^2)$, or "Order n squared" and as you can imagine, it's not very efficient and is considered quite inelegant.

Analysis

The big problem is this: if we have 1000 numbers in our list, we'll have $1000 * 1000 = 1,000,000$ operations! That's bad.

If you remember only one thing from this chapter, let it be this: *there is almost always a better way!* In fact, this is one of those things where once you see it, it's pretty hard to *not see it* again. Nested loops working over the same collection - always $O(n^2)$.

These problems appear often in coding interviews as well, where the simplest answer is to brute force your way through an " n^2 " solution, trying to figure out how it's possible to do things better.

Which you might be wondering right now... so let's play Big Job Interview!

Me: well, this solution works, true, but I'm wondering if we can do better? Is there a way we might be able to improve on your $O(n^2)$ solution?

You: possibly –

Me: sure, chicken. If it makes you feel better to be told an answer as opposed to figuring it out yourself... then...

You: passive-aggressive bullying isn't cool Rob

Me: sorry... I don't know the answer either...

In a few sections we'll revisit this problem and see if we can make it better.

Refining To $O(\log n)$

We've been relying on a brute force approach to work with our array up to this point which works, but as I keep saying it's 1) inefficient and 2) inelegant. We want to do both of those things so we don't just look the part of super stellar programmer!

For the next task, let's search through our array for a given number. If you don't have a CS degree (having a few search algorithms at the ready) – how would you go about such a thing. Once again: yes, we have eyes but a program

doesn't, so we need to have a systematic (and deterministic) way of finding a given value in a collection.

The simplest way to go about this process is to think about the *time complexity* of the operation, which is what we've been doing. Instead of offhandedly saying "we'll have to search the array top to bottom" or "just iterate until you find the number", we can now use a simpler, more direct term: $O(n)$.

If we use a brute force approach, we just loop over the entire array *once* to find the number we're looking for. This means that in the worst case scenario (which is how we think when we use Big-O) we have to perform one operation per input.

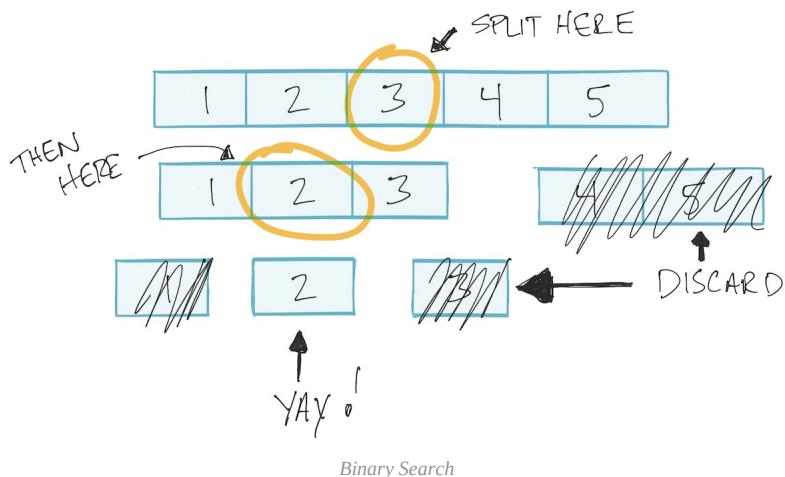
Can we do better than a linear $O(n)$? It turns out that yes, we indeed can if we make a few assumptions.

Assumption 1: we've noticed that the list is sorted; can we assume that it is sorted for this exercise? For the purpose of illustration, I'll say *yes, we can*.

Assumption 2: I didn't mention anything about *in-place*; can we use more than just the given array? *Sure, why not.*

Given these two assumptions, we can use a clever algorithm called *binary search*. We'll get into the details later on, but this algorithm is performed in a set of simple steps where we split the list in half, discard the half we don't need, then go on splitting until we have the value we're looking for.

Let's see if we can use this approach to find the number 2 in our array:



As you can see, by splitting the list and discarding the half without our target value each time, **we're able to find the number we're looking for in far fewer operations**.

This type of algorithms is called "divide and conquer" and works on the mathematical principle of *logarithms*. If you're like me (just a year ago) you haven't had to think about logarithms in quite a few years and, perhaps, you've forgotten how they work.

Let's take a quick tangent to review how they're helping us here.

Quick Logarithm Review

In essence, logarithms help make working with exponents a bit easier. Let's do some quick math and hopefully things will be a bit clearer. What if I asked you to figure out what x is here:

$$x^3 = 8$$

To answer this you simply need to take the cube root of 8 to get 2. Simple enough! What if we changed things up now:

$$2^x = 512$$

This makes things a bit harder, doesn't it? How would you solve for x here? If you're a math person you already know the answer; but if you're like me and have forgotten most of your math the answer wouldn't exactly pop out at you. What we need here are *logarithms*.

If we used logarithms we could rewrite this equation like this:

$$\log_2(512) = x$$

This equation says "log base 2 of 512 equals some number x ". The key here is the "base 2" part. It means that we're *thinking in 2s*, and we want to know *how many times to multiply 2 to arrive at 512*. That's all logarithms do! How do we solve this equation? Duh! With a calculator of course!

The important part here is the recognition that logarithms basically deal with splitting a given number into some type of *base*. The default base is 10, which means you should be able to guess at the answer of this expression:

$$\log(100) = x$$

The question you're asking here is *how many times do you multiply 10 by itself to arrive at 100?* The answer is 2!

OK, so let's get back to the problem at hand. We're splitting a list of numbers in half, continually, until we get down to single value that we want. The key to this is the name of the algorithm itself: *binary search... splitting things in half*.

We're dealing with 2s which means our logarithm will be base 2. So this operation can be expressed like this:

$$\log_2(n) = x$$

Which brings us to an interesting thing about Big-O which I'm repeating now and I'll be repeating again later: **the exact math doesn't matter.** What matters is that this is logarithmic - nothing else. We can even do away with the base and x - we just care that the operation is $O(\log n)$.

Analysis

If we were to use an $O(n)$ scan on our list we would have had to perform $n = 5$ operations in order to find our target number 2. By splitting the list and evaluating the number at the split (thus discarding the half without our number), we were able to bring the number of operations down to 3.

You might be thinking "hey wait a minute! Splitting *and* evaluating is 2 operations!" and this is true - but they are 2 $O(1)$ operations so we just consider the entire thing to $O(1)$.

Cutting from 5 operations down to 3 isn't terribly exciting. But what if we had 100 total elements? If you whip out a calculator and run some math, you'll find that:

$$\log_2(100) = 7$$

OK that's not exactly true. The number is something a bit more like 6.643856. But we can reason through this in a simpler way.

We're programmers and have a reasonable facility for base 2 multiplication, hopefully. We know that multiples of 2 are 2, 4, 6, 8, 16, 32, 64 and finally 128. The number we want is somewhere between 64 and 128, so we *round up* to be safe. **This gives us 7.**

Why do we care? Because we can see how our search algorithm scales! We upped our inputs by a factor of 20, but our operations only went up by 4.

That is why we care about doing this entire exercise! We can prove, with math now, that our approach here is better than scanning the entire list for a given number.

Rethinking $O(n^2)$ With $O(n \log n)$

A few sections ago I asked you to find duplicate numbers in our list of numbers and we decided to iterate over every number (which is $O(n)$) and then do it once more to see if our current number was repeated.

This put an $O(n)$ operation inside of another $O(n)$, which we can simplify to $O(n * n)$ which is $O(n^2)$, which is a Bad Thing.

But *wait a minute...* didn't we just use a nifty algorithm to find a number within our list that was much more efficient than our brute force linear scan? Yes!

So here's a quiz: *what is the Big-O* of using our binary search inside of our linear scan? Let's break it down.

Iterating and Searching

We have an $O(n)$ as before with our array iteration – we know that much. We now have a single operation within the iteration that is $O(\log n)$. Putting the two together we have $O(n * \log n)$ which we can just think of now as $O(n \log n)$.

So: in terms of *time complexity*, we have a better algorithm here! It will provably scale better than our original brute force approach as well!

```
const nums = [1, 2, 3, 4, 5];
const searchFor = function(items, num){
    //use binary search!
    //if found, return the number, otherwise
    //return null. We'll do this in a later chapter...
}

const hasDuplicates = function(nums){
    for(let num of nums){
        //let's go through the list again and have a look
        //at all the other numbers so we can compare
        if(searchFor(nums, num)){
            return true;
        }
    }
    //only arrive here if there are no dups
    return false;
}
```

Finding duplicates with binary search

Common Uses

You see $O(n \log n)$ complexity in a number of algorithms, but they are most common when doing sorts. We'll get into sorting algorithms later on, but the most commonly used (because they're the most efficient) are Quicksort and Merge Sort. Each is $O(n \log n)$ with a few variations which we'll discuss in a later chapter.

Going Even Further

Our interviewer is digging our style right now and decides to up the game, one more time.

What if I told you that you could assume there was only one duplicate number?

Well now... that's an interesting development. Let's consider everything that we just did and see how we can bring this $O(n \log n)$ algorithm down just a bit more.

The first question we should ask ourselves is *can we do some math* to avoid an iteration? A few sections ago we used Gauss' trick to sum a series of integers, which could help us here. To do that, we'll need to perform three operations:

- **Find out what the sum should be** for an array without duplicates. For this all we need is the highest number in the array past to us. Since we know it's sorted, that's the last element and we know this is a constant time thing at $O(1)$.
- **Find out the actual sum.** This will necessarily be an $O(n)$ scan.
- **Subtract the non-duplicate sum from the actual sum** and that's the number we're looking for! This is also $O(n)$.

If we're right, this will bring our algorithm down to $O(n)$.

Here's our new algorithm:

```
const findDuplicate = function(ary){  
    //sum what we're given  
    let actualSum = 0;  
    //our O(n) scan  
    ary.forEach(x => actualSum += x);  
    //get the last item  
    const lastItem = ary[ary.length - 1];  
    //create a new array  
    const shouldBe = lastItem * (lastItem + 1) / 2;  
    return actualSum - shouldBe;  
}  
const ary = [1,2,3,4,4,5];  
const duplicate = findDuplicate(ary); //4
```

Nice work! Our interviewer is impressed!

A Quick Break for Math Geeks

One of the YouTube channels I love watching is Numberphile. They tackle all kinds of interesting math problems and concepts, creating 10 minute videos illustrating the results.

One day, while researching the math tricks you just read about (specifically Gauss' trick) I stumbled on this video, which asks the question:

What is the sum of $[1 + 2 + 3 + 4 + \dots \infty]$?

What do you think the answer is? You would *think* it would be a pretty large number if not infinity itself. But then... the answer turns out to be -1/12.

No, seriously. Fair warning: *you will get lost* down this rabbit hole...

Oh Please No: $O(n!)$

The array that we've been working with is, as you've probably been able to tell, a simplified version of the number set we're going to be working with in a few chapters in this book, which is:

```
var soUsers = [23, 4, 42, 15, 16, 8, 3];
```

Now that you know the basics of time complexity, we're ready to turn you lose on something a bit more complex.

In short: we think a few of our employees are teaming up to answer questions on Stack Overflow under a single account name. I won't give away any names, but it rhymes with Bon Sweet.

Here's what we need you to do: we'll supply you with a dump of all the questions and answers from Stack Overflow as well as a listing of each of the above users' IP addresses (so we know who's answering what). We need you to go through these questions and answers and find which grouping of our users is the best at answering the most popular questions.

If they're going to answer questions on our dime, we have to make sure they do well!

Analysis?

My head hurts even thinking about this. For fans of tough problems out there, you might have recognized elements of The Traveling Salesman Problem? If I did my example right, that is.

We talked about NP-Hard problems in the last part of this book, and here we have one, which I hope you recognized as a *combinatorial optimization*.

We know that problems in Exp have the characteristic of 2^n complexity, which is also $2n!$. It actually takes a bit of work to accidentally create a factorially complex algorithm – but you'll know it once you try to run it over any data set of appreciable size.

Another way that I've come to recognize $O(n!)$ is if your solution calls for n nested loops. If this is you, you're fired.

Thinking in Big-O

By now you should be able to equate certain basic operations in code to a given Big-O. If you practice this for a bit, you'll be able to quickly spot patterns and, ideally, improve them if you know the right algorithm to do so.

Specifically:

Random access to a given element in a collection is always $O(1)$, depending on how the list is indexed. Arrays, for instance, allow you to access elements randomly if you know their index. HashSets allow you to access if you know what the value is (the hashed value is the index). Dictionaries allow you random access if you know what they key is, and so on. These *types* of operations are *always* $O(1)$ which means if they are combined with other Big-O, they will remain static or *constant time*.

List iterations are always $O(n)$. If you need to evaluate every item in a list for a given algorithm, this means it

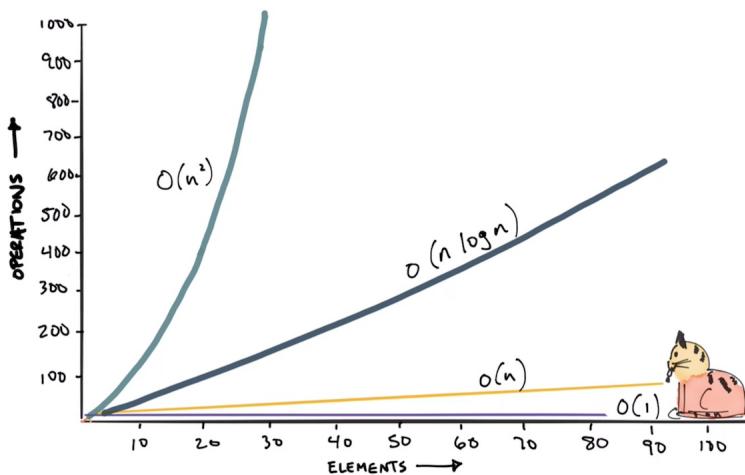
will at least be $O(n)$. Sometimes you can get around this with some trickery, which I'll discuss later on.

Nested Loops on the same collection are always at least $O(n^2)$. Loops within loops ... sometimes necessary, but can usually be improved by thinking about data structures (which we'll do later on).

Divide and Conquer is always $O(\log n)$. The very act of dividing a list into smaller sublists is logarithmic. If you have an $O(1)$ operation once the list is split apart, then the Big-O for the entire operation is $O(1 * \log n)$ which is just $O(\log n)$.

Iterations that use divide and conquer are always $O(n \log n)$. Think about looping a list and then executing some algorithm to search for list value or, possibly, to run some kind of sorting.

If you solve a problem by adding another nested loop for every input that you have: that's $O(n!)$ which is bad and you should probably find another job!



Space Complexity vs. Time Complexity

You may have noticed that I've been using the term *time complexity* a lot in this chapter, as that's what we've been focused on: *how long will it take to do a given operation given n inputs*. In data analysis speak this is called a *dimension* and is just a way to think about how complex an algorithm is.

There is another dimension that is also important: *space*. In other words: **what are your algorithms resource requirements?**

The same type of Big-O classifications still apply in that we still refer to things such as $O(n)$ or $O(1)$ "space", but the meaning is somewhat different. Space where? The answer is: *it doesn't matter*. "Computer resources somewhere" is all we care about.

However. There is a bit of reasoning we should be able to apply so we can go a bit deeper.

When a program executes it has two ways to "remember" things: the *heap* and the *stack*. I'll get into the details of each later on, but for right now just consider the stack to be the thing we're worried about. It's the thing that remembers the variables *in the scope of the currently executing routine*. When a variable is declared in a block of code, it's stored on the stack. When a block of code goes out of scope, the variables are removed from the stack. Sometimes the scope is the current block, other times it's the current function or procedure.

Why do we care? The short answer is that **you can easily run out of resources before you run out of time** depending on how you've written your program. The simplest way to think about this is the dreaded "Stack Overflow Exception" which simply means you're executing some kind of loop that has used up every last bit of space on the stack. This can (and often does) happen with a recursive routine as each value remains on the stack until all functions have executed.

Working with strings is another way to cause yourself *space complexity* problems. For example, if you use a

loop to build a string, your space complexity might be as bad as $O(n * m)$ where n is the number of iterations and m is the length of the string. Not so bad if you're just building out memes, but if you're trying to evaluate string patterns in a book... that could be bad.

CHAPTER SIX: DATA STRUCTURES

Now that we understand how to quantify complexity, let's look into the various ways that we can work with data in order to increase efficiency. As I mentioned in the beginning of the previous chapter: *algorithms, complexity and data structures work hand in hand.*



Arrays are simple to work with but make life hard if you're trying to find a particular value. Dictionaries help with that, but are slightly more involved. Hash Sets (or Hash Tables) are lovely because the hash of their value is also their key, so accessing data is done in constant time and is $O(1)$ if you know the value you want. From there things get quite a bit more involved.

Some languages offer a plethora of choices when it comes to working with data, but all of these choices are variations of (or small improvements upon) a standard, core set of data structures that every programmer should know. That's what we'll take a look at in this chapter.

Array

You know about arrays, I'm sure. Unless you work in JavaScript every day, however, it's unlikely you use them much. They're very lightweight and simple to use, but there is also weirdness.

When you declare a non-dynamic array, you must specify its size upfront. This size cannot change. In most modern languages we don't think about this – but in some older languages this is still the case.

Arrays hold values which are referenced by an index. They allow very fast random access, which means you can access any value from an array using a $O(1)$ routine as long as you know the index.

Arrays are typically bound to a specific length once they are created. If you need to add an item to an array, the language you're working in will typically copy the original plus whatever value you want to add to approximate dynamically changing the array's size.

This can be a costly operation, as you can imagine. Let's see why.

Resizing

Many languages allow you to dynamically resize an array (Ruby and Python, e.g.) while other, more strict languages, do not (C# and Java e.g.). Ruby and Python allow for dynamic arrays which are basic arrays, but with some magic behind the scenes that help during resizing. C# and Java have different structures (like Lists) built specifically for expanding and shrinking.

Arrays are allocated in adjacent blocks of memory when they are created. There is no guarantee that additional memory can be allocated adjacent to an array if you need to add an element, so it becomes necessary to copy and rebuild the array at a new memory location, which is a costly process.

Strings, for instance, are simply arrays of characters. If you append a string with another string value in C#, for instance, an array copy/rebuild needs to happen. This is why working with strings in a loop is typically not a good idea.



Arrays in JavaScript

JavaScript, however, is an interesting case. Arrays in JavaScript are just objects with integer-based keys that act like indexes. They are instantiated from the Array prototype, which has a few “array-like” functions built into it.

Here is a description from Mozilla:

Arrays are list-like objects whose prototype has methods to perform traversal and mutation operations. Neither the length of a JavaScript array nor the types of its elements are fixed. Since an array's length can change at any time, and data can be stored at non-contiguous locations in the array, JavaScript arrays are not guaranteed to be dense; this depends on how the programmer chooses to use them. In general, these are convenient characteristics; but if these features are not desirable for your particular use, you might consider using typed arrays.

Hey, it's JavaScript.

Why Choose an Array?

Arrays are the simplest data structure with the least amount of rules. Bits of data are stored *in contiguous memory*, so they are also have the smallest footprint of any data structure.

If all you need is to store some data and iterate over it, arrays are a fine choice, especially if you know the indices of the items you're storing. Random, $O(1)$ access to values in an array is also a great reason to choose an array over other, more “ceremonial” data structures.

The final reason to keep arrays fresh in your mind is The Big Job Search. Coding interviews will almost *always* involve working with an array of values at some level (usually integers), so understanding their restrictions (and advantages) is quite important.

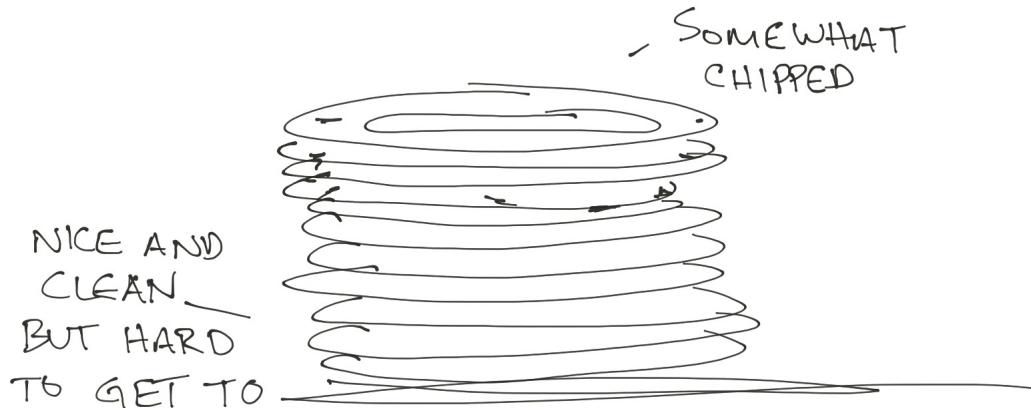
Stack

I discussed the idea of a stack in the previous chapter when describing a Pushdown Machine. It's just like an array, but it has a few restrictions:

- You can't access items randomly using an index
- You can only add and retrieve items in the stack from the "top"
- It has three explicit methods: Push, Pop, and Peek.

A term that is often used for the nature of a stack is "Last In, First Out" or LIFO. You *push* an item onto a stack, *peek* to see the value that's currently on top of the stack and then *pop* that item off of the stack.

The simplest way to think about the stack is by visualizing one of my big pet peeves: a stack of plates. The ones on the bottom never get used because we always grab the plates on the top!



Why Choose a Stack?

Stacks are useful in quite a number of surprising ways. They are perfect when you need to know the very last value that was seen in a loop. You might want to know this when you're:

- Reversing a string
- Traversing a graph or a tree
- Checking an opening/closing structure of some kind (such as balanced parentheses in a sentence)

Honestly, the biggest reason you want to know what a stack is (and does) is for interviews. They come up often! You might not use them regularly in your current job, but you never know when they might fit something you're trying to do.

Queue

A queue is also like an array, but with some additional rules:

- You can't access values randomly using an index
- You can only add values in one end and retrieve them from the other
- It has two explicit methods: *Enqueue* and *Dequeue*

A queue queues things in a queue, and is described as "First In, First Out" or FIFO. You *enqueue* an object into the "end" of the queue and *dequeue* it from the "front".

It's likely you've used queues often so I won't labor the point for too long. It's also a word that's kind of tough to type so I'll keep this section short.

An interesting thing about a queue is that you can create one with two stacks. This is one of those interview questions that you'll likely need to know at some point in your career! How would you go about doing this?

The mechanism is pretty obvious: you have an "in" stack and an "out" stack – the tough part is deciding at which point you'll move items from the "in" stack onto the "out" to avoid collisions. If you do it one at a time then things can easily get out of order! Imagine calling *enqueue* 5 times, then *dequeue* once, then *enqueue* again. Your stacks would be a mess!

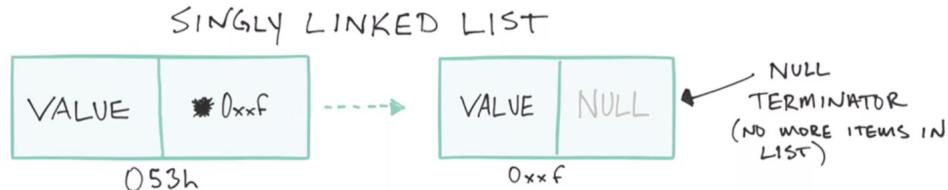
How would you avoid this? I'll leave it to you...

Linked List

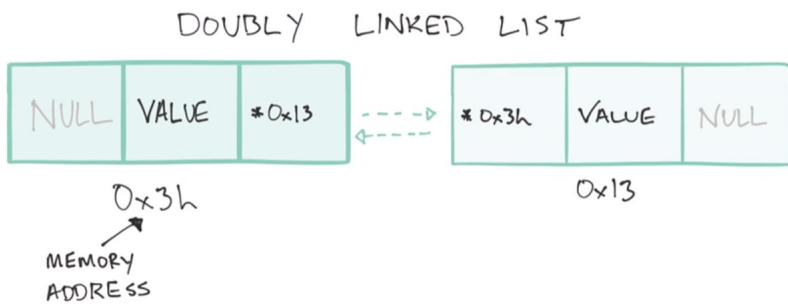
There are two types of linked lists: singly and doubly linked. Both of these are graphs (which we'll discuss in a minute), which means you can think of them as two dimensional structures with associations.

A singly linked list consists of a set of "nodes" in memory, that have two elements:

- The value you want to store
- A pointer to the next node in line



A doubly linked list is exactly the same, but contains an additional pointer to the previous node.



Operations

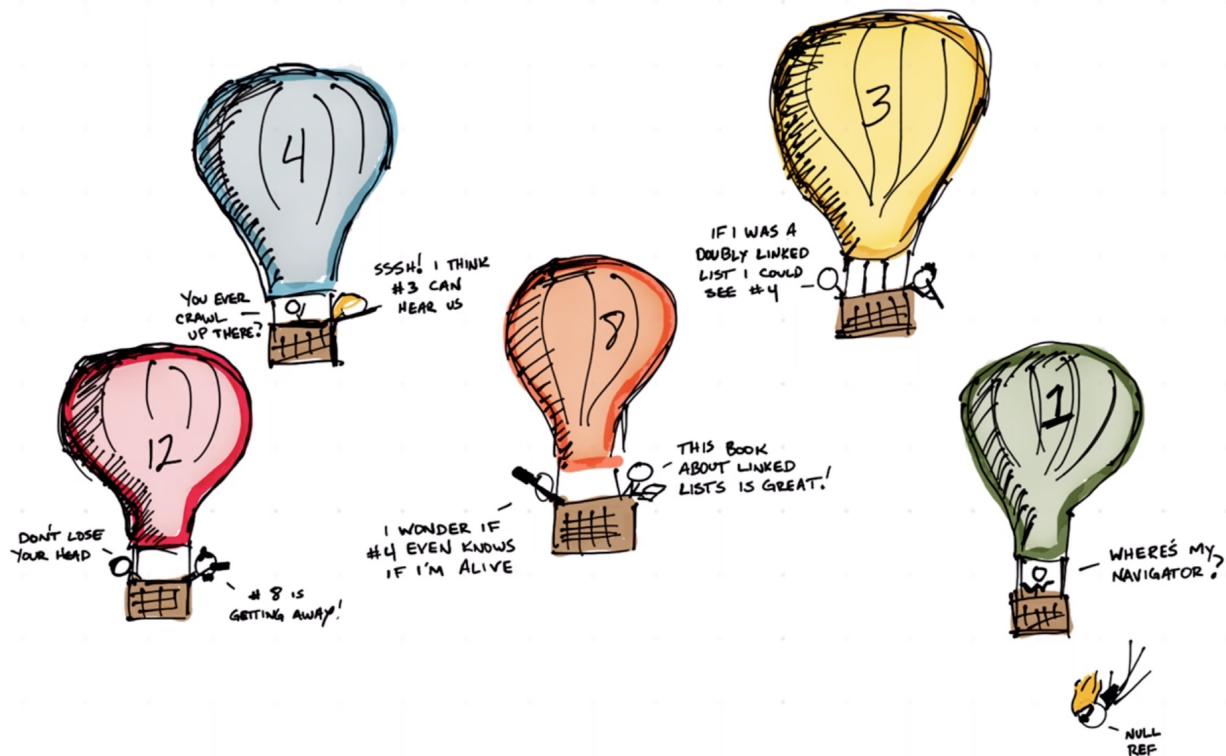
Nodes in a linked list don't need to reside next to each other in memory, and therefore can grow and shrink as needed.

Their loose structure allows you to insert values into the “middle” of the list by simply resetting a few pointers. Same for deleting a node in a linked list.

The downside to linked lists is that finding an item means you have to traverse the list to get to the item you want, which is an $O(n)$ operation, always. Arrays, on the other hand, allow $O(1)$ access if you know the index.

Linked lists are null-terminated, which simply means if a node’s pointer is null, then that signifies the end of the list.

The first item in a linked list is called the head, the rest of the list is, typically, called the tail. Some languages refer to the tail being the last item in the list – there is no hard definition so if you hear about “tailing the list” just think about the end of it and hope for some context.



Why Choose a Linked List?

The primary reason to choose a linked list over something like an array is *simplicity* and the ability to grow and shrink as needed.

Working with them can be a bit weird, but also kind of interesting in that they are very self contained and lightweight:

```

public class LinkedListNode
{
    public int Value { get; set; }
    public LinkedListNode Next { get; set; }

    public LinkedListNode(int value)
    {
        this.Value = value;
    }
}
var a = new LinkedListNode(1);
var b = new LinkedListNode(2);
var c = new LinkedListNode(3);

a.Next = b;
b.Next = c;

```

Iteration is best thought of as “traversal” because you don’t really know when a linked list will end – so you end up using something like a **while** loop:

```

var thisNode = a;
while(thisNode != null){
    //do something
    //...
    //traverse
    thisNode = thisNode.Next;
}

```

Hash table

Arrays are fast for reading data, linked lists are good for writing data and having more flexibility. A hash table is a combination of the two.

A hash table stores its data using a computed index, the result of which is always an integer. The computation of this index is called a hash function and it uses the value you want to store to derive the key:

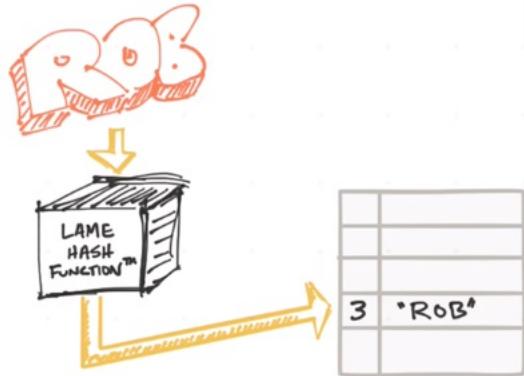
```

var value = "Rob";
function myLameHashFunction(val){
    return val.length; //3
}

```

Most modern programming languages have a hash function you should use; don’t create your own. I’ll go into why in just a second.

Once you have a key, you can store your value. The key/value pair is often referred to as a bucket or slot.



Speed

Hash tables are great when you want quick access to certain values. Because their value is also their index, hash table reads are $O(1)$. This can be complicated, however, if a hashing function is overly complex. You will typically leave this to whatever framework or language you're using, so assuming $O(1)$ is fine.

Similarly, adding values to a hash table does not involve (typically) any traversals – you just hash the value and create the key. This leads to really good performance, however in the real world this doesn't happen that often.

Even the best hashing algorithms will create duplicate keys (called collisions) if the data size is large enough. When this happens, your reading and writing can be reduced to $O(n/k)$, where k is the size of your hash table, which we can just reduce to $O(n)$.

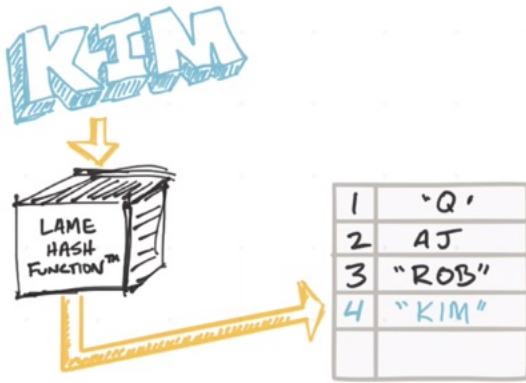
Collisions will likely happen in any hash table implementation, so most implementations have built-in ways of handling these. Let's take a look at two common ways to deal with collisions: open addressing and separate chaining.

Open Addressing

Open addressing resolves a collision by finding the next available slot, and dropping the value there.

Let's assume I'm using my Lame Hash Function with a new name: "Kim". The key produced will be a 3 since my Lame Hash Function only uses the length of the value instead of something more intelligent. The key produced (3) will collide with my existing entry for "Rob", which is also a 3. Using open addressing to resolve the collision, "Kim" will be added to the first slot at the end, which is 4.

To find "Kim" in the table, open addressing dictates that we do an $O(n)$ scan from index 3 (where "Kim" should be) and then work our way down until we find the value.



This presents some interesting problems – what if “Doug” and “Dana” want to join our list? They can, but they’re indexes will be 5 and 6. It spirals from there.

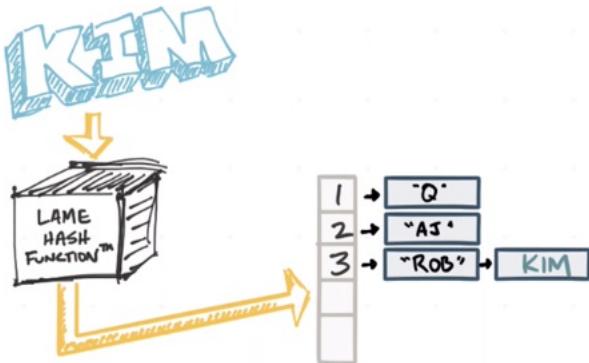
This is called *clustering* and it can quickly turn a lovely, fast hash table into a simple $O(n)$, not fast hash table.

Separate Chaining

Separate chaining involves combining two data structures that we’ve already looked at: arrays and linked lists.

When a key is hashed, it’s added to an array that points to a linked list. If we have only one value in our hash table for a given key, then we’ll have a linked list with only a single element.

However, in the case of key 3, the linked list can expand easily and accommodate “Kim” as well:



Why Choose a Hash table?

Hash tables are one of the most common data structures you’ll find in modern languages because they are fast. As long as the hashing algorithm is comprehensive and capable, the hash table will be $O(1)$.

One of the main drawbacks, however, is again the hashing algorithm. If it’s too complex then it will take longer to run, and you will still have $O(1)$ read/write but it could actually be slower than executing an $O(n)$ over an array, for instance.

A dictionary is exactly like a hash table except it has a unique key for accessing a given value. This has an advantage over hash tables in that you *can't have a dictionary with the same key* - so collisions are not something you need to worry about.

Given that a key is guaranteed to be unique, dictionaries provide $O(1)$ access to any element, as long as you know what the key is. The downside is that dictionaries can be a bit larger than hash tables and therefore have increased *space complexity*. Most of the time this isn't something you need to worry about.

Dictionaries are ubiquitous in programming, even more so than arrays and hash tables. If you're a JavaScript developer you might be wondering about this assertion – after all, the only type of list you get to work with is an array... which isn't exactly true.

Under the covers, your array is really a dictionary with integer keys, as is a plain old object:

```
var nums = [23,4,42,15,16,8,3];
//this array is really a dictionary
//which is also used to represent an object
nums[0] // 23;
nums[1] // 4;
nums[2] // 42;

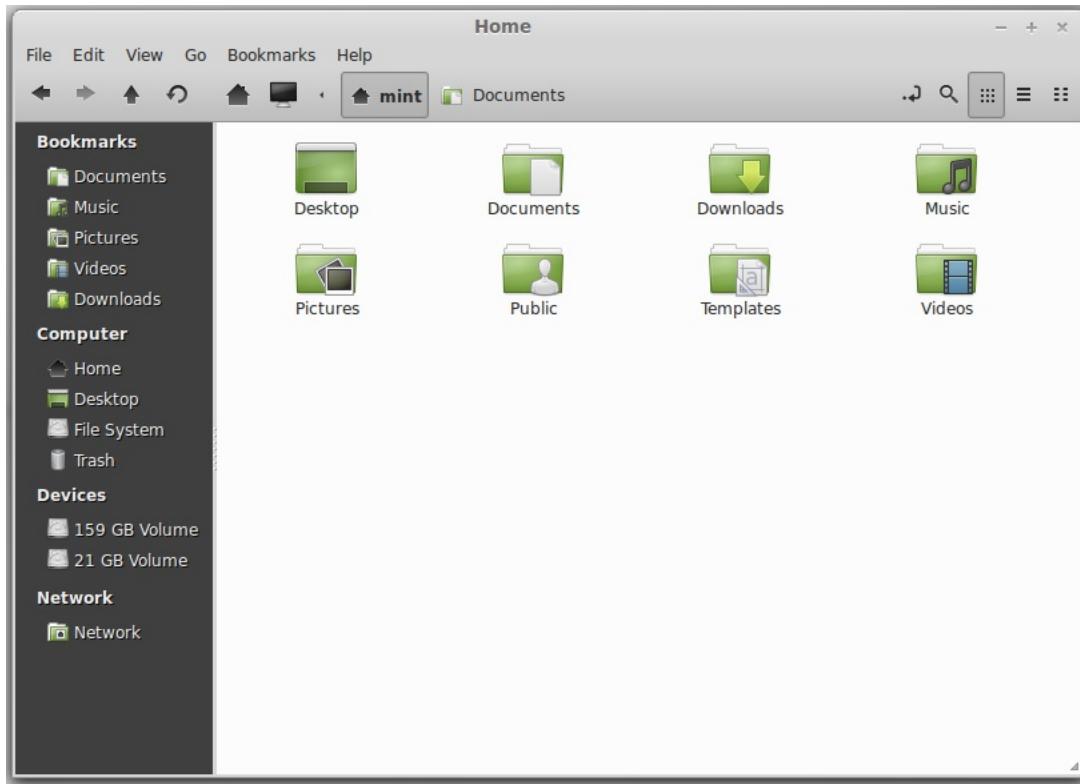
var book = {title: "The Imposter's Handbook"};
book["title"] //The Imposter's Handbook
```

Trees

Trees are another type of graph (which is a collection of associated nodes that I'll get into in just a minute) with a hierarchical structure. As opposed to linked lists, which are linear, trees can have 0 or more ancestor nodes.

Every tree has a single root node. Every child of the tree descends from this root node and has *only one parent*.

You work with tree data structures every day: your computer's file system is a tree:

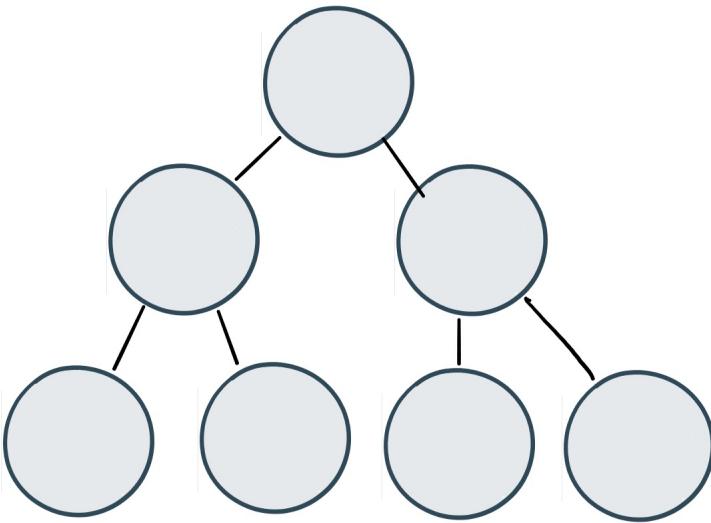


Using a tree data structure usually involves traversing using recursion, which can be expensive in terms of *space complexity*. Every time you recurse a child node, the stack has to store the values for that scope and you could run out of space.

There are other ways to traverse a tree by iterating using a Queue and a Stack - I'll get to that in the Algorithms chapter, specifically Breadth-first and Depth-first traversals.

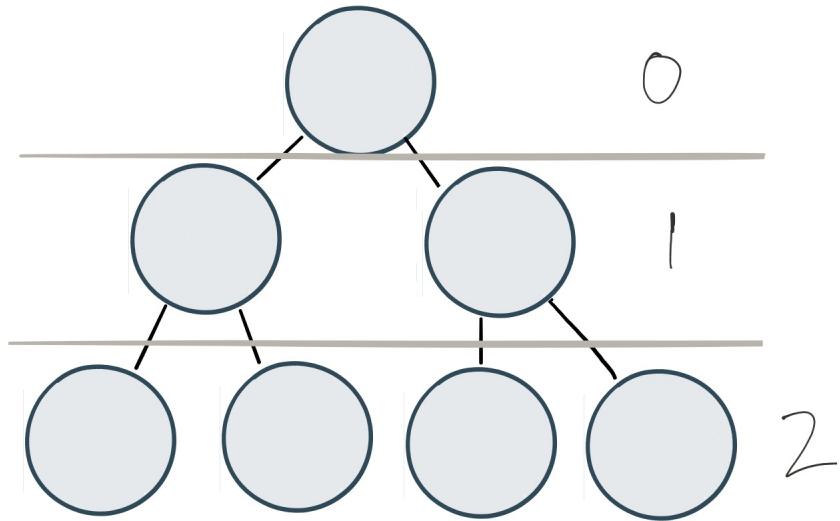
Binary Tree

A binary tree is something you've likely seen. It's a type of graph, but with some rules applied, which are straightforward: each node can have 0, 1, or 2 *child* nodes, but only 1 parent:



Binary trees can represent a number of interesting things, including *decisions*. Each node represents a given state that is the result of a yes or no decision.

Another interesting aspect of a binary tree is that each level of the tree represents a *logarithmic value*:



At level 0 we have $\log_2(0) = 1$ Node. At level 1 we have $\log_2(1) = 2$ and finally $\log_2(2) = 4$ nodes.

You can also visualize the *divide and conquer* search algorithm we used in the last chapter with a binary tree. Each split operation represents another level of a binary tree.

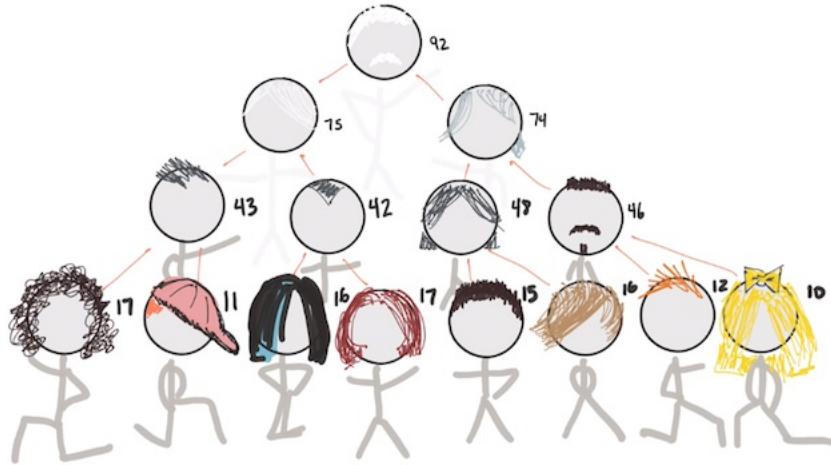
Heap

A heap is an inverted tree that is a set of interconnected nodes that store data in a particular way, which is called the *heap property*.

Every child node belongs to a parent node that has a greater priority (or value, whatever) – this is called a max heap. A min heap is the opposite: every parent has a value less than all of its children.

Put another way: if we're dealing with a max heap then every node on a top level has a greater value than every node on the next level down.

We can see this in a recent family picture of mine, where we all posed based on our ages:

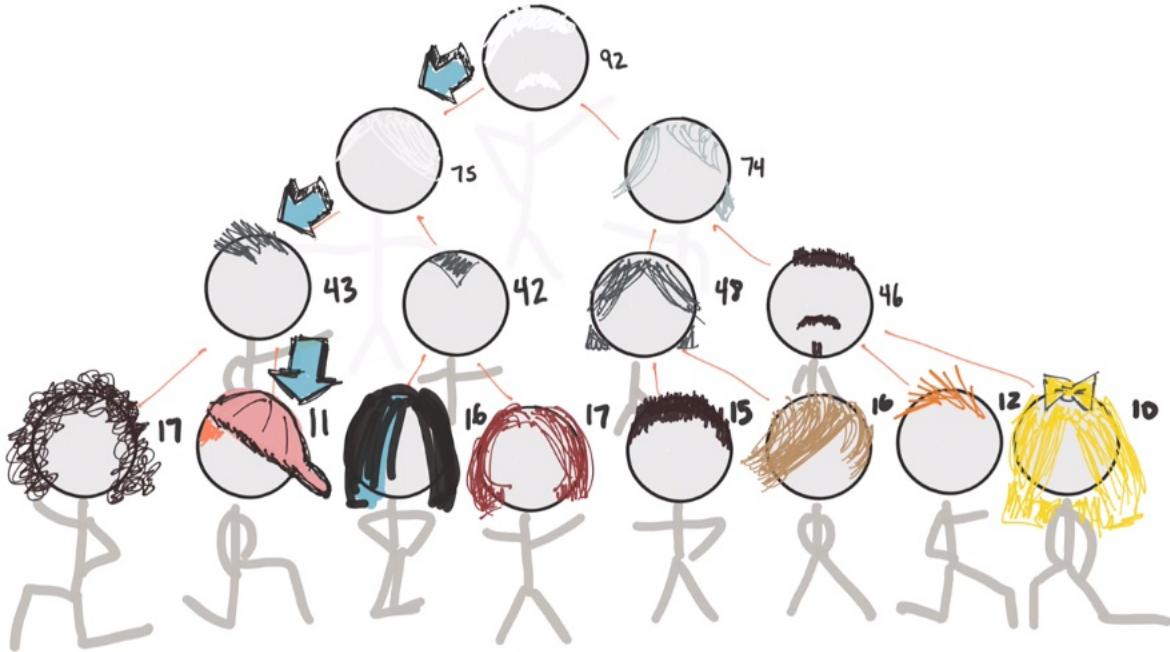


We couldn't afford a camera so we had to draw it out

Usage

A heap (and its variants, which I'll get to later on) can be used in any algorithm where ordering is required. Arrays are random, and allow random access to any element within them. Linked lists can change dynamically but finding something within them is $O(n)$ (linear); heaps are a bit different.

You can't do $O(1)$ random access and a single node knows nothing about its children. This means you need to do some type of traversal to find what you're looking for. Given the structure of the tree, however, finding things is considerably easier than with a linked list:



We found Tommy, age 11, easily here. However we could easily have had to traverse over a few times if he was on the end there next to Jujubee, age 10.

So what are heaps good for then? It turns out they're wonderful if you're doing comparative operations – something like “I need all people in this heap with an age greater than 23”. Doing that in a linked list would be quite slow – same with an array and a hash as no order is implied.

Why Choose a Heap?

Heaps are used in data storage, graphing algorithms, priority queues and sorting algorithms. In many languages you'll find data structures that will give you the same kind of ordered structure if you choose your keys wisely.

For instance: in C# you could use a `SortedList` or a `SortedDictionary`, storing objects with parent/child associations.

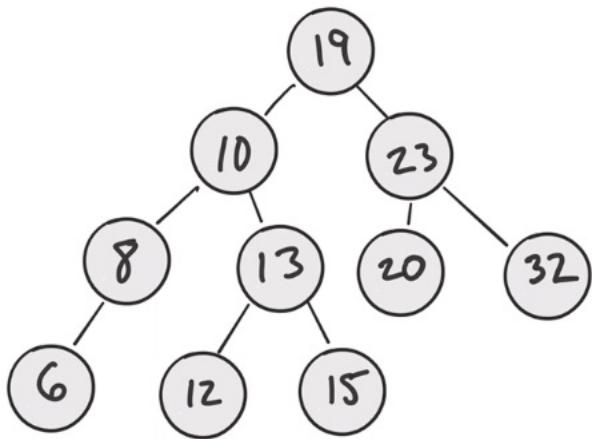
The primary advantage of heaps is performance. They're very fast when it comes to searching as they are pre-sorted and don't involve the extra step of an $O(\log n)$ split – it's just a traversal.

Binary Search Tree

A binary search tree is just like a heap in that it is organized based on the values of the nodes, with one major addition: there is also a left to right value priority.

The rules are:

- All child nodes in the tree to the right of a root node must be greater than the current node
- All child nodes in the tree to the left must be less than the current node
- A node can have only two children



Binary Search Tree

The advantage of a binary search tree is, obviously, searching. It's very easy to find what you're looking for as you'll see down below. The downside is that insertion/deletion can be time consuming as the size of the tree grows.

For instance, if we remove 13 from the tree above, a decision needs to be made as to which node will ascend and take its place. In this case I could choose 15 or 12. This operation seems simple on the face of it, but if 15 had children 14 on the left and 16 on the right, some reordering would need to happen.

Searching A Binary Search Tree

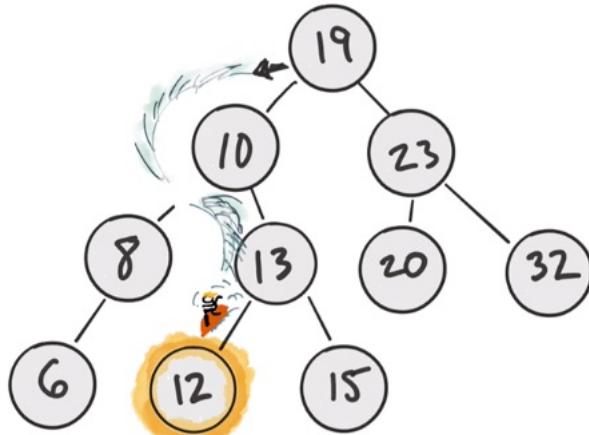
If you know the rules, finding a value in a BST can be done recursively, or by using breadth-first or depth-first search (which we'll get to in a minute):

- If the value is less than the current node you're on, go to the left child node.
- If the value is greater, go to the right child node.

Do this until you find what you're looking for.

For example, here we need to find the number 12. Our root node is 19, so we traverse to the left because $12 < 19$. When we reach 10, we traverse right, because $12 > 10$.

Finally we come to 12, so we go to the left again and arrive at 12.



Why Choose a Binary Tree?

Binary search trees are very useful, especially if you remove the restriction that each node can only have two children and then make it self-balancing.

This is called a b-tree and is the foundation for many data storage systems and indexing algorithms.

Many of the algorithms we'll be discussing are built into database engines and, by now, you should be able to reason through their utility for whatever it is you're trying to do.

For instance – PostgreSQL allows you to specify what type of index you want to apply to a given table and columns. Two of the most popular are HASH and BTREE.

Can you reason through what each type of index would be good for? We know that Hash Tables are $O(1)$ average complexity for reading and writing single elements. Binary search trees are great for finding things.

Here is the [documentation](#) from PostgreSQL for BTREE:

By default, the CREATE INDEX command creates B-tree indexes, which fit the most common situations... B-trees can handle equality and range queries on data that can be sorted into some ordering. In particular, the PostgreSQL query planner will consider using a B-tree index whenever an indexed column is involved in a comparison using one of these operators: <, <==>= >

Here is the HASH index documentation:

Hash indexes can only handle simple equality comparisons. The query planner will consider using a hash index whenever an indexed column is involved in a comparison using the = operator.

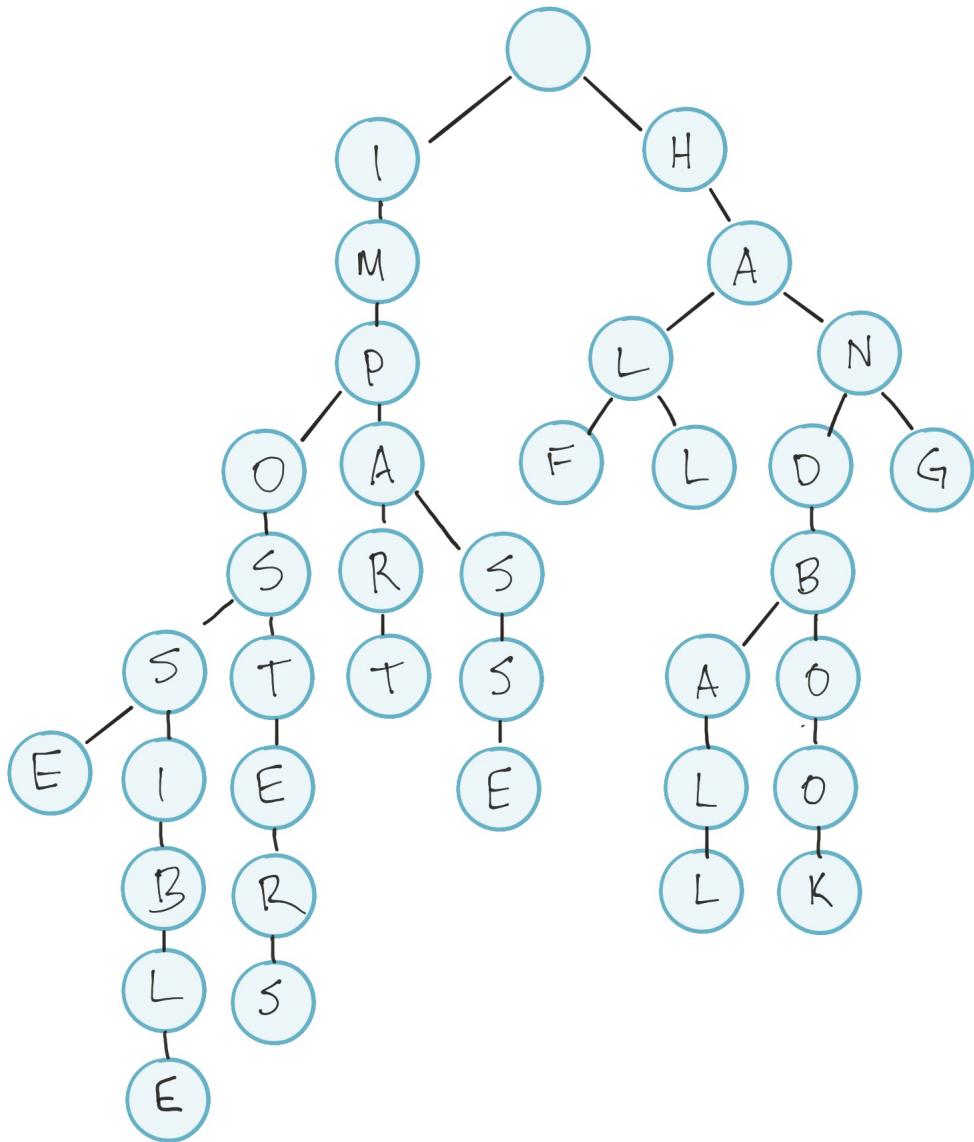
Makes perfect sense. If you find yourself looking up email addresses, for instance, a HASH index might suit you better than a BTREE.

If you're looking for ranges, a BTREE is preferable. It's also fast enough to do equality – but hashes are faster.

Digital Tree (or Trie)

A digital tree (or “trie”, which most pronounce “try” even though the term comes from [retrieval](#)) is a specialized tree used in searching, most often with text. In many cases it can outperform a binary search tree or hash table, depending on the type of search you're doing.

Tries allow you to know if a word (or part of a word) exists in a body of text. The easiest way to understand a trie is to see one, so here goes:



This try has an empty root node, and from there letters are added as child nodes. The power of a trie is evident with the prefix "imp" - in this section of the trie there are 6 distinct words represented by 20 total nodes:

- Imp
 - Imposter (and Imposters)
 - Impossible
 - Impart
 - Impasse

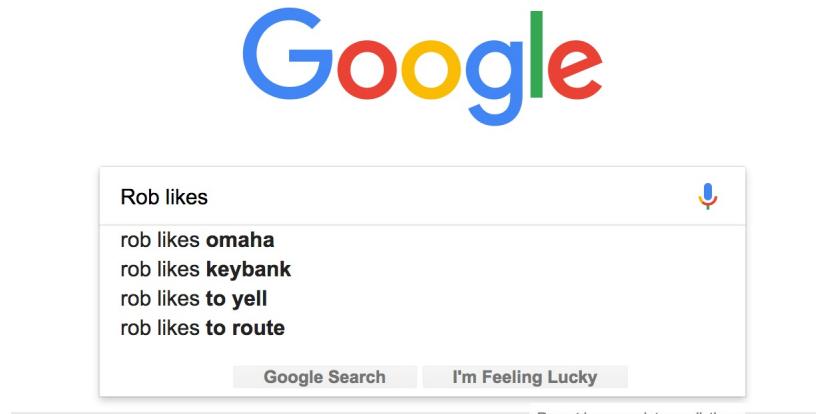
The advantages of this structure are speed and space. Finding a word in this structure is $O(m)$, where m is the length of the word you're trying to find.

Tries also have a major advantage when it comes to *space complexity*. Common prefixes are reused so repetition

within the structure is kept at a minimum.

Finally, the major advantage of a trie is that you're able to search the structure for *partial matches* based on a prefix. This kind of thing is great for word completion.

Have you ever wondered how code completion works in your favorite editor, or how Google can do the below so quickly?

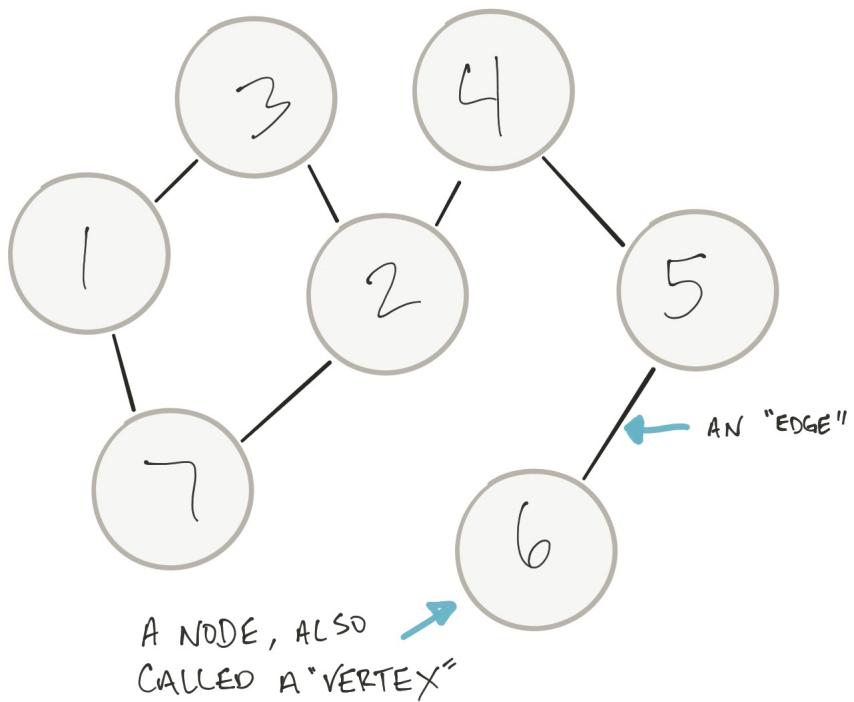


That's right: using tries.

Graphs

I've been mentioning graphs quite a lot in this chapter and it's time, finally, to dive in and get to know these data structures a bit more.

Graphs are one of the most useful and most *used* data structures in computer science. In short, a graph is a set of values that are related in a pair-wise fashion. Again, the easiest way to understand this (if it's not intuitive) is to see it:



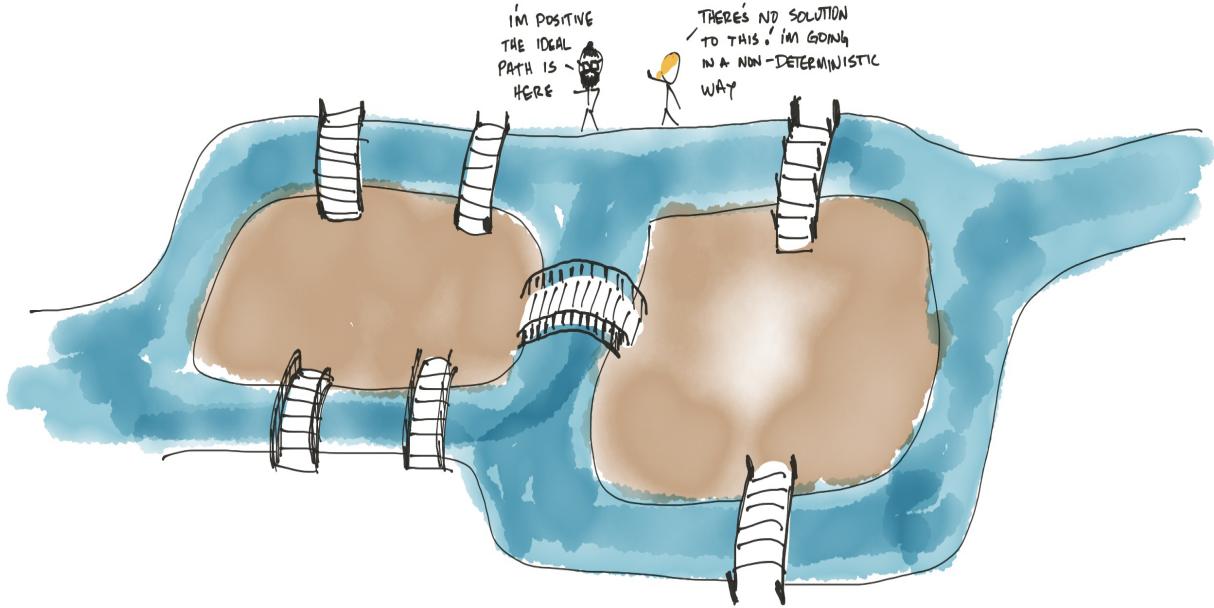
As you're probably figuring: graphs can describe quite a few things. Let's take a look at a few.

Graph Theory and The Bridges of Königsberg

If you took calculus in school, you probably learned about the origins of graph theory with Leonhard Euler's [Seven Bridges of Königsberg](#) problem:

The city of Königsberg in Prussia (now Kaliningrad, Russia) was set on both sides of the Pregel River, and included two large islands which were connected to each other and the mainland by seven bridges. The problem was to devise a walk through the city that would cross each bridge once and only once, with the provisos that: the islands could only be reached by the bridges and every bridge once accessed must be crossed to its other end. The starting and ending points of the walk need not be the same.

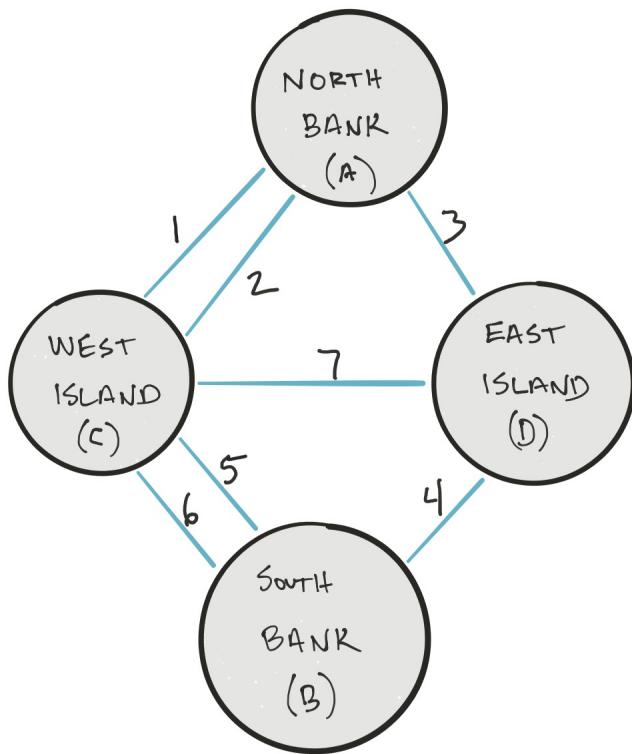
It might be easier to visualize the problem ... which gives me another reason to draw!



The Seven Bridges of Königsberg.

In trying to solve this problem, Euler reasoned that the route on land was not important, only the sequence in which the bridges were crossed. Given this, he could relate the land to the bridges in an abstract way, using the idea of nodes that are accessible by an edge.

In other words, *a graph*.



Now that we have our graph, we can restate the problem in terms of a graph:

Can you visit vertices A, B, C and D using edges 1 through 7 only once?

Before reading on, take a second and see if you can solve the problem just tracing your finger across the page. Or draw it out on a paper yourself and see if you can trace a line using a pencil or pen, visiting nodes A through D using edges 1 through 7 only once.

In mathematical terms, a *simple path* access every vertex. A *Euler path* will access every edge just once – that's the one we want, a Euler path.

The Solution

There is an algorithm we can use to solve this problem, which is to determine the number of degrees each vertex has and apply some reasoning. A degree is how many edges a given vertex has, by the way.

Euler reasoned that a graph's degree distribution could determine whether a given edge must be reused to determine the path. His proof, in short, is that a graph must have either zero or two vertices with an odd degree in order to have a Euler path (a path which visits each edge just once).

Let's tabulate the seven bridges graph:

VERTEX	DEGREES
A	3
B	3
C	5
D	3

Here we have four vertices with odd degrees, which tells us that there is no Euler path and, therefore, that the Seven Bridges Problem has no solution.

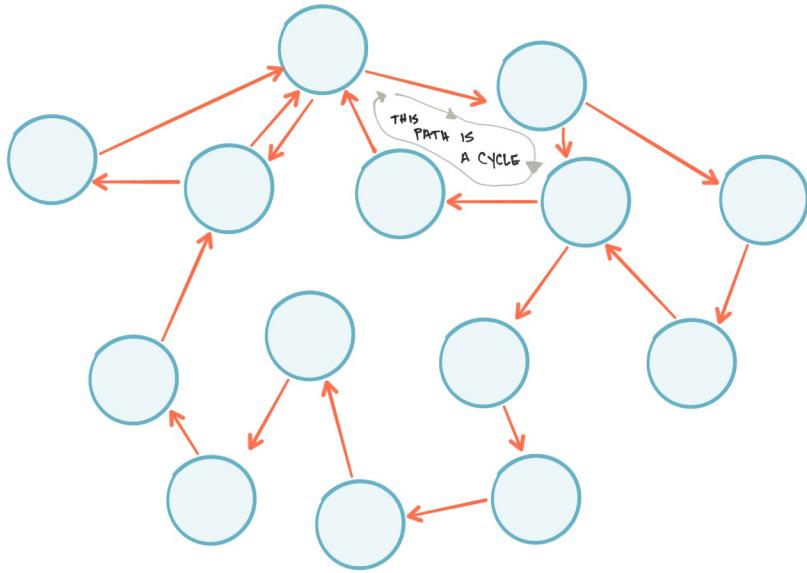
So What?

I wouldn't blame you if you're wondering right now why I've included the Seven Bridges Problem in this book as well as a discussion of Euler. These seem to be more math-related than anything in computer science, don't they?

In your next job interview you may very well be asked how you would solve some algorithmic problem that you (hopefully) will recognize as graph-based. Fibonacci, Traversal, Balancing, or a Shortest Path problem – if you can spot a graph problem you'll have a leg up on the question.

Directed And Undirected Graphs

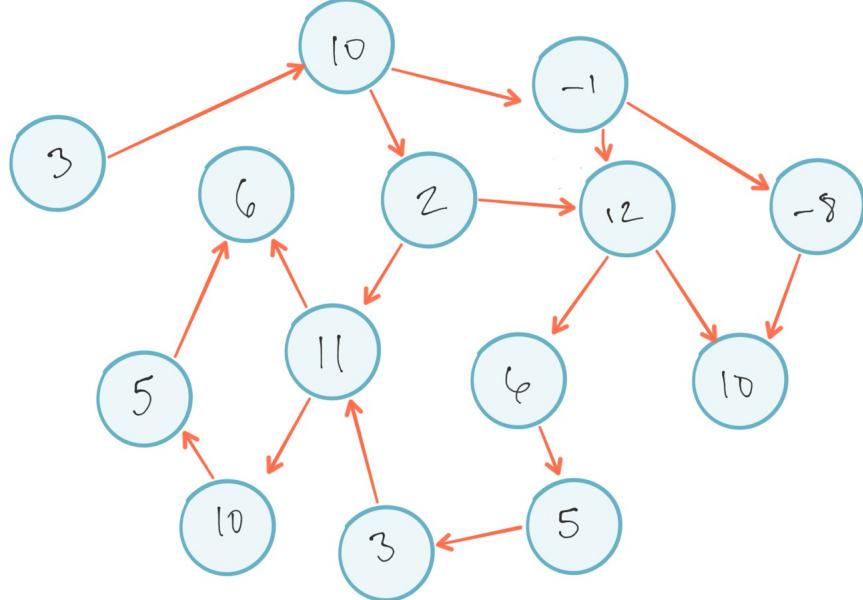
There are different types of graphs, as you can imagine. One which you'll want to be familiar with is a directed graph, which has the notion of direction applied to its edges:



These types of graphs are useful for describing traffic flow of some kind, or any system in which movement is not bi-directional in every case. There is also an undirected graph which you can think of as a series of two-lane highways that connect towns in a countryside. Travel is bidirectional between each town and not directed along a given path.

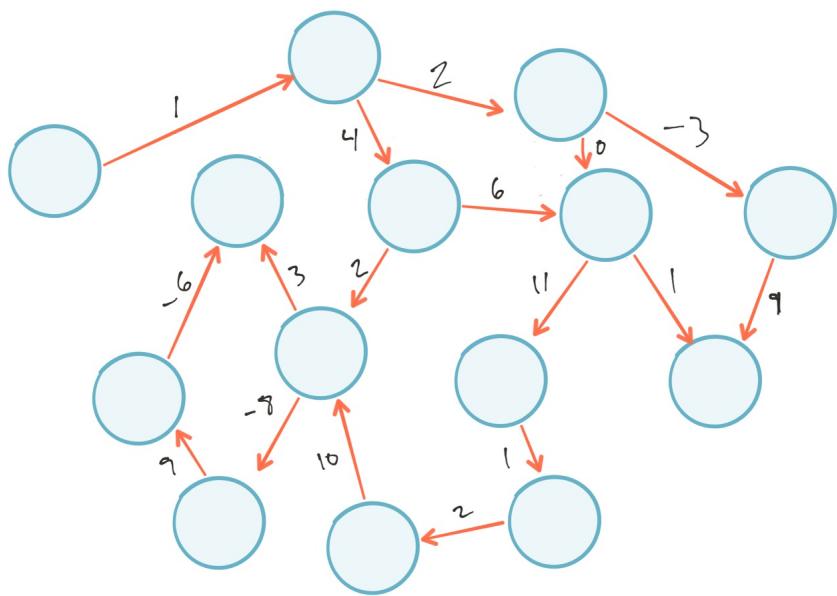
Weighted and Unweighted Graphs

Values can be applied to various aspects of a graph. Each vertex, for instance, might have a weight applied to it that you'll want to use in a calculation of some kind:



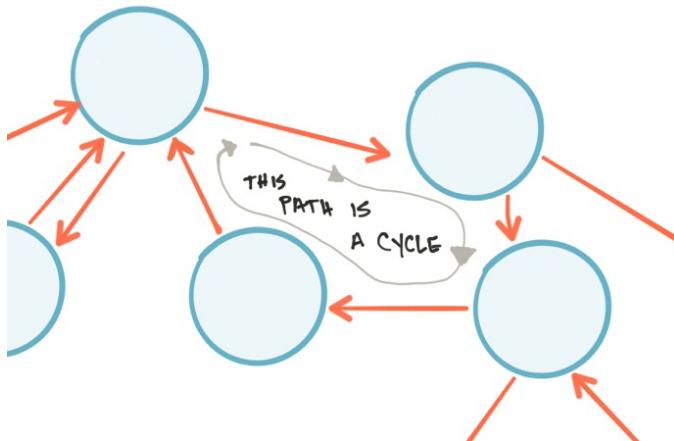
Maybe you're going on a trip, trying to figure out the most efficient way to see the cities you like the most.

There is also an edge-weighted graph, which is useful for calculating optimal paths:



Cycles

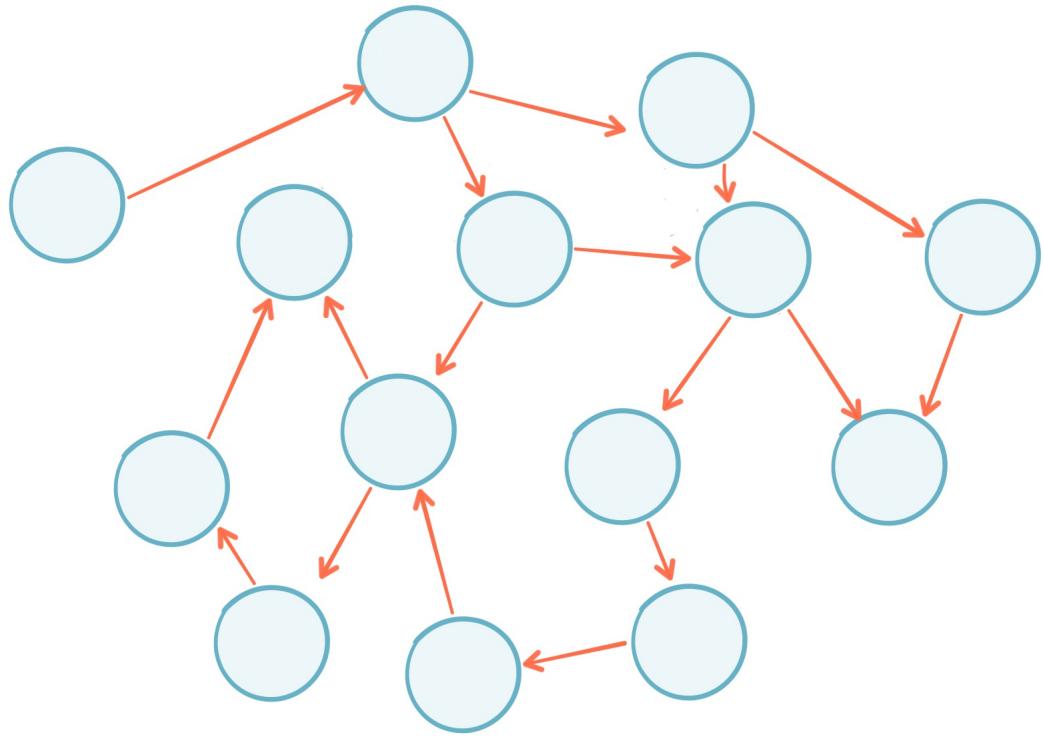
When you have vertices connected in a circular fashion, it's called a cycle. We can see one of these as part of our first graph above:



Cycles are fairly common in directed graphs. If a graph doesn't have a cycle, however, it has a special name.

Directed Acyclic Graphs (DAGs)

If we redraw the graph above with edges that don't cause any cycles, we'll have a directed acyclic graph, or DAG:



DAGs are useful for modeling information and processes, like decision trees, algorithmic processing or data flow diagrams. Many of the data structures that we've been playing with up until now have been DAGs, including all of the trees and linked lists.

CHAPTER SEVEN: SIMPLE ALGORITHMS

You don't need to know how to write a sorting or searching algorithm from scratch, frameworks do that for us. You do, however, need to know *how they work* because 1) it's likely you will be asked some details about them during interviews and 2) understanding their complexity could be the difference between keeping and losing your job!



In 2014 I visited my friend Frans Bouma at his home in The Hague, Netherlands. While there he showed me the multiple books he had on his shelf and I asked him:

Which ones are your favorite? If you had to pass just one to a developer starting today, which would it be?

He didn't hesitate at all:

Know your algorithms. Pick any of these here.

So I did just that. I'll admit, getting through all of these was tough, but it was incredibly fun!

The Code

You can find the code used in parts of this section up at [my github repo](#). In addition, you [can buy 17 video walkthroughs](#) of the algorithms you see in this chapter and others [from here](#).

The videos cover some new ground as well, and use languages other than C#.

Bubble Sort

Let's start with the simplest sorting algorithm there is: bubble sort. The name comes from the idea that you're "bubbling up" the largest values using multiple passes through the set.

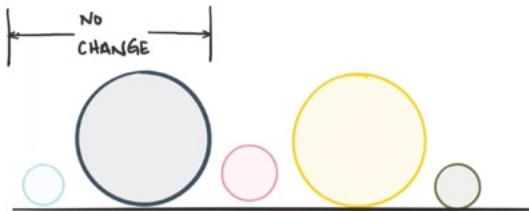
So, let's start with a set of marbles that we need to sort in ascending order of size:



A reasonable, computational approach to sorting these marbles is to start on the left side and compare the first

two marbles we see, moving the larger to the right.

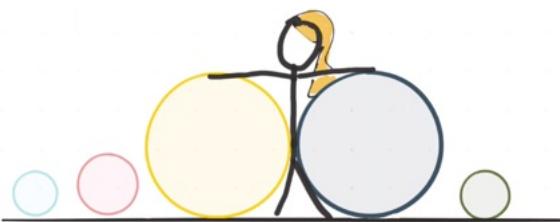
As you can see, the smaller marble is already on the left so there's no change needed:



Then we move on to the next two, which are the dark blue and the pink, switching their positions because pink is smaller than dark blue:



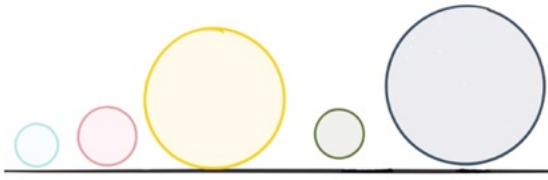
The same goes for yellow and dark blue, although an argument could be made that the author's drawing skills don't make it clear that dark blue is slightly larger.



The last two are simple: the green marble is much smaller than the dark blue, so they switch positions as well.



OK, we're at the end of our first pass, but the marbles aren't sorted yet. The green is out of place still.

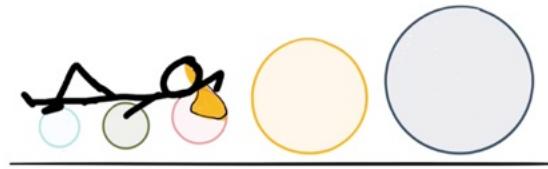


We can fix this by doing another sorting pass. This one will go a bit faster because blue and red are in order, red and yellow are in order, but green and yellow are not – so we make that switch:



Not too hard to figure it out from here. The green ball needs to move 1 more time to the left, which means one more pass to sort the marbles – making 3 passes in total.

Eventually, we get there:



Bubble sorts are not efficient, as you can see.

JavaScript Implementation

Implementing bubble sort in code can be done with a loop inside a recursive routine. That sentence right there should raise the hairs on the back of your neck! Indeed, bubble sort is not very efficient (as we'll see in a minute).

Here's one way to implement it:

```

//the list we need to sort
var list = [23,4,42,15,16,8];

var bubbleSort = function(list){
  //a flag to tell us if we need to sort this list again
  var doItAgain = false;
  //loop variables
  var limit = list.length;
  var defaultNextVal = Number.POSITIVE_INFINITY;

  //loop over the list entries...
  for(var i = 0; i < limit; i++){
    //the current list value
    var thisValue = list[i];
    //the next value inline, which we'll default to a really high number
    var nextValue = i + 1 < limit ? list[i + 1] : defaultNextVal;

    //is the next number lower than the current?
    if(nextValue < thisValue){
      //if yes, we switch the values
      list[i] = nextValue;
      list[i + 1] = thisValue;
      //since we made a switch we'll set a flag
      //as we'll need to execute the loop again
      doItAgain = true;
    }
  }
  //recurse over the list if the flag is set
  if(doItAgain) bubbleSort(list);
};

bubbleSort(list)
console.log(list);

```

Executing this code with Node we should see this:

```
[ 4, 8, 15, 16, 23, 42 ]
```

Looks sorted to me!

Complexity Analysis

As you can see, we're using recursion as well as a for loop, which should set off some alarms. If you recall from the chapter on Big-O, nested loops operating on the same list almost always means $O(n^2)$ and this algorithm is no exception to that, even if we're using recursion.

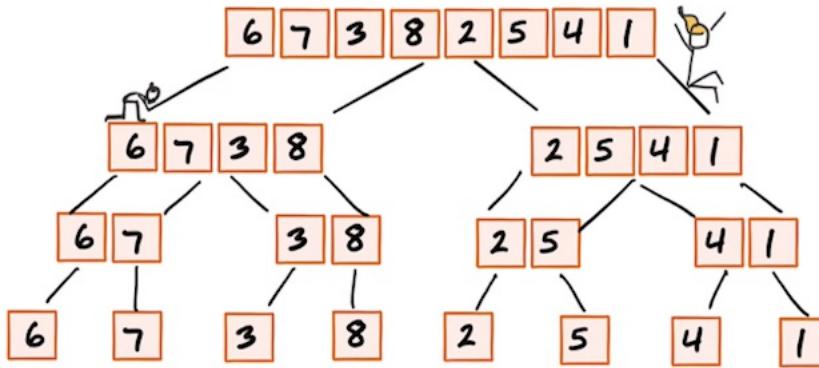
The use of recursion also means we're potentially taking up $O(n)$ space as well and opens up the possibility of a stack overflow exception given enough items to sort.

Merge Sort

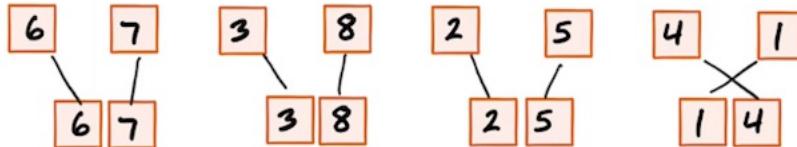
Merge Sort is one of the most efficient ways you can sort a list of things and, typically, will perform better than most other sorting algorithms. In terms of complexity, we're using a *divide and conquer* approach, which should tip you off that this is going to be at least $O(\log n)$. Once we divide the array, we need to sort the items which is going to be an $O(n)$ operation since we need to address each item. That means this algorithm's complexity is $O(n \log n)$.

Merge Sort works by splitting all the elements in a list down to smaller, two-element lists which can then be sorted easily in one pass. The final step is to recursively merge these smaller lists back into a larger list, ordering as

you go - this is the $O(\log n)$ part:

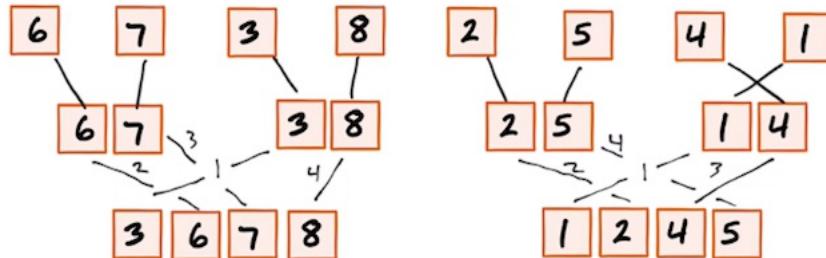


Now we need to merge the lists. The rules for this are simple: compare the first elements of adjacent lists, the lowest one starts the merged list – this is the $O(n)$ part:



This is straightforward with lists of one element being combined into lists of two elements. But how do we match up lists of two elements?

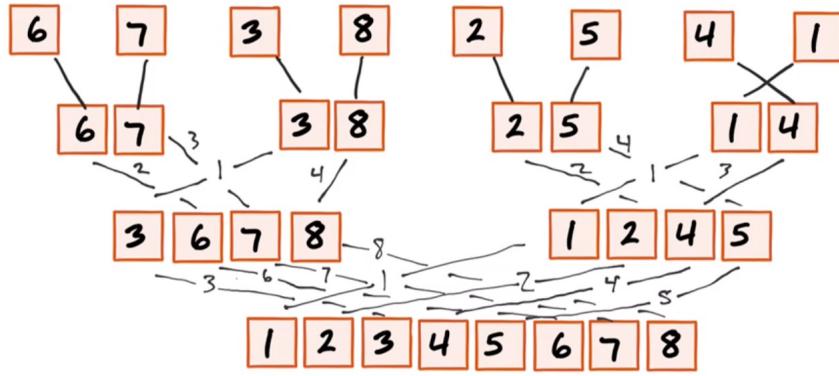
The same way. When combining the [6,7] list with the [3,8] list, we compare the 3 with the 6 – the 3 is smallest so it goes first. Then we compare the 6 with the 8 and the 6 is smaller, so it goes next. Finally we compare 7 and 8 and add them accordingly:



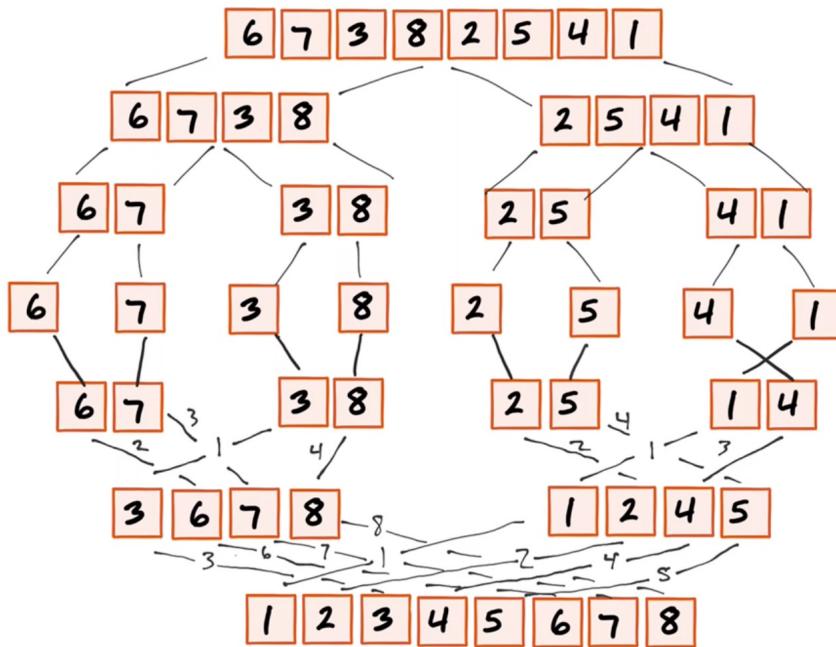
Now, you might be thinking “wait a minute – how do we know that a smaller number isn’t sitting to the right of the 6? Wouldn’t that mess up the sort?” That’s a good question.

It’s not possible to have a lower number to the right of any element in a merged list – when the [6,7] list was created we sorted it. This is the power of Merge Sort: the leftmost numbers are always smaller which gives us a lot of power.

OK, so now we continue on in the same way, merging the final lists of 4. We start on the left-hand side of each list, comparing the values, and adding the lowest to the merged list first:



And we're done! Here's the full operation, in case you'd like to see it top to bottom:



JavaScript Implementation

Implementing merge sort in code is a bit tricky. You need to have two dedicated routines, one for splitting the list and one for merging.

The first step is to recursively split the list:

```

var list = [23,4,42,15,16,8,3];

var mergeSort = function(list){
  var left=[], right=[];
  //if there's only one item in the list
  //return
  if(list.length <= 1) return list;

  //cut the list in half
  var middle = list.length / 2;
  var left = list.slice(0, middle);
  var right = list.slice(middle, list.length);

  //recursively run through the splits
  //left and right will be separated down to single elements
  return merge(mergeSort(left), mergeSort(right));
};


```

In this routine we're just splitting whatever list comes in right down the middle. If the list only has one entry, we're returning. This prevents the recursive call on the last line from blowing up.

Next is our merge function:

```

var merge = function(left, right){
  var result = [];
  //if the left and right lists both have elements
  //run a comparison
  while(left.length || right.length) {
    //if there are items in both sides...

    if(left.length && right.length) {
      //if the first item on left is
      //less than right...
      if(left[0] < right[0]) {
        //take the first item on the left
        result.push(left.shift());
      } else {
        //otherwise take the first item
        //on the right
        result.push(right.shift());
      }
      //if the right list is empty and
      //the left is not...
    } else if (left.length) {
      result.push(left.shift());
    } else {
      //there are items remaining on the right
      result.push(right.shift());
    }
  }
  return result;
};

console.log(mergeSort(list));

```

This routine takes two lists and compares their leftmost values. If one of the lists is empty then the left-most value from the other list is appended as the result.

Running this we get:

[3, 4, 8, 15, 16, 23, 42]

Hurrah!

Quicksort

Quicksort is a divide and conquer algorithm that uses a pivoting technique to break the main list into smaller lists. These smaller lists use the pivoting technique until they are sorted. The complexity of quicksort, however, is not constant because the pivot (as you're about to see) is determined at random and the partitioning of the list can put the algorithm at a disadvantage.

In the worst case, quicksort is $O(n^2)$ when the pivot is the smallest or largest element in the list. In the best case it's $O(n \log n)$, like merge sort. We'll discuss this a bit more at the end of this section; for now let's get to know this algorithm.

There are two ways to implement quick sort. Let's go over how the algorithm works and then I'll discuss the different implementations.

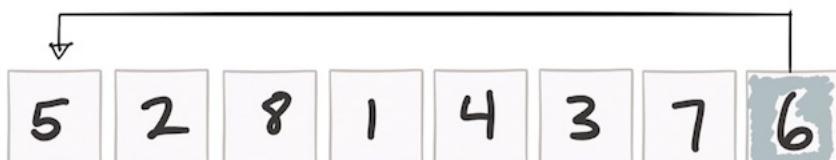
We'll start with a set of 8 elements (sorry, no cats or marbles this time).



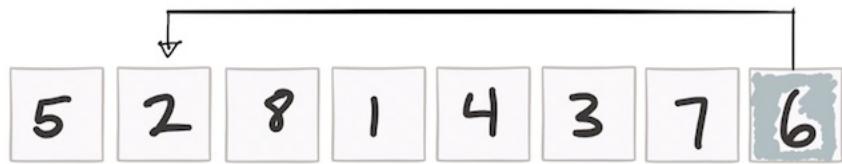
We need a pivot, so we'll choose the very last element of the list (by convention).



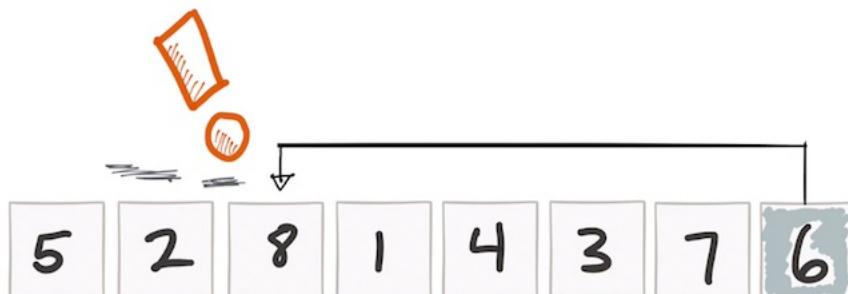
The next step is to partition our list so that all elements in the list that are less than our pivot are in a separate partition to the left, and all the elements greater are in a partition to the right. There are various ways to do this, but the simplest is to start at the beginning of the list – in this case a 5 – and if it's smaller we'll leave it in place.



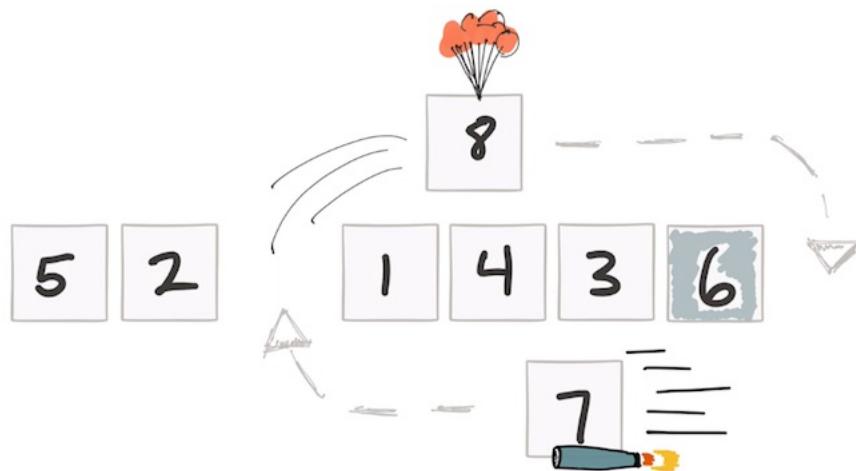
The next element is a 2, so we'll leave that there as well.



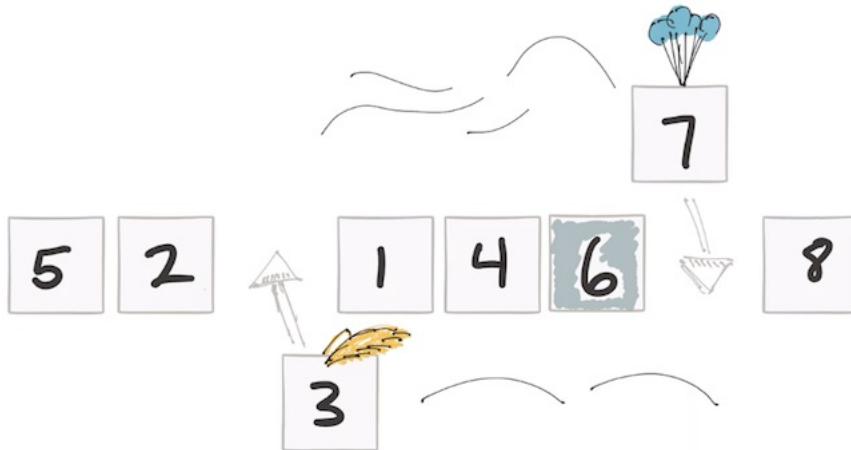
Then we come to an 8, which is greater than our pivot, which means we need to move it.



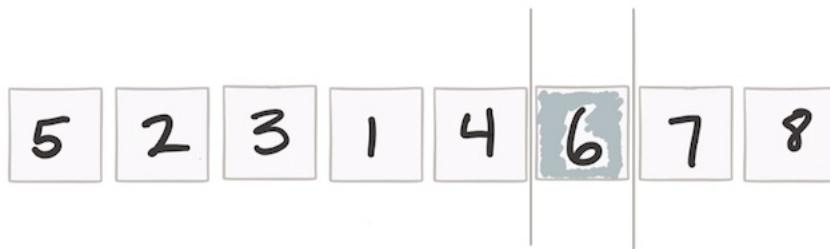
I'll pick it up and move it behind the pivot as it's a larger number, and I'll move the pivot down one position. There's already a number there – a 7 – so I'll move that to where the 8 was.



Now the 7 is the next thing to evaluate. It's greater than 6 so I'll do the same maneuver, switching the 3 and the 7, putting the 7 in the position my pivot was just in.



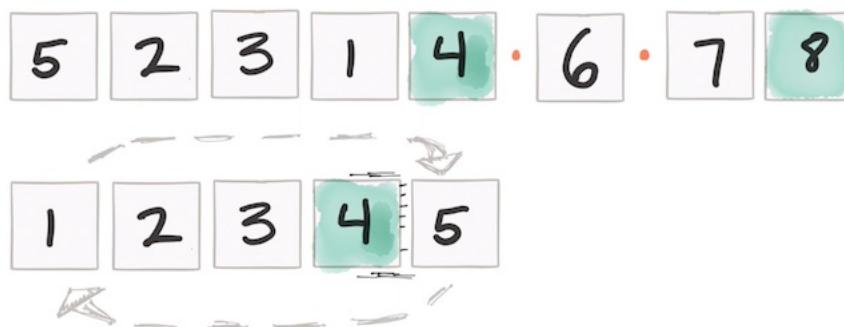
The next elements are 1 and 4 – which are both smaller than 6. This means we’re done with our first partitions!



We now have two new lists – the numbers less than 6 on the left and the numbers greater than 6, on the right. Six itself in its final position, so we’ll ignore it for now.

Now we pick a pivot for the new lists in exactly the same way – the last element in each list. This means that 8 is the pivot on the right, 4 is the pivot on the left.

The neat thing here is that there are no numbers greater than 8 in our list on the right, so there’s nothing we need to do.



The list on the left, however, can be separated in exactly the same way we did before. We’ll compare once again from the first position – it’s a 5 so we’ll stick it behind our 4, and move the number to the left of the 4 (a 1) to 5’s old position.

And just like that – we’re done!

Avoiding a Mess

If our list was presorted for whatever reason (and it turns out that yes, this does happen) then our sorting routine here would take a very, very long time. We'd have to split and order the list for every element, turning the complexity to $O(n^2)$.

To get around this we can select our pivot intelligently. This requires an initial step where you find the median value of a list and then make that the pivot. From there the sorting operation will usually beat out merge sort because the sorting happens *in-place*, which means we have decreased *space complexity* and fewer overall operations.

JavaScript Implementation

Again we'll use recursion to help us split the list into partitions. The one tricky thing here is that we need to remove the pivot from the evaluation, which is commented inline:

```
var list = [23,4,42,15,16,8,3];

function quicksort(list) {
    //recursion check. If list is empty or of length 1, return
    if (list.length < 2) return list;

    //these are the partition lists we'll need to use
    var left = [], right = [];

    //default the pivot to the last item in the list
    var pivot = list.length - 1;

    //set the pivot value
    var pivotValue = list[pivot];

    //remove the pivot from the list as we don't want to compare it
    list = list.slice(0, pivot).concat(list.slice(pivot + 1));

    //loop the list, comparing the partition values
    for (var i = 0; i < list.length; i++) {
        if(list[i] < pivotValue){
            left.push(list[i])
        }else{
            right.push(list[i]);
        }
    }

    //do it all again
    return quicksort(left).concat([pivotValue], quicksort(right));
}

console.log(quicksort(list));
```

Selection Sort

Selection sort is likely the simplest possible way to sort a list. It's how you and I might think of telling a duck to do it if a duck had hands and needed to get a job as a coder at Amazon.

This algorithm works by scanning a list of items for the smallest element, and then swapping that element for the one in first position. This continues with the remaining items until the list is sorted.

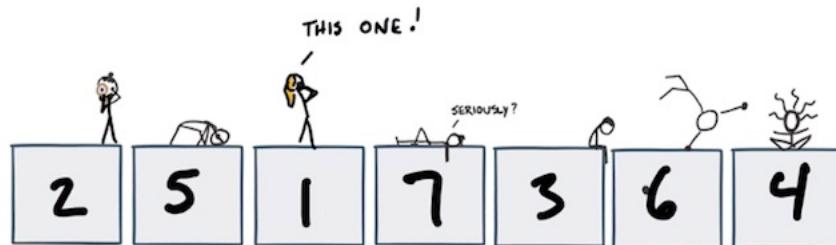
The complexity of selection sort ranges from $O(1)$ when the list is already sorted to $O(n^2)$ when the list is presorted in the reverse order you want. The $O(1)$ complexity makes this an interesting choice if there's a chance of

sorting a pre-sorted list (or a list of equal objects) – which happens fairly often.

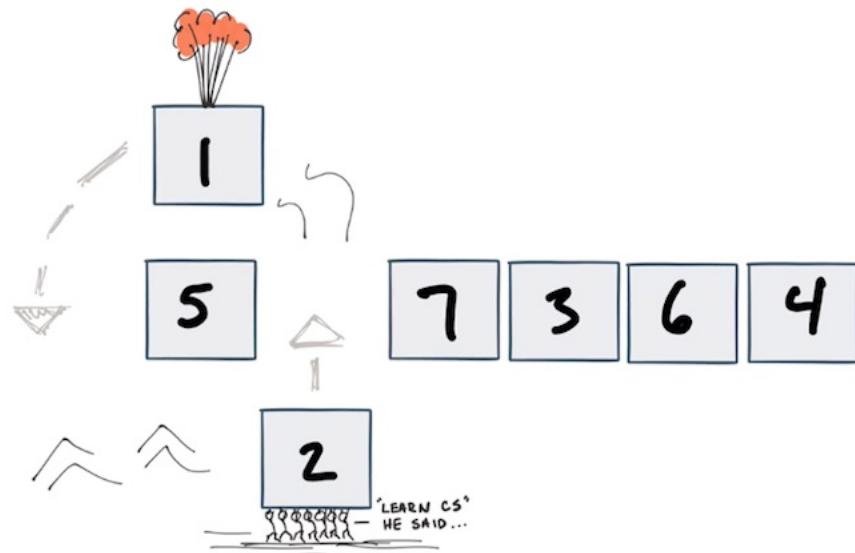
We'll start out with an unsorted list of 7 elements:



The first task is to find the lowest number, which (in the worst case scenario) is a linear scan:



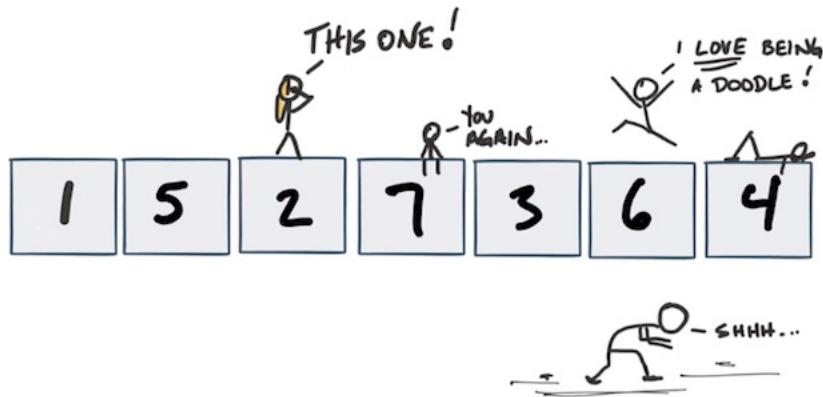
Once we've found the lowest element, we swap it with the first element in the list, as we know this is where the lowest element belongs.



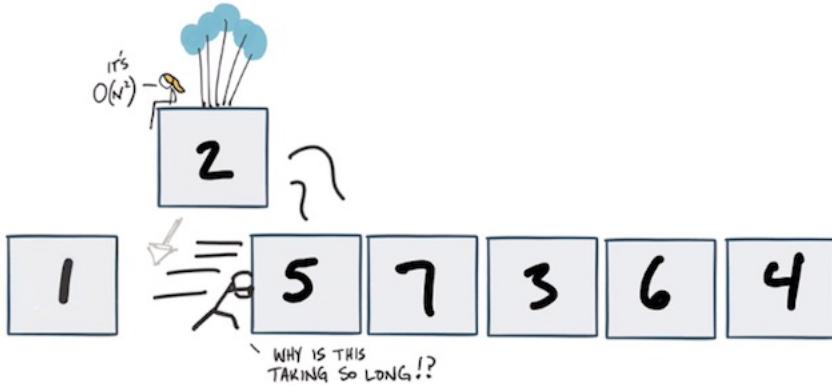
Then we do it all over with the remaining items. In this case the next lowest element is a 2.



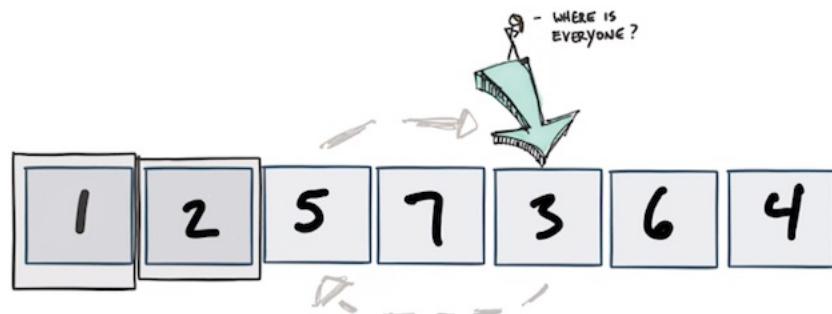
So we swap the two with the element that was in second position.



Rinse, repeat. We keep scanning the list for the lowest item and swapping it for the items that aren't in their final position (shaded):



Not a very efficient operation. My doodles have rebelled, apparently – but our last operation is swapping the 4 and the 7 and then we're done!



Implementation in JavaScript

This one is straightforward – no need for recursion. We can use two loops: the first will loop over our list, the second will loop forward for every step:

```

var list = [23,4,42,15,16,8,3];

function selectionSort(list) {

    for (var i = 0; i < list.length - 1; i++) {
        //default the min value to the first item in the list
        //all we need do is track the index for now
        var currentMinIndex = i;
        //loop over the list, skipping the currentMinIndex
        for (var x = currentMinIndex + 1; x < list.length; x++) {
            //if the current list item is less than the current min value...

            if (list[x] < list[currentMinIndex]) {
                //reset the index
                currentMinIndex = x;
            }
        }
        //has the index changed?
        if (currentMinIndex != i) {
            //if yes, switch the values in the list
            var oldMinValue = list[i];
            list[i] = list[currentMinIndex];
            list[currentMinIndex] = oldMinValue;
        }
    }
    return list;
}

console.log(selectionSort(list));

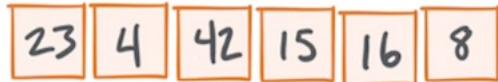
```

Heap Sort

Heap sort is a bit like selection sort in that it moves unsorted data to a sorted “partition” selectively. The difference, however, is that it uses a heap to do so. If you recall, a heap is a tree structure where parent nodes in one level are either greater than (max heap) or less than (min heap) child nodes in descendent levels.

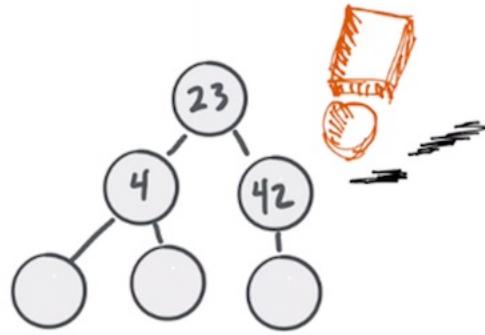
Heap sort is $O(n \log n)$, however it has an advantage over quicksort of being $O(n \log n)$ in the worst case scenario, whereas quicksort in the worst case is $O(n^2)$.

We have an unsorted list of numbers, as always.



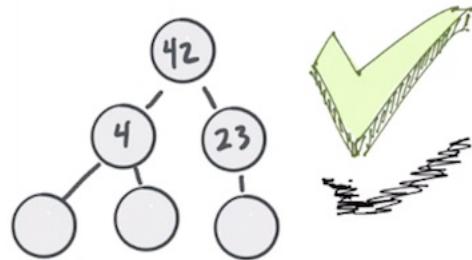
The first step is to move these numbers into a heap:

23	4	42	15	16	8
----	---	----	----	----	---



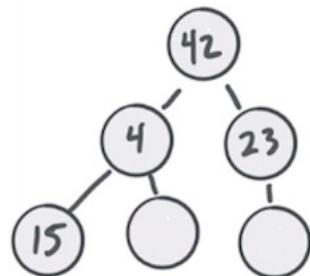
As we go along, however, we realize that we would violate the heap rules if 42 was to be placed before 23, so we swap them.

23	4	42	15	16	8
----	---	----	----	----	---

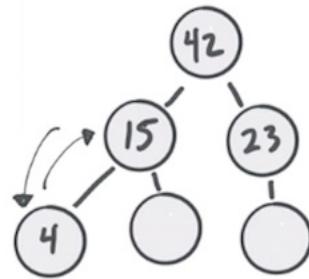
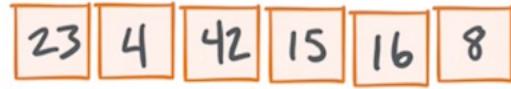


Now we have a valid heap.

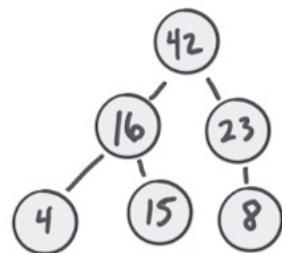
23	4	42	15	16	8
----	---	----	----	----	---



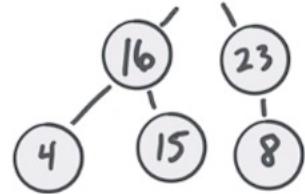
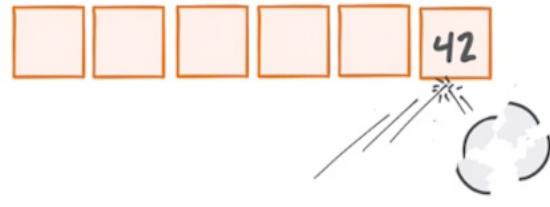
In the same way, 15 and 4 would be in violation – so we swap them.



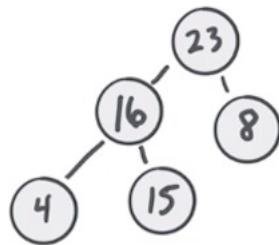
We end by adding 16 and 8, swapping positions for the 15 and 16 so we avoid violations. We now have a completed heap from our original list.



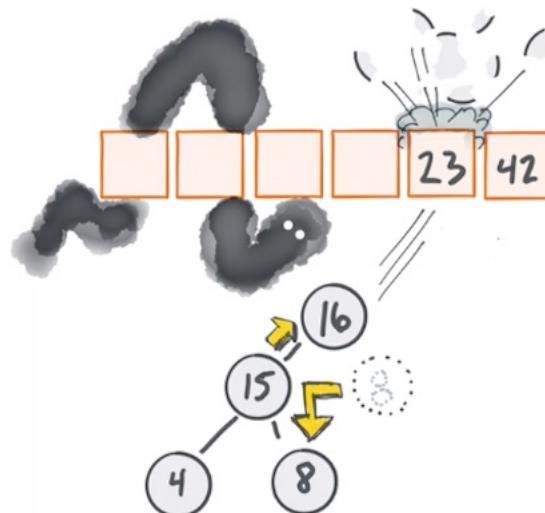
At this point we move on to the final stage, which is systematically removing the root of the heap, putting it back into the list.



As we add items back to our list, the heap is adjusted to make sure it remains valid. Here, we need to move 23 to the root as it's larger than its sibling, 16. 8 tags along for the ride as it's still in a valid position.



Next, we add 23 back to our list, and then elevate 16 to the root, as it's larger than 8. However, this means that 8 doesn't have a parent, so we move it over and place it under the 15.



We keep going in this way until all the nodes in our heap are gone.



Binary Search

Binary search is a fascinating thing. At first glance it seems rather ridiculous – the list you're searching over has to be sorted first! This isn't so nuts if you store your data in a binary tree (a BTREE) which we've discussed already.

Again, this algorithm is *divide and conquer* which should give a clue – I hope it does because we've already had a look at it in the Big-O chapter!

The deal is this: you split a list of sorted items and decide, from there, whether the item you're looking for is in the left or right list. You can decide this accurately because the list is sorted.

You then split that list and check to see if values on the right or left of the split are greater, lesser, or equal to the value you're searching for.

Keep going until you find what you want.

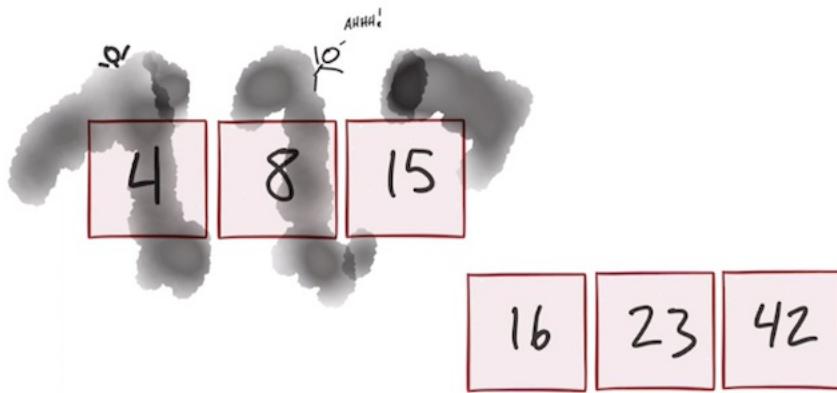
Let's take a look.

We have an ordered set to play with. How it became ordered is a mystery – as are the numbers – which we will figure out later in the book.

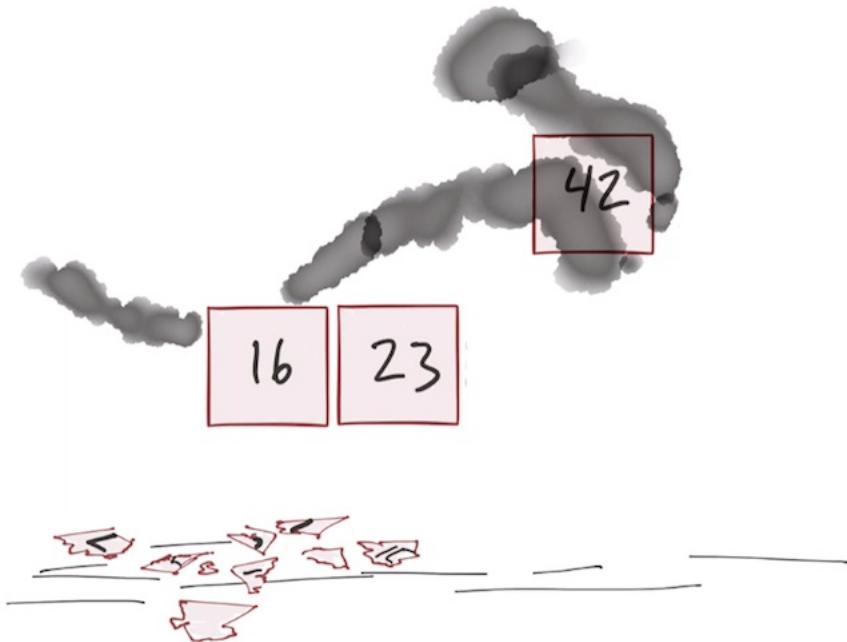


What we need to do is to find the number 23. Now, you and I both have eyes, but a computer doesn't – so we need to give it an intelligent way to find the number 23.

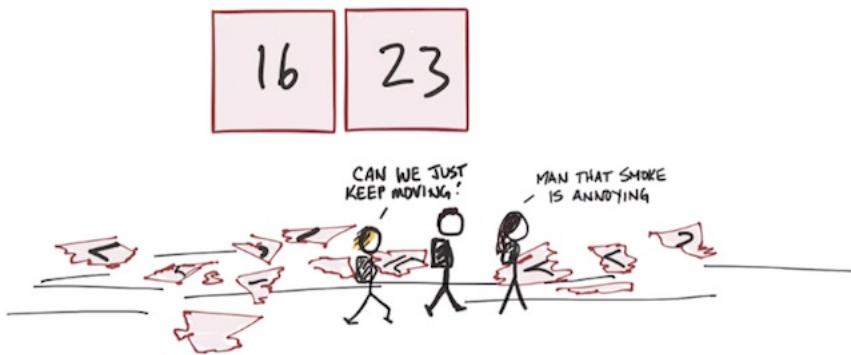
To do this, we'll split the list in two, right down the middle, and remove the entire half that is less than, or greater than, our number. In this case it's the side with 15 and below.



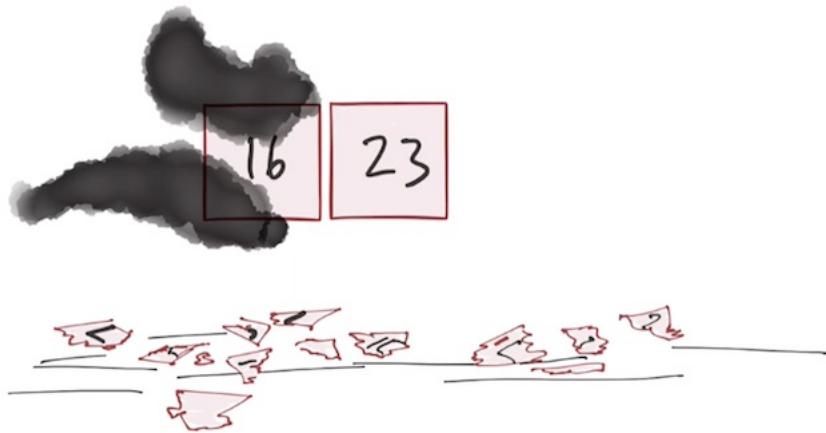
Since our list is sorted, we know the number we're looking for is in the remaining set. So, let's do the same thing. We'll split that list down the middle (or close to it).



The 42 is on the right – which means it and everything after it can be removed.



Now we're left with only two elements, which again we split down the middle. Evaluating each side, we see that the number we're looking for is on the right! We can remove the 16 and we're done!



JavaScript Implementation

Looking through a list by splitting it in half continually is fairly straightforward. In this routine we'll search for a value and return its index:

```

var list = [1,2,3,4,5,6];

function binarySearch(list, lookFor) {
    var min = 0, max = list.length -1;
    var middle;

    while (min <= max) {
        //find the middle of the list
        middle = Math.floor((min + max) / 2);

        //if we just happen to land on it...
        if (list[middle] === lookFor) {
            return middle;
        }
        else {
            //the list is sorted, so reset the list according to
            //position
            list[middle] < lookFor ? min = middle : max = middle;
        }
    }
    return -1;
}

console.log(binarySearch(list,3));

```

Graph Traversal

We've learned how to split lists apart using binary search, but how do we search over something a bit more complex, such as a graph? As you've noticed, many of the data structures we've been working with are based on graphs, so traversing them properly is quite important.

Before we go on, it's important to understand that you can carry out these operations in one of two ways: using *recursion* or using *iteration*. One is clever and one is costly – thankfully I'm referring to the same method: *recursion*.

Every time a recursive function is executed, the variables that the function uses remain in scope and are stored on the stack. I know I've brought this up before, but it's worth repeating: *recursion is interesting and clever, but can also be costly*.

I point this out because *you will be asked something like this in an interview someday*. Recursion and space complexity are not friends – so come prepared to understand why. Let's go a bit deeper.

Using Recursion

For the sake of variety, let's do this with C#. Here is a basic representation of a node in a binary tree:

```

public class BinaryNode
{
    public int Value { get; }
    //child nodes
    public BinaryNode Left { get; set; }
    public BinaryNode Right { get; set; }

    public BinaryNode(int value)
    {
        Value = value;
    }
    public bool IsLeaf(){
        return this.Left == null && this.Right == null;
    }
    public void Traverse(BinaryNode node){
        if(node.Left != null){
            //do some kind of calculation and then...
            Traverse(node.Left);
        }
        if(node.Right != null){
            //do some kind of calculation and then...
            Traverse(node.Right);
        }
        //Leaf node!
        if(this.IsLeaf()) {
            //we're done!
        };
    };
}

```

A node instance has a value as well as two possible descendants, **Left** and **Right**. The **Traverse** method allows you to start from any node in the tree and then traverse through each node, doing something exciting and amazing.

There are a number of ways to run calculations with this class. We can hard code what we need done in the **Traverse** method (searching for a value, for instance) or we can pass along some kind of callback, perhaps a **Func<T, bool>**. For the sake of keeping things at a pace for this book, I'm going to sidestep implementation details here and simply wave my arms and say "don't worry about this for right now".

The reason why is that this code, while clever, is also a bit inefficient. I already discussed the stack overflow problems above, but what if you wanted to control what *type* of search you wanted to perform? In other words, what if you wanted to *go deep* into the tree as you knew the thing you're looking for was likely at the lower levels of the tree?

Conversely: what if you wanted to *go wide* instead? Recursion can be tweaked to address these things, but then you're altering your class to accommodate what should really be the responsibility of an entirely different bit of code altogether.

Using Iteration

Instead of filling up our stack, what if we used a simple loop? I'll use C# again for this, but hopefully you can translate it to whatever language works for you.

Let's strip out the **Traverse** method and go with a more cohesive **BinaryNode** class:

```

public class BinaryNode
{
    public int Value { get; }

    //child nodes
    public BinaryNode Left { get; set; }
    public BinaryNode Right { get; set; }

    public BinaryNode(int value)
    {
        Value = value;
    }

    public bool IsLeaf(){
        return this.Left == null && this.Right == null;
    }
}

```

Now let's fill out our nodes:

```

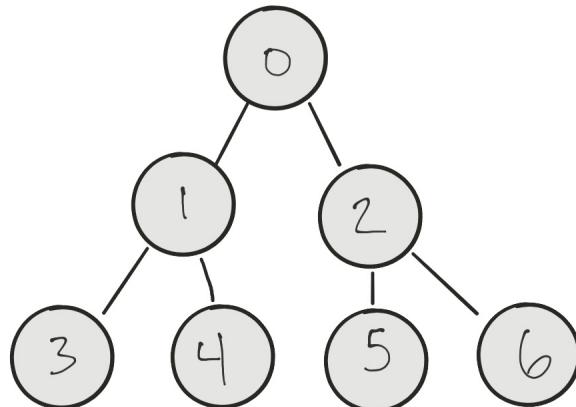
//create the tree structure..
var rootNode = new BinaryNode(0);
rootNode.Left = new BinaryNode(1);
rootNode.Right = new BinaryNode(2);
rootNode.Left.Left = new BinaryNode(3);
rootNode.Left.Right = new BinaryNode(4);
rootNode.Right.Left = new BinaryNode(5);
rootNode.Right.Right = new BinaryNode(6);

```

This will give us a two level tree, which will be good enough to traverse. But how? It turns out there are two ways: *Breadth-first* and *Depth-first*. These types of traversals are typically referred to as “breadth-first search” and “depth-first-search”.

Depth-first Search

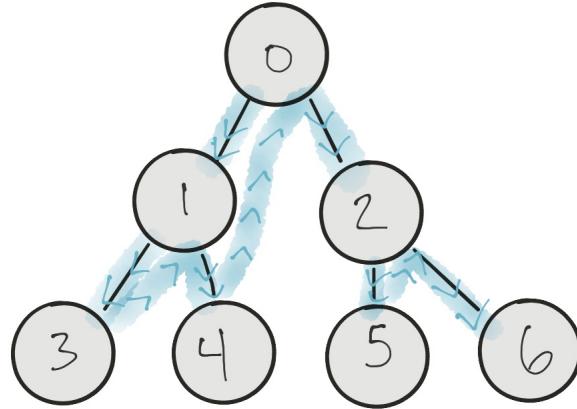
Let's start by visualizing our graph:



With a depth-first search, we want to go as deep as possible into the graph from either the left or the right. Traditionally it's done from the left.

The idea is to go as deep as you can and then backtrack your way back up and then over to the right until the traversal of the tree is done.

With my crude drawing skills, it might look something like this:



We start at the root, proceed to 1 and then 3. Then backtrack to 4, back up and over to 2, then 5 then 6.

To implement this we need a data structure that will keep track of the very next descendant as well as the current sibling. We will always follow the next descendant, and if it doesn't exist, we'll go with the last saved sibling. If both of those don't exist we're done with our tree.

For this, we can use a stack. Here's some basic code, once again in C#:

```
public class DFS{
    public void Traverse(BinaryNode root){
        var stack = new Stack<BinaryNode>();
        BinaryNode thisNode = null;
        //push the root onto the stack
        stack.Push(root);
        while(stack.Count > 0){
            thisNode = stack.Pop();
            Console.WriteLine(thisNode.Value);
            //push right then left
            if(thisNode.Right != null){
                stack.Push(thisNode.Right);
            }
            if(thisNode.Left != null){
                stack.Push(thisNode.Left);
            }
        }
    }
}
```

All we need to do is to keep track of how many nodes there are left to traverse. We know that by using a stack and pushing values onto it if they exist down below.

Let's load it up and run it!

```

public class DFS{
    public void Traverse(BinaryNode root){
        //...
    }
    public void Run(){
        var rootNode = new BinaryNode(0);
        rootNode.Left = new BinaryNode(1);
        rootNode.Right = new BinaryNode(2);
        rootNode.Left.Left = new BinaryNode(3);
        rootNode.Left.Right = new BinaryNode(4);
        rootNode.Right.Left = new BinaryNode(5);
        rootNode.Right.Right = new BinaryNode(6);
        this.Traverse(rootNode);
    }
}

```

The result, when we execute this, is right on the money:

```

0
1
3
4
2
5
6

```

“Going deep” on a graph is useful if you’re interested in the nature of parent/child relationships.

You could also look through a binary tree and see if it is in fact a binary *search* tree, with all of the node values set as required. We know from reading the last chapter that a binary search tree requires that every child node to the right of a given node must have a greater value; every child node to the left must have a lesser one.

Is depth-first the best choice for this? You might check a lot more nodes than you need to! Let’s have a look at an alternative: *breadth-first*.

Breadth-first

With breadth-first search we need to track the nodes of a tree in each level before traversing to the next level. This means we need to track every node and its children *in order*.

For this we can use a queue:

```
public class BFS{
    public void Traverse(BinaryNode root){
        var queue = new Queue<BinaryNode>();
        BinaryNode thisNode = null;
        //push the root onto the queue
        queue.Enqueue(root);
        while(queue.Count > 0){
            //remove the next item from the queue
            thisNode = queue.Dequeue();
            //do whatever calc
            Console.WriteLine(thisNode.Value);
            //add the children into the queue
            if(thisNode.Left != null){
                queue.Enqueue(thisNode.Left);
            }
            if(thisNode.Right != null){
                queue.Enqueue(thisNode.Right);
            }
        }
    }
}
```

If we load it up just like we did in the last section and run it, we will have a completely different result:

```
0
1
2
3
4
5
6
```

We're traversing the graph one level at a time, which makes great sense if we're trying to evaluate all of the parent nodes *first*, which we would be doing if we're validating a binary search tree.

CHAPTER EIGHT: ADVANCED ALGORITHMS

I wasn't sure about using the terms "advanced" and "simple" for these chapters on algorithms. I don't want you to think these are harder to implement or to figure out – they're not (necessarily).

The reason I chose these terms is that it's more likely you'll find yourself getting paid to implement a shortest path algorithm or some type of sieve at some point in your career. It's *unlikely* that you'll be hired to implement bubble sort.

I also wanted to provide a framework for you to create your own algorithms for solving the more interesting problems you work on every day. This, to me, is the most important part of this chapter, so we'll start there.



Dynamic Programming

No, this section is not about Ruby, Python, JavaScript, etc. Dynamic programming is a way to solve a problem using an algorithm in a fairly prescribed way. It sounds complicated, but it's anything but.

Dynamic programming gives us a way to elegantly create algorithms for various problems and can greatly improve the way you solve problems in your daily work. It can also help you ace an interview.

Definition

Let's start with a quick definition so we know what dynamic programming is and how it works. At its core, dynamic programming is simply solving an optimization problem by guessing in a systematic way. It's almost laughable to think about dynamic programming in terms of this definition, but as you'll see it turns out to be rather powerful.

To use dynamic programming, the problem you're solving must be:

- **An optimization problem.** We saw one of these in Chapter 1 (the Bin Packing Problem) when I tried to optimize storage for my daughter's things.
- **Dividable into subproblems.** With dynamic programming you can recurse over and solve in order to solve the larger, objective problem.
- **Have an optimal substructure.** That's a mouthful, but what it means is that the subproblems you solve must be complete unto themselves. In other words, if you solve subproblems x, y and z in order to solve objective problem A, then the solutions to x, y and z should be sufficient on their own to solve A. You don't need to use x plus some other algorithm.
- **Reducible to P time through memoization.** Some of the problems you can solve with dynamic programming are solvable in exponential time (like Fibonacci), however this can be reduced to P time through memoization. Another fun word, but you can think of this basically as "caching" the answers to the subproblems and then applying.

I wouldn't blame you if you're underwhelmed at this point. The name "dynamic programming" seems pretty bland, and the underlying techniques more than a little vague. You'll understand it well in a few sections as we solve some problems with it, I promise.

Before we get there, it's important to understand where the name came from and why dynamic programming even exists.

Origins

This is a funny story (emphasis mine):

An interesting question is, ‘Where did the name, dynamic programming, come from?’ The 1950s were not good years for mathematical research. We had a very interesting gentleman in Washington named Wilson. He was Secretary of Defense, and he actually had a pathological fear and hatred of the word, research. I’m not using the term lightly; I’m using it precisely. His face would suffuse, he would turn red, and he would get violent if people used the term, research, in his presence. You can imagine how he felt, then, about the term, mathematical. The RAND Corporation was employed by the Air Force, and the Air Force had Wilson as its boss, essentially. Hence, I felt I had to do something to shield Wilson and the Air Force from the fact that I was really doing mathematics inside the RAND Corporation. What title, what name, could I choose? In the first place I was interested in planning, in decision making, in thinking. But planning, is not a good word for various reasons. I decided therefore to use the word, ‘programming.’ I wanted to get across the idea that this was dynamic, this was multistage, this was time-varying—I thought, let’s kill two birds with one stone. Let’s take a word that has an absolutely precise meaning, namely dynamic, in the classical physical sense. It also has a very interesting property as an adjective, and that is it’s impossible to use the word, dynamic, in a pejorative sense. Try thinking of some combination that will possibly give it a pejorative meaning. It’s impossible. Thus, I thought dynamic programming was a good name. It was something not even a Congressman could object to. So I used it as an umbrella for my activities.

If you want to read Richard Bellman's original paper on dynamic programming, [you can do so here](#).

There you have it: **the name means nothing**. The dynamic programming design process, however, is behind some of the most powerful algorithms we know of. We'll see those in the next section.

The best way to see its power, however, is to just do it. So let's! We'll use dynamic programming to help us get through a job interview.

Fibonacci

You knew Fibonacci was going to come up in this book at some point didn't you! Well, here it is. I'm using it here because it's the simplest way to convey the dynamic programming process. Also: you will be asked how to solve Fibonacci at some point in your career, and you're about to get three different approaches!

Which leads right to a great opening point: *our jobs are about solving problems*. When you go to these interviews, they mostly want to see how you would go about solving something complex. As it turns out, the Interviewing For Dummies book says that Fibonacci is a great question for just that case.

Definition

So let's start with a definition, just in case you don't know or remember what a Fibonacci Sequence is:

A series of numbers in which each number (Fibonacci number) is the sum of the two preceding numbers. The simplest is the series 1, 1, 2, 3, 5, 8, etc.

Lovely. Why do we care about these numbers? These numbers (and the algorithm we're about to discuss) underpin [nature's symmetry](#):

The Fibonacci numbers are Nature's numbering system. They appear everywhere in Nature, from the leaf arrangement in plants, to the pattern of the florets of a flower, the bracts of a pinecone, or the scales of a pineapple. The Fibonacci numbers are therefore applicable to the growth of every living thing, including a single cell, a grain of wheat, a hive of bees, and even all of mankind.

If you divide each successive number by itself (so: 5/3, 8/5...) you converge on a fascinating number called [phi](#):

What makes a single number so interesting that ancient Greeks, Renaissance artists, a 17th century astronomer and a 21st century novelist all would write about it? It's a number that goes by many names. This "golden" number, 1.61803399, represented by the Greek letter Phi, is known as the Golden Ratio, Golden Number, Golden Proportion, Golden Mean, Golden Section, Divine Proportion and Divine Section. It was written about by Euclid in "Elements" around 300 B.C., by Luca Pacioli, a contemporary of Leonardo Da Vinci, in "De Divina Proportione" in 1509, by Johannes Kepler around 1600 and by Dan Brown in 2003 in his best selling novel, "The Da Vinci Code."

Absolutely fascinating stuff. Our interviewer, however, is waiting patiently for us to come up with an algorithm for calculating a Fibonacci Sequence to the nth position – so let's get to it!

The Painful Way

The interviewer has asked us a standard question:

How would you derive a Fibonacci sequence up to a given position?

In other words, if we're given a value of 10, the interviewer will want to see the first 10 Fibonacci numbers. We can solve this (and more!) using dynamic programming.

The first step is to break the problem down into smaller problems (called subproblems) that we can solve. If we're trying to derive a Fibonacci sequence to the 10th position, we can do it with pen and paper like this:

- The first number in the Fibonacci sequence is 0
- The second number is 1
- The third number is $0+1=1$
- The fourth number is $1+1=2$

And so on. This would answer the interviewer's question (about the sequence) but it wouldn't show them what they're after: our ability to solve a problem programmatically. We can do this with the next step in dynamic programming: recursively solve the subproblems until the objective problem is solved.

It's easiest if we see some code at this point. Here's my Fibonacci solver implemented in JavaScript:

```

var calculateFibAt = function(n){
  if(n < 2){
    return n;
  }else{
    return calculateFibAt(n-2) + calculateFibAt(n-1);
  }
}

for(var i = 0; i<=10; i++){
  console.log(calculateFibAt(i));
}

```

Running this (using Node):

```

0
1
1
2
3
5
8
13
21
34
55

```

Great! By the way I tried four times to write this from memory and *completely failed*. You would think this little routine would be embedded in my mind but ... oh well. If it took you a few times to come up with it don't feel bad! Recursive programming takes some getting used to.

The code in this routine works, is straightforward, and is standard interview fare. We're feeling happy about ourselves at this point, when the interviewer says:

Talk to me about the complexity of this routine in terms of time and also space...

Uh-oh... time for some Big-O! Good news for us is the that we started this part of the book off with a discussion of Big-O and our brains are burning right now, thinking "right... wasn't there something bad about using recursion?"

Complexity Analysis

Let's start off by thinking in terms of *time complexity*. How long will it take to run our recursive algorithm? You can see why you might get fired by changing the loop value to 1000. In short: this routine scales horribly. To see this, let's add a counter to the function:

```

var fibCount=0;
var calculateFibAt = function(n){
    fibCount = fibCount+1;
    var calc;
    if(n < 2){
        return n;
    }else{
        return calculateFibAt(n-2) + calculateFibAt(n-1);
    }
}

for(var i = 0; i<=10; i++){
    var fib = calculateFibAt(i);
    console.log("The Fibonacci number at position " + i + " is " + fib + "; It took
}

```

When we run this:

```

The Fibonacci number at position 0 is 0; It took 1 calls to fib to get here
The Fibonacci number at position 1 is 1; It took 2 calls to fib to get here
The Fibonacci number at position 2 is 1; It took 5 calls to fib to get here
The Fibonacci number at position 3 is 2; It took 10 calls to fib to get here
The Fibonacci number at position 4 is 3; It took 19 calls to fib to get here
The Fibonacci number at position 5 is 5; It took 34 calls to fib to get here
The Fibonacci number at position 6 is 8; It took 59 calls to fib to get here
The Fibonacci number at position 7 is 13; It took 100 calls to fib to get here
The Fibonacci number at position 8 is 21; It took 167 calls to fib to get here
The Fibonacci number at position 9 is 34; It took 276 calls to fib to get here
The Fibonacci number at position 10 is 55; It took 453 calls to fib to get here

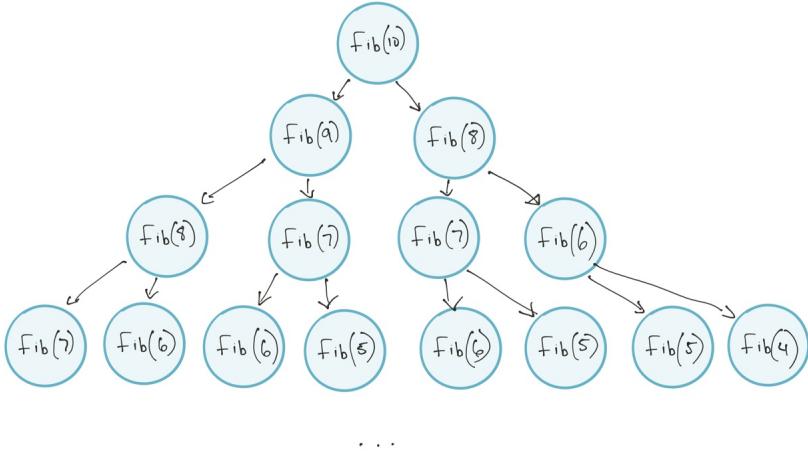
```

453 calls! Good grief! As you can see, the number of calls to our routine goes up more than exponentially with each additional input:

- Input 0 resulted in 1 call
- Input 1 resulted in 2
- Input 2 resulted in 5 calls
- Input 6 resulted in 59
- Input 10 resulted in 453

If you really want to have some fun, let it run for 15 minutes and see how many calls it takes to calculate **calculateFibAt(32)**... it's 18,454,894!!!

Another way to think about this is from the top down. Here we can visualize the complexity for calculating the Fibonacci number in 10th position using a graph:



This is horribly inefficient. Look how many times $\text{fib}(6)$ and $\text{fib}(7)$ are called! The interviewer seems happy with our answer and then asks us:

Tell me about the space complexity...

Right. We learned from the last section that a recursive routine will push values onto the stack repeatedly; once for every single call of the current function. If you do this enough you'll run into a stack overflow exception, which is bad.

As mentioned before: *recursion and space complexity aren't friends.*

Sounds good. So how would you improve this routine?

This is where we get to the next step of dynamic programming. We can reduce the time and space complexity of our algorithm by using memoization. We can do this because the solution to each subproblem is optimal, meaning that it can stand alone and we don't need anything else to use its value.

In Big-O terms, we can use the memoized solution in linear time, $O(n)$ where n is the position we're interested in, which will speed things up tremendously. What about *space complexity*? Can you figure out a way to do this in constant space? I'll talk about that in the next section.

The Faster Way

Memoization is simply caching. In more formal terms it's remembering the solution to a subproblem so you don't have to calculate it again recursively. This only works if the subproblems are in an optimized substructure. You can think of that as a large graph (like the one above), where you can simply replace $\text{fib}(6)$ with the number 8. That's darn optimal if you ask me.

To accomplish this, I'll store the results of our loop in some kind of data structure; the question is *which one?* We've learned about a whole mess of them in a previous chapter... which would be the best?

All we need to do is to remember some values in memory and then to iterate over them. If this is all you need, *don't overthink it!*

Since we know that Fibonacci numbers start with 0 and 1, I can use those seeds to calculate the remaining

numbers:

```
var fibFaster = function(n){
  var sequence = [0,1];
  for(var i=2; i<=n; i++){
    sequence.push(sequence[i -2] + sequence[i-1]);
  }
  return sequence;
}

console.log(fibFaster(10));
```

The result, when run:

```
[ 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55 ]
```

Perfect. But what about *time complexity*? This is simple to reason through, but before we do let's run this faster Fibonacci routine 1000 times. You should see that it returns almost instantly.

Complexity Analysis

We want to find the fibonacci numbers up to a given number n , which means we'll need to perform some operation for each n . If you recall from a few chapters back, iterating over a collection is always $O(n)$. Accessing a value from an array using its index is always $O(1)$, so our total operation here is $O(n * 1)$ which is $O(n)$.

But what about the *space complexity*? We're not using recursion so that means we won't be potentially overloading the stack, which is a good thing. We have a few variables and an array for every n number that we're evaluating, so our space complexity is also $O(n)$.

You might be thinking "hey wait a minute you have a loop variable in there too!" and yes, that's true, but with Big-O you're more concerned about the *nature* of algorithm. In this case it's simply $O(n)$.

Can we do better here? Yes. Well... sort of. Right now we're returning an array of fibonacci numbers up to the n th number. We could ask our interviewer at this point if they're interested in the whole sequence or just the n th number?

Just the nth number will do.

Perfect. This means we can now use a *greedy algorithm*, which is a term you should remember for interviews. Let's take a small diversion (again).

Greedy Algorithms

A greedy algorithm does what's best at that moment. Put in math terms:

A greedy algorithm is an algorithmic paradigm that follows the problem solving heuristic of making the locally optimal choice at each stage with the hope of finding a global optimum.

Say what?

Let's do this with code, then I'll see how well I can explain the idea. We'll start by redoing our **fibFaster** function:

```

var fibConstantSpace= function(n){
  let twoFibsAgo = 0, oneFibAgo=1, currentFib=0;
  for(var i=2; i<=n; i++){
    currentFib = twoFibsAgo+oneFibAgo;
    twoFibsAgo = oneFibAgo;
    oneFibAgo = currentFib;
  }
  return currentFib;
}

```

No more arrays required! What we're doing here is simply "remembering" *only what we need to remember*, essentially "locally optimizing" our decision making by setting these values to variables and forgetting the rest.

That's a greedy algorithm in practice. Our interviewer likes this, but has a question:

Can you think of a way to make this a bit more flexible for calling code? Right now you're just returning the nth fibonacci number; what if I wanted to do something else in my calling code?

If this happens to you in an interview, take it as a good sign. You've nailed the question so far! They're happy with your response and are likely just wanting to see how much better you are than the question allows.

We're using JavaScript, so ideally the answer is jumping out at you. Most languages support the idea of *callbacks*, something that will yield control of the current iteration/operation. We can use that here to *yield* the currentFib value back to the calling code:

```

var fibConstantSpace= function(n, fn){
  let twoFibsAgo = 0, oneFibAgo=1, currentFib=0;
  //make sure to output the first two
  if(fn){
    fn(0);
    fn(1);
  }
  for(var i=2; i<=n; i++){
    currentFib = twoFibsAgo+oneFibAgo;
    twoFibsAgo = oneFibAgo;
    oneFibAgo = currentFib;
    if(fn) fn(currentFib);
  }
  return currentFib;
}

```

There we go! The best of both worlds: we can report back each number so the calling code can do whatever it wants or we can just return the final value.

Use In Heuristics

Greedy algorithms can be useful when solving some very complex problems. A few chapters ago we took a look at some very tough, NP-Hard optimization problems, one of which was The Traveling Salesman Problem.

One way to approximately solve this problem is using a heuristic (that's a fancy word for "rule of thumb") called **Nearest Neighbor**:

The nearest neighbour algorithm was one of the first algorithms used to determine a solution to the traveling salesman problem. In it, the salesman starts at a random city and repeatedly visits

the nearest city until all have been visited. It quickly yields a short tour, but usually not the optimal one.

In other words: *there's no master plan here*. Nearest Neighbor just looks at the next cheapest city and goes there. This is a classic *greedy algorithm*.

Another greedy solution is finding your way out of a maze. You just put your right hand on the nearest wall and keep walking until you're out. Not the *optimal* solution, but it *will* solve the problem.

For a final example: consider Agile Development. Teams gather quickly in the morning so everyone's aware of what's going on with everyone else. Adjustments are welcomed and deployment rapid - it's all about quick adaptation to the changes in the development process.

Now, if you were to step back and look at a software project as a series of decisions which you could represent on a graph (which you can), then Agile is, itself, a greedy algorithm! It might not be the *optimal* solution to success for a project, but it is a solution! You're simply doing the next, closest thing that makes the most sense to the team and client without a master (waterfall) plan in place.

Calculating Primes

We've made it to the second round of our interview! Good work. This next interviewer looks trickier, however, and is readying up her question with a sly look in her eye.

Tell me how you would calculate all the prime numbers in a given range.

This problem has been solved (rather elegantly I might add) over 2200 years ago by a chap named Eratosthenes. His algorithm is so elegant that just about anyone can understand it. In mathematical terms, his algorithm uses a sieve to filter and extract the numbers in a set:

Sieve theory is a set of general techniques in number theory, designed to count, or more realistically to estimate the size of, sifted sets of integers. The primordial example of a sifted set is the set of prime numbers up to some prescribed limit X .

Before you read further: *how would you solve this problem?* Also: would you consider this dynamic programming? We'll answer the latter question at the very end.

The Sieve of Eratosthenes

The first step is to create the set of numbers you're interested in. We'll set our range limit (n) to be 100 ($n=100$) and I'll use a table to represent the set:

	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	100

Start with the number 2 (which is prime) and mark all the multiples of 2 in the set, as they (by definition) are not prime:

	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	100

█ = 2's

Move to the next number in the set (3). It's not marked, which means it's prime! Now we need to mark off all multiples of 3 which are, by definition, not prime:

	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	100

= 2's
 = 3's

We next hit 4, which is marked. It's not prime and, as it turns out, all of its multiples have been marked because of 2, so we skip it. Next is 5, which is not marked (therefore it's prime), and we mark off all its multiples as before, unless they're already marked:

	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	100

= 2's
 = 3's
 = 5's

We can skip 6, 8 and 10 because all of their multiples were covered by 2. 9 is covered by 2s and 3s so we can skip that as well. That leaves marking off 7 as our last effort:



The remaining unmarked numbers are all primes! Simple stuff. There's one other interesting thing here, however: notice the list of numbers that we were able to mark off on the right?

Yep. Primes. As we discussed in a previous chapter, prime numbers underly all kinds of sequencing algorithms, including generating every integer in our number system except for (of course), the prime numbers themselves.

OK, enough drawing! Let's implement this now with some code.

Break It Down

Our main problem (as defined by our interviewer) is to generate a list of primes within a given range. To do this, we need to:

- Create the range
- Calculate our iterator
- Remove the multiples of the given index
- Spit out a final list

Let's implement this in the simplest way possible (again, using JavaScript):

```

const sieve = (n) => {
  //the grid of numbers from 2 to n
  //you could use an array here too
  var grid = {};

  //load the grid
  for (var i = 2; i <= n; i++) {
    grid[i]={marked: false};
  }

  //figure out how many multiples we'll need
  const limit = Math.sqrt(n);

  //loop the multipliers...
  for (var i = 2; i <= limit; i++) {
    //calculate the multiples...
    for(var x = i + i; x <= n; x += i){
      //mark them on the grid
      grid[x].marked = true;
    }
  }
  var out =[];
  //return only those marked as true
  for (var i = 2; i <= n; i++) {
    if(!grid[i].marked) out.push(i);
  }
  return out;
};

console.log(sieve(100));

```

This works, and if you bump the list size to 10,000 you can see that we don't necessarily have a problem with performance. That said I'm sure some of you will find some optimizations!

Complexity Analysis

This is a straight up list iteration so the Big-O on this is $O(n)$. Our interviewer is happy with our answer, but that sly look is coming back! She casually says:

Let's say that space isn't a concern; I'd like to see you implement this with head/tail recursion...

We've given her a very standard answer to this problem, and she's onto us! And what the heck is *head/tail recursion*?

You can split a list into two parts: the *head* (which is the first item in the list) and the rest of the list itself, which is called the *tail*. If you do this, you can recurse over the list easily and run some calculations. This kind of thing is common in functional programming (which we'll get to in a later chapter) but, if you don't know it, you will in a minute!

Once Again, With More Recursion

In functional programming, head/tail recursion is a way to iterate over a list in a functional way. The head is simply the first value of the list, the tail is the rest of the list. We can implement head and tail functions in JavaScript in this way:

```

Array.prototype.head = function(){
  if(this.length > 0) return this[0];
  else return null;
}
Array.prototype.tail = function(){
  if(this.length > 0) return this.slice(1, this.length);
  else return null;
}

```

Simple enough. If you're not familiar with JavaScript, the **Array.prototype** declarations here are simply extending the **Array** type, so that we can use the head and tail functions directly.

*Note: if you're a JavaScript person you're rightfully scratching your head thinking “why is he reimplementing **Array.forEach**?” A perfectly lovely question, with a simple answer (hopefully): I want to show recursion explicitly. You should know how to do this.*

Now that we have the notion of head and tail, we can recurse into our list easily:

```

var recurseInto = function(list, fn){
  var h = list.head();
  var t = list.tail();

  if(fn && h) fn(h);
  if(t) recurseInto(t, fn);
}

```

Here I'm grabbing the head and tail values. If there's a head and a callback function, I'll call it, passing along the head. If there's a tail I'll call `recurseInto` one more time, passing along the tail and the callback. This is straight up functional programming! Now all I need to do is to call it directly instead of nesting many for loops:

```

var sieveTwo = function(n){
  var list = [];

  //load the list, defaulting to "true" as unmarked
  for(var i = 2; i <= n; i++){
    list[i]={index : i, value: true};
  }
  //figure out our multiplier limit
  var limit = Math.sqrt(n);

  for(var i = 2; i <= limit; i++){
    recurseInto(list, function(val){
      //is it eligible?
      if(val.index > i && val.value){
        //is it a multiple of i?
        if(val.index % i === 0) val.value = false;
      }
    });
  }

  //build an output
  var out = [];
  recurseInto(list, function(val){
    if(val.value) out.push(val.index);
  });
  return out;
}

console.log(sieveTwo(100));

```

To get all of this to work properly I had to use an object to hold the index as well as the boolean flag value. I have to do it this way since I'm working with recursion without an iteration index. Aside from that, the results are the same.

Give it a whirl! See which one you like better.

Was This Dynamic Programming?

What do you think? Consider:

- Was this an optimization problem?
- Did I break the problem into subproblems?
- Did those subproblems have an optimal substructure?
- Did we memoize the subproblems to speed things up?

We did break things down into subproblems and you could mistake our initial list for a memoized table ... but no, **this is not dynamic programming**.

Why not?

In short, a sieve problem like this is just working over a set of numbers, filtering out (sequentially) values that we don't need for the result. We're not collecting a list of solutions to subproblems to solve a greater objective problem. We did use recursion, but that's about it.

Finally, while our initial list might look like a memoized table, it is not. Memoization is the process of writing down and remembering a solution. A pre-filled list of integers is not that.

Bellman-Ford

One of the most fascinating uses of graphs is in the optimization of path traversal, which can be used in a vast number of calculations.

As mentioned in the previous chapter, graphs can be used to represent all kinds of information:

- A network of any kind. Social (friends) or digital (computers or the internet), for example
- A decision tree
- Contributions from members of any kind to a cause of any kind
- Atomic interactions in physics, chemistry or biology
- Navigation between various endpoints
- If you apply weighting to the edges or vertices, you can run useful calculations for just about anything. One of the most common is finding the shortest path between two vertices.

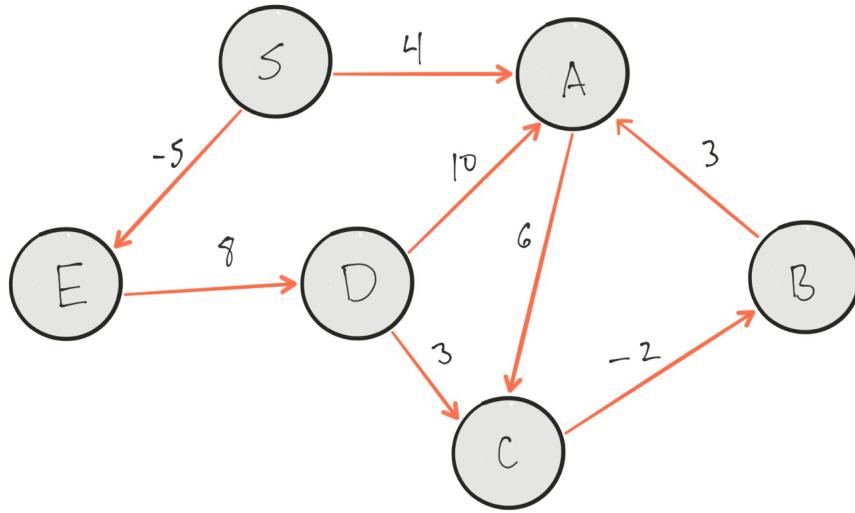
There are numerous algorithms to touch on at this point, but I have to round this chapter out by discussing the two you should be aware of: Bellman-Ford and Dijkstra. In this section we'll discuss Bellman-Ford; Dijkstra comes next.

The Problem Graph

This algorithm is named after Richard Bellman (the same person who wrote about dynamic programming) and Lester Ford Jr. It shares a lot with Dijkstra's algorithm (which we'll see next), but has one major advantage: it can

accommodate negative edges.

Let's see how it works. Consider this graph:



This is an edge-weighted, directed graph. We want to calculate the shortest paths between our source vertex S and the rest of the vertices, A through E. We can do this using dynamic programming.

Now, if you were to stare at this and I told you what your task was (to calculate the smallest cost between S and the rest of the vertices), you would probably get overwhelmed! I know I did.

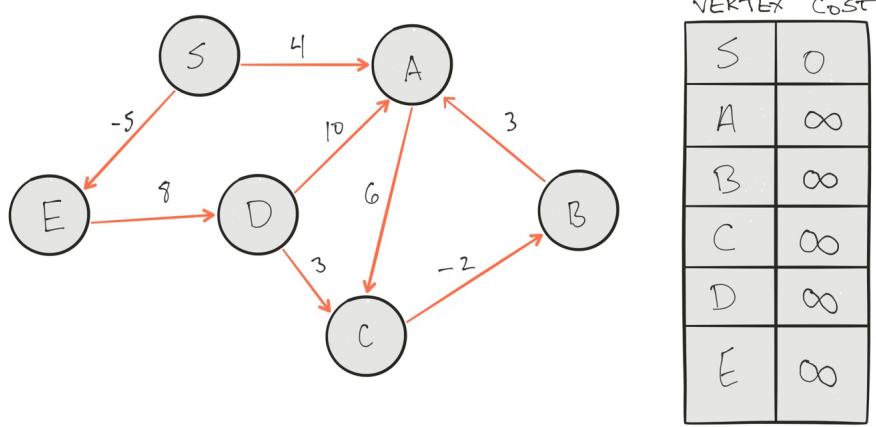
The good news is that we have the Bellman-Ford algorithm, which makes this calculation rather fun. Rather than spend many words on it, let's do it visually, then implement it with some code.

We will be using memoization to keep our calculations in P time, so let's setup a memoization table that we can update as we run our calculations. We know that there are only 6 vertices, so we only need to track, at most, 6 total costs. The only cost we know is the base cost: $\delta(S, S) = 0$

Note: if you're not familiar with that squiggle, δ , it's a "delta", which indicates a difference. This equation states that the difference (or "cost") from S to S is 0.

For the rest of the vertices, we don't know so we'll set them all to infinity. Why infinity? Simply because the calculation we do later will be based on finding the smaller value, and infinity guarantees that the initial state will not remain.

Here's our setup:



Great. The plan is to calculate the cost of each outgoing edge, for each vertex in our graph. This will be 6 total calculations which we'll repeat in iterations. The number of iterations i you need to perform when using Bellman-Ford is: $i=|V|-1$.

Why is this? You'll see in a second, but the crux of it is that we're calculating the distances between all nodes and remembering the smallest ones. This is called relaxation: we start with the largest values we can think of (infinity) and then slowly relax the costs through an iterative calculation.

OK, enough words, let's do this.

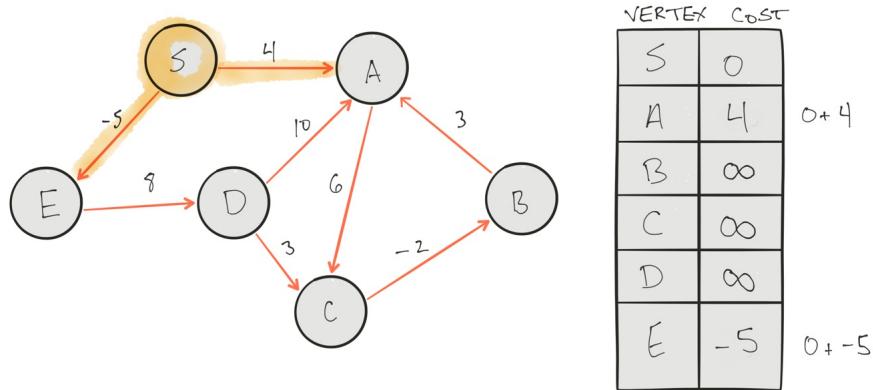
Iteration 1

We'll start with our source, S, and move clockwise around the graph. From S we have two outgoing edges with the values 4 and -5. These costs (C_{sa} and C_{se}) associate S with A and E so we'll add them to our table using this calculation:

$$C_{sa} = \delta(S, S) + \delta(S, A)$$

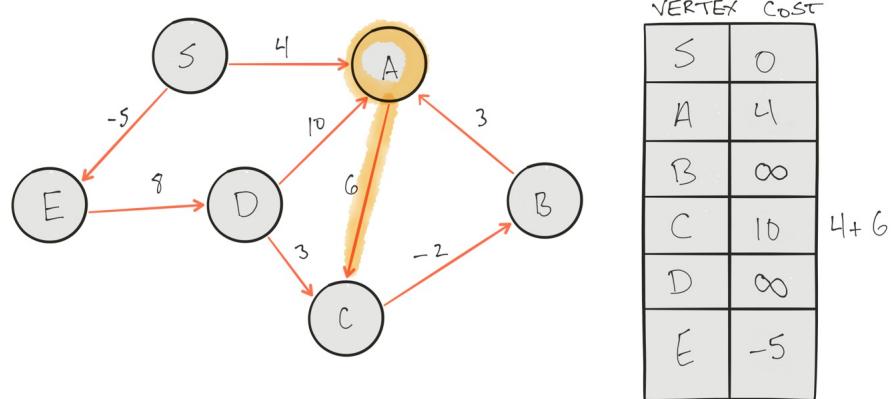
$$C_{se} = \delta(S, S) + \delta(S, E)$$

You can visualize this on the graph itself. We're using the cost of the highlighted edges below and adding them to the initial cost of S to S (which is 0):



If this seems complicated, let it wash over you. As we move through this it will make more sense.

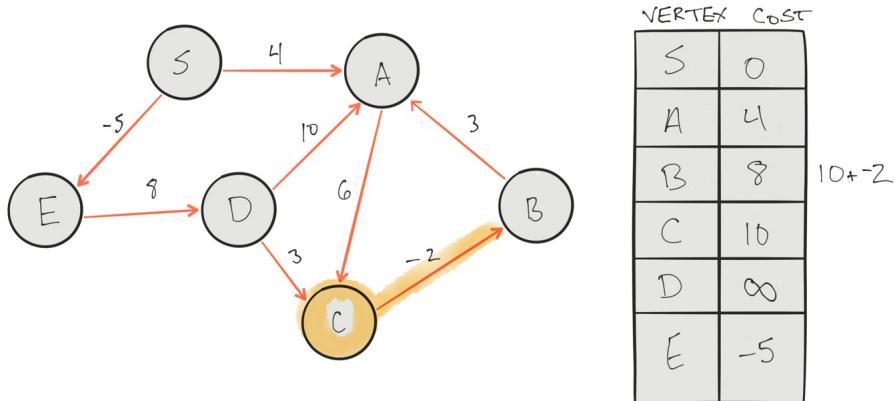
OK, we're still in our first iteration of calculating the costs between S and the rest of the vertices. Let's now move to A and calculate the outgoing edges from A:



A only has a single outgoing edge with a cost of 6, so we take that cost and add it to A's current cost (which is 4). This gives us a value of 10, which we add to our memo table.

Now we move on to B, which has a current cost of infinity. This means that we don't have a path from S to B yet, which means we can't calculate it in this iteration. So we skip it.

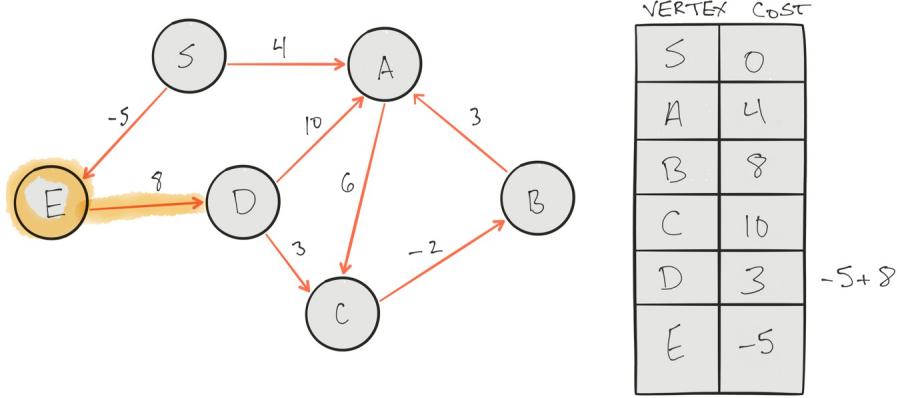
Next up is C, which does have a current cost of 10. It also has an outgoing edge to B which is good news as we'll be able to use that in the next iteration. For now, let's update our table:



The current cost of the path to C is 10, so adding -2 to that (which is the cost of the edge between C and B) gives B a current cost of 8.

Now we're up to D. What do you think we do here? The current cost is set to infinity, which means we haven't calculated a path to it yet, so we skip it.

Finally we close off this iteration by calculating the outgoing edges from E:

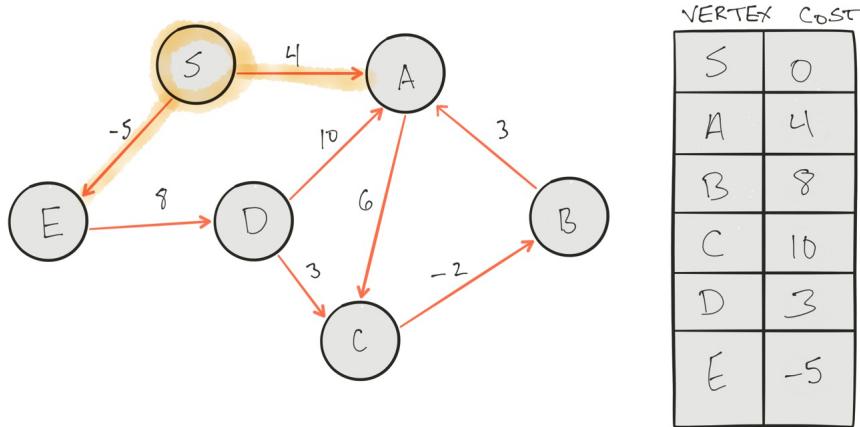


By now you should be able to reason what will happen. We know that E has a current cost of -5, so we add that to the cost of the edge between E and D, which is 8. This gives a cost from S to D equal to 3.

Great! We now have costs for all the vertices! It's time to do another iteration to see if we can relax these values a bit more.

Iteration 2

Let's step through this a bit quicker this time. We'll start with S again, our source, and note that the cost of the outgoing edges does not improve on the costs we've recorded (4 and -5). This makes sense as these are single edges and there really is no way to improve the costs here:



If we move to A, again the only outgoing edge we can evaluate is A to C, which is 6. The cost from S to C remains the same at 10, so there's no improvement here and we can move on.

Next up is B, which now has a cost associated with it so we can use it in our calculations. The current cost to get to B is 8, so getting from B to A is 11 total, which is not an improvement over A's current cost of 4.

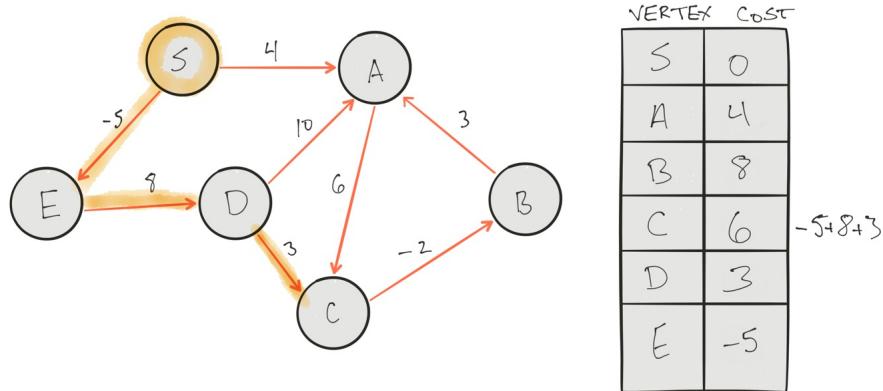
Then we come to C, which is 10. We've already calculated C to B as 8 and there's no improvement here, so we can move on again.

Seems a bit boring, doesn't it? We keep skipping things – but that's OK! It's about to get a bit more exciting

when we consider D, our next vertex.

The current value of D is 3 and has two outgoing edges to A and C. The value of the edge from D to A is 10, so adding the current cost of D (3) to this edge cost of 10 would be 13. This doesn't reduce A's cost so we leave it. But what about edge D to C?

This produces a cost of 6, which means we can lower C's cost in our memo table to 6:



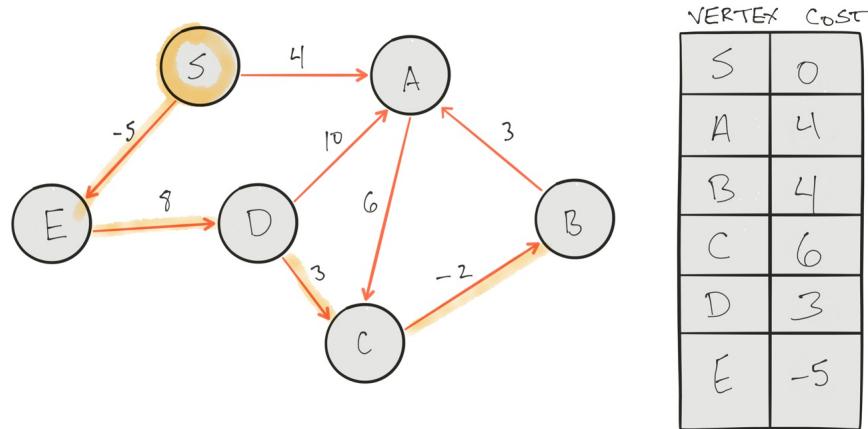
The excitement! Can you feel it! Now, as you might imagine, changing C like this means we'll be able to change more values in our table – but that will have to wait for the next iteration.

Speaking of, we're back to E now and there are no changes we can make here, so we'll skip on to the next iteration.

Iteration 3

We skipped S and A last time and that's still the case this time: there are no changes we can make to improve our current costs for these vertices. Same for boring old B, again. We can't improve the cost of A, which is its only outgoing edge.

C, however, makes things interesting. The current cost of C is now 6, which means we can reduce the cost of B to 4:



Oh my that's so exciting isn't it! From this point we go to D and E just like before, skipping them as there's no way to improve their costs.

We can also use some more reasoning here. Since B was the only change we made from the last iteration, we can just evaluate its outgoing edge to see if it will improve A. Since B's current cost is 4 and A's current cost is 4, then no improvement will happen if we add the cost of the edge from B to A, which is 3.

So we're done! The shortest paths for each vertex are now calculated.

Implementing With JavaScript

I'm sure you'll find some ways to improve this code, which is great! I've made it a little more verbose for clarity, which I think will help if you're still having problems getting your head around this algorithm.

```
//define the vertices - these can just be string values
var vertices = ["S", "A", "B", "C", "D", "E"];
```

Next up is our memoization table. This can be a simple JavaScript object:

```
//our memoization table, which I'll set to an object
//defaulting as described
var memo = {
  S : 0,
  A : Number.POSITIVE_INFINITY,
  B : Number.POSITIVE_INFINITY,
  C : Number.POSITIVE_INFINITY,
  D : Number.POSITIVE_INFINITY,
  E : Number.POSITIVE_INFINITY
}
```

Now we need to define the graph itself. For this I simply need to track which vertices are involved and the costs associated with the relationship:

```
//this is our graph, relationships between vertices
//with costs associated
var graph = [
  {from : "S", to : "A", cost: 4},
  {from : "S", to : "E", cost: -5},
  {from : "A", to : "C", cost: 6},
  {from : "B", to : "A", cost: 3},
  {from : "C", to : "B", cost: -2},
  {from : "D", to : "C", cost: 3},
  {from : "D", to : "A", cost: 10},
  {from : "E", to : "D", cost: 8}
];
```

Looking good! At this point you should be able to reason how we'll use these data structures to iterate over our graph. In short, we need to:

- Iterate over the vertices array
- Using the current vertex from our vertices array, we select the outgoing edges from the graph array.
- Once we have the outgoing edges, we run a quick calculation using our memo object. The calculation is straightforward: we take the cost of the current vertex and add it to the cost of the current outgoing edge. If that value is less than the cost of the current edge, we update the memo for the current edge.

Translating that word salad to code:

```
//represents a full iteration of Bellman-Ford on our graph
var iterate = function(){
    //the return value
    var doItAgain = false;
    //check each vertex
    for(fromVertex of vertices) {

        //get the outgoing edges for this vertex
        var edges = graph.filter(function(path){
            return path.from === fromVertex;
        });

        //iterate over the edges and set the costs
        for(edge of edges){
            //calculate the cost of the outgoing edge
            const potentialCost = memo[edge.from] + edge.cost;
            //if it's less than what's in our memo table for the
            //target vertex..
            if(potentialCost < memo[edge.to]){
                //update the memo table
                memo[edge.to] = potentialCost;
                //since we made an update, we'll want to
                //iterate again
                doItAgain = true;
            }
        });
    };
    return doItAgain;
};
```

Great! Now all we need to do is to iterate over our iterate function and we're good to go:

```
for(vertex of vertices){
    if(!iterate()) break;
}
console.log(memo);
```

You'll notice that I'm only iterating to `vertices.length - 1`. Can you reason why? If you run this code (using Node), you should see:

```
{ S: 0, A: 4, B: 4, C: 6, D: 3, E: -5 }
```

These are the exact values of our exercise above. Nicely done!

Analysis and Summary

This is dynamic programming in action. Dividing the objective problem (finding the shortest paths) into smaller problems (calculating the costs between each vertex). We then recursed over the smaller problems (the `iterate` function) and used memoization to derive the answer.

There is room for improvement, however. Think about the process we went through in the first section. I only needed 3 total iterations to derive the shortest paths. The code I wrote, however, required 5. How would you optimize this?

In addition, I'm not using recursion in a code sense. I am calling the same function repeatedly, but there's probably a tighter, cleaner way to do this using a true recursive function. Can you see how?

Finally, here's something to ponder: does the order in which we evaluate the vertices matter? If yes, why? If no ... why not? Rearrange the code a bit and see if you come up with a different answer than you see here.

The Bellman-Ford algorithm is quite effective, as we can see, but it can also take a long time to run. In terms of complexity, the algorithm runs in $O(V * E)$ time, where V is the number of vertices and E is the number of total edges. It can work well for simple graphs, but for more complex (or dense) graphs, it is not the most efficient algorithm.

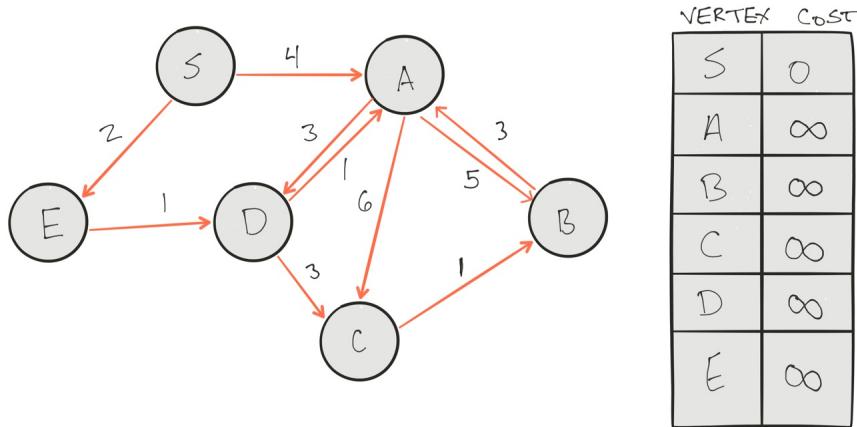
Djikstra

In the last section we iterated over a simple graph using Bellman-Ford to find the shortest paths from a single vertex (our source) to all other vertices in the graph.

The complexity of Bellman-Ford is $O(|V|E)$, which can approximate $O(n^2)$ if every vertex has at least one outgoing edge. In other words: *it's not terribly efficient*.

Dijkstra's algorithm requires only one iteration, however and has a complexity of $O(|V| \log V)$, which is much more efficient. Let's see why.

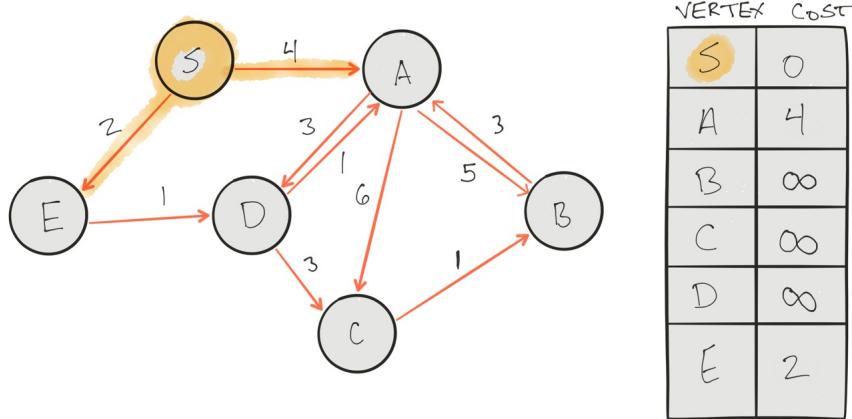
As with Bellman-Ford, we'll use a directed, weighted graph with 6 vertices. In addition, we'll setup a memo table with our source S set to 0 and the rest of the vertices set to infinity:



There is a difference here, however, and it's critical! Dijkstra doesn't work with negative edge weights! I have adjusted this graph so that we don't have any negative weights, as you can see. Specifically the edges between S and E as well as C to B. In addition I've added a few edges to show that the algorithm will scale easily regardless of the number of edges involved.

Starting At The Source

The first step is to evaluate the source, S. We do the same thing as before, with Bellman-Ford, where we tally up the cost of the outgoing edges and store them in the memo table. In addition, we'll mark S as complete:



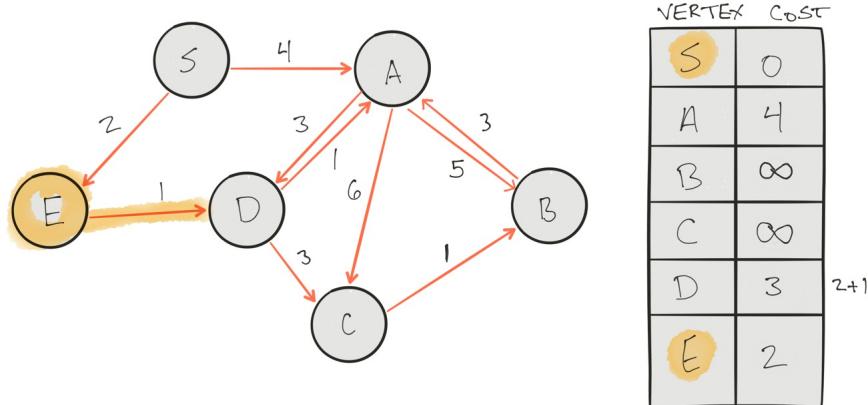
This is another way that Dijkstra differs from Bellman-Ford: each vertex is visited only once.

The next vertex is chosen using these rules:

- It must not have been visited previously
- It has the smallest cost of the remaining unvisited vertices In our case, this would be vertex E.

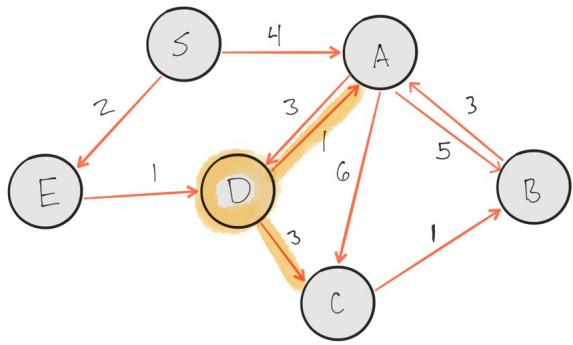
Traversing Each Vertex

The next vertex that we'll visit is E. The cost to reach E is 2, which is less than 4, and E has not been visited previously:



We then update the memo table, setting D to 3 (which is the cost of S to E plus the cost of E to D). We then mark E as visited.

The next vertex in our table that is unvisited with the lowest cost is D, so we calculate that next:



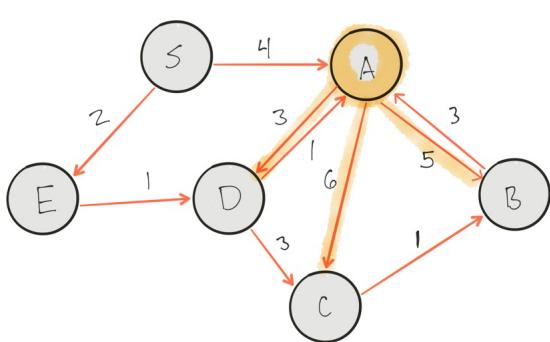
VERTEX	COST
S	0
A	4
B	∞
C	6
D	3
E	2

$3+1$

$3+3$

D has two outgoing edges: to A and to C. The cost of A would be $3+1$ which is 4 and no improvement, so we leave A as it is in our table. C has no cost, so we update it to 6, which is the current value of D (3) + the cost to get to C.

The next vertex we'll visit is A. It has the least cost and has not been visited before:



VERTEX	COST
S	0
A	4
B	9
C	6
D	3
E	2

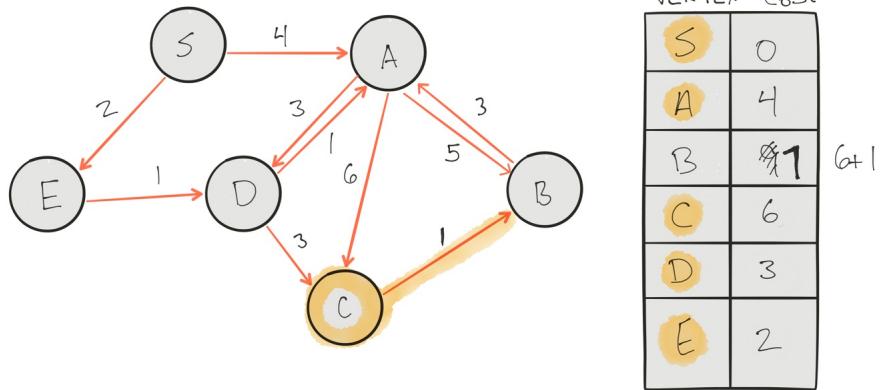
$4+5$

~~$4+6$~~

~~$4+3$~~

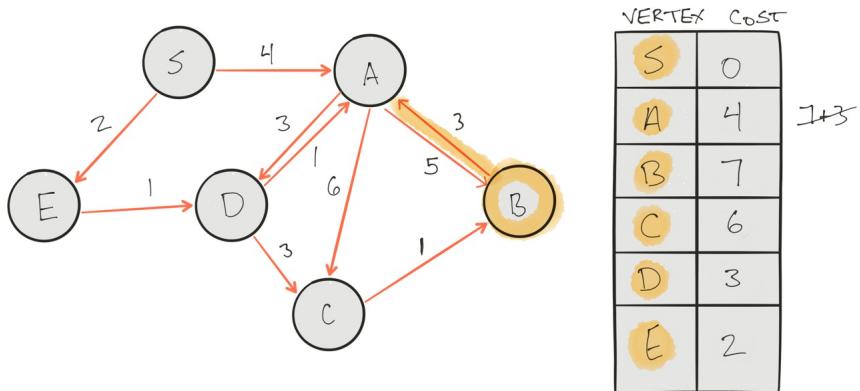
The current cost of A is 4 and the edges going out of A go to vertices D, C and B. The cost of reaching D through A is 7, which is not an improvement of D's current cost of 3, so we leave it. Same with C: reaching C through A does not reduce C's cost, so we leave it. B, however, is still infinity so we'll set it to 9, which is A's cost plus the cost to reach B, which is 5.

C is the next vertex we'll choose as its current cost is 6 and it hasn't been visited:



C has only one outgoing edge: to B. The current cost of C is 6 and adding 1 to it would be less than the current cost of B, so we'll update B's cost to 7.

There is only one vertex left, which is B:



The only candidate for improvement is A, with a current cost of 4. This is less than the tentative cost of B to A, which is 10, so we leave it as it is. We're done!

All of our vertices are visited, so we're done! We can be sure we calculated the shortest path by using our Bellman-Ford code from the last section to make sure it matches, and it does!

Implementing With JavaScript

The code you're about to see is probably a lot more verbose than you might write it. I like clarity, mostly for my own sake, because you can bet I'll be looking over this page and this code quite a few times in the future! I want to remember what I was thinking.

To start with, let's alter our implementation of Bellman-Ford in the last section to have a memo table with a little more smarts:

```

class MemoTable{
    //the constructor for our ES6 class
    constructor(vertices){
        //the source vertex
        this.S={name: "S", cost: 0, visited: false};
        this.table = [this.S];
        //loop over the passed-in vertices and
        //initialize
        for(var vertex of vertices){
            this.table.push({name: vertex, cost: Number.POSITIVE_INFINITY, visited: false});
        }
    }
    //returns all unvisited vertices
    getCandidateVertices(){
        return this.table.filter(entry => {
            return entry.visited === false;
        });
    };
    //returns the next unvisited vertex with least cost
    nextVertex(){
        const candidates = this.getCandidateVertices();
        //return the lowest
        if(candidates.length > 0){
            return candidates.reduce((prev, curr) => {
                return prev.cost < curr.cost ? prev : curr;
            });
        }else{
            return null;
        }
    };
    //sets the cost of a given vertex in our table
    setCurrentCost(vertex, cost){
        this.getEntry(vertex).cost =cost;
    };
    //marks the vertex as visited
    setAsVisited(vertex){
        this.getEntry(vertex).visited = true;
    }
    //returns a single entry in our table
    getEntry(vertex){
        return this.table.find(entry => entry.name === vertex);
    }
    //returns the cost associated with reaching
    //a given vertex from the source
    getCost(vertex){
        return this.getEntry(vertex).cost;
    }
    //for display purposes
    toString(){
        console.log(this.table);
    }
}
const vertices = ["A", "B", "C", "D", "E"];
var memo = new MemoTable(vertices);

```

I've added the logic for retrieving a given entry as well as updating its values. In addition I've added logic for determining the next vertex to traverse to based on the rules of Dijkstra that we saw above.

Next, we have our graph:

```
//this is our graph, relationships between vertices
//with costs associated
const graph = [
  {from : "S", to :"A", cost: 4},
  {from : "S", to :"E", cost: 2},
  {from : "A", to :"D", cost: 3},
  {from : "A", to :"C", cost: 6},
  {from : "A", to :"B", cost: 5},
  {from : "B", to :"A", cost: 3},
  {from : "C", to :"B", cost: 1},
  {from : "D", to :"C", cost: 3},
  {from : "D", to :"A", cost: 1},
  {from : "E", to :"D", cost: 1}
];

```

This array represents the visual graph we worked with above. Now we just need to evaluate our graph using our **MemoTable** functionality:

```
const evaluate = (vertex) => {

  //get the outgoing edges for this vertex
  const edges = graph.filter(path => {
    return path.from === vertex.name;
  });

  //iterate over the edges and set the costs
  for(var edge of edges){
    //the cost of the edge under evaluation
    const currentVertexCost = memo.getCost(edge.from);
    const toVertexCost = memo.getCost(edge.to);

    //the proposed cost from S to the current vertex
    const tentativeCost = currentVertexCost + edge.cost;

    //if it's less, update the memo table
    if(tentativeCost < toVertexCost){
      memo.setCurrentCost(edge.to, tentativeCost);
    }
  };

  //this vertex has been visited
  memo.setAsVisited(vertex.name);

  //get the next vertex
  var next = memo.nextVertex();

  //recurse
  if(next) evaluate(next);
}

//kick it off!
evaluate(memo.S);
memo.toString();
```

The code looks a bit familiar – it follows what I did (for the most part) with Bellman-Ford but this time I've added a few tweaks to accommodate setting a vertex as visited, and I'm also using recursion off the memoization table instead of a loop.

What do you think? Can you improve this without losing its clarity? Let's run it and make sure it works. Again, using Node:

```
[ { name: 'S', cost: 0, visited: true },
  { name: 'A', cost: 4, visited: true },
  { name: 'B', cost: 7, visited: true },
  { name: 'C', cost: 6, visited: true },
  { name: 'D', cost: 3, visited: true },
  { name: 'E', cost: 2, visited: true } ]
```

Right on! That's the exact answer we got above!

CHAPTER NINE: COMPILATION

Compilation is a set of five discrete steps: lexical analysis, parsing, semantic analysis, optimization and code generation. Every compiler does this exact same thing, but some are more capable than others. The main differences between compilers is in the optimization and code generation steps.

A stack is a way to store program data and is a *per thread process*. The stack is very fast and is used for *value types*. The heap is a *per application process* and is used for *reference types*. You can make your application more efficient by understanding how these two interact.

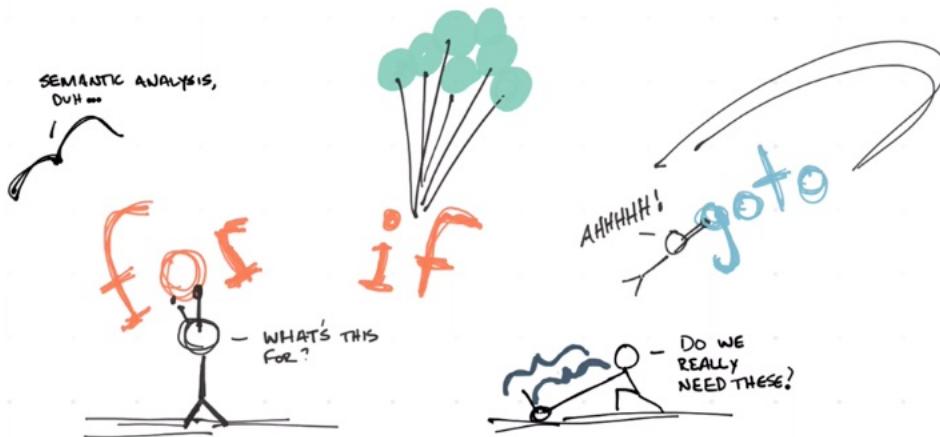
Garbage collection is the process of freeing up application memory (the heap) automatically. This is not a free process and requires system resources. There are a number of strategies used for garbage collection (GC), but tracing is the most widely used.



How a Compiler Works

A compiler simply brings different things together and makes them one thing. In a programming sense, a compiler is a program that reads the code you write and generates something that the runtime engine can execute.

Compilation can happen on command – for instance using a make file. It can happen just in time (JIT) or at runtime with a thing called an interpreter.



They can compile code to different languages (CoffeeScript to JavaScript, e.g) or down to byte code ... or something in between. C#, for instance, compiles to an intermediate language (MSIL), which is then compiled to native code upon execution. This has the advantage of portability – meaning that you could create a compiler for different platforms (Windows 32-bit, 64-bit, Mac, Linux, etc) without having to change the code.

But how do these things work? Let's take a look at a high level.

The Compilation Process

Compilation in a computer is just like compilation in your head when you read these words. You're taking them

in through your eyes, separating them using punctuation and white space, and basing the meaning of those words on emphasis.

These words are then turned into meaning in your mind, and at some point sink into your memory ... causing (hopefully) some type of action on your part.

A compiler does the same things, but with slightly larger words. The main compilation steps are:

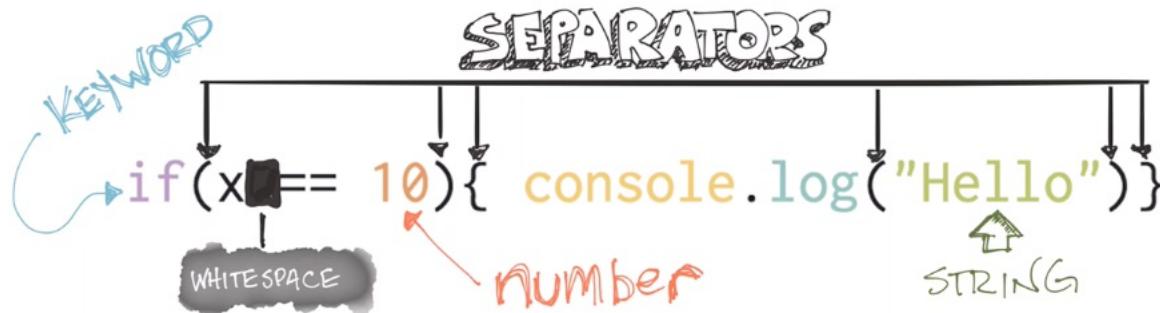
- Lexical Analysis
- Parsing
- Semantic Analysis
- Optimization
- Code Generation

Every compiler goes through these steps.

Lexical Analysis

Lexical Analysis simply analyzes the code that's passed to the compiler and splits it into tokens. When you read this sentence, you use the whitespace and punctuation between the words to "tokenize" the sentence into words, which you then label.

A compiler will scan a string of code and separate it into tokens as well, labeling the individual elements along the way:



The program within the compiler that does this is called the lexer. So, using our code sample, the lexer will generate these pseudo code tokens (using tuples):

```
{keyword, "if"}  
{paren, }  
{variable, "x"}  
{operator, "=="}  
{number, "10"}  
{left-brace, }  
...
```

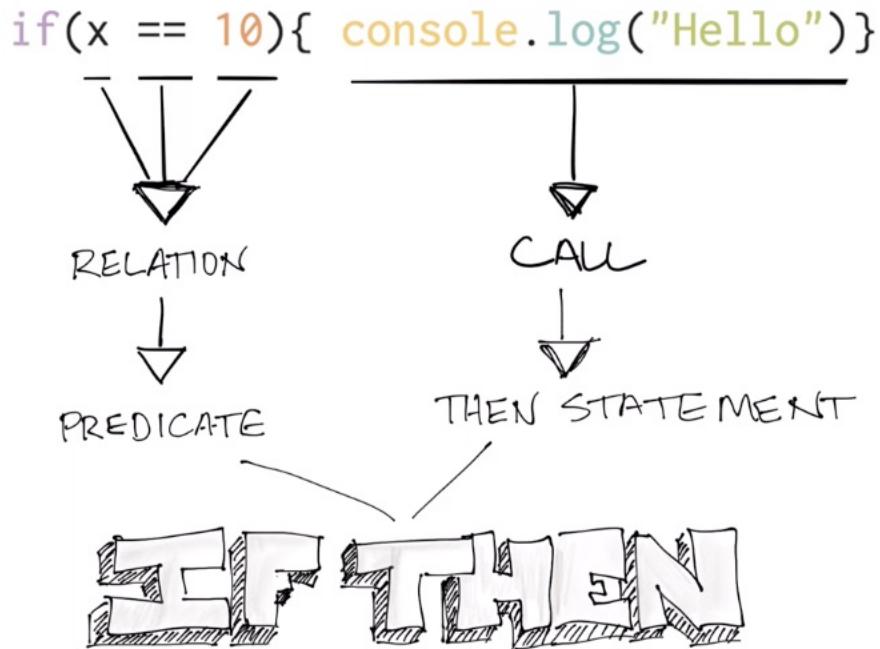
The thing being analyzed by the lexer is the lexeme. A token is what's produced from this analysis. Yay for more random words to know!

Parsing

After the lexer has tokenized the incoming code string, the parser takes over, applying the tokens to the rules of the language – also known as a *formalized grammar*.

Pushing this back into the realm of written language: the words you're reading now are tokenized using whitespace and punctuation – the next process is to parse their meaning and, essentially, what they're supposed to mean.

A parser analyzes the tokens and figures out the implicit structure:



We know, at this point, that we have an if statement. This is just one line of code, of course, however our entire codebase would be parsed in exactly this way: turning tokens into language structures.

Semantic Analysis

Now we get to the fun part. Semantic Analysis is where the compiler tries to figure out what you're trying to do. We have an if block for our code, but can the predicates be resolved to truthy/falsey expressions? Do we have more than one else associated with it? Consider this sentence:

Kim and Jolene want to go to her party

We can reason that “Kim and Jolene” are the subjects, “to go” is the verb and “party” is the indirect object. But who is her? When a compiler goes through semantic analysis, it has to reason through the same thing. For instance:

```
var x = 12;
var squareIt = function(x){
    return x * x;
};
y = squareIt(x);
console.log(y);
```

What will y evaluate to? If you’re thinking 144 – you’d be right. Even though we’ve reused x here, the

JavaScript interpreter figured out what we meant, even though we tried to confuse it.

```
var x;
console.log(x);
x = "Hello!";
```

If you're Brendan Eich (creator of JavaScript) and it's 1995 and you have [10 days to create a language](#) – what choices do you make when it comes to semantic analysis?

I'm sure many of you are thinking “hang on – JavaScript is not a compiled language – it's interpreted” and you'd be correct. They still go through the same steps.

I bring up JavaScript because it's precisely these decisions, made at the semantic analysis level, that have caused developers so much confusion over the years. If you've used the language, you'll know what I mean.

Lexical Scoping

Most object-oriented languages are lexically scoped, which means the scope of a variable extends to the block of code that contains it. So, in C#, for instance:

```
public class MyClass {
    public MyClass(){
        var x = 100;
    }
}
```

The scope of **x** in this case is only within the constructor. You cannot access **x** outside of it.

Let's change the code so we can rely on lexical scoping to make **x** available to all properties and methods in **MyClass**:

```
public class MyClass {
    int x;
    public MyClass(){
        x = 100;
    }
}
```

All we needed to do was to declare the variable within the **MyClass** code block.

JavaScript, however, does things differently. Scopes in JavaScript are defined by function blocks. So, strictly speaking, JavaScript is sort of lexically scoped.

Consider this code:

```
if(1 === 1){
    var x = 12;
}
console.log(x); //12
```

If lexical scoping was enforced here we should see **undefined**. But what happens in C#? Let's try it:

```
public class ScopeTest
{
    [Fact]
    public void TheScopeOfXIsLexical()
    {
        if(1 == 1){
            var x = 12;
        }
        Console.WriteLine(x);
    }
}
```

If I run this test, I get the expected response:

```
error CS0103: The name 'x' does not exist in the current context
```

The reason for the difference? A different semantic analysis for the C# compiler vs. the JavaScript interpreter.

There are more issues that I'm sure you're aware of: hoisting, default global scope for variables with **var**, confusion about **this**. There are many things written about these behaviors and I don't need to go into them here. I bring all of "this" up simply to note the choices made by semantic analysis.

Optimization

Once the compiler understands the code structures that are put in place, it's time to optimize. Of all the steps in the compilation process, this is typically the longest.

There are almost limitless optimizations that can occur; little tweaks to make your code smaller, faster, and use less memory and power (which is important if you're writing code for phones).

Let's optimize that previous sentence: compiler optimization produces faster and more efficient code. Reads the same doesn't it? The same meaning, anyway – compilers don't care about creative expression.

Which is a very important point. Modern languages are leaning more on syntactic niceties and shortcuts (aka "*syntactic sugar*"). This is great for developers like me, who enjoy reading code where the intent is clear. It's not so great for the optimizer.

Ruby and Elixir are prime examples of this. A number of constructs in these languages are optional (parentheses for example), and a compiler must work through various syntactic shortcuts to figure out the instructions. **This takes time**.

Does a `meaningful_variable_name` need to be 24 characters long. Not to the compiler, which will often rename the variable to something shorter. List operations as well – often you'll see arrays substituted for you, behind the scenes.

How about this rule – does this rule make sense?

X=Y*0::X=0

Basically, any time you see a number multiplied by 0, replace it with 0. Seems to make sense ... but ...

```
public class MultiplyingByZero
{
    [Fact]
    public void UsingNaN()
    {
        double x = Double.NaN;
        Console.WriteLine(x * 0); // NaN
    }
}
```

Yeah that won't work because **NaN** $\neq 0$. It will work if **x** is an integer, however.

Compiler optimizations are at the center of “what makes a language good” – something we’ll get into more, later on.

Code Gen

Our code has been parsed, analyzed, and optimized – we’re ready for output! This is one of the simpler steps: *we just need to package it up and off we go*.

Intermediate Language Output

A number of platforms out there will compile code text files into an intermediate structure, which will be further compiled later on. Code that runs in a virtual machine (like Java) will do this, and later compile it down to machine-level code.

C# does this as well, and compiles into an actual second language called IL (or MSIL). People don’t typically read or write IL directly (unless you’re [Jon Skeet](#)) – but it is possible to dig down into it to see what’s going on.

When a C# program is run for the first time, the IL is recompiled down to native machine code, which runs quite fast. It’s also JIT compiled every time it’s run, after that. The term JIT stands for *Just In Time* – which usually means “last possible moment”

Getting The JITters

When discussing multistep compilation, you’ll often hear other developers talk about “the JITter” and what it will do to your code.

In some cases it will output code optimized for an executable that will then run it. Other times it will produce native byte code that runs at the processor level.

Interpreters

Dynamic languages, like Ruby, JavaScript, or Python, typically use a different approach. I say typically because multiple runtime engines have been created that will produce byte code for you from each of these languages. More on that in a later section.

You can think of these languages as “scripting languages” – in other words, they are scripts that are executed by a runtime engine every time they are invoked.

Ruby, for example, has the MRI (Matz’s Ruby Interpreter) which will read in a Ruby file, do all the things discussed above, then execute things as needed. You have to go through the full set of compiler steps *every time a routine needs to be executed*. This is not as fast as native code, obviously, which has already been analyzed, parsed and optimized.

JavaScript is interpreted on the fly in the same way as earlier versions of Ruby, depending on where you use it. The browser will load in and compile any script files referenced on a page, and will hold the compiled code in memory unless/until you reload the browser.

Node works in the same way: it will compile your source files and hold them in memory until you restart the

Node runtime.

As mentioned, Ruby and Python have alternative runtimes that can be used. We'll discuss those in the next chapter.

LLVM

One of the biggest names in the compilation space (if not the biggest) is LLVM:

The LLVM Project is a collection of modular and reusable compiler and toolchain technologies. Despite its name, LLVM has little to do with traditional virtual machines, though it does provide helpful libraries that can be used to build them. The name “LLVM” itself is not an acronym; it is the full name of the project.

So you can use LLVM “stuff” to build your own compiler. One that you've probably heard of is CLANG which is an LLVM-based compiler for C.

Native Gem/Module Pain

You know when you go to install a Node module or Ruby gem and you get some nasty message about native bindings or build tools required? The reason for this is that some of these modules use C libraries that need to be compiled to run. You see this a lot with database drivers, for instance, which want to be super fast.

With Node, you'll usually see a reference to node-gyp rebuild, which is a module specifically created for compiling native modules.

This can cause headaches, especially when you have other developers working on Windows. The workaround, typically, is to install Visual Studio, referencing C++ build tools during the install. Installing these things on a Windows Server is one of the most frustrating things about using Node on Windows in production.

GCC

The GCC project is from GNU:

The GNU Compiler Collection includes front ends for C, C++, Objective-C, Fortran, Java, Ada, and Go, as well as libraries for these languages (libstdc++, libgcj,...). GCC was originally written as the compiler for the GNU operating system.

GCC and LLVM do the same things... sort of. They're both compiler toolchains that also provide compilers for C, C++, Java and Go (among others).

The main differences have been speed (with GCC typically being a bit faster) and licensing. GCC uses the GPL and LLVM uses a license based on MIT/BSD which is a little more permissive.

Apple moved from the GCC to LLVM back in 2010/2011 and it caused a lot of code to break, especially for Rails developers. Those darn native gems!

Stack Vs. Heap

I remember reading about Erlang a number of years ago, and the author started discussing memory usage – specifically how the Erlang VM supports individual processes with private heaps. He went on to explain how garbage collection speeds were amazing and oh wow isn't that neat.

I had no idea what he was talking about. And, to be honest, I'm still a bit hazy. Stack vs. heap discussions come

and go in my life, and every time the subject comes up I can feel certain parts of my memory slowly grind to life. It takes a minute or two to get the basics sorted... and honestly I really should just know this stuff cold.

Understanding how runtime engines use memory to run your code is critical. Trusting the garbage collection process is, generally speaking, not always a good idea.

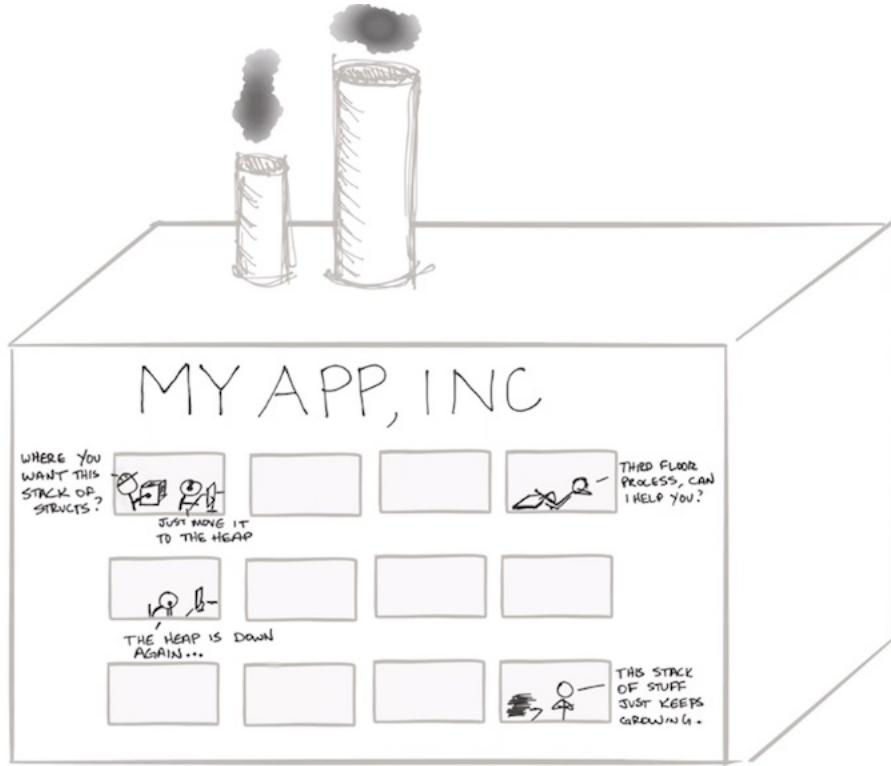
It's likely you already know this, so bear with me. Writing things down (and drawing pictures) cements things in my head. If you feel like you know this cold, however, skip ahead to garbage collection.

To get right to it:

- A stack is created **per thread or process** – whatever construct is executing your code.
- A heap is created **per application**

The stack is faster than the heap for storage and access mainly because of its structure (last-in, first-out linear storage), whereas the structure of the heap is a bit more complex.

Memory that lives on the heap requires careful management so it doesn't bloat and cause problems. Memory allocated on the stack is freed up when it's no longer needed.



The Stack

So, think back to a Pushdown Machine and what made it different from a Finite State Machine. It was a thing called a stack. The stack was simply a place to put values that can be used during state transitions. It's the same deal for your code.

A stack is created per thread or process and is Last In, First Out or LIFO. This is a very important idea to understand because it will affect how you write your code in the future.

Imagine you're walking down the street with a friend who's on his way to an interview. He's late so he's walking briskly down the street and you've decided to tag along to help out and offer moral support (you're a good friend!)

He's trying to finish tying his tie when he passes a bagel cart and realizes he didn't eat breakfast – so he buys it and hands it to you so he can finish tying his tie. When he's done, he asks for the bagel back and eats it.

He then grabs a cup of coffee, some napkins, and a doughnut and, again, hands them to you while he tucks in his shirt (the doughnut, the napkin, then the coffee). He **eatsTheDonut()** and thanks you for not taking a bite (I totally would have), grabs the napkin from you and **wipesHisMouth()** and hands it back to you, and finally gets the coffee and **sipsItCasually()**. When he's done he'll probably **wipesHisMouth()** again so you need to hold onto that napkin for him.

In this scenario: you're the stack, holding things for your friend, the thread. Since the thread can only do one thing at a time a LIFO stack system works perfectly. When values on the stack (in your hands) go out of scope (are eaten or thrown away) then they're just gone. There's nothing that you or the thread need to do.

The values on the stack are held in static memory, which means they reside in a given memory location (your arms) and are commonly referred to as value types – meaning they don't represent something, they are something.

The heap, however, uses a pointer to dynamic memory – something that can grow or shrink based on garbage collection. The heap is your application's memory, and things that live here are things you write and then reference all over your application – also known as reference types.

The Heap

As you and our friend walk down the street he starts telling you about dinner last night with his partner, who calls right then believe it or not. Our friend is still a bit late and trying to get the doughnut crumbs off his face so he hands the phone to you, where you get to have a conversation with the partner:

```
//when you see a "new" keyword, you see the heap in action
var conversation = new ThePartner.PhoneCall();
```

You know the partner pretty well; at least the public things. You're not too sure how **ThePartner** works in general... they're kind of abstract.

The conversation is nice enough, and as it goes on **ThePartner** tells you all about **List<Friend>()** they went out with last night and how much fun everyone had at their **FavoriteClub**.

At that moment your friend trips and you put the phone in your pocket, telling **ThePartner** to hold on a second. Turns out your buddy scuffed up his shoes badly and you have to grab a cab quickly to head back to his place.

Unfortunately you forgot about the phone call! Which is OK as the battery runs out on the phone and dereferences the call for you – this town's code is fully managed.

OK that was a bit geeky, but in essence that's what is going on:

When you use value types, you're using the stack. Reference types use the heap.

Why should you care? **Because the stack is typically faster.** But why? This is where we get into garbage collection, which is no doubt something you've heard of. We'll get into this in the next section, but for now understand that a garbage collector cleans up the heap for you.

This isn't free, and it requires overhead and processing time – which means it's slower. It can also be buggy, which is a pain.

I might get in trouble for saying this because compiler nuts out there will remind me that the .NET CLR's garbage collection (and that of a tuned up JVM) is so good that it's almost indistinguishable from just using the stack.

I'll sidestep all of that for now and let you make up your own mind on the subject.

Garbage Collection

Garbage collection (GC) is the process of cleaning up the heap. It's a facility of managed languages such as C# (and other CLR-based languages), Java (and anything that runs on the JVM) as well as dynamic languages such as Ruby, Python and JavaScript (and many others).

Some languages, like Objective-C, do not have garbage collection directly. This was a choice Apple made to keep their phones as fast as possible. You can, if you want, implement Automatic Reference Counting (ARC) which is a feature of LLVM.

When you write code in Objective-C, you allocate memory as you need and specify pointers vs value types explicitly. When you're done with the variable, you deallocate it yourself.

Correction on ARC

In previous versions of this book I made the claim that “most developers don’t use ARC”, which is false. I made this claim based on a conversation I had 5 years ago with an iOS developer friend. At that time Objective-C developers needed to know how to manually manage memory, and some didn’t trust ARC to work perfectly.

This has changed. Use of ARC is standard these days. A comment from a reader (Adrian Kosmaczew) sums it well:

... after ARC was introduced in 2010 (I was in the room during the WWDC in San Francisco when they announced it) Apple rebuilt pretty much every single app in Mac OS X with it, and ended up removing the garbage collection (which had been added to Objective-C in Mac OS X Leopard, back in 2007) a few years later. iOS never had the GC for the reasons of performance you mention in your book; that is correct. Having a “mark and sweep” GC was too much to handle for the small CPU of the original iPhones, and as such we devs ended up using “retain” and “release” calls to make sure objects didn’t die before we had a chance to use them.

The original Objective-C runtime (originally designed for the NeXTSTEP operating system in 1989) used a simpler GC scheme based in resource counting. What ARC does is to perform a static analysis of the source code and to insert the calls to “retain” and “release” automatically where needed. The final binaries are 100% compatible with older versions of OS X and iOS, because there’s no runtime library needed.

Now the Objective-C GC is 100% deprecated, and everybody in the “Apple galaxy” uses ARC these days. There are only two known cases of retain cycles one has to take care of when using it (that is, from a lambda to an object holding it, and from an object to another through strong references) but apart from that, it works beautifully well. Actually, the concept was quite revolutionary, it trades a bit of a longer compilation cycle with binaries that literally never leak.

There are a number of things to know about GC:

- **It's not free.** Determining what objects and memory are subject to collection requires overhead and can slow things down
- **It's undecidable.** This is as Turing illustrated with the Halting Problem: it's just not possible to know if a process will complete, therefore it's not possible to know if memory will ever not be needed

- **It's the focus of a lot of amazing work.** Speeding up GC means speeding up the runtime, so language vendors spend much time on it.

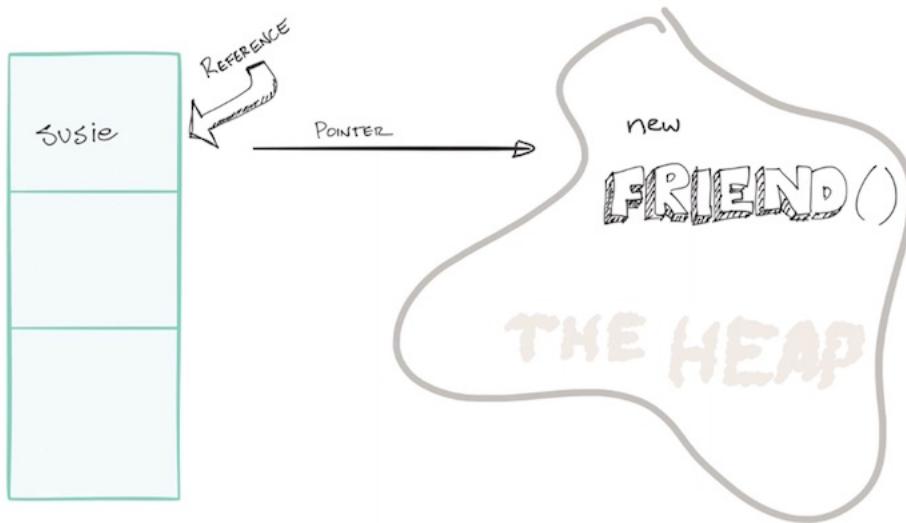
Saying that GC “cleans up the heap” is not going to satisfy the curious, so let’s dive into it a bit.

As with all things computer science, there are multiple ways to go about reducing the memory size of the heap and they all center on probability analysis. Some are sophisticated, some are very outdated and buggy.

The most common strategy for GC is tracing, to the point that unless you say otherwise, people will assume this is what you mean.

Tracing

Tracing is the most widely used method of garbage collection. When you create an object on the heap, a reference to it exists on the stack, as we’ve discussed.



As your program runs, the garbage collector will take a look at a set of “root” objects and their references, and trace those references through the stack. Objects on the heap can refer to other objects, so the trace can be much more complex than what is represented here.

As the trace runs, the GC will identify objects that are not traceable – meaning they aren’t reachable by other objects that are traceable. The untraceable objects are then deallocated and the memory freed.



The advantages of tracing are:

- It's accurate. If objects aren't referenced they are targeted for deallocation and that's that.
- It's easy. You just have to find the objects!
- The disadvantages are:
 - When will the GC get around to executing?
 - What happens when there are multiple threads? The stack for one thread may not have any references, but what about the other threads running?

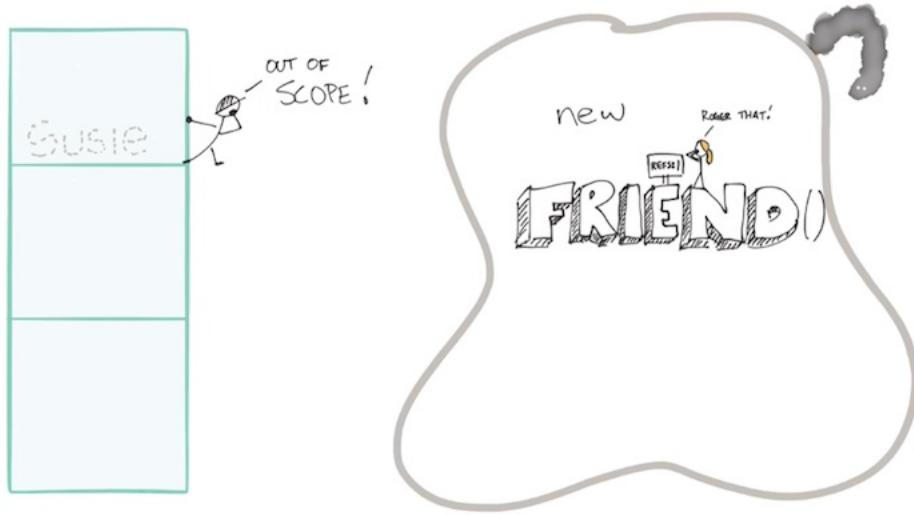
Obviously, it gets tricky. Within the tracing strategy there are various algorithms for marking and chasing down memory to be cleaned up, with different levels of speed and complexity. Both Java and .NET use tracing for GC, both of them implementing various flavors of generational (or ephemeral) implementations.

This is a really deep topic, full of algorithms, probability and statistical analysis ... and I could fill up many pages on the smallest details. Alas I have the rest of the book to write and so I need to clip the discussion here.

If you want to know more, Google away!

Reference Counting

Reference counting works almost exactly like tracing, but instead of running a trace, an actual count of references is made for each object. When the count goes to 0, the object is ready to be deallocated.



The advantages to reference counting are:

- It's simple. Objects on the heap have a counter which is incremented/decremented based on references to that object (from both the heap and the stack) rather than a trace algorithm.
- There's less guesswork as to "when" GC will happen. When local reference variables fall out of scope, they can be decremented right away. If that count goes to 0 then GC can happen shortly thereafter
- The asymptotic complexity (Big-O) of reference counting is O(1) for a single object, and O(n) for an object graph

Sounds simple and rather obvious doesn't it? What about this code (pseudo code):

```
public class Customer{
    public int ID {get;set;}
    //...

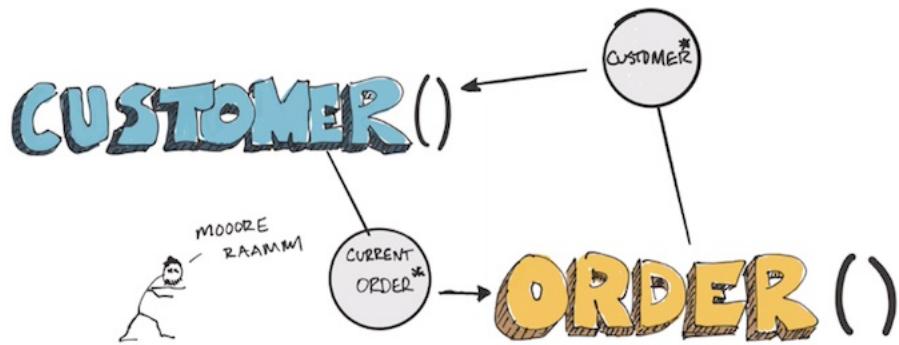
    public Order CurrentSalesOrder{get;set;}
    public Customer(int id){
        //fetch the current order from the db
        this.CurrentSalesOrder = db.getCurrentOrder(id);
    }
}

public class Order{
    public int ID {get;set;}
    //...

    public Customer Buyer{get;set;}
    public Order(GUID key){
        this.Customer = db.getCustomerForOrder(key);
    }
}
```

This code is quite common. What we have here is a circular reference on the heap. You could create an instance of the Customer and along with it comes a reference to an Order ... which has a reference back to the Customer.

The reference count for these objects will always be > 1 , so they are, in effect, little memory bombs eating up RAM.



Compile-time

I mentioned before that you can use Automatic Reference Counting (ARC) when working with XCode and Objective-C. Let's explore that a bit more.

If you have a sophisticated enough compiler it should be able to analyze what variables need what memory, where, and (possibly) for how long. Without it, Objective-C developers have had to manage memory themselves, allocating and deallocating within the code (much like C).

The XCode compiler analyzes the code written and decides the memory use upfront, freeing the runtime from the overhead of GC. It does this by using the notion of strong, weak, and unowned objects, each with an explicit level of protection that would, again, take me pages to explain properly (along with some goofy drawings).

The essence is this: using strong variables keeps them in memory longer, based on other strong references they're related to. If you have a weak reference to an object it tells the compiler "no need to protect the referenced object for longer than now". Finally with unowned you're making sure that a strong reference won't keep an object alive for longer than you want – so you explicitly say "make sure this object goes away".

If you want to know more about this, there's [a great article here](#) that discusses the ideas in terms of the human body:

A human cannot exist without a heart, and a heart cannot exist without a human. So when I'm creating my Human here, I want to give life to him and give him a Heart. When I initialize this Heart, I have to initialize it with a Human instance. To prevent the retain cycle here that we went over earlier (i.e. so they don't have strong references to each other), your Human has a strong reference to the Heart and we initialize it with an unowned reference back to Human. This means Heart does not have a strong reference to Human, but Human does have a strong reference to Heart. This will break any future retain cycles that may happen.

PART 3: PRACTICE

Resources

You can find the code used in parts of this section up at [my github repo](#). In addition, you [can buy 17 video walkthroughs](#) (of algorithms and the design patterns you see in this chapter) [from here](#).

The videos cover some new ground as well, and use languages other than C#.

CHAPTER TEN: OO PATTERNS

People have been writing code in object-oriented languages for a long time and, as you might guess, have figured out common ways to solve common problems. These are called design patterns and there are quite a few of them.

In 1994 a group of programmers got together and started discussing various patterns they had discovered in the code they were writing. In the same way that the Romans created the arch and Brunelleschi created a massive dome – the Gang of Four (as they became known) gave object-oriented programmers a set of blueprints from which to construct their code. The Gang of Four are:

- Erich Gamma
 - Richard Helm
 - Ralph Johnson
 - John Vlissides
- 

This entire chapter will be argumentative. I hate to say it, but there's just no escaping it: *how we build software is still evolving*. You might disagree with what you read here, which is fine. I'll do my best to present all sides – but before we get started please know this: *I'm not arguing for or against anything*.

These concepts exist; you should know they exist. You should know why people like or dislike them and, what is more important, what they argue about when they argue. And oh how they do.

Which is healthy! We need to suss out different approaches – to think in different ways about problems. Most importantly: *we need to know if others have already done so!*

In here be dragons. Every section is primed for detonation ... so grab your popcorn, do some meditation, and let's get into this.

When We Started Thinking In Objects...

The idea of “objects” in computer program design has been around since the late 50s and early 60s, but it was the release of Alan Kay’s Smalltalk in the 70s that really started to push the idea forward.

In the 80s the notion of “purely object-oriented languages” started cropping up and in the 90s it went through the roof when Java hit the scene.

We live in the aftermath of Java’s eruption.

Computer programs used to be a set of instructions, executed line by line. You could organize your code into modules and link things expertly :trollface: with **GOTO** statements. This had an amazing simplicity to it that kept you from over-thinking what it is you were trying to create.

This changed with the spread of object-oriented programming (OOP). People began building more complex applications, and the industry needed a different way of thinking about a program. We no longer write programs, we architect them.

For most developers, thinking in objects is natural. We write code that represents a thing and we can align that thinking with solving certain problems for a business.

This might come as a bit of a surprise, but there are a growing number of developers who are becoming less and less interested in OOP. For most developers, this is all they know.

Many, however, are questioning this.

You might be wondering: why are you bringing this up? The answer is simple: I want to challenge your assumptions before we even get started talking about software design. Everything that lives in this chapter has sprung from OOP. Not everyone believes any of this is a good thing.

One of the best opinion pieces that I've ever read, opposing OOP, is entitled [Object Oriented Programming is an expensive disaster which must end](#) by Lawrence Krubner. If you can allow yourself to get past the inflammatory title, it's worth every minute of your time. At least to get the juices flowing.

He brings up [Alan Kay's original vision](#):

The mental image was one of separate computers sending requests to other computers that had to be accepted and understood by the receivers before anything could happen. In today's terms every object would be a server offering services whose deployment and discretion depended entirely on the server's notion of relationship with the servee.

As Lawrence later points out, this is a strikingly apt description of the Actor Model in Erlang, which is a functional language.

This next statement is arguable, and it might make you angry. It might make you want to dismiss everything you've read thus far and maybe skip ahead to another chapter ... which is fine it's your book. I do hope you'll at least consider pondering Lawrence's main point, which I believe he puts together rather well:

My own experience with OOP involves long meetings debating worthless trivia such as how to deal with fat model classes in Ruby On Rails, refactoring the code into smaller pieces, each piece a bit of utility code, though we were not allowed to call it utility code, because utility code is regarded as a bad thing under OOP. I have seen hyper-intelligent people waste countless hours discussing how to wire together a system of Dependency Injection that will allow us to instantiate our objects correctly. This, to me, is the great sadness of OOP: so many brilliant minds have been wasted on a useless dogma that inflicts much pain, for no benefit. And worst of all, because OOP has failed to deliver the silver bullet that ends our software woes, every year or two we are greeted with a new orthodoxy, each one promising to finally make OOP work the way it was originally promised.

This is what we do when we implement software design patterns. From one perspective: *we waste precious time on worthless trivia*. From another perspective *we act disciplined and build for the future*.

Does the answer lie somewhere in the middle? To be honest with you: I don't know. You can't build something "half way" and expect it to work right. In many ways you commit, or you don't do it. Building software is a rather precision process and doing it well requires rigor.

Let's wander through some of the principles that Lawrence discusses, and see if we can make sense of it all.

The Code

The code for the examples you will read in this chapter can be [downloaded from Github](#). I recommend doing this if you want to play along, as copy/paste out of this book messes up the formatting.

Creational Patterns

When working with an OO language you need to create objects. It's a simple operation, but sometimes having some rules in place will help create the correct object, with the proper state and context.

Constructor

Most OO languages have a built-in way of creating an instance of a class. Here's one in C#:

```
1 public class Thing {  
2     //the constructor  
3     public Thing(){  
4         //set some class variables, etc  
5     }  
6 }  
7 var thing = new Thing(); //a constructor in C#
```

Here's a constructor in Ruby:

```
thing = Thing.new
```

Other languages, like JavaScript, require you to use a specific construct:

```
var Thing = function(){  
    //body  
}  
var thing = new Thing();
```

You can invoke this function directly, but if you use the new keyword it will behave like a constructor. This is very important if you're keen on creating an object in a valid state.

Factory

Sometimes instantiating an object can be rather involved, and might require a little more clarity. This is where the Factory Pattern comes in.

For instance: our Customer class might have some defaults that we want set:

```

1 public class Order{}
2
3 public class Customer
4 {
5     public string Id { get; set; }
6     public string Name { get; set; }
7     public string Email { get; set; }
8     public string Status { get; set; }
9     public List<Order> Orders { get; set; }
10
11    public static Customer FromDefaults ()
12    {
13        var customer = new Customer { Status = "unregistered", Name = "Guest" };
14        customer.Orders = new List<Order> ();
15        return customer;
16    }
17
18    public static Customer FromExisting (IDictionary values)
19    {
20
21        var customer = new Customer ();
22        //populate the values on the class, validating etc.
23
24
25        return customer;
26
27    }
28 }
```

Our **Customer** class isn't really all that complex, but you do gain some clarity by calling **Customer.FromDefaults()**. This can become important as your code base grows because it's not terribly clear what's going on if you simply use new **Customer()**.

For very complex class construction you could create a dedicated factory class. You see this often in Java. For instance, we could pull the instantiation logic completely out of our **Customer** and into a **CustomerFactory**:

```

1 public class Customer{
2     public string Id {get;set;}
3     public string Name {get;set;}
4     public string Email {get;set;}
5     public string Status {get;set;}
6     public List<Order> Orders {get;set;}
7 }
8
9 public class CustomerFactory
10 {
11     public Customer FromDefaults ()
12     {
13         var customer = new Customer { Status = "unregistered" , Name = "Guest" };
14         customer.Orders = new List<Order> ();
15         return customer;
16     }
17
18     public Customer FromExisting (IDictionary values)
19     {
20         var customer = new Customer ();
21         //populate the values on the class, validating etc.
22
23
24         return customer;
25     }
26 }
27
28 var customerFactory = new CustomerFactory();
29 var customer = customerFactory.FromDefaults();

```

Again: this is a bit simplistic. You can do a lot more with a factory class, such as deciding which object to create altogether. Your application might have the notion of an **Administrator** that inherits from **Customer**:

```

public class Administrator : Customer {
    //specific fields/methods for admins
}

```

You can use a variation of the Factory Pattern (called the Abstract Factory Pattern) to decide whether an **Administrator** should be returned or just a **Customer**:

```

1 public class CustomerFactory
2 {
3     public Customer FromDefaults ()
4     {
5         var customer = new Customer { Status = "unregistered", Name = "Guest" };
6         customer.Orders = new List<Order> ();
7         return customer;
8     }
9
10    public Customer FromExisting (IDictionary values)
11    {
12        if (values.Contains ("Email")) {
13            if (values ["Email"].ToString () == "admin@example.com") {
14                var admin = new Administrator ();
15                //populate values
16                return admin;
17            } else {
18                var customer = new Customer ();
19                //populate the values on the class, validating etc.
20
21                return customer;
22            }
23        } else {
24            return null;
25        }
26    }
27 }
28 }
29
30 var customerFactory = new CustomerFactory();
31 var customer = customerFactory.FromDefaults();

```

This pattern is useful but it can spiral on you if you are really into patterns. Consider this question on [StackOverflow](#):

What is a good name for class which creates factories? (FooFactoryFactory sounds silly imo)?

This [happens with C#](#) as well:

I make extensive use of the interface-based Typed Factory Facility in Windsor, but there are times when I must pass a lot of arguments to a factory around with the factory itself. I'd much prefer to create a factory factory with these arguments so that I don't need to muddy up the constructors of objects more than I need to.

The term “FactoryFactory” has become a bit of a joke in dynamic and functional language circles, throwing some shade at excessively zealous OOP programmers. Personally I think I might take a second and think through what I’m doing before using anything with a name like **FooFactoryFactory**, but then again I’m not building massive, monolithic applications either.

The point is: you now know the pattern, perhaps be reasonable with its use.

Builder

The Factory pattern can only do so much until it becomes too convoluted. This usually happens with very complex objects. Many developers consider this a “code smell” (when you find yourself needing it, it means there’s

a simpler way). Overly complex objects are ripe for bugs and, typically, means you've probably over-thought your solution.

There are times, however, that a Builder makes sense. Consider a class that .NET developers use all the time: **System.Text.StringBuilder**.

Strings are immutable in C#, so if you try to build a string from many string fragments, you can run into the memory problem as seen here:

```
1 public class NaiveStringBuilder {
2     IList<string> _strings;
3     public NaiveStringBuilder(){
4         _strings = new List<string>();
5     }
6     public void Append(string val){
7         _strings.Add(val);
8     }
9     public override string ToString(){
10        var result = "";
11        foreach (var s in _strings) {
12            // a new string is built each time
13            result = result + s + " ";
14        }
15        return result;
16    }
17 }
18
19 var naiveBuilder = new NaiveStringBuilder();
20 naiveBuilder.Append("This");
21 naiveBuilder.Append("could be");
22 naiveBuilder.Append("very long");
23 naiveBuilder.Append("and blow up");
24 naiveBuilder.Append("your program...");
25 var result = naiveBuilder.ToString(); //BOOM
```

If you ever find yourself writing a string concatenation routine in a loop, stop. It's a memory problem just waiting to happen.

The good news is that the C# team contemplated this and decided to help out, using the Builder pattern with **System.Text.StringBuilder**:

```
var goodBuilder = new System.Text.StringBuilder();
goodBuilder.Append("This ");
goodBuilder.Append("won't ");
goodBuilder.Append("blow up ");
goodBuilder.Append("my program ");
var result = goodBuilder.ToString(); //yay!
```

If you're curious about how the **StringBuilder** works, you can view the source code online. There's a lot going on in there! The thing to take away, however, is that an instance of an object (**System.String**) is being built for us in a very specific way to avoid problems. This is what the Builder Pattern is good for.

There is a more elegant way of doing this, however...

Method Chaining

Instead of calling **stringList.Add("...")** or using a **StringBuilder** directly, you can encapsulate what you're doing into a class that uses a fluent interface, otherwise known as Method Chaining:

```

1 public class Message
2 {
3     System.Text.StringBuilder _stringBuilder;
4     public Message (string initialValue)
5     {
6         _stringBuilder = new System.Text.StringBuilder ();
7         _stringBuilder.Append (initialValue);
8     }
9     public Message Add (string value)
10    {
11        _stringBuilder.Append(" ");
12        _stringBuilder.Append(value);
13        return this;
14    }
15    public override string ToString ()
16    {
17        return _stringBuilder.ToString ();
18    }
19 }
20
21 var message = new Message("Hello")
22     .Add("I might be")
23     .Add("a really long string")
24     .ToString();
25 //Hello I might be a really long string

```

Singleton

A Singleton is a class that only allows one instance of itself. It's not an easy thing to do correctly and many blog posts have been written about the perils of Singletons and threading or multiple processes.

You should know the pattern, however. Here's a rather naive one in C#:

```

1 public class SingleThing
2 {
3     //single instance holder
4     private static SingleThing _instance;
5     //disallow calling constructor directly
6     protected SingleThing () { }
7     //access to the instance
8     public static SingleThing Instance ()
9     {
10        if (_instance == null) {
11            _instance = new SingleThing ();
12        }
13        return _instance;
14    }
15 }

```

The problem with this code is that it will likely work fine most of the time. Until it gets used more and the Instance method is called simultaneously and a nasty collision happens. Or if, more likely, someone decides to use your code in a threaded environment.

The problems with the Singleton (in C#, at least) have been explored by some very bright people, [including Jon Skeet](#). The referenced article explores 5 or 6 different ways you can do it!

Interestingly, in JavaScript land, Node only uses Singletons for its moduling system. It can do this because Node is single threaded and mounts each module when it is used the first time.

Structural Patterns

Code needs to have some structure, so we use things like methods, modules, and classes. As code becomes more complex, these modules and classes might also need some structure to reduce confusion and unneeded complexity. That's what we'll take a look at here.

Adapter

The Adapter Pattern is all about making one interface work with another. You see them used often with data access tools (ORMs) where the abstracted query interface needs to work against different databases, such as PostgreSQL, SQL Server, etc.

For instance, we might create a way of working with a database that we really like, so we'll abstract it into a basic interface of some kind:

```
1 public abstract class GroovyQuery{
2     //groovy interface
3     //find
4     //fetch
5     //save
6 }
7
8 public class GroovyPostgresAdapter: GroovyQuery{
9     //implements groovy interface for PostgreSQL
10 }
11
12 public class GroovySQLServerAdapter: GroovyQuery{
13     //implements groovy interface for SQL Server
14 }
```

You just need to pick the correct adapter for the database you're working with. This can be done manually or by way of configuration and some higher-level patterns which we'll see later on.

Bridge

The Bridge Pattern is quite subtle and tends to look a lot like the Adapter Pattern, but it's one step up the abstraction curve. You use the Bridge Pattern when your abstraction gets complicated enough that you need to split things out.

People really like our **GroovyQuery** tool and we want to add a feature: document queries. It turns out that you can store JSON happily in PostgreSQL and also in SQL Server – so we decide to implement a document API that handles parsing and so on:

```
public abstract class GroovyQuery{
    //groovy interface
    //Find
    //Fetch
    //Save
    public abstract T GetDocument<T>();
    public abstract T SaveDocument<T>();
    public abstract IList<T> FetchDocuments<T>();
}
//etc
//etc
```

This is a very interesting idea! The problem is that we now have to go and implement it for every adapter. Unless we abstract the document interface and bridge it to our **GroovyQuery** as in [Listing 12.7](#):

```

1 public abstract class GroovyQuery{
2     //groovy interface
3     //Find
4     //Fetch
5     //Save
6     public IDocumentQueryable Documents();
7     //etc
8 }
9
10 //a document query interface
11 public interface IDocumentQueryable{
12     T Get<T>();
13     T Save<T>();
14     IList<T> Fetch<T>();
15 }
16
17 //implementation of the document query interface for
18 //relational systems.
19 public class RelationalDocumentQueryable : IDocumentQueryable{
20     GroovyQuery _adapter;
21     public RelationalDocumentQueryable(GroovyQuery adapter){
22         this._adapter = adapter;
23     }
24     //implement Get, Save, Fetch
25 }
26
27 //our SQL Server adapter
28 public class GroovySqlServerAdapter: GroovyQuery{
29     public GroovySqlServerAdapter(){
30         this.Documents = new RelationalDocumentQueryable(this);
31     }
32     //implement Get, Save, Fetch
33 }

```

The neat thing about this new structure is we can change our **IDocumentQueryable** interface and the implementation, without breaking any of our adapters.

Composite

The Composite Pattern deals with parent-child relationships that are composed to create a whole object. They can grow and shrink dynamically and child objects can move between parents.

Our GroovyQuery tool is really picking up steam! People are really happy with it, mostly because we have a cool document abstraction they can use next to your typical ORM interface. The problem is we need more speed!

It turns out that the some of the drivers we've been using don't implement connection pools – basically a set of 10 or so open connections to the database that we keep alive so we don't need to take the time establishing a connection for each query.

We can create our own incredibly naive implementation using the Composite Pattern:

```

1 public class Connection
2 {
3     public bool CheckedOut { get; set; }
4     public Connection (string connectionString)
5     {
6         //connect
7     }
8     public void Close ()
9     {
10        //close the connection
11    }
12 }
13
14 public class ConnectionPool
15 {
16     public IList<Connection> Pool;
17     public ConnectionPool (string connectionString)
18     {
19         this.Pool = new List<Connection> ();
20         for (var i = 0; i < 10; i++) {
21             this.Pool.Add (new Connection (connectionString));
22         }
23     }
24     public void Checkout ()
25     {
26         //grab a list of connections which aren't checked out
27         //return the first
28     }
29     public void Checkin ()
30     {
31         //tick the boolean
32     }
33     public void Drain ()
34     {
35         foreach (var connection in this.Pool) {
36             connection.Close ();
37         }
38         this.Pool = new List<Connection> ();
39     }
40 }

```

I hesitated to show a **ConnectionPool** example as I'm sure many of you will be poking holes in it (as you should)! Pooling is a hard thing to do and I don't recommend writing your own. I include it here because it's a real-world example that's easily understood (as opposed to the mind-numbing Foo and Bar nonsense you see everywhere).

If the **ConnectionPool** goes away, so do all the connections. If there are no children (in other words the **IList<Connection>** is empty, there is no **ConnectionPool**. The parent and children work together to provide functionality.

If you work in an IDE (such as Visual Studio or Eclipse) – each of the UI elements you see is a component that has a parent. This, again, is the Composite Pattern.

Decorator

The Decorator Pattern adds behavior to an object at runtime. You can think of it as “dynamic composition”.

We could use the Decorator Pattern as an alternative to the Bridge Pattern above for our **GroovyQuery** engine:

```

1 public abstract class GroovyQuery
2 {
3     //groovy interface
4     public abstract T GetDocument<T> ();
5     public abstract T SaveDocument<T> ();
6     public abstract IList<T> FetchDocuments<T> ();
7
8     public IDocumentQueryable Documents;
9     //etc
10 }
11
12 public interface IDocumentQueryable
13 {
14     T Get<T> ();
15     T Save<T> ();
16     IList<T> Fetch<T> ();
17 }
18
19 //implementation of the document query interface for
20 //relational systems.
21
22 public class RelationalDocumentDecorator : IDocumentQueryable
23 {
24     GroovyQuery _adapter;
25
26     //Find, Fetch, and Save use the _adapter passed in
27     public RelationalDocumentDecorator (GroovyQuery adapter)
28     {
29         this._adapter = adapter;
30     }
31     //implement Get, Save, Fetch for Documents below
32 }
```

With our **RelationalDocument** Decorator we're able to "decorate" the **GroovyQuery** base object with the ability to work with JSON documents.

Facade

A Facade hides implementation details so clients don't have to think about it. We can use a Facade for our **GroovyQuery** to pick an adapter for the calling code, so they don't need to worry about how to wire things together as you can see here:

```

1 using System;
2 using System.Collections.Generic;
3
4 namespace Facade
5 {
6
7     //abstract base class
8     public abstract class GroovyQuery
9     {
10         public GroovyQuery (string connectionString) { }
11     }
12
13     //implementation for PostgreSQL
14     public class PostgreSQLQuery : GroovyQuery {
15         public PostgreSQLQuery(string connectionString)
16             : base(connectionString){}
17     }
18
19     //implementation for SQL Server
20     public class SQLServerQuery : GroovyQuery
21     {
22         public SQLServerQuery (string connectionString)
23             : base(connectionString) { }
24     }
25
26
27     //a simple class that hides the selection details
28     public class QueryRunner
29     {
30
31         string _connectionString;
32
33         //Find, Fetch, and Save use the _adapter passed in
34         public QueryRunner (string connectionString)
35         {
36             _connectionString = connectionString;
37         }
38
39         public void Execute ()
40         {
41
42             GroovyQuery runner;
43             if (_connectionString.StartsWith ("postgresql://",
44                 StringComparison.InvariantCultureIgnoreCase)) {
45                 runner = new PostgreSQLQuery (_connectionString);
46             } else if (_connectionString.StartsWith ("sqlserver://",
47                 StringComparison.InvariantCultureIgnoreCase)) {
48                 runner = new SQLServerQuery (_connectionString);
49             } else {
50                 throw new InvalidOperationException ("We don't support that");
51             }
52             //execute with the runner
53         }
54
55     }
56 }
```

Flyweight

In the initial versions of **GroovyQuery** we decided it would be very useful to introspect our database whenever a write needed to happen (insert or update query). We did this because knowing more about each table (data types, primary key fields, column names, and default values) would be extremely helpful in crafting up a very usable API.

Unfortunately this became very slow when the system came under load, so we opted to implement the Flyweight Pattern.

Now, when **GroovyQuery** starts up, it runs a single query that introspects every table in our database, and then loads up a series of small objects that can be used throughout our application:

```
1 //our Flyweight class
2 public class Table{
3     public string Name {get;set;}
4     public string PrimaryKeyField {get;set;}
5     //column and data type information...
6 }
7
8 public abstract class GroovyQuery{
9     //the API as we've come to know it
10    List<Table> _tables;
11    public void Initialize(){
12        _tables = new List<Table>();
13        //query the database for meta information
14        //load up the _tables list
15    }
16 }
```

Now, whenever we run an insert or update query we can reuse one of the **Table** instances we have in memory, avoiding the need to make a special query call on each write operation. This pattern can scale reasonably well to quite a large number tables, and helps to scale our app to thousands of write operations per second.

As a quick aside: this is precisely the pattern I have used in two of the data access libraries I put together: [MassiveJS](#) and [Moebius](#).

Behavioral Patterns

We've figured out various ways to create our **GroovyQuery** class as well as how to enable functionality by structuring things a certain way. Now let's see how we can use patterns to simplify how clients can use our **GroovyQuery** operations.

Chain of Responsibility

We've decided to implement validations for our User and have to orchestrate a bit of an approval chain. We can use the Chain of Responsibility Pattern for this, which is focused on moving data through a set of handlers.

There Is A Better Way

Moving objects through a process chain can be subject to many high-level patterns that are, frankly, much better than this one. I'm showing you this example because you should know the pattern – but when it comes to validations there are better ways to do this.

The first thing to do is to create an abstract handler class:

```
//our handler class
public abstract class UserValidator
{
    protected UserValidator Successor = null;
    public void SetSuccessor (UserValidator successor)
    {
        this.Successor = successor;
    }

    public abstract void Validate (User user);
    public void HandleNext (User user)
    {
        if (user.IsValid && this.Successor != null) {
            this.Successor.Validate (user);
        }
    }
}
```

This handler will allow us to define our **Validate** routines as well as any successors that we might have. Now we can create individual validations:

```
public class NameValidator : UserValidator
{
    public override void Validate (User user)
    {
        user.IsValid = !String.IsNullOrEmpty (user.Name);

        if (user.IsValid) {
            user.ValidationMessages.AppendLine ("Name validated");
        } else {
            user.ValidationMessages.AppendLine ("No name given");
        }
        HandleNext (user);
    }
}
```

```
public class AgeValidator : UserValidator
{
    public override void Validate (User user)
    {
        user.IsValid = user.Age > 18;
        if (user.IsValid) {
            user.ValidationMessages.AppendLine ("Age validated");
        } else {
            user.ValidationMessages.AppendLine ("Age is invalid - must be over 18");
        }
        HandleNext (user);
    }
}
```

One of the great things about using Chain of Responsibility is that we can formalize our validations into classes that target a single use case, rather than writing a ton of validation code onto our **User**.

Speaking of, let's create our **User** class and orchestrate the validations:

```

public class User
{
    public string Name { get; set; }
    public int Age { get; set; }
    public bool IsValid = false;
    public System.Text.StringBuilder Logs;

    public User ()
    {
        this.ValidationMessages = new System.Text.StringBuilder ();
        this.ValidationMessages.AppendLine ("Pending save");
    }

    public void Validate ()
    {
        var nameCheck = new NameValidator ();
        var ageCheck = new AgeValidator ();
        nameCheck.SetSuccessor (ageCheck);

        //kick it off
        nameCheck.Validate (this);
    }
}

```

The **IsValid** and **ValidationMessages** properties will let us know if we can save this user, and what's happened during our validation process.

The user's **Validate** routine is where the orchestration is at. We instantiate the name and age validators, and then decide which goes first. We could use Method Chaining here – but I think this is clear enough.

Now we just need to kick it off:

```

var user = new User { Name = "Larry", Age = 22 };
user.Validate();
user.IsValid //true
user.ValidationMessages.ToString() //Pending Save, Name validated, Age validated

var user2 = new User { Name = "Larry", Age = 16 };
user2.Validate();
user2.IsValid //false
user2.ValidationMessages.ToString() //Pending Save, Name validated, Age is invalid

```

As I mention, there are more elegant ways to do this, but this is simple to understand, easy to test and gets the job done.

Command

The Command Pattern formalizes requests from one API to the next.

Our data access tool, **GroovyQuery**, is all about writing and reading data from the database. It does this by creating SQL statements that our adapter then executes. We could do this by passing in a SQL string and a list of parameters – or we could formalize it into a command:

```

public class QueryParameter
{
    public QueryParameter (string name, string value)
    {
        this.Name = name;
        this.Value = value;
    }
    public string Name { get; private set; }
    public string Value { get; private set; }
}
public interface IQueryCommand
{
    string SQL { get; set; }
    IList<QueryParameter> Parameters { get; set; }
    IDbCommand BuildCommand ();
}
public class QueryCommand : IQueryCommand
{
    public string SQL { get; set; }
    public IList<QueryParameter> Parameters { get; set; }
    public IDbCommand BuildCommand ()
    {
        //return a command that can be executed
        //...
    }
}

```

This is our new **GroovyQuery** class:

```

public class GroovyQuery
{
    //the API
    //...

    public IDataReader Execute (IQueryCommand cmd)
    {
        //build the command and execute it
        var dbCommand = cmd.BuildCommand ();
        //...

    }
}

```

One thing about formalizing a request like this is that we can scale it to specific needs:

```

public class CreateUserCommand : QueryCommand
{
    public CreateUserCommand (string name, string email, string password)
    {
        this.SQL = @"insert into users(name, email, hashed_password)
values(@1, @2, @3);";

        this.Parameters = new List<QueryParameter> ();
        this.Parameters.Add (new QueryParameter("@1", name));
        this.Parameters.Add (new QueryParameter("@2", email));
        this.Parameters.Add (new QueryParameter("@3",
            SomeHashingAlgorithm (password)
        ));
    }
    private string SomeHashingAlgorithm (string val)
    {
        //some solid hashing here...

        return "";
    }
}

```

A little naive, perhaps, but this command encapsulates what it means to add a **User** to our system for SQL and the parameters required.

Mediator

We want to formalize our document storage capabilities, however adding methods and abstractions to our **GroovyQuery** will make the API more complex, which goes against some programming principles we'll discuss in a later chapter.

In short: simplicity is our goal. We want our class abstractions to do one thing and to do it well.

Let's formalize our document storage idea with the Mediator Pattern. A Mediator is simply a class that sits between two other classes, facilitating communication. It's often used in message-based applications, but we can use a simplified version here:

```

1 //our Mediator
2 public class DocumentStore{
3     GroovyQuery _adapter;
4     public DocumentStore(GroovyQuery adapter){
5         _adapter = adapter;
6     }
7     public T Save<T>(T item){
8         //parse and save the object
9     }
10    public T Get<T>(){
11        //pull the record, dehydrate
12    }
13    public IList<T> Fetch<T>(){
14        //pull the list, dehydrate
15    }
16    string Dehydrate<T>(T item){
17        //turn the object into JSON
18    }
19    T Hydrate<T>(string json){
20        //resolve
21    }
22 }

```

Here, we're mediating between our database adapter and any class type of **T**. The adapter doesn't need to know anything at all about **T**, and **T** knows nothing about the adapter.

Observer

The Observer Pattern facilitates event-based programming. You use this pattern whenever you wire up events in a language like C# or JavaScript (using the **EventEmitter** in Node or listening to DOM events in the browser).

Many frameworks have the mechanics for observation already built in, but let's take a look at how we can construct an observer by hand by adding methods to our **GroovyQuery** that get fired when certain events occur. These are commonly referred to as callbacks:

```
1 public interface IListener
2 {
3     void Notify<T> (T result);
4     void Notify ();
5 }
6 public abstract class GroovyQuery
7 {
8     //API methods etc
9     //...
10
11    public IList<IListener> Listeners { get; set; }
12    public GroovyQuery ()
13    {
14        //constructor stuff
15        //...
16
17        this.Listeners = new List<IListener> ();
18    }
19    public virtual IDataReader Execute (IDbCommand cmd)
20    {
21        //the execution stuff
22        //notify all listeners
23        foreach (var listener in this.Listeners) {
24            listener.Notify ();//optionally send along some data
25        }
26    }
27 }
```

There are other ways to do this in C# – namely using virtual methods that inheriting classes can implement directly.

State

The State Pattern changes an object's behavior based on some internal state. Often this is done by creating formalized state classes.

We want to know what current state our **QueryCommand** is in – if it's new, succeeded, or failed. Let's create some classes that tell us this. We'll start with the **QueryCommand** and a base class for the state:

```
public class QueryCommand
{
    public QueryState State { get; set; }
    //...

    public QueryCommand ()
    {
        this.State = new NotExecutedState (
            "Query has not been run"
        );
    }
    //pass execution off to the state bits
    public T Execute<T> ()
    {
        return this.State.Execute<T> (this);
    }
}
public abstract class QueryState
{
    protected string Message { get; set; }
    public QueryState (string message)
    {
        this.Message = message;
    }
    public abstract T Execute<T> (QueryCommand cmd);
}
```

Next we implement each possible state as an explicit class:

```

public class SuccessState : QueryState
{
    public SuccessState (string message) : base (message) { }
    public override T Execute<T> (QueryCommand cmd)
    {
        throw new InvalidOperationException (
            "This query already executed successfully"
        );
    }
}
public class FailState : QueryState
{
    public FailState (string message) : base (message) { }
    public override T Execute<T> (QueryCommand cmd)
    {
        throw new InvalidOperationException (
            "This query already failed execution"
        );
    }
}
public class NotExecutedState : QueryState
{
    public NotExecutedState (string message) : base (message) { }
    public override T Execute<T> (QueryCommand cmd)
    {
        try {
            //run query execution... and if it works
            cmd.State = new SuccessState (
                "Query executed successfully"
            );

        } catch (Exception x) {
            //on error
            cmd.State = new FailState (x.Message);
        }
        //return query results
    }
}

```

Notice in this code how the actual execution is handed to the **QueryState**? This probably seems a bit counterintuitive, but if you think of it as a formalized state of the **QueryCommand** it makes more sense. It also allows you to forego a big switch statement.

Strategy

The Strategy Pattern is a way to encapsulate “doing a thing” and applying that thing as needed. Code is the easiest way to explain this pattern, as it’s quite simple and useful.

Our document query capability is working well, but it turns out that SQL Server has excellent support for XML, and some users have asked that we support that along with JSON storage.

We can do this using the Strategy Pattern. We’ll start by defining our **GroovyQuery** class and the interfaces we need:

```

public interface IDocumentQueryable
{
    T Get<T> ();
    T Save<T> ();
    IList<T> Fetch<T> ();
}

public abstract class GroovyQuery
{
    //groovy interface
    public abstract T GetDocument<T> ();
    public abstract T SaveDocument<T> ();
    public abstract IList<T> FetchDocuments<T> ();

    public IDocumentQueryable Documents;
    //etc
}

public interface IStorageStrategy
{
    T Hydrate<T> (string document);
    string Dehydrate<T> (T item);
}

```

Next we implement two strategies for storing documents, one for XML and the other for JSON:

```

public class JsonStorageStrategy : IStorageStrategy
{
    public string Dehydrate<T> (T item)
    {
        //turn the object into JSON, return the JSON
    }
    public T Hydrate<T> (string json)
    {
        //resolve from JSON
    }
}

public class XmlStorageStrategy : IStorageStrategy
{
    public string Dehydrate<T> (T item)
    {
        //turn the object into XML
    }
    public T Hydrate<T> (string xml)
    {
        //resolve from XML
    }
}

```

Finally we implement the **DocumentStore** with the needed strategies:

```

public class DocumentStore
{
    GroovyQuery _adapter;
    IStorageStrategy _parser;
    public DocumentStore (GroovyQuery adapter)
    {
        _adapter = adapter;
        _parser = new JsonStorageStrategy ();
    }
    public DocumentStore (GroovyQuery adapter, IStorageStrategy parser)
    {
        _adapter = adapter;
        _parser = parser;
    }
    public T Save<T> (T item)
    {
        var document = _parser.Dehydrate (item);
        //parse and save the object
    }
    public T Get<T> ()
    {
        //pull the record, dehydrate
        //get the results
        return _parser.Hydrate<T> (result);
    }
    /**
     */
}

```

In The Real World...

Many of you will likely notice that I left a few patterns out of the above list – namely the Visitor Pattern, Memento, Template, etc. These are useful patterns to know about, but their use is rather rare.

For instance the Visitor Pattern – it's useful if you're parsing tree structures (like Expression Trees in C#) but in everyday code, this is kind of rare. For me, at least.

Also: as you implement patterning as we've done here, the code you write tends to become more generalized and you end up writing a lot more of it just to do a simple operation. This is not what these patterns are for.

A design pattern should make things simpler. If you implement one, have a look at your code before and after, and see if it makes more sense or less.

CHAPTER ELEVEN: OO PRINCIPLES

As you build applications using the patterns we learned in the previous chapter, you begin to see some common side effects.

For instance: the Strategy, Adapter, Mediator and Bridge Patterns lead you to think a little bit more about better ways to manage dependencies between classes. You also begin to create rules and reasons why code should even exist in the first place.

Obviously, this is not a small topic. In this chapter we'll discover the key principles you should understand, who came up with them, and why.

Coupling and Cohesion

You've likely heard these terms before, they're thrown around a lot and have fairly straightforward definitions:

- **Cohesion** applies to how well you've thought out the concepts (and concerns, for that matter) of your application. In other words: how related are the functions of each module or class? You put your **Membership** code into a **Membership** module and your **User** code in a **User** model. The functionality here is cohesive (meaning the ideas bond together logically).
- **Coupling** is kind of the opposite of cohesion. When you couple two or more things, their separate notion becomes one. In our code above we had an example where **Membership** created a **newUser** during the registration process. This coupled **Membership** to User. If we moved/renamed/got rid of the notion of a User we'd have an error in our **Membership** code. This is tight coupling.

You want high cohesion, low coupling. Your classes and modules should make sense for isolating ideas, and not rely on each other to exist. This is the goal, at least.

These ideas were invented in the 60s by Larry Constantine and later formalized in a white paper called [Structured Design](#) (Yourdon and Constantine, 1979):

For most of the computer systems ever developed, the structure was not methodically laid out in advance – it just happened. The total collection of pieces and their interfaces with each other typically have not been planned systematically. Structured design, therefore, answers questions that have never been raised in many data processing organizations.

It's a fascinating and easy read, and I highly suggest it.

Separation of Concerns

Separation of Concerns is about slicing up aspects of your application so they don't overlap. These are typically thought of (by developers) as horizontal concerns (they apply to the application as a whole): such as user interface, database, and business logic. The term can equally (and confusingly) be applied to more abstract ideas, such as authentication, logging and cryptography.

Finally there are vertical concerns, which deal with more business-focused functionality such as Content Management, Reporting, and Membership.

Some Opinion

I had a discussion once with a developer who suggested I separate SQL from my data access code so I can have a “cleaner separation of concerns”.

Another time it was suggested to me that dividing my .NET code into separate library projects was a great way to separate the concerns of my application, instead of having all those files together in one place.

Ruby on Rails (which is responsible for the spread of the term) famously suggested that the Model View Controller approach they used was a great “separation of concerns” as it decoupled data access and business logic from HTML. The reality is it did exactly the opposite. A Rails view is HTML strewn with artifacts (which are Models) created in a Controller that deal directly with data access.

There is no separation there. Ruby on Rails version 3.0 tried to get there with more generic implementations (so called Railties), but this ended slowing everything down.

The term “Separation of Concerns” honestly doesn’t mean anything anymore. But it used to, and I’ll devote the rest of the chapter to that meaning.

The Origin Of The Term

The origin of the term comes from this quote from one of my favorite computer science people, [Edsger W. Dijkstra](#):

Let me try to explain to you, what to my taste is characteristic for all intelligent thinking. It is, that one is willing to study in depth an aspect of one’s subject matter in isolation for the sake of its own consistency, all the time knowing that one is occupying oneself only with one of the aspects. We know that a program must be correct and we can study it from that viewpoint only; we also know that it should be efficient and we can study its efficiency on another day, so to speak. In another mood we may ask ourselves whether, and if so: why, the program is desirable. But nothing is gained — on the contrary! — by tackling these various aspects simultaneously. It is what I sometimes have called “the separation of concerns”, which, even if not perfectly possible, is yet the only available technique for effective ordering of one’s thoughts, that I know of. This is what I mean by “focusing one’s attention upon some aspect”: it does not mean ignoring the other aspects, it is just doing justice to the fact that from this aspect’s point of view, the other is irrelevant. It is being one- and multiple-track minded simultaneously.

Every application we build is composed of a vertical subset of processes and rules that try to solve a business need. An eCommerce application will have a sales aspect, a membership aspect, accounting, and fulfillment. These are concerns of the application.

Can we apply this type of thinking to more horizontal ideas? In other words, can we study in depth the notion of logging? Or data access? I’m sure friends of mine would argue that I have, indeed, done the latter many times! Studying these horizontal aspects of our application might make the application better, but I don’t think it will make it more correct.

Now this is where we come up against the weight of history and a little trick that every politician knows: **If you say something long enough it become true**. I think it’s the same with the phrase “separation of concerns”. It’s reminiscent of the phrase “I could care less” or “hone in on”. These phrases make no sense at all, but for some reason popular American English vernacular has twisted them to mean something and, as time goes on, they get adopted.

I think the same is true with Separation of Concerns. It’s a catchall phrase which means “I’m trying to do the right thing”, whatever that thing may be. Perhaps it’s an effort at file organization or using one of Fowler’s enterprise patterns (which we won’t be discussing in this book) to abstract away a part of your application … at this

point *it doesn't matter*.

So: when in conversation and someone invokes this trite little phrase, perhaps ask them for some detail. After a few years of doing this you should have some fun tales to share.

YAGNI and DRY

I remember when I started learning Ruby. I loved the simplicity of the language as well as its dynamic design which, I know, many people dislike. You had to have some rigor and much care when building programs with Ruby because you didn't have a compiler and static type checking.

This was freeing, and it was also a little scary.

Part of this rigor was learning a new set of jargon. YAGNI (You Aint Gonna Need It) and DRY (Don't Repeat Yourself) – neither of which came from the Ruby community – started becoming more popular precisely because the practices you needed to adopt to write good Ruby code leaned squarely on you, as developer, rather than your tooling.

More (and better) testing replaced compiler checks. Test-driven Development (TDD, which we'll get to in a later chapter) helped in this regard as well – forcing you to justify the code you needed to write with a set of tests. In other words: if a test didn't mandate some code's existence, you didn't need that code.

Don't Repeat Yourself is something that most developers understand. Duplicated code is difficult to maintain. With large applications, this is almost impossible – but it is an important idea to keep a focus on.

If you have an application with 300 classes and 50,000 lines of code, you're bound to have some kind of duplication. The trick is to spot it and, hopefully, to simplify your future by abstracting it in some way.

The Wrong Abstraction

In January of 2016 Sandi Metz [wrote a great article](#) about the perils of focusing too much on DRY. The punchline being:

... duplication is far cheaper than the wrong abstraction... prefer duplication over the wrong abstraction...

The article sprang from [a talk she gave](#) at RailsConf in 2014.

The main idea is that you don't want to create abstractions solely for the sake of avoiding repetition; it needs to fit your overall approach. If you can't fit it, just let the duplication be.

Tell, Don't Ask

Another Rubyism that I quite like came from Ruby's inspiration: Smalltalk. Whenever you invoke a method on a Ruby class you send it a message. You tell that instance that you need it to do something, or that you need some data back of some kind.

If you ask an object instance a question, then you'll need to know something about that object or its state, which breaks the notion of encapsulation.

If you think back to our **GroovyQuery** from a previous chapter, imagine this as our API:

```

using System;
using System.Data;

public class GroovyQuery
{
    public bool IsCommandValid ( IDbCommand cmd)
    {
        //logic
    }
    public bool IsConnectionAvailable ()
    {
        //check connection pool to see if one is ready
    }
    public IDataReader Execute ( IDbCommand cmd)
    {
        //execution
    }
}

```

To use this API effectively I would need to ask two questions and finally get around to telling the class what to do (**Execute**). In short: I need to know way more about the API than is needed.

A better way to do this is by moving a few things around:

```

using System;
using System.Data;
public class GroovyQuery2 //Telling
{
    bool CommandIsValid ( IDbCommand cmd)
    {
        //logic
    }
    bool ConnectionIsAvailable ()
    {
        //check connection pool to see if one is ready
    }
    public IDataReader Execute ( IDbCommand cmd)
    {
        var commandIsValid = CommandIsValid (cmd);

        if (ConnectionIsAvailable () && commandIsValid) {
            //execution
        } else {
            throw new InvalidOperationException ("Can't run this query");
        }
    }
}

```

The responsibility for deciding whether the query can run is now within **GroovyQuery**, which is where it should be.

Law Of Demeter (or: Principle of Least Knowledge)

The Law of Demeter (LoD, or “Deep Dotting”) is an offshoot of loose coupling. In short: you shouldn’t have to “reach through” one object to get to another. This can be further nuanced to mean you shouldn’t have to reach deeply into one object to do the thing you need to do.

Let’s examine both.

Our **Membership** system is working well, but some users aren't behaving themselves so we need to give them a bit of a timeout. With our first go we decide to drop a Suspend method on User because we're telling them they're suspended:

```
public class DB
{
    public User GetUser (int id)
    {
        //call to the DB, getting record
        //returning an empty user for now
        return new User ();
    }
}

public class User
{
    public String Status { get; set; }
    public void Suspend ()
    {
        this.Status = "suspended";
    }
}

public class Membership
{
    DB _db;
    public Membership ()
    {
        _db = new DB ();
    }
    public User GetUser (int id)
    {
        //get the user
        return _db.GetUser (id);
    }
}
```

To suspend a user we need to access them from the Membership module and then suspend them:

```
var membership = new Membership();
membership.GetUser(1).Suspend();
```

This is a violation of LoD. We had to reach through **Membership** to get to the **User**. You might be wondering ... so what?

It's a subtle point, sure, but the more you think on it the more you realize how you're muddying the principles we've been reading about.

In essence: *we've punched a hole in our membership abstraction* by dividing the responsibility for changing the user between two different classes. Cohesion is breaking down and coupling is going up.

It doesn't make sense to involve **Membership** at all here, except for the fact that we need to get at the **User**. So let's make a choice, and it's a simple one: **Membership** has the responsibility of adding and retrieving users (aka changing them) so let's have it update the user's status as well:

```

public class DB
{
    public User GetUser (int id)
    {
        //call to the DB, getting record
        return new User ();
    }
    public void Save (object item)
    {
        //save to DB
    }
}

public class User
{
    public String Status { get; set; }
}

public class Membership
{
    DB _db;
    public Membership ()
    {
        _db = new DB ();
    }
    public User GetUser (int id)
    {
        //get the user
        return _db.GetUser (id);
    }

    public void SuspendUser (int id)
    {
        var user = this.GetUser (id);
        user.Status = "suspended";
        _db.Save (user);
    }
}

```

Some developers will focus on “dot counting”, claiming that the use of too many dots is, all by itself, a violation of LoD. Sometimes it is, sometimes not.

Consider this API:

```
var liTag = new Html.Helpers.HtmlTags.Lists.ULTag().LiTag();
```

This API is kind of ridiculous, I must say. However it’s an organizational choice and not necessarily a violation of LoD. This could be horrible namespacing for all we know! Or it could be a lack of imagination. We don’t know the inner workings of the Html helper library (and trust me, you don’t want to), so it’s not exactly accurate to call for a violation just by looking at dots.

I was reading my friend [Phil Haack’s blog](#) while researching this subject, and he had a great quote from [Martin Fowler](#):

I’d prefer it to be called the Occasionally Useful Suggestion of Demeter.

I hate to leave you with vagary, but hopefully you can see how “deep dotting” and LoD aren’t always the same thing.

Dependency Injection

One way to loosen up your code is to send in the dependencies that a class needs through its constructor. The best way to see this is with some code.

Our **Membership** class is using the database to retrieve and save a **User**:

```
public class Membership{
    DB _db;
    public Membership(){
        _db = new DB();
    }
    public User GetUser(int id){
        //get the user
        return _db.GetUser(id);
    }
    public void SuspendUser(int id){
        var user = this.GetUser(id);
        user.Status = "suspended";
        _db.Save(user); //Coupling
    }
}
```

This couples the **Membership** class to the **DB** class which is responsible for data interactions. We've read *The Imposter's Handbook*, so we know that coupling is bad – but how can we change this?

The simple answer is to inject the dependency through the constructor rather than to invoke it in place:

```
public class Membership
{
    DB _db;
    public Membership (DB db)
    {
        _db = db;
    }
    public User GetUser (int id)
    {
        //get the user
        return _db.GetUser (id);
    }
    public void SuspendUser (int id)
    {
        var user = this.GetUser (id);
        user.Status = "suspended";
        _db.Save (user); //Coupling
    }
}
```

Now our class doesn't need to know how to instantiate **DB**, which is one step in the right direction. There's still a bit too much coupling, however as our **Membership** class cannot be used unless a **DB** instance is passed in.

Let's see how we can loosen this up a bit more.

Interface-based Programming

Many languages support the idea of interfacing with an ability, rather than a type itself. With C# and Java these are called Interfaces. With languages such as Swift and Elixir this is done with Protocols. For our purposes I'll use interfaces, so translate as you need.

The goal of working with interfaces is to describe an ability of your application. In our case all we care about is that we can retrieve and save a record from a data store. It's important to keep this interface light because doing so will actively increase cohesion and drive down coupling:

```
public interface IDataStore {
    public void Save<T>(T item);
    public T Get<T>(int id);
    public IList<T> Fetch<T>();
}
```

This is a good start. We can now use our new interface:

```
public interface IDataStore
{
    void Save<T> (T item);
    T Get<T> (int id);
    IList<T> Fetch<T> ();
}

public class Membership
{
    IDataStore _db;
    public Membership (IDataStore db)
    {
        _db = db;
    }
    public User GetUser (int id)
    {
        //get the user
        return _db.Get<User> (id);
    }
    public void SuspendUser (int id)
    {
        var user = this.GetUser (id);
        user.Status = "suspended";
        _db.Save (user);
    }
}
```

Much better. We still have coupling to the notion of an **IDataStore**, but it's unavoidable at this point (unless we want to work directly with eventing, but that's probably overkill). We can now implement an **IDataStore** to do all kinds of things for us, such as:

- Store data in a relational system
- Store data in a NoSQL system
- Store data directly in memory for testing purposes

Using interfaces like this is a cornerstone of object-oriented programming. Injecting them, as we're doing here, is a great way to keep your code isolated.

It does come at a price.

Inversion of Control

As you build out your application, paying attention to interfaces and dependency injection, you will start to see the number of dependencies for a given class begin to spiral a bit out of control. The best way to see this is with

some code.

In the real world, our **Membership** class will probably need quite a few external dependencies:

- A hashing library for password storage
- An email library for sending a new user a note
- A rules module for accepting new users
- A logger module for logging

This means our constructor is going to grow:

```
public class Membership{
    IDataStore _db;
    IEmailer _email;
    ICrypto _crypto;
    ILogger _logger;
    IRulesEngine _rules;
    public Membership(IDataStore db,
                      IEmailer email,
                      ICrypto crypto,
                      ILogger logger,
                      IRulesEngine rules){
        _db = db;
        _email = email;
        _crypto = crypto;
        _logger = logger;
        _rules = rules;
    }
    public void Register(IRegisterable user){
        //validations etc
        if(_rules.CanRegister(user)){
            user.Status="Registered";
            user.HashedPassword = _crypto.HashPassword(user.Password);
            _db.Save(user);
            _email.SendWelcome(user);
            _logger.Info("New user added: " + user.Email);
            return user;
        });
    }
}
```

This is nuts. Every time we want to use **Membership** we'll need to create instances of its dependencies which, themselves, likely have dependencies of their own we'll need to create (and then inject). This is simply sweeping the dependency coupling somewhere else.

This is where Inversion of Control comes in. With Inversion of Control you have a separate mechanism (called a “container”) which is responsible for creating and injecting all the dependencies you need and then giving those injected objects to you when you need it.

Here is some pseudo code for an IoC container modeled after my friend Nate Kohari's excellent [Ninject Project](#):

```
//our app start
public void Main(){

    Container container = new Container();
    container.Bind<IMembershipStore>().To<PostgreSQLAdapter>();
    container.Bind<IEmailSender>().To<MailgunSender>();
    container.Bind<ILogger>().To<Log4Net>();
    container.Bind<ICrypto>().To<SuperCryptoThingy>();
    container.Bind<IMembership>().To<Membership>();

    //get an instance of Membership
    var membership = container.Get<IMembership>().InSingletonScope();

}
```

We have our interfaces mapped to concrete implementations in a single place. If we ever need to change anything, we just change it here.

When we need an instance, we simply need to access our container, which needs to be global to our application. The container then orchestrates the instantiation of the classes we want. A bonus to this is that we can set a “scope” on the object. For instance in the example above I’m setting the lifecycle to a Singleton.

You can do many other things with Inversion of Control containers – and they are quite useful.

Is This Really Useful?

You might be getting the sense that we’re creating a bit of a “meta” programming system here, where object instantiation is removed from the language constructs themselves and into this separate...mechanism of our own creation.

This is where we start getting subjective. There are quite a few developers out there who see patterns like the above as flaws in the language or, more broadly, as flaws in object-oriented programming itself. We started off this entire part of the book with a quote from Lawrence Krubner:

I have seen hyper-intelligent people waste countless hours discussing how to wire together a system of Dependency Injection that will allow us to instantiate our objects correctly. This, to me, is the great sadness of OOP: so many brilliant minds have been wasted on a useless dogma that inflicts much pain, for no benefit.

Now that we understand a bit more about dependency injection and inversion of control containers – do you think you’ll waste “countless hours”? To be honest: yes, *I have*. But it was my fault.

Keeping your containers happy and working properly is not as simple as it seems. As your application grows and becomes more complex, it becomes easier to find yourself creating circular dependencies. For instance: we might decide to create an **ILogger** implementation that saves logs to a database. We decide to reuse our **IDataStore**, which requires an instance of **ILogger** which requires an instance of **IDataStore**...

These problems, as you might be sensing, typically have to do with application design rather than object-oriented programming. Which seems to be a recurring problem in our industry.

If a language, platform or framework leads you down a snarled path of bad design, it’s usually your fault. Or is it?

I’d like to leave this chapter with a great quote from my friend Gary Bernhardt, which he offered during his amazing talk [The Birth and Death of JavaScript](#):

The behavior that you see a tool being used for is a behavior that tool encourages

It's easy to dismiss recurring structural problems as ignorance on the programmer's part. If the same problem occurs throughout the development community, however, is it really a problem of ignorance?

I don't have an answer.

SOLID

In object-oriented programming circles it's almost impossible to escape Uncle Bob's SOLID principles:

- Single Responsibility Principle (SRP)
- Open/Closed Principle
- Liskov Substitution Principle
- Interface Segregation Principle
- Dependency Inversion Principle

Uncle Bob is the primary name behind these principles, but Michael Feathers and Bertrand Meyer should also get some credit for naming and refining some of them.

Single Responsibility

A class should have a single responsibility to your application or, as Uncle Bob puts it, *a single reason to change*. This sounds simple, but it's a little tricky to grasp.

Most applications have a class for a **User**:

```
public class User
{
    public string Name { get; set; }
    public string Email { get; set; }
    public string Status { get; set; }
    public User ()
    {
        this.Name = "Guest";
        this.Status = "Anonymous";
    }
}
```

This class changes when the data about the **User** changes. But how does that information actually change?

Let's say we want to register the user into our system. We could do something like this:

```

public class User
{
    public string Name {get;set;}
    public string Email {get;set;}
    public string Status {get;set;}

    public User()
    {
        this.Name = "Guest";
        this.Status = "Anonymous";
    }

    public void Register(string name, string email)
    {
        this.Name = name;
        this.Email = email;
        this.Status = "Registered";
        //save to the DB or something else
    }
}

```

This class is now doing two things: describing a **User** based on some data and registering a user into the system. We can keep with SRP if we move the registration responsibility off to another class:

```

public class User
{
    public string Name { get; set; }
    public string Email { get; set; }
    public string Status { get; set; }
    public User ()
    {
        this.Name = "Guest";
        this.Status = "Anonymous";
    }
}

public class Membership
{
    public User Register (string name, string email)
    {
        //validations etc
        var newUser = new User { Name = name,
                               Email = email,
                               Status = "Registered" };
        //save to the DB or something else
        return newUser;
    }
}

```

Open/Closed

This principle may seem rather obvious to you, but at the time it was created and refined (late 80s, early 90s) it addressed a real problem.

Open/Closed says that a class or module should be open for extension, closed for modification. Or, put another way: let people override/extend your code without needing to modify it.

Let's say we're working in Node and we need to install a module from NPM. We run our npm install command and a node_modules directory appears with our module.

It turns out that we find a bug in one of the methods! We could go and fix the bug directly inside of node_modules, but that would violate Open/Closed! The good news is that the developer left the module extensible, so we can go in and override the bug with a fix. This module was open for our extension, but closed for our modification.

With object-oriented languages, this kind of thing is solvable if you allow users of your classes and modules to inherit and extend key bits of functionality that you're providing.

Liskov Substitution

Liskov is a subtle principle, but can catch some very serious bugs that creep into your application. The principle has to do with inheritance and how inheriting objects behave. It says:

if S is a subtype of T , then objects of type T may be replaced with objects of type S

Let's say we have a User and Administrator class:

```
public class User{
    //...
}

public class Administrator : User{
    //...
}
```

My program should work if I pass **Administrator** to any routine that expects a User.

The classic example of breaking LSP is the **Square** and **Rectangle** analogy:

```
public class Rectangle
{
    int _height;
    int _width;

    public virtual void SetHeight (int height)
    {
        _height = height;
    }
    public virtual void SetWidth (int width)
    {
        _width = width;
    }
}
public class Square : Rectangle
{

    public override void SetHeight (int height)
    {
        this.SetHeight (height);
        this.SetWidth (height);
    }
    public override void SetWidth (int width)
    {
        this.SetHeight (width);
        this.SetWidth (width);
    }
}
```

This code is obviously confusing in that we're coding our way around a problem in our implementation. While a square, in reality, is a rectangle (if we're talking math) – implementing it as a **Rectangle** in our code causes us to do some weird things.

Moreover, passing a **Square** around as if it were a **Rectangle** could cause some very strange things to happen in our video game code later on.

Interface Segregation

The Interface Segregation Principle is all about targeted, simple API creation so code is easy to use and implement. Rather than create a small set of large interfaces, favor a larger set of smaller, more generic interfaces.

Let's think about registering our **User** again:

```
public class User
{
    public string Name { get; set; }
    public string Email { get; set; }
    public string Status { get; set; }
    public User ()
    {
        this.Name = "Guest";
        this.Status = "Anonymous";
    }
}
public class Membership
{
    public User Register (string name, string email)
    {
        //validations etc
        var newUser = new User { Name = name, Email = email, Status = "Registered" };
        //save to the DB or something else
        return newUser;
    }
}
```

This works fine, but the notion of a **User** is bound to our **Membership** class. We could solve this with an interface (something like **IUser**), or we could lean on ISP and focus on what's really needed:

```

public interface IRegisterable
{
    string Name { get; set; }
    string Email { get; set; }
    string Status { get; set; }
}
public class User : IRegisterable
{
    public string Name { get; set; }
    public string Email { get; set; }
    public string Status { get; set; }
    public User ()
    {
        this.Name = "Guest";
        this.Status = "Anonymous";
    }
}
public class Membership
{
    public IRegisterable Register (IRegisterable user)
    {
        //validations etc
        user.Status = "Registered";
        //save to the DB or something else
        return user;
    }
}

```

We could extend this notion further with **IAuthenticatable** as well, if we want. Abstracting your code like this helps hide implementation details – which means making changes in the future becomes a lot easier.

Dependency Inversion

Dependency Inversion is all about loosening up the relationship between classes and modules in your code. The definition is a bit wonky:

- High-level modules should not depend on low-level modules. Both should depend on abstractions
 - Abstractions should not depend upon details. Details should depend upon abstractions
- You know, sometimes I wonder why programmers even bother trying to define things. Let's see if we can break this notion apart.

Our **Membership** module is a high level module, and we need to depend on some low-level modules in order for things to work.

For instance: we need to save our **IRegisterable** object to the database. We could go the route of creating a new database connection right in our **Register** code, if we want:

```

public interface IRegisterable
{
    string Name { get; set; }
    string Email { get; set; }
    string Status { get; set; }
}
public class User : IRegisterable
{
    //...implement interface
}
public class Membership
{
    public IRegisterable Register (IRegisterable user)
    {
        //validations etc
        user.Status = "Registered";
        //save to the DB or something else
        var db = new PostgreSQLAdapter ();
        db.Save (user);
        return user;
    }
}

```

But this is coupling or binding our **Membership** class directly to our **PostgreSQLAdapter**. This will solve our problem now, but in the future we will likely come across some problems:

- We might want to switch data access tools later on, maybe moving from our homespun routines to an ORM
- We might move to a document database, or something hosted (like Amazon's Dynamo DB)
- Our adapter might change the way it's constructed – perhaps moving to a Factory pattern or the like. Or maybe we'll have a Mediator...

In short: our **Membership** class knows way too much about the construction and execution of our adapter. Let's invert this dependency, shall we?

```

public interface IMembershipStore
{
    void Save (object item);
}

public interface IRegisterable
{
    string Name { get; set; }
    string Email { get; set; }
    string Status { get; set; }
}
public class PostgreSQLAdapter : IMembershipStore
{
    //implement interface
}
public class User : IRegisterable
{
    //implement interface
}
public class Membership
{
    IMembershipStore _store;
    public Membership (IMembershipStore store)
    {
        _store = store;
    }
    public IRegisterable Register (IRegisterable user)
    {
        //validations etc
        return _store.Save (user);
    }
}

```

Now we're free to change our storage approach however we like in the future. Our high level module doesn't depend on a lower-level one (our database bits) and the abstractions we're using don't depend on the details of the implementation.

Dependency Inversion vs. Dependency Injection

These are not the same thing, though they sound alike and, in many cases, look alike. Dependency Inversion is simply structuring our code to work with interfaces in a particular way (as described above). Dependency Injection is how these interfaces are provided to the classes that need them.

CHAPTER TWELVE: FUNCTIONAL PROGRAMMING

Functional Programming: some people love it and claim it's the only way to write software. Others see it as a fad and roll their eyes.

Let's come away from those extremes. Functional Programming is based in a foundational concept which we learned about a few chapters ago: *Lambda Calculus*. It's not magical, nor is it something you should ignore because you're amazing. It is simply something you need to understand.

Like most things in Computer Science, functional programming is full of jargon, idioms and practices that, at first, might be a bit opaque. If you take the time, however, to let it soak in... functional programming can change the way you write software if you come from an object-oriented background.



A Change In Thinking

I learned a functional language two years ago; one that I love using: *Elixir*. It has changed the way I think about programming. It's the only functional language I know, so I'll be using it to do some of the demos you're about to see.

For others, I'll be using JavaScript/ES6. I chose this language because 1) most programmers know it at least a little bit and 2) it has some functional characteristics - at least to the level where you can get your point across.

We'll cover four main topics in this chapter:

- Immutability (things that can't/don't change)
- Purity (functions rely only on what they're given)
- Side Effects (functions only operate on data they're given)
- Currying (breaking big functions down to little ones)

I also threw in a very brief discussion about functors and monads at the very end, something I barely understand but I think is critical to have a think on.

Off we go...

Immutability

The first word that comes to mind whenever you hear "functional programming" is usually "immutable". As I am sure you know, it means "not changeable" in plain English, but how does that translate to programming? Moreover: *who cares?*

If you're an object-oriented (OO) programmer you're used to creating classes and then instantiating them. You set their properties and send them messages which change their properties or tell you something about themselves. This is not possible with functional programming.

There are no objects, no classes, no "state" as you might think of with OO. There are only functions which transform things. You give a function what it needs, it hands you back what you want (hopefully). This might sound

incredibly confining, but there are practices that go along with this idea that make it rather compelling. Let's take a look.

If you want to follow along, you should have Elixir (version 1.3 or so) installed, as well as Node (version 6+).

Simple Immutability with Elixir

Consider this assignment:

```
friend = %{name: "Clara"}  
IO.inspect friend  
#%{name: "Clara"}
```

This is a **map** in Elixir and it looks a lot like an object in JavaScript or a hash in Ruby. The difference is you can't do this:

```
friend = %{name: "Clara"}  
friend.name = "Mike"  
  
*** (CompileError): cannot invoke remote function friend.name/0 inside match  
[Finished in 0.159s]
```

The error message is a bit odd as it has to do with the way Elixir tries to match values (pattern matching), but essentially it means “you can't change Clara's name to Mike”. So what do we do?

We ask the **Map** library to update the map. This is going to look a bit weird, but once I explain it more you'll hopefully understand:

```
friend = %{name: "Clara"}  
friend = Map.put(friend, :name, "Mike")  
IO.inspect friend  
#%{name: "Mike"}
```

I do not blame you if you're completely underwhelmed at this point... and more than a little confused. There's just no graceful way to go about this, so I'm throwing you in the deep end straight away.

Here's what just happened:

- We created a **map** with a name key set to “Clara”
- We asked the **Map** library to put the value “Mike” in Clara's place
- The **Map** library gave us back a *completely new map* ... sort of
- The **friend** variable was *rebound* to the new **Map**

Elixir provides some helpful features which you don't find in more “strict” functional languages. For instance: in Erlang (a functional language that Elixir is based on) you can't do what we just did (rebinding **friend**). You would need to create a whole new variable - something like **renamed_friend** to hold the result of the **Map.put/2** operation. Elixir is less strict, so it *looks* like we're changing the **friend** variable.

A natural thought that you might be having at this point is *wait a minute, if I need to update a map, I need a whole new one? Isn't this horrible for memory?* This is a very good question! The short answer is: *no, it's not*. This is because Elixir simply uses a pointer under the covers. The initial **friend** map is still there and the updated map points back to it. This keeps things rather light and you don't fall into the immutability traps that you have in other languages (strings with .NET, for example). Furthermore: this kind of thing is only possible with a functional

language.

That's the technical teardown, let's discuss something a bit more intangible: *this code is just ugly*. I would agree with you there too. Fortunately what I wrote is *not* idiomatic. I said at the beginning that functional programming is all about *transforming* data through a set of functions. Let's rewrite our code to support that idea:

```
%{name: "Clara"} |> Map.put(:name, "Mike") |> IO.inspect
```

Much better. That toothy thing is the “pipe” operator and works in much the same way as its Unix counterpart. The result of a given function is piped into a following one as its first argument, and so on.

Before we end our introduction to immutability, I want to underscore the notion of rebinding. Let's change our code a bit, and I'll reintroduce our **friend** variable:

```
friend = %{name: "Clara"}
friend |> Map.put(friend, :name, "Mike") |> IO.inspect
IO.inspect friend
```

Running this you'll see that our **friend** variable wasn't changed at all. **Map.put/2** simply passed back a new map with the updated key, which we piped into **IO.inspect** directly. Thus we have two different outputs:

Formalizing Data with Structs

I don't want to leave you with the impression that data in a functional language like Elixir is just tossed around without any formalized rules. There are no classes in Elixir, but their close relatives, “Structs”, do provide some of their utility.

Let's formalize our code, making our intentions clearer and giving our friends some structure:

```
defmodule Friend do
  defstruct name: "Clara", age: 0
end

defmodule Immutability do
  def change_name(friend, new_name) do
    Map.put(friend, :name, new_name)
  end
  def get_friend do
    %Friend{}
  end
end

Immutability.get_friend |> IO.inspect
```

This looks better don't you think? Structs allow you to set defaults for your data, and give it a prescribed structure. Elixir also gives you some shorthand for updating a struct:

```
...
friend = Immutability.get_friend
%{friend | name: "Mike"} |> IO.inspect
IO.inspect friend
```

Once again, we get the same result as we did with **Map.put/2** above as the rebinding returns an immutable new

Friend struct:

```
{name: "Mike"}  
{name: "Clara"}  
[Finished in 0.22s]
```

Transforming Data

If you're new to functional programming ideas than none of this is probably convincing. It takes a while to get into the functional groove, so let's write some more code.

Three things to focus on for this section are:

- We are transforming data by passing it through a set of functions
- We like smaller functions
- We can treat functions the same as values

The latter comes straight from Lambda Calculus. In fact all of this does - that's where functional programming has its roots in case that wasn't obvious. In addition, things work better overall if we focus on smaller, clearer functions.

To see this in action, let's do some math, shall we?

```
defmodule Ops do
  def square_it(num), do: num * num
  def double_it(num), do: num + num
  def root_it(num), do: :math.sqrt(num)
  def print_it(num), do: IO.inspect num
end

4 |> Ops.square_it
|> Ops.double_it
|> Ops.root_it
|> Ops.print_it

# 5.656854249492381
# [Finished in 0.274s]
```

What we have done here is fairly typical in the functional world. Small, concisely-named functions are arranged as needed, and we can simply pass some data along. This type of approach can work with anything, you just have to change your thinking a bit.

Discussion: Why Would You Do Any of This?

Right now you might be thinking about how certain things might be solved with functional programming, maybe translating problems you work on every day? It's an interesting thing to ponder, but from my experience it takes a good two weeks to really hit that AHA! moment... at least for me.

What you need is a *reason to care about it all* in order to pull you through the process of changing your thinking. Let's see if I can help with that.

First: *no, functional programming is not perfect nor the answer to everything*, of course. But it does greatly expand your ability to think through a solution!

The first advantage of functional programming is its immutability. Quite a few bugs in OO are caused by the

state of something not being correct, or what you wanted. You might have quite a few tests that show that yes, indeed, this code *should work* provided X, Y and Z conditions are met, but then condition Σ comes along and screws the whole thing up!

If you write code that depends on anything going on outside of it, you're prone to these types of errors. It's the same reason most developers loathe global variables - one change and it ripples throughout your program, causing things to break.

Functional programming is different in that, ideally, a function should receive everything it needs in order to do its job, when asked. There are all kinds of safeguards you can attach to these functions to prevent them from being called if the data isn't correct. As opposed to being a burden, it's quite freeing! In order to get it right, however, you have to change your thinking, which takes time and effort.

Speaking of effort, let's get back to it. We'll pick up this conversation again later.

A Real Example: A Shopping Cart

Let's compare and contrast a functional style vs. an OO style, and I'll do that by creating a **ShoppingCart** in Elixir and also in JavaScript.

Let's start with JavaScript:

```
class Cart{
  constructor(){
    this.items = [];
  }
  addItem(item){
    //make sure it's a proper item
    //and then...
    this.items.push(item);
  }
}
const cart = new Cart();
cart.addItem({sku: "SOCKS", price: 12.00})
cart.addItem({sku: "MORESOCKS", price: 18.00})
```

Lovely. A very, very simple cart but it captures the idea: we have a class, an instance and we change the state of our instance by adding items to it.

Elixir is a bit different. As a first pass you might be tempted to do something like this:

```
defmodule Cart do
  items = []
  def add_item(item) do
    %Cart{items: item}
  end
end
```

Which wouldn't work. Functional languages don't have the notion of state, so holding onto an **items** array would cause an error. We need to pass everything to the **Cart** that it needs to operate, *including the items*:

```

defmodule Cart do
  def add_item(items, item) do
    items ++ item
  end
end

items = []
Cart.add_item(items, %{sku: "SOCKS", price: 12.00}) |> IO.inspect
IO.inspect items

# %{sku: "SOCKS", price: 12.00}
# []

```

This works and, once again, underscores the notion of rebinding. Let's tweak this to look more like functional programming:

```

defmodule Cart do
  def add_item(items, item) do
    items ++ item
  end
end

[] |> Cart.add_item(%{sku: "SOCKS", price: 12.00}) |> IO.inspect

# %{sku: "SOCKS", price: 12.00}

```

Much cleaner. Still a bit... wonky however. We're passing around arrays and maps without knowing what's going on. Let's clean this up a bit.

When you're forced to do without state and objects, you begin to think in terms of data and things that happen to that data. In our example we're using OO thinking by trying to represent a **Cart** as an object, which isn't very functional of us. Instead we should think about it in terms of data moving through a *process*.

The **Cart** is our data and our process is, more accurately, described as **Shopping**:

```

defmodule Cart do
  defstruct items: [], total: 0, count: 0
end

defmodule Shopping do
  #use pattern matching to guarantee data we need
  #the _ ignores the data, we just want the pattern
  def add_item(%Cart{} = cart, %{sku: _, price: price} = item) do
    %{cart | items: cart.items ++ [item], total: cart.total + price, count: cart.count + 1}
  end
  #get from the DB?
  def get_cart, do: %Cart{}
  #save to DB?
  def save_cart(cart), do: cart
end

# add an item to the cart
Shopping.get_cart
|> Shopping.add_item(%{sku: "SOCKS", price: 12.22})
|> Shopping.save_cart
|> IO.inspect

```

I threw a bit more Elixir at you in this go, hope you don't mind. Of particular note is how I'm able to use pattern matching to guarantee the data this function needs, which is a **Cart** and a map with a **sku** and a **price**.

From end to end, I'm able to retrieve a **Cart** from somewhere (assume it's a database for now), add an items to it and then put it back. It's clear what this process is by reading the code and the only change of state is that of our database, which is acceptable in functional programming realms. Barely.

Side Effects and Purity

There are two terms that you hear often in discussions about functional programming: *purity* and *side effects*. Both of these terms stem from the idea of immutability.

The whole notion of interacting with a system outside the scope of the function you're in is called a "side effect" - something that happens as a result of your function being invoked. Working with a database, for instance, is referred to as a "necessary side effect" because you're changing the state of something outside the scope of your function.

The more you do this, the less "pure" your code is. *Purity* is not a term dedicated to functional programming - it refers to the level of interaction any code has with the outside world. You can write "pure" code in OO programming just as you can with functional, it just happens to have a little more focus in the functional world.

Functional programmers like purity. It's easier to test a function that doesn't change behavior based on some external setting or function call. It's also easier to debug. The downside is that you end up with more functions that are (typically) a lot smaller than what you might write in OO land.

There's a way to work with those, too, as we're about to see.

Currying

Functions, functions, functions. They're everywhere! Organizing a program full of them can feel overwhelming, especially when you're just starting out with functional programming.

Let's take a look at one of the very first practices you'll want to take advantage of: *currying*. Currying is the act of using smaller, single arity functions in a chain rather than a larger function with multiple/complex arguments. It's easier to understand using code.

Consider date night with your partner:

```
const dateNight = (who, what, where) => {
  return `Out with ${who} having fun ${what} at ${where}`;
};
```

This function takes 3 arguments (or, in more programmy speak, has an arity of 3). If we split these functions into a smaller set of chained, single arity functions, we're currying:

```
const nightOut = who => what => where => {
  return `Out with ${who} having fun ${what} at ${where}`;
};
```

Looks a bit weird doesn't it? Especially if you're not used to lambdas. But how might we use this function? Like this:

```
const funTime = nightOut("Dancing")("wife")("Club 9");
console.log(funTime);
//Out with wife having fun Dancing at Club 9
```

Does this look familiar to you? It should, this is the *exact* thing we did when applying values to functions in Lambda Calculus a few chapters back.

I know what you're thinking: *calling functions like this is a bit goofy*, and I would agree. This is where the idea of *partial application* comes in:

```
const dancing = nightOut("Dancing");
const dancingWithWife = dancing("wife");
const funTime = dancingWithWife("Club 9");
console.log(funTime);
//Out with wife having fun Dancing at Club 9
```

This is the power of currying. Let's see something a bit more concrete, however.

A Curried SELECT Query

I worked with a very nice, very opinionated developer recently who loves her object-oriented programming. I was trying to explain how I thought it would be fun to add some more functional ideas to the data access code we were writing, and she was rather resistant. Finally she said "just show me what you mean".

So I did. I started out with a curried function set for building a basic select query:

```
const selectQuery = table => (where, params) => order => limit => {
};
```

At this point I simply needed to evaluate what was sent in to each function call:

```
const selectQuery = table => (where, params) => order => limit => {
  const whereClause = where ? ` where ${where}` : "";
  const orderClause = order ? ` order by ${order}` : "";
  const limitClause = limit ? ` limit ${limit}` : "";
  if(params.length > 0) params = [params];
  const sql = `select * from ${table}
    ${whereClause}${orderClause}
    ${limitClause}`; //wrapped this for readability
  return {sql: sql, params: params};
};
```

I've written a lot of code for building SQL statements, and this is probably the smallest I've ever created. Using this is even more fun:

```
const usersQuery = selectQuery("users");
console.log(usersQuery()());
//{sql: "select * from users", params: []}
```

You probably wouldn't want to create a query in that way, however, and this is where partial application comes in. You could create a module (let's call it a repository for fun) which built up the following:

```

const usersQuery = selectQuery("users");
const allUsersQuery = usersQuery();
const usersByEmail = email => usersQuery("email = $1", email)();
console.log(usersByEmail("test@test.com"));
//{sql: "select * from users where email = $1", params: ["test@test.com"]}

```

I think this is rather interesting, and a very nice way to reuse functionality. Let's round this out by plugging in [pg-promise](#) and having it actually do something. I'll be using the [Chinook test database](#) for this:

```

const pgp = require('pg-promise')();
const db = pgp("postgres://localhost/chinook");
//our selectQuery as before
const selectQuery = table => (where, params=[]) => order => limit => exec => {
  const whereClause = where ? ` where ${where}` : "";
  const orderClause = order ? ` order by ${order}` : "";
  const limitClause = limit ? ` limit ${limit}` : "";
  if(params.length > 0) params = [params];
  const sql = `select * from ${table}${whereClause}${orderClause}${limitClause}`;
  const query = {sql: sql, params: params};
  //should we execute?
  if(exec){
    db.many(sql, params).then(res => exec(res)).catch(err => console.log(err))
  }
  //if not, just return the query
  else return query;
};

const albumQuery = selectQuery("album");
const albumSearch = title => albumQuery(`title LIKE $1`, `%${title}%`)();
const rockAlbums = albumSearch("Rock");
//fire it up!
rockAlbums(res => {
  console.log(res);
})

```

If you're playing along and run this code, you should see all the albums in the Chinook database with the term "Rock" in the title.

When I showed this code to my friend, her response was wonderful:

oh my god, that's equal measures ingenious and horrifying :D Being able to reuse selectQuery, albumQuery, etc almost seems like a twisted take on models at an arbitrary scope...

Twisted indeed. Personally I quite like this approach!

A Very Brief Discussion About Functors and Monads

There's a great line from Douglas Crockford about monads:

In addition to it being useful, it is also cursed and the curse of the monad is that once you get the epiphany, once you understand - "oh that's what it is" - you lose the ability to explain it to anybody.

From ["Monads & Gonads"](#), December 2012

Very, very true. That's why I want to keep this discussion as brief as possible and invite you to do your own investigation. You should have a little bit of an itch in your brain right now, wondering *just how am I supposed to orchestrate things beyond simply currying my way to insanity?*

This is where structures like *functors* and *monads* (among others) come in.

When you're slinging functions around as we have, you might want to have a higher level of abstraction for working with them. Something that might *wrap* those functions and handle certain situations for you the way you want.

One thing in particular that you might want to do is to interrogate/iterate over some values, such as an array or a struct like a **user**. You could write a loop or do some type of recursion, or you could orchestrate the effort using a set of functions.

A popular way of doing this is to create a function which allows you to map values from a given object:

```
class Monkey{
  constructor(val){
    this.__value = val;
  }
  map(fn){
    return fn(val);
  }
}
```

This is pretty simplistic, but hopefully you get the idea. This simple construct has a pretty lofty name: it's a *functor*. It's only purpose in life is to run **map** over things that are *mappable*.

Functors have big brothers that will implement some logic on top of that mapping, namely returning some alternate values based on the data contained in your mappable object. For instance you might want to access a **name** field on your **user** object, but don't want bad things to happen if there is no **name** field.

This is where monads come in, and where I start gracefully backing my way to the door. They're easy to understand in concept, and when you first encounter them it can be a little anticlimactic. I don't know why this is. Monads really aren't terribly scary, but it's like some beautiful gateway to programming hell: if you understand it you're doomed. I've found Crockford's rule above to be spot on.

Let's see if you agree with me.

Assume that we're allergic to **if** statements. This is one lovely oddity that happens when you start working with functional languages. There are so many interesting constructs and ways of doing things (like pattern matching) that you find yourself avoiding complex conditional trees.

Our **selectQuery** is pretty good and I feel happy about it, but we wouldn't be allowed into Functional Club with this code:

```
const selectQuery = table => (where, params = []) => order => limit => exec => {
  const whereClause = where ? ` where ${where}` : "";
  const orderClause = order ? ` order by ${order}` : "";
  const limitClause = limit ? ` limit ${limit}` : "";
  if(params.length > 0) params = [params];
  const sql = `select * from ${table}${whereClause}${orderClause}${limitClause}`;
  const query = {sql: sql, params: params};
  // ...
};
```

Monads exist to handle just this case! Let's create the simplest one: the **Maybe** monad:

```

class Maybe{
  constructor(val){
    this.__value = val;
  }
  isNothing(){
    return (this.__value === null || this.__value === undefined);
  };
  map(fn){
    return this.isNothing() ? Maybe.of(null) : Maybe.of(fn(this.__value));
  };
  val(){
    return this.isNothing() ? "" : this.__value;
  };
}
Maybe.of = (val) => new Maybe(val);

```

That's not so scary is it! See what I mean? Ah the temptation... let's see if we can confuse ourselves a bit.

This class exists to do one thing: *handle conditional values* in an orchestrated, functional way. When you call **map()** you say “here’s a function, use it if there’s a value”. The trick, however, is that if there *is* a value for your monad then **map** hands you back *another* **Maybe** monad with its value set as the result of your mapped function call.

This means we can move data through a chained set of functions without worrying about **null** or **undefined** values. How fun.

We can now use our **Maybe** monad to evaluate our **whereClause**:

```
const whereClause = Maybe.of(where).map(w => `where ${w}`).val();
```

If **where** has a value then our lambda will be called and we’ll get back another **Maybe** with its current value set to a where clause. We can then call **val()** if we just want this value directly, which we do. If where isn’t set, we’ll get back an empty string.

We can apply this to the other arguments as well:

```
const whereClause = Maybe.of(where).map(w => `where ${w}`).val();
const orderClause = Maybe.of(order).map(o => `order by ${order}`).val();
const limitClause = Maybe.of(limit).map(l => `limit ${limit}`).val();
```

We could, at this point, build our SQL statement as before by building a string with template literals, but where’s the fun in that! Let’s use our **Maybe** monad one more time:

```
const sql = Maybe.of(`select * from ${table}`)
  .map(sql => `${sql} ${whereClause}`)
  .map(sql => `${sql} ${orderClause}`)
  .map(sql => `${sql} ${limitClause}`)
  .val();
return {sql: sql, params: params};
```

Let’s step through this. We start out with an initial value that we want to work with, in this case it will be our **select * from \${table}** statement. We move this value through a successive chain of functions, where it is applied to a function that is passed in via **map**. That value is then passed on to the next in the chain. If you’re thinking that this looks a lot like Elixir’s **|>** (pipe) operator, you’d be correct.

Note: if you’re a functional programming person and you see a better way to do this, please do let me know!

Here's the final code:

```
const selectQuery = table => (where, params = []) => order => limit => {

  const whereClause = Maybe.of(where).map(w => `where ${w}`).val();
  const orderClause = Maybe.of(order).map(o => `order by ${o}`).val();
  const limitClause = Maybe.of(limit).map(l => `limit ${l}`).val();

  const sql = Maybe.of(`select * from ${table}`)
    .map(sql => `${sql} ${whereClause}`)
    .map(sql => `${sql} ${orderClause}`)
    .map(sql => `${sql} ${limitClause}`)
    .val();

  return {sql, params};
};

const albumQuery = selectQuery("albums");
const albumSearch = title => albumQuery("title LIKE $1", `%${title}%`());
const rockAlbums = albumSearch("Rock");
console.log(rockAlbums());
//{
//  sql: 'select * from albums where title LIKE $1  ',
//  params: '%Rock%'
// }
```

Wahoo!

Going Further With Monads

If you're interested in reading more about monads, I found a wonderful blog post from [James Sinclair](#) which is not only well-written, it's also gorgeous! I was able to stay with the explanation up until the last third of the article, which I've reread about 12 times now.

I also found an excellent example [on Github](#) which I found approachable and easy to read. I asked the author, Kim Hogeling, if it would be OK if I included these examples as I thought they were outstanding and he graciously agreed.

Here they are. Consider the task of doing some simple math on numbers in an array, like this one:

```
const sales = [183, 229, 331, 453];
```

This is all well and good, but if you're using an array that doesn't have type enforcement, you might end up with something like this:

```
const sales = [183, 229, 331, 453, "NOTHING"];
```

This would break things badly and would likely cause an error. By wrapping your functionality with a monad, however, you can gracefully handle this situation. Here's Kim's monad for handling this situation:

```

const isSomething = a => {
  return (
    a !== null &&
    a !== undefined &&
    (typeof a !== 'number') ||
    (typeof a === 'number' && !isNaN(a))
  );
}

class Maybe {
  constructor(a) {
    return isSomething(a) ? new Just(a) : new Nothing()
  }
}
//added by me because it's standard - RC
Maybe.of = (val) => return new Maybe(val);
class Just {
  constructor(a) {
    this.val = a
  }
  value() {
    return this.val
  }
  // I named the Left's and Right's functor .fmap
  // instead of .map to prevent confusion with Array.prototype.map
  fmap(fn) {
    return new Just(fn(this.val))
  }
}

```

The **Maybe** monad works with the **Nothing** monad:

```

class Nothing {
  value(fallbackValue) {
    return fallbackValue
  }

  fmap(fn) {
    return this
  }
}

```

These are used together with the **Either** monad:

```

class Either {
  constructor(a, b) {
    // take value out of functors before rewrapping into Left/Right
    if (a && typeof a.value === 'function') {
      a = a.value(null)
    }
    if (b && typeof b.value === 'function') {
      b = b.value(null)
    }
    return isSomething(a) ? new Left(a) : new Right(b)
  }
}

class LeftOrRight {
  constructor(a) {
    this.val = a
  }
  value() {
    return this.val
  }
}

class Left extends LeftOrRight {
  fmap(fn) {
    return new Left(fn(this.val))
  }
}

class Right extends LeftOrRight {
  fmap(fn) {
    return new Right(fn(this.val))
  }
}

```

We can use these monads together to run calculations without having to write a bunch of conditional logic and validation checks:

```

const times = nums => nums.reduce((a, b) => a * b, 1)
const double = a => a * 2
const quarter = a => a / 4
const pretty = a => `The value is ${a}!`

// Composing the functions together
const timesDoubleQuarterPretty = arr =>
  Maybe.of(times(arr))
    .fmap(double)
    .fmap(quarter)
    .fmap(pretty)

// a bad array
const broken = [1, 2, 'x', 3, 4]
const result1 = timesDoubleQuarterPretty(broken)
// Nothing

// A better array
const intact = [1, 2, 3, 4]
const result2 = timesDoubleQuarterPretty(intact)
// Just 'The value is 12!'

// And here we can simply use Either to pick the valid one
// Either returns the Left or Right functor which both have the .value method
const result = Either(result1, result2) // Right 'The value is 12!'
result.value() // 'The value is 12!'

```

If nothing else, hopefully you understand monads at a very, very high level. If not, please don't give up! I encourage you to follow the links above and write some code; see if your brain changes shape a little.

CHAPTER THIRTEEN: DATABASES

There is an established way to design a transactional database: following the rules of normalization. This is the process of essentially turning a single, large spreadsheet into a set of related tables.

Transactional systems can back our applications, but analytical systems power decision making. Data warehouses and data marts store the data and OLAP systems provide the analysis infrastructure.

Distributed database systems use the notion of *horizontal scaling* (adding more machines) vs *vertical scaling* (increasing machine size and power). Distributed systems are much different than traditional “big machine” databases in that they have to balance tradeoffs with data consistency, availability and network problems.

Social media is driving the need to store gigantic, petabyte-sized data stores. This movement is called Big Data and it's often difficult to grasp the sheer size of it all.



This chapter might confuse you. What do databases have to do with computer science?

I suppose the correct answer to that question is *not much*. The real-world answer is the opposite: *you need to know this stuff*.

I was lucky enough to get to know databases well when I started coding, and I think it's contributed directly to whatever success I've had in my career. Your application is grand and I'm sure the code you right is amazing; but *the database is where the gold is*.

Every time I mention that to fellow programmers get a mixed reaction. Some agree wholeheartedly, others laugh in my face and condescendingly explain IP and how VCs don't fund databases.

A business is built on customers. Customers drive sales, sales pay your wages. There really is nothing else! Your code will probably be replaced in the next 3-5 years, but the data will remain – likely the database system too!

I know far too many developers that don't know the first thing about database design, SQL, and the various systems out there. This chapter is for them.

Normalization

A relational database consists mainly of a bunch of tables. These tables contain rows of data, organized by columns. There has to be a method to this madness, and it's called normalization.

Database normalization is all about controlling the size of the data as well as preserving its validity. Back in the 70s and 80s you simply did not have disk space, but companies did have a ton of data that would fill it up. Database people found that they could reduce the size of their database and avoid data corruption if they simply followed a few rules.

Before we get to these rules, let me just add upfront that it's truly not that complicated once you grasp the main ideas. As with so many things computer science related, the jargon can be rather intense and off-putting, making simple concepts sound hard. We'll slowly work our way up to it.

Finally: rules are meant to be broken. In fact a DBA will break normalization quite often in the name of performance. I'll discuss this at the end.

First Normal Form (1NF): Atomic Values

We've decided to buy a food truck and make tacos and burritos for a living. We have our ingredients and drive up to our favorite street corner, opening our doors for business.

When the orders come in we put them into a spreadsheet – knowing we'll need to deal with it later on:

email	name	order_id	items	price
joe@example.com	Joe Tonks	1	Pollo Burrito, Diet Coke	\$12.50
jill@example.com	Jill Jones	2	Carne Asada, Sprite	\$14.50

Let's move this spreadsheet into the database, altering it as we go.

First normal form (1NF) says that values in a record need to be atomic and not composed of embedded arrays or some such. Our items are not atomic – they are bunched together. Let's fix that.

EMAIL	NAME	ORDER-ID	ITEMS	TOTAL
JOE @ EXAMPLE.COM	JOE TONKS	1	POLO BURRITO DIET COKE	\$10.50
JILL @ EXAMPLE.COM	JILL JONES	2	CARNE ASADA Sprite	\$14.50



EMAIL	NAME	ORDER-ID	ITEMS	TOTAL
Joe @ EXAMPLE.com	Joe Tonks	1	Pollo Burrito	\$8.00
Joe @ EXAMPLE.com	Joe Tonks	1	Diet Coke	\$2.50
Jill @ EXAMPLE.com	Jill Jones	2	Carne Asada	\$12.00
Jill @ EXAMPLE.com	Jill Jones	2	Sprite	\$2.50

Lovely. Our new orders table is 1NF because we have atomic records. You'll notice that things are repeated in there, which we'll fix in a bit. Our next task is to isolate the data a bit.

Second Normal Form (2NF): Columns Depend On a Single Primary Key

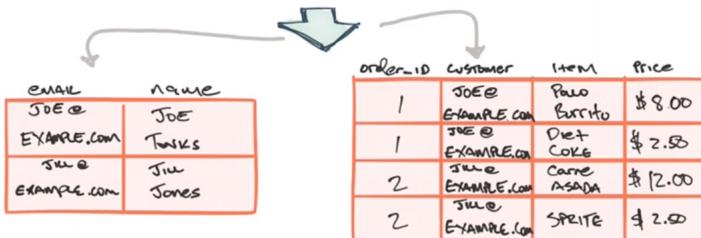
Now that we're in 1NF, we can move on to 2NF because part of being in 2NF is complying with 1NF. Our task to comply with 2NF means that we need to identify columns that uniquely define the data in our table.

An order is a customer buying something. In our case the email field uniquely identifies a customer, and the order_id field uniquely identifies what they've ordered. You put that together and you have a sale – which could

uniquely identify each row in our table.

The problem we have is that name does not depend on the order_id and items and price have nothing to do with email – this means we're not in 2NF. To get to 2NF we need to split things into two tables:

email	name	order_id	item	total
joe@example.com	joe tonks	1	pollo burrito	\$ 8.00
joe@example.com	joe tonks	1	diet coke	\$ 2.50
jill@example.com	jill jones	2	carne asada	\$ 12.00
jill@example.com	jill jones	2	sprite	\$ 2.50



Much better, but not quite there. Our customers table looks good but our orders table is a bit of a mess.

Third Normal Form (3NF): Non-keys Describe The Key And Nothing Else

3NF says that every non-primary field in our table must describe the primary key field, and no field should exist in our table that does not describe the primary key field.

Our orders table has repeated values for the key (which is a no-no) and the price of each item has nothing to do with the order itself. To correct this we need to split the tables out again:

The diagram illustrates the process of normalizing data. It starts with two flat tables:

email	name
JOE@EXAMPLE.COM	JOE TONKS
JILL@EXAMPLE.COM	JILL JONES

order_id	customer	item	price
1	JOE@EXAMPLE.COM	POLLO BURRITO	\$ 8.00
1	JOE@EXAMPLE.COM	DIET COKE	\$ 2.50
2	JILL@EXAMPLE.COM	CARNE ASADA	\$ 12.00
2	JILL@EXAMPLE.COM	SPRITE	\$ 2.50

A pink arrow points from these tables down into a cloud containing the text "3NF". Below the cloud, the data is shown in four normalized tables:

The four normalized tables are:

SKU	PRICE	NAME
PB	\$ 8.00	POLLO BURRITO
DC	\$ 2.50	DIET COKE
CA	\$ 12.00	CARNE ASADA
S	\$ 2.50	SPRITE

Customer	email	name
1	JOE@EXAMPLE.COM	JOE TONKS
2	JILL@EXAMPLE.COM	JILL JONES

ID	ORDER_ID	SKU
1	1	PB
2	1	DC
3	2	CA
4	2	S

ID	CID	DATE
1	1	04/03/2016
2	2	04/03/2016

Relationships are indicated by arrows: an arrow from the 'Customer' table to the 'order_items' table, and another arrow from the 'order_items' table to the 'orders' table.

We now have 4 tables to support our notion of an order:

- **customers** holds all customer data
- **orders** holds meta data related to a sale (who, when, where, etc)
- **order_items** holds the items bought
- **products** holds the items to be bought

You'll notice, too, that I replaced email with an integer value, rather than using the email address. I'll explain why in just a second.

In The Real World: The Normalization Process

Normalizing a database requires some practice. As programmers, hopefully you understand how to model classes and objects. It's almost the same process that we just went through: *what attributes belong to which concept?*

It's at this point that I get to tell you (with a sinister giggle) that the rules of normalization are more of a guideline, not necessarily law. A well-normalized database may be theoretically sound, but it will also be kind of hard to work with.

We managed to move a fairly simple spreadsheet with 2 rows and 5 columns into a 4 table structure with 3 joins and multiple columns! We only captured a very, very small fraction of the information available to us and our Taco Truck.

It's very easy to build a massively complex database with intricate lookups, foreign keys, and constraints to support what appear to be simple concepts. This complexity presents two problems to us, right away:

- Writing queries to read and write data is cumbersome and often error-prone
- The more joins you have, the slower the query is

The bigger the system gets, the more DBAs tend to cut corners and denormalize. In our structure here it would be very common to see a total field as well as item_count and embedded customer information.

To show you what I mean – have a look at the structure for the StackOverflow posts table:

Name	Type	Length	Decimals	Dimen...	Not Null	Key
id	int4	0	0	0	<input checked="" type="checkbox"/>	
post_type_id	int4	0	0	0	<input type="checkbox"/>	
accepted_answer_id	varchar	255	0	0	<input type="checkbox"/>	
parent_id	int4	0	0	0	<input type="checkbox"/>	
created_at	varchar	255	0	0	<input type="checkbox"/>	
deleted_at	varchar	255	0	0	<input type="checkbox"/>	
score	int4	0	0	0	<input type="checkbox"/>	
view_count	varchar	255	0	0	<input type="checkbox"/>	
body	text	0	0	0	<input type="checkbox"/>	
owner_user_id	int4	0	0	0	<input type="checkbox"/>	
owner_display_name	varchar	255	0	0	<input type="checkbox"/>	
last_editor_id	varchar	255	0	0	<input type="checkbox"/>	
last_editor_display_name	varchar	255	0	0	<input type="checkbox"/>	
last_edit_date	varchar	255	0	0	<input type="checkbox"/>	
last_activity_date	varchar	255	0	0	<input type="checkbox"/>	
title	varchar	255	0	0	<input type="checkbox"/>	
tags	varchar	255	0	0	<input type="checkbox"/>	
answer_count	varchar	255	0	0	<input type="checkbox"/>	
comment_count	int4	0	0	0	<input type="checkbox"/>	
favorite_count	varchar	255	0	0	<input type="checkbox"/>	
closed_date	varchar	255	0	0	<input type="checkbox"/>	
community_owned_date	varchar	255	0	0	<input type="checkbox"/>	

Figure 11.5: StackOverflow posts table.

Notice the `view_count`, `owner_display_name`, and `_count` fields? Also the `last_editor_id` field?

The count fields are calculated and don't belong in a table, theoretically speaking. The `owner_display_name` and `last_editor_id` fields should be foreign keys that link to an **authors** or **users** table of some kind – and they do with `last_editor_id` and `owner_id`. Querying this massive table using the required joins, however, would be way too slow for what they need.

So they denormalized it. Many businesses do – it just makes things faster and simpler.

Primary Key Selection

In our example here I decided to change the primary key of our customers table from email to an integer field called id. Why?

As the database grows, we're going to add tables. Many of these tables will reference our customers table using the email primary key. The net result of this will mean our customer's email address sprinkled throughout the database.

When they go to change their email address, an updating nightmare ensues. I made this exact mistake once, and I knew better!

So: choose a primary key that will not change or has little to no semantic meaning.

The other problem with having `email` as the primary key is that the database engine will index that field and will try to order it (this is called a “clustered” index which we’ll talk more about later). Indexes take time to maintain and they also take up space.

This is why you see so many integer-based keys.

Preparing For Distribution

In large organizations (or super-hot big time startups), there comes a time when “spreading the load” needs to happen, for one reason or another. You might want a redundant and highly available database, so you would choose replication (mirroring data between two servers).

You might want super-fast response times so you would shard your database based on the tables soaking up the most data. For instance StackOverflow would probably shard their posts over 4 or 5 different servers – each server holding a subset of the data.

For either situation, an integer-based key could cause problems. For these cases a GUID (Globally Unique Identifier) is usually a good choice – but it’s hard on a database. GUIDs aren’t sequential in nature, so a clustered index is meaningless. They’re also very large, so they take up space.

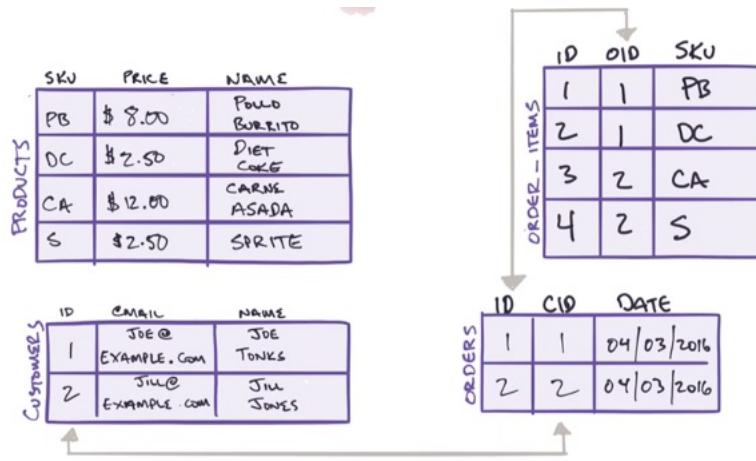
Not All Database Engines

If you run SQL Server you’ll likely want to tell me that you can do a new `sequential_id` as a default GUID – and that does work. SQL Server also has a lovely, crafty way of dealing with GUIDs.

What some companies have started doing is to have a centralized “identifier” service – something like [Twitter Snowflake](#). These are sequential, long integers that will grow for a very long time.

More Real World: Is This Schema Correct?

Let’s take a look at the final schema we came up with:



While it is theoretically correct, there is a problem with being historically correct. For instance, if you come to my Taco Truck and buy some Carne Asada, I’ll have a record of it stored happily in my orders table.

When I run my sales queries at the end of the month, your sale will be in there, adding \$12.00 to the total. In July of this year I have \$6800 in total sales! Wahoo!

Sales have gone well, and being a good capitalist I decide I’m going to charge \$15.00 for Carne Asada from now

on. I'm really proud of myself and so I run July's sales reports one more time so I can print them out – I want to see that money rolling in!

Hmmm. The numbers are off for some reason. It used to say \$6800 for July, but now it says \$7300! What happened?

We've made a rather critical mistake with our design, here. One that you see constantly. The deal is that **order_items** is what's known as a "slowly-changing historical table". The data in this table is not transactional, it's a matter of record.

So what do you do if you don't want to change the past? We'll discuss that in the next section.

OLAP and OLTP

99% of the databases you and I work in are considered "OLTP": Online Transaction Processing. This type of system is based on performance – many reads, writes and deletes. For most applications this is appropriate.

At some point, however, you're going to want to analyze your data, which is where "OLAP" comes in: *Online Analytical Processing*. These systems are low-transaction systems that change little, if at all, over time apart from nightly/weekly loads. These systems power data warehouses and support data mining.

The structure of each system varies quite a lot. OLTP systems are relational in nature and are structured using the rules of normalization discussed in the last chapter.

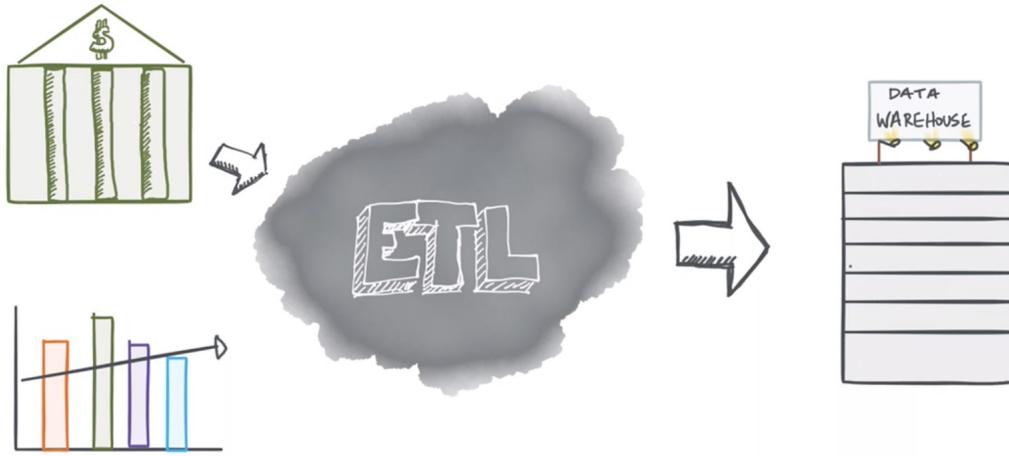
OLAP systems are heavily denormalized and are structured with dimensional analysis in mind. Building these systems can take hours and usually happens on a nightly basis, depending on the need.

Let's start where OLTP ends and OLAP begins...

Extraction, Transformation, and Loading (ETL)

My accountant, who's also my best friend, has the same thing to say to me every year when preparing my taxes: "trash in, trash out". That's his warning to me as I prepare my account statements to bring to him.

This is a form of ETL that people do every year (especially in the US): pull all their financial data from their banks, savings, investments etc. and compile it in a single place – maybe Excel. They go through it at that point, making sure it all adds up. Each bank statement reconciles and there are no errors.



You do the same with analytical systems. The first step is to extract the information you want from your OLTP system (and/or other sources) and comb through it for any errors. Maybe you don't want null sales totals or anything tagged test.

You then transform as required. Reconciling customer information with your CRM system so you can add history data, account numbers, location information, etc.

Finally you load the data into your system, which is usually another database that has a special layout. The system I'm most familiar with (and spent years supporting) is Microsoft's SQL Server Analytical Services (SSAS) so I would usually extract the data from one SQL Server database to another.

They also had a built-in transformer that worked with VBScript, of all things! I used it sometimes but more often than not it would fail. We later moved to a system called Cognos that was a gigantic pile of XML pain.

Today, you can perform quite complicated ETL tasks efficiently and simply by using a set of simple scripts. These can be as simple as shell scripts or, more commonly, quite complex using a programming language like Python or Ruby. Python's speed and popularity make it a very common choice for ETL.

Data Marts and Warehouses

You'll often hear these terms used interchangeably, but they're two very different things. A data warehouse is like a filing cabinet in your office or at home where you keep all of your financial information: statements, tax documents, receipts, etc. Hopefully you keep this organized so it's easy to sift through and put in a form that your accountant can understand – such as an Excel spreadsheet.

There are other things we can do with this data; maybe we want to know how much we spent on groceries so we can budget properly next year. We might want to calculate how much our life savings could be if we had any ... stuff like that. For now, however, we need to get some data to our accountant because it's tax season here in the US, so we load the data into Excel with a targeted use case: Accounting.

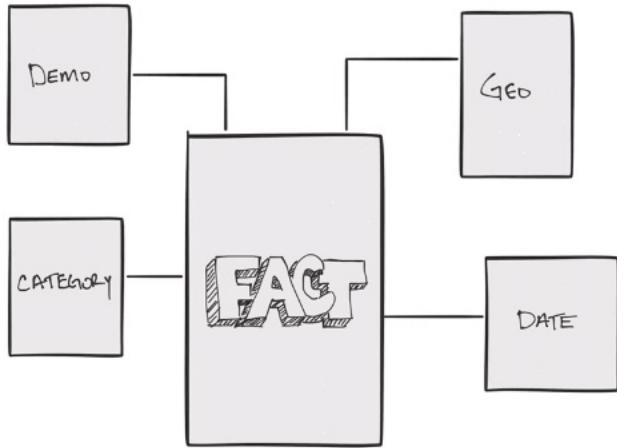
This is a data mart. A place (typically focused/targeted) that can answer questions. We could use the Data Warehouse for this (sending the accountant our shoe box full of receipts and statements) but that would take her a really long time and she wouldn't be happy.

Data Mart Schemas

The Excel spreadsheet is an apt way to think about how data is stored in a data mart: flattened. You might have a

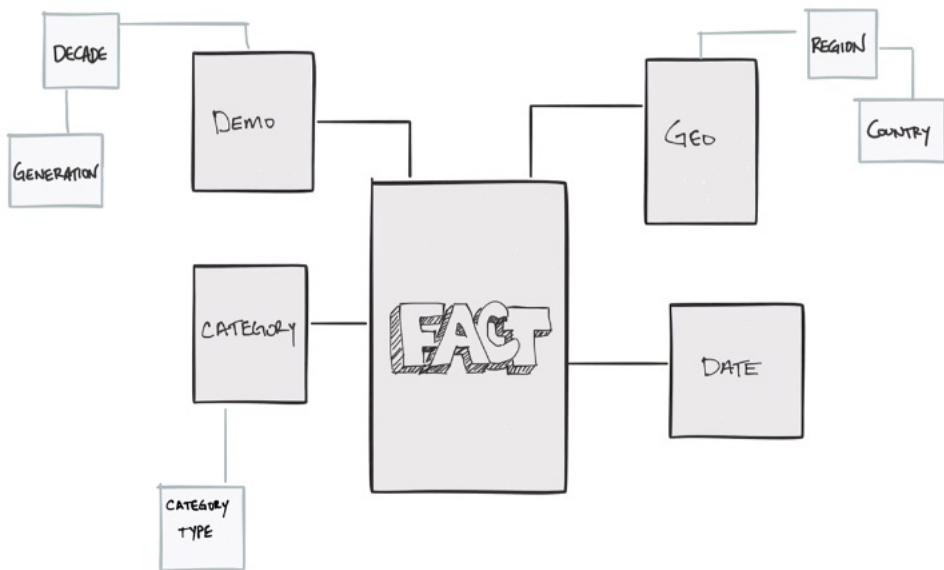
few joins in there, but the fewer the better because processing the data mart, which is typically millions and millions of records, will slow down with more joins.

What does this look like, however? The one you'll see most often is the star schema:



In the center is a table called “the fact table” which represents a single fact you want to report on. For my accountant this would be a singular transaction of some kind: deposit, debit, adjustment, etc.

A *snowflake* schema is the same, but the dimension tables themselves have more dimensions.



Dimensions

The fact table has keys which link off to dimensional look up tables, which you can use to roll up the data for specific queries. They're called dimensions because they present a different way of rolling up the data. For sales (or anything relating to people) these are typically:

- Categories of some kind
- Time (week, month, quarter, year)
- Geography (city, state, country)
- Demographic (gender, age, race)

Selecting dimensions is much harder than it seems. They need to be fundamentally different from each other, or you'll be rolling up on data that has crossover meaning.

You see this sometimes with schemas that confuse demographic data with geographic data – typically with the region that a person is from. I had a hour-long discussion with a client once about the meaning of "Southerner" in their data.

It might not seem like a big deal, but making sure the data can be cross-checked is absolutely critical.

With a data mart it's possible (and common) to query on multiple dimensions at once. If we had left "Southerner" as a bit of demographic information, we would have had conflicting questions and answers:

- Show all sales for both men and women located in the Southern United States
- Show all sales for both men and women who are "Southerners"

I have friends in Hawaii who call themselves "Southerners". I have Hawaiian friends who live in Louisiana. What are we learning with these questions? *Analytics is difficult*. The point is: pick your dimensions with care and make sure you involve the people who are using the reports you'll generate.

Bad Dimensions

The "Southerner" problem (as it became known) is rather intangible and it takes some experience with data to be able to spot reporting issues like that.

Others are far easier to spot – such as "double-labeling", which happens all the time and is rather infuriating.

As a programmer I hope you have a blog where you share your ideas. If you do, it's likely you have a way of tagging your posts with small, contextual keywords (tags).

Let's do a counting query to find out how many comments your blog has for the tag **opinion** vs the tag **humor** (if you have such things... if not let's pretend). It's a simple enough query because, as it turns out, you only have 3 posts with 5 comments apiece:

- "Data Analysis is Silly" tagged **opinion**; 5 comments
- "Hadoop Honeybadger" tagged **opinion** and **humor**; 5 comments
- "A DBA Walks Into a Bar..." tagged **humor**; 5 comments

So you run these queries:

```

select count(1) as comment_count from posts
inner join comments on post_id = comments.id
inner join posts_tags on posts_tags.post_id = posts.id
inner join tags on posts_tags.tag_id = tags.id
where tags.tag = 'humor'

--comment_count
---
--10

select count(1) from posts
inner join comments on post_id = comments.id
inner join posts_tags on posts_tags.post_id = posts.id
inner join tags on posts_tags.tag_id = tags.id
where tags.tag = 'opinion'

--comment_count
---
--10

```

Simple enough. But then you remember reading *The Imposter's Handbook* which mentioned cross-checking your rollup queries for accuracy, so you do:

```

select count(1) from comments;

--comment_count
---
--15

```

Uh oh. Ten humor posts + 10 opinion posts does not equal 15!

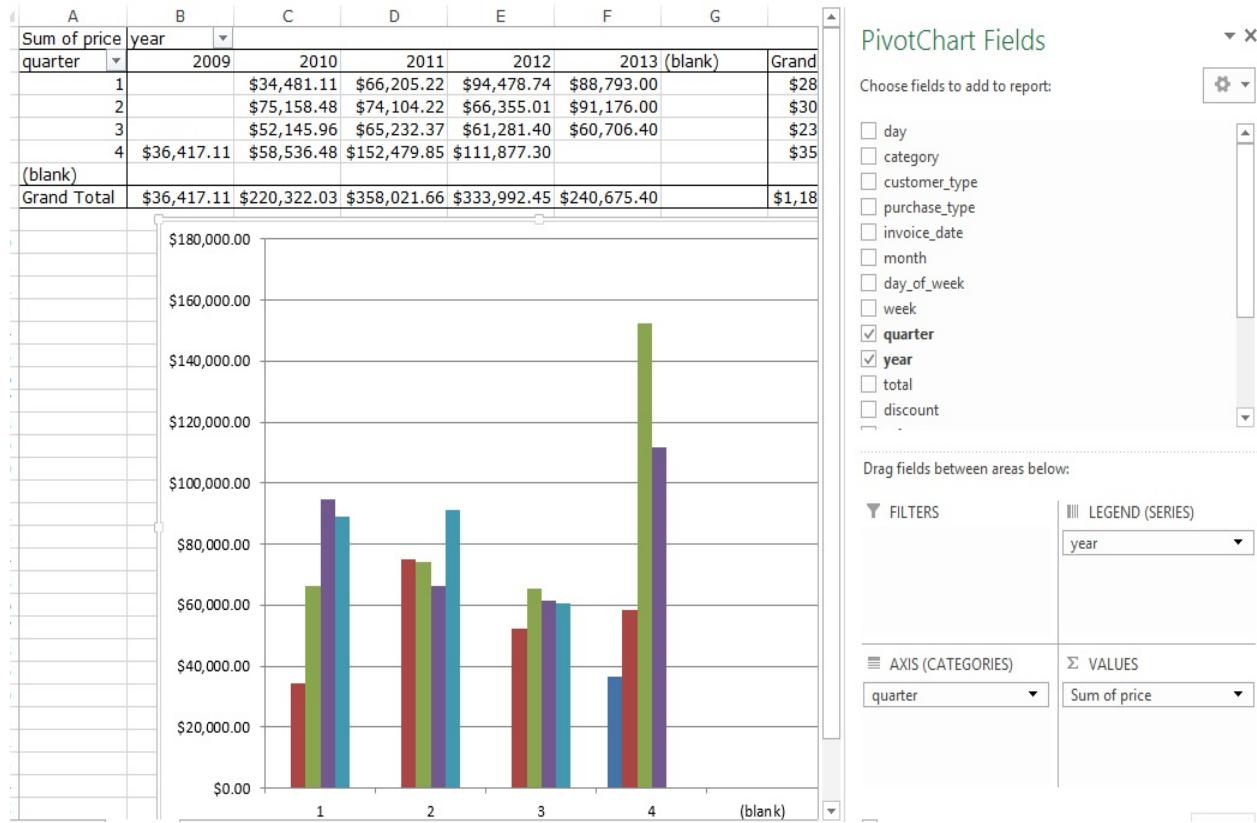
Now you might be thinking “of course it doesn’t” and that cross-checking like this is not accurate! My answer to that is “tell it to the product specialist who wants a sales rollup on various product tags”.

Right now, across the world, sales reports are in this “semi error” state. You cannot do rollup queries that involve many to many categorizations and expect to keep your job. Even if you add warnings! The numbers will suggest a reality that’s not there.

By the way: cross-checking like this is all part of ETL. Bad data should never make it into your data warehouse/data mart.

Analyzing a Data Mart

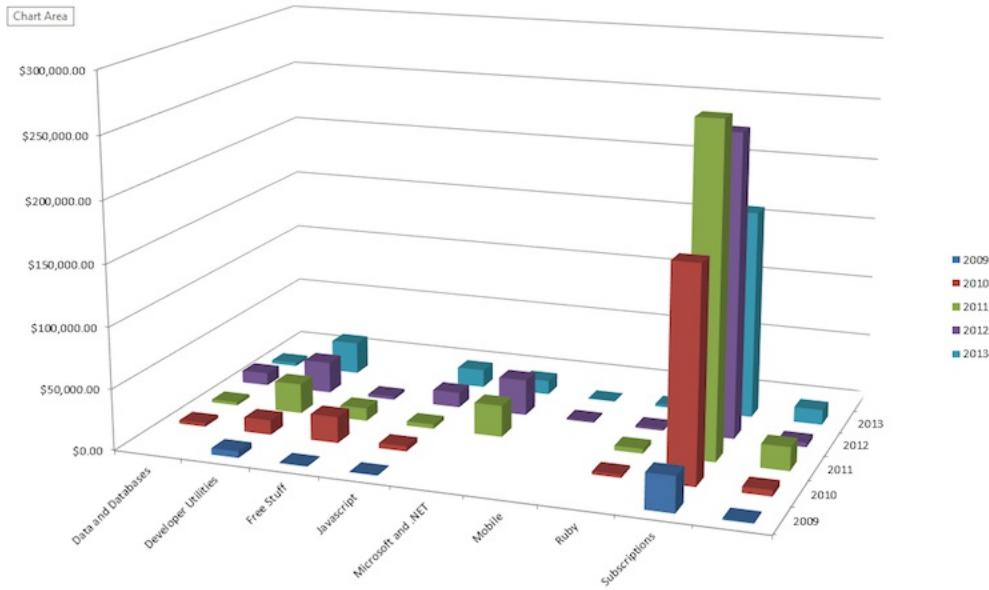
OK, you’ve gone through your data and have decided it’s clean of weirdnesses (good job!) and imported it into your data mart. How do we analyze it? The simplest way is with a Pivot Table and/or a Pivot Chart. It’s likely you’ve seen these in action – here’s some sample data in Excel:



The idea with a pivot table is that you can move dimensions around, rolling your facts up in various interesting ways.

The axes of the graph is a perfect example of why a dimension is called a dimension: the x dimension is often time and the y dimension is often the sum of sales.

What if you wanted to visualize sales by category over time? You just need to add another dimension to the graph – and thank goodness most people in the world can understand things in three dimensions:



Your boss likes this report a lot! It's interesting to give your data a "surface" as we're doing here because, in a way, you can feel what's going on. Now your boss wants to see this data with some demographic information – for instance the buying patterns between men and women.

That requires a fourth dimension. How in the world would you do that! Well, without getting into a physics discussion – you can treat time as a fourth dimension – which can work really well. Unfortunately for me, this book only works in two dimensions (with the illusion of a third), so I can't show you a moving time-graph ... but close your eyes and see if you can imagine a three dimensional graph slowly changing over time...

The neat thing is that time is that one of your dimensions so you can lift that to the fourth axis and watch sales by category and gender change slowly.

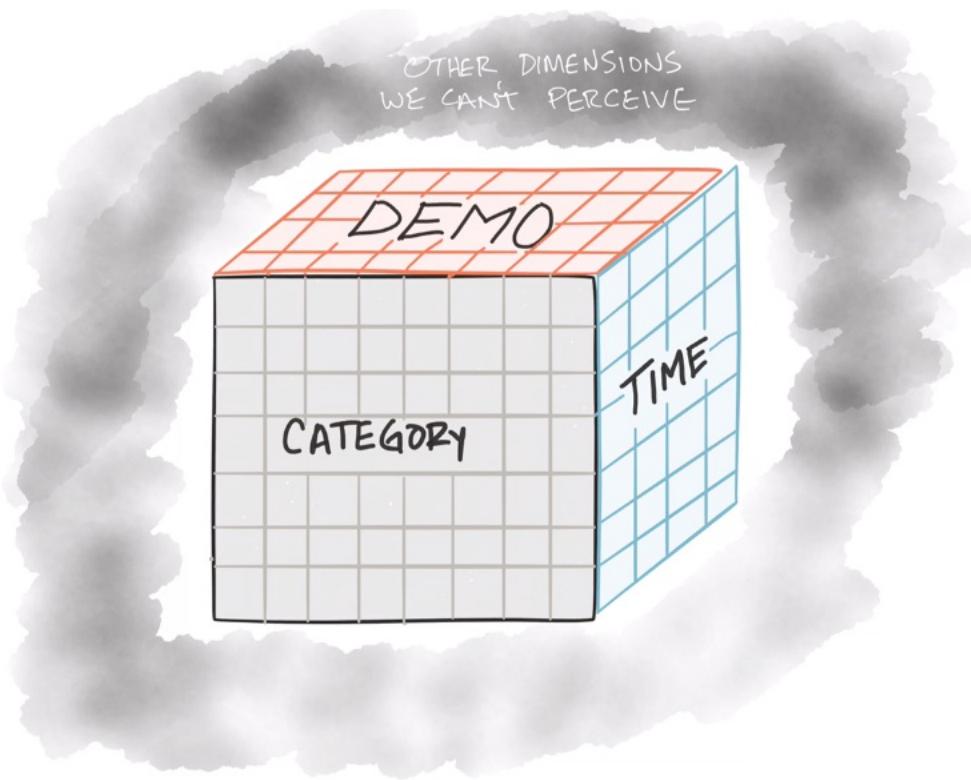
If your boss asks for more axes on this report you need to ask for a raise.

Using an OLAP Cube

Pivot tables work well over data structured in a flat way. If you have more than a few thousand rows, however, things can get mighty slow.

This is where a structure called an OLAP Cube comes in. You tell it about your fact table, the dimensions you're using and their hierarchy, and then begin processing.

An OLAP cube is simply a bunch of pre-calculated data. It's called a "cube" because the data is described typically in three dimensions, and as I mention above people can't really conceive more than four dimensions anyway. Time is usually one dimension, some type of categorization is another, and customer demographic is usually the third. Any more than that and things just get weird.



When you view data along a dimension, you're viewing a slice of the cube and you usually do this with a pivot table of some kind. Excel, for example, will hook up to an OLAP cube.

Pre-calculating data like this makes OLAP cubes very fast at preparing ad-hoc reports over millions of rows of historical data, but that comes at the cost of preprocessing the data. For this reason data marts that act as the source of an OLAP cube should be structured in a very specific way ... and this will sound counterintuitive.

Fact Tables and Indexes

Your fact table should not have a primary key or indexes of any kind, for that matter. Inserting data into a table with an index means the database needs to update the index whenever data is inserted, which takes time.

Ideally you've already vetted and cleaned your data during ETL and you trust it – so no keys or indexes. A fact table can grow to billions of records – can you imagine the index sizes on that!

Favor a Star Schema

Joins are slow, so denormalize your dimensional look up tables for speed of processing. Building an OLAP cube with millions of facts can take hours depending on the number of dimensions you're rolling up on.

Date rollups are the easiest thing to contain. For instance your boss might think she wants a weekly sales report – but that's adding 52 additional slices to the OLAP structure – and every other dimension will need to be pre-calculated based on those 52 weeks times however many years.

This will move a three hour processing run to an overnight run easily. So push back, if possible, or consider building an additional cube.

Distributed Database Systems

It's 2008 (or thereabout) and you're realizing that processors aren't really getting any faster. Buying more RAM

used to solve all kinds of problems but lately you're finding that 12G is really all you need; the processor has become the bottleneck for your database.

In 2010 you had two processors. In 2012 you have four – all for the same price. Today if you pay enough money you can have 32 ... THIRTY TWO CPU cores on that lovely machine in the sky.

Can your database take advantage of each of those cores? THAT is the question for this chapter.

Multiple CPUs means that many things can be processed at once. This means the software has to be able to *do things in parallel* without freaking out. Parallel processing is not a simple topic – especially concerning data.

Imagine 3 things happening in parallel, each on a different core:

- **user 3002** changed their password
- **user 3002** updated their profile picture
- **user 3002** canceled their account

What happens first? Is there a priority here ... and if so what is it based on? What happens if core #1 goes offline for 30 milliseconds because of network trouble?

Very Smart People have focused specifically on these problems over the last 10 or so years and come up with some interesting ways of doing things. **Parallel processing is where things are going** because that's where the hardware is taking us.

A Shift In Thinking

When computer science people tried to figure out data storage back in the 70s and 80s (aka: databases), they did so with two primary constraints in mind:

- **Storage capacity:** hard drives were not cheap so they had to focus on ways of storing data that would be extremely efficient with hard drive space. This led to column names like "UUN1" and hard-core adherence to normalization.
- **Memory and processing speed:** computers were simply slower, so storage needed to be optimized for read efficiency as well as overall size.

This is what most developers (including myself) "grew up" with. You used a relational engine to store data and you created your tables, keys, etc. in a very particular way.

NoSQL systems have been around since the 60s, but it was only in the late 90s that the development community really started to pay attention. Then, right around 2010, big software (Amazon, Facebook, Google, etc) began to see the advantages of using NoSQL systems.

There was one advantage, however, that stood out above the rest: distribution. Simply put: **it's easier (technically and economically) to build distributed databases with a NoSQL system than it is to run a few, gigantic servers with oceans of RAM and disk space.** Smaller servers are cheaper, and you spread your risk, mitigating data loss and disaster recovery.

You can scale horizontally with relational systems such as PostgreSQL and SQL Server – the problem, however, is that these systems need to remain ACID Compliant, which is a problem in distributed systems. They don't like to work in a parallelized way.

ACID-compliance means that you have certain guarantees when writing data in a transaction. In summary form, each transaction will be:

- **Atomic.** This means that a single transaction happens, or it doesn't. There is no concept of a "partial" transaction

- **Consistent.** The entire database will be aware of a change in the data whenever a transaction is completed. In other words: the state of the database will change completely, not partially.
- **Isolated.** One transaction will not affect another if they happen at the same time. This has the basic appearance of transactions being queued in a single process.
- **Durable.** When a transaction concludes it concludes. Nothing can change the data back unless another transaction changes the data back to the way it was. In essence: there is no undo.

Distributed systems are much different from this, and rely on a different rule set entirely.

CAP Theorem

In 1998 Eric Brewer theorized that distributed processing of any kind can only provide two of the following three guarantees at any given time:

- **Consistency.** The same meaning as with ACID above; the state of the database will change with each transaction.
- **Availability.** The distributed system will respond in some way to a request.
- **Partition tolerance.** A distributed system relies on a network of some sort to function. If part of that network goes offline (thus “partitioning” the system), the system will continue to operate.

So far, this has proven to be true. Sort of. In 2012 Brewer [wrote a followup](#) which suggested “picking two of three” can be misleading:

...the “2 of 3” view is misleading on several fronts. First, because partitions are rare, there is little reason to forfeit C or A when the system is not partitioned. Second, the choice between C and A can occur many times within the same system at very fine granularity; not only can subsystems make different choices, but the choice can change according to the operation or even the specific data or user involved. Finally, all three properties are more continuous than binary. Availability is obviously continuous from 0 to 100 percent, but there are also many levels of consistency, and even partitions have nuances, including disagreement within the system about whether a partition exists.

Modern distributed database systems are addressing exactly this. RethinkDB is a prime example (full disclosure: it’s one of my favorite distributed databases and I love it. Further disclosure: during the writing of this book the company behind RethinkDB closed its doors. The project, however, is open source and will continue on).

You can choose the level of consistency you want on a per table or per query basis. Meaning that you choose whether you want an ack (acknowledgment of write to the entire system) or you can just trust the system to do it when it gets around to it.

In addition, you can architect your database on a table by table basis to enhance which of the three you want.

This can be really confusing, so let’s dive into each of these ideas (as well as the jargon for each) using RethinkDB as an example system (because it’s what I know).

Enhancing A and P with Eventual C

You’ve heard of eventual consistency, it’s a buzzword that makes many ACID-loving, relational DB people freak out. I was one of them.

The idea is a simple one: a write operation (think of it as an insert into query) is handed to the system and you receive an ack immediately, meaning that the system has received it and will write it to disk when it can.

At this point (before the write hits the disk), the database is not consistent. You can visualize this by thinking of a database with three nodes. One node receives the write request, which means the other two nodes are not consistent.

This inconsistency is risky. If the power goes out for some reason, the write will be lost, which could be a bad thing depending on the data.

The benefit, however, rests squarely on the benefits of distributed systems in the first place: parallel processing. The system is available to handle more operations, so many things can happen at the same time, making the system extremely fast.

If the node handling the write goes offline (due to a netsplit, or network partition) it doesn't matter (assuming it still has power) because the write is queued, and will remain queued, until the rest of the system is brought back online and consistency is achieved.

These systems are called “AP systems” (generally) and are built to handle gigantic loads with guaranteed uptime. Facebook’s Cassandra and Riak from Basho are prime examples of these systems.

An Alternative to ACID: BASE

Systems that focus on availability and partition tolerance comply with BASE:

- Basically Available
- Soft state
- Eventually consistent

If you’ve grown up with ACID systems, as I have, this idea might sound nightmarish. ACID is all about data paranoia: *protect it at all costs*. Make sure it’s written to disk in isolation.

BASE, on the other hand, is the opposite of paranoid – which kind of makes sense when you consider the idea of strength in numbers: the more machines, the wider the risk is spread.

ACID systems, on the other hand, are typically “big machine” systems. You scale these systems up (bigger hardware) as opposed to out (more hardware).

The problem for ACID, however, is there is only so big you can get. BASE systems can scale massively simply by adding more nodes. I’m not saying this is easy by any stretch, running these systems is very complicated – but it’s doable and is being done in very large companies. We’ll discuss this more in the chapter on Big Data.

Assessing and Mitigating The AP Risk

You’re in a meeting where the CTO has just announced that your fictional company, Red:4 Aerospace, is moving its orbiter telemetry data over to a distributed system. She’s heard about CAP and needs a risk analysis – and she’s looking at you.

What do we tell her in our analysis?

The first thing is that we’ll gain a ton of processing power, which is good. Customers might have to wait a few extra seconds (depending on load), but if things get too slow we can just add a few more nodes!

This has the dual advantage of mitigating a netsplit. If we’re strategic about our nodes and data centers we should be able to survive just about any outage.

But... what about the data? *There is the possibility of data loss*, always, when you make the AP tradeoff. Losing data can mean losing business (or worse) and to people (like me) who live and breathe the idea of Good Data, this is a hard subject to muse on clearly.

This is where we begin staring out the window as we remember various Hacker News threads on the

unreliability of MongoDB which, in many developers' minds, means all distributed NoSQL systems. We remember the many blog posts we've read that ultimately resolve to "we lost data because we did NoSQL wrong"...

Giving Into Paranoia By Leaning On C

You've turned in your report on AP systems and the CTO now wants to know what other options are. Most (if not all) of the distributed database systems out there support partition tolerance well – it's just a matter of choosing availability or data consistency.

Do you want your system to stay up? Or the data to be correct? RethinkDB and MongoDB lean towards the latter – they are CP systems. By default, RethinkDB will only acknowledge a transaction when the data is persisted to disk.

Both MongoDB and RethinkDB are configurable so that you can tweak consistency settings the way you want and need. You can make some tables more AP if you want as the data allows.

The more you get into the nuances of CAP and how modern NoSQL databases handle it, the more confusing things get. Rapidly. As with every topic in this book I could fill chapters and chapters with detail. Instead I'll leave the CAP discussion here and suggest you read more about how RethinkDB and MongoDB are put together.

It's time to turn our attention to the mechanisms for tweaking availability and consistency.

Best Of Both Worlds: Applying A and C Where Needed

Distributed databases specialize in handling very large volumes of data. They do this by "spreading the load" between individual database nodes.

Not Only NoSQL

I've been focused on NoSQL systems like Cassandra, MongoDB and RethinkDB but relational systems can be clustered into a distributed system as well. Namely my favorite database engine: PostgreSQL. This is how Instagram does it and they've written extensively about maintaining a PostgreSQL distributed cluster.

The rest of this section is about distributed systems in general, unless otherwise specified.

Your CTO is quite happy that the distributed system chosen by the company can handle both AP and CP, and she wants you to come up with strategies for data coming from the orbital probe that's due to arrive in orbit around Jupiter 8 months from now.

Let's shape our distributed system using sharding and replication.

Increasing A With Sharding

We will have telemetry data coming in at a very high rate, and we need to process this data continually to be sure our orbital calculations are correct and that our probe is where it's supposed to be.

A fast database can hold the majority of *current data* in RAM. By current data I mean the stuff that we care about. With a demo database, this might come to a few megabytes.

The data generated by most consumer-focused websites remains in the < 1Gb realm, which means that sharding/replication will have little effect.

Another good friend of mine, Rob Sullivan is a PostgreSQL DBA who fields some very interesting questions from developers he runs into at conferences and cafes. Recently he was asked how he would suggest sharding a database system that just hit 5Gb total data.

His answer (paraphrased):

... they were trying to do all of this on the cheapest Heroku database possible, and having issues

because they didn't want to pay for the appropriate tier... talking themselves into a complete rearchitecture because of a price list...

There is, and will always be, confusion about when and mostly if you should shard your database.

Knowing When To Shard

The outcome: when you run out of RAM and it's cheaper to add an additional machine vs. scaling up to a bigger VM or server.

Let's take a closer look at this.

Here's a price breakdown for [Digital Ocean](#) as of summer, 2016:

Choose a size

\$5/mo \$0.007/hour	\$10/mo \$0.015/hour	\$20/mo \$0.030/hour	\$40/mo \$0.060/hour	\$80/mo \$0.119/hour	\$160/mo \$0.238/hour
512 MB / 1 CPU 20 GB SSD Disk 1000 GB Transfer	1 GB / 1 CPU 30 GB SSD Disk 2 TB Transfer	2 GB / 2 CPUs 40 GB SSD Disk 3 TB Transfer	4 GB / 2 CPUs 60 GB SSD Disk 4 TB Transfer	8 GB / 4 CPUs 80 GB SSD Disk 5 TB Transfer	16 GB / 8 CPUs 160 GB SSD Disk 6 TB Transfer
\$320/mo \$0.476/hour	\$480/mo \$0.714/hour	\$640/mo \$0.952/hour			
32 GB / 12 CPUs 320 GB SSD Disk 7 TB Transfer	48 GB / 16 CPUs 480 GB SSD Disk 8 TB Transfer	64 GB / 20 CPUs 640 GB SSD Disk 9 TB Transfer			

The price of each machine goes up according to the RAM. Double the RAM, double the price. So here's the question: *would bumping to the top instance they have (\$640/month) be the same as 4 of the \$160/month?*

In short: **no**. A single machine is simply easier to maintain, especially when it comes to databases. You're much better off just upgrading until you can't – let's see why.

Sharding

The top of the line DigitalOcean VM has 64G of RAM, most of which you can probably use for your database cache. You also have 20 CPU cores at your disposal: this is a fast machine.

Unfortunately the telemetry data is projected to have a current data size of about 250Gb – this is data we'll need to actively write and query extremely quickly.

We don't have access to a machine with this much RAM at our provider, so we'll need to distribute the load across multiple machines. In other words: shard.

There are two ways we can do this with modern databases:

- **Let the machine decide how.** Modern databases can use the primary keys to divide the data into equal chunks. In our case, we might decide to go with 6 total shards – so our database system will divide the primary keys into groups of 6 based on some algorithm
- **We decide how.** Our telemetry data might lend itself to a natural segregation in the same way a CRM system might divide customers by region or country – our telemetry data could be divided by solar distance, attitude, approach vector, etc. If the majority of our calculations only need a subset of this data, sharding it logically might make more sense.

The less we have to mess with a sharding situation, the better – so I would suggest letting the machine decide unless you’re utterly positive your sharding strategy won’t change.

Once you’ve decided the strategy, it’s a matter of implementing it. Some systems make this very easy for you (RethinkDB, for instance) and others require a bit of work (PostgreSQL requires an extension and some configuration, MongoDB is a bit more manual as well).

If everything works well, you should be able to up A – your throughput and processing capacity – dramatically, at the price of C, consistency.

Other data, however, needs some additional guarantees in a distributed system.

Replication

We’re receiving and processing telemetry data constantly, and every now and then it might cause us to alter the current mission plan just a little.

Given that we’re using a distributed system, we need to be sure that the mission plan the probe is following is up to the minute and guaranteed to be as accurate as possible – even if there is a massive power outage at our data center.

So we replicate the data across the nodes in our cluster.

We have data centers in the US, Europe, Asia and the Middle East – all of which can communicate with the probe throughout the daily rotation of the Earth. This data doesn’t change all that often – perhaps a few hundred times per day – so we can implement replication at the database level.

Whenever data changes in our replicated tables, we can guarantee that:

- The data will, at some point, propagate if a netsplit occurs. In other words: if we write to a node in the US and our US datacenter goes down, the write will happen eventually when the US datacenter comes back online.
- The system will compensate and rebalance for the loss of a node, ensuring the data will be as consistent (and available) as possible.

In The Real World...

As I mention above: distributed systems will be the future because this is where the hardware is going. It’s important to look past the “relational vs. nosql” debate because it simply misses the point: you need to start thinking horizontally.

Even if you and your company choose to go with a hosted solution such as Compose – understanding what’s happening behind the scenes is critical.

Finally, if you’re like me and you really love PostgreSQL but you also see the need to start doing things in a distributed way, you might want to have a look at the company that Craig Kerstiens works for: [Citus Data](#). It’s an extension system on top of PostgreSQL that allows you to build real time, distributed PostgreSQL systems.

Big Data

Most applications generate low to moderate amounts of data, depending on what needs to be tracked. For instance: with Tekpub (my former company), I sold video tutorials for a 5 year period. The total size of my database (a SQL dump file) was just over 4Mb.

Developers often over estimate how much data their application will generate. My friend [Karl Seguin](#) has a great quote on this:

I do feel that some developers have lost touch with how little space data can take. The Complete Works of William Shakespeare takes roughly 5.5MB of storage.

Millions vs. Billions vs. Trillions

A megabyte (1 million bytes) seems so small, doesn't it? A gigabyte (a billion bytes) seems kind of skimpy for RAM and we all want a few terabytes on our hard drives. What do these numbers really mean outside of computers? Consider this:

- A million seconds is almost 12 days
- A billion seconds is just over 31 years
- A trillion seconds is 317 centuries Many people simply do not grasp just how much bigger a terabyte is vs. a gigabyte. Let's do this again with inches:
- A million inches is almost 16 miles
- A billion inches is almost 16,000 miles, or a little over half way around the earth
- A trillion inches is 16 million miles, or 631 trips around the planet When considering the data generated by your application, it's important to keep these scales in the back of your mind.

A Truly Large Data Store: Ancestry.com

In 2006, Ancestry.com added every census record for the United States, from 1790 to 1930.

That's a lot of data. Ancestry.com tracks many people, as you can imagine. How big was their archive back then?

The project added 540 million names, increasing the company's genealogical database to 600 terabytes of data.

Storing This Data In Your Closet

If you had an extra \$32,000 lying around after your seed round of funding, you could buy 15 x 40Tb drives and store all of this information in your closet today. In 2006 this amount of data was quite impressive and would have cost a fortune. Today ... not so much. Don't get me wrong: \$32,000 is a whole lot of money but it's pocket change for big companies needing to store tons of data.

In 2013 Information Week [wrote an article](#) about Ancestry.com and how it stores its data. This, friends, is a massive growth in data:

A little over a year ago [2012], Ancestry was managing about 4 petabytes of data, including more than 40,000 record collections with birth, census, death, immigration, and military documents, as well as photos, DNA test results, and other info. Today the collection has quintupled to more than 200,000 records, and Ancestry's data stockpile has soared from 4 petabytes to 10 petabytes.

A petabyte is 1000 terabytes – or 1000 trillion bytes of data. If we translate that into seconds it would be almost 32 million years.

Computer Weekly [wrote a fascinating article](#) on visualizing the petabyte, with some amazing quotes from industry experts:

... Michael Chui, principal at McKinsey says that the US Library of Congress "had collected 235 terabytes of data by April 2011 and a petabyte is more than four times that."

Wes Biggs, chief technology officer at Adfonic, ventures the following more grounded measures... One petabyte is enough to store the DNA of the entire population of the US – and then clone them, twice.

Data analysts at Deloitte Analytics also put on their thinking caps to come up with the following... Estimates of the number of cells in a human body vary, but most put the number at approaching 100 trillion, so if one bit is equivalent to a cell, then you'd get enough cells in a petabyte for 90 people – the rugby teams of the Six Nations.

A petabyte is **huge**. You might be wondering why I'm throwing these statistics at you? It has to do with the title of this chapter.

So... What Is Big Data?

Social media is driving the idea of Big Data:

Big data is a term for data sets that are so large or complex that traditional data processing applications are inadequate. Challenges include analysis, capture, data curation, search, sharing, storage, transfer, visualization, querying, updating and information privacy. The term often refers simply to the use of predictive analytics, user behavior analytics, or certain other advanced data analytics methods that extract value from data, and seldom to a particular size of data set.

It's a buzzword, sure, but there is meaning behind it. Companies like Google, Facebook and Twitter are generating gigantic amounts of data on a daily basis. Back in 2008 Google [was processing over 20 petabytes of data per day](#):

Google currently processes over 20 petabytes of data per day through an average of 100,000 MapReduce jobs spread across its massive computing clusters. The average MapReduce job ran across approximately 400 machines in September 2007, crunching approximately 11,000 machine years in a single month.

While researching this chapter I stumbled on an interesting post from FollowTheData.com, which outlined [how much data was processed](#) by certain organizations on a daily basis back in 2014:

- The US National Security Administration (NSA) collects 29 petabytes per day
- Google collects 100 petabytes per day
- Facebook collects 600 petabytes per day
- Twitter collects 100 petabytes per day
- For data storage, the same article states:
- Google stores 15,000 petabytes of data, or 15 exabytes
- The NSA stores 10,000 petabytes
- Facebook stores 300 petabytes

These numbers are rough estimates, of course. No one knows about Google's storage capabilities outside of Google, but Randall Munroe (of xkcd fame) decided to [try to deduce](#) how much data Google could store for one of his What If? articles using metrics like data center size, money spent, and power usage:

Google almost certainly has more data storage capacity than any other organization on Earth... Google is very secretive about its operations, so it's hard to say for sure. There are only a handful of organizations who might plausibly have more storage capacity or a larger server infrastructure.

A fascinating read. Please take a second and have a look – but be warned! You will likely get lost in all the *What If?* posts.

Simply put: relational systems are just not up to the task, for the most part. The reason for this is a simple one: processing this data needs to be done in parallel, with multiple machines churning over the vast amounts of data. This is the most reliable way to scale a system like Google's that generates gigantic quantities of data every day: just add another data center.

How do you process this kind of information in parallel, however? This is where systems like [Hadoop](#) come in:

Hadoop is an open-source software framework for storing data and running applications on clusters of commodity hardware. It provides massive storage for any kind of data, enormous processing power and the ability to handle virtually limitless concurrent tasks or jobs.

Hadoop was born from efforts at Yahoo!, and then turned into an open source project that any organization can download and install.

Hadoop partitions your data using its dedicated file system, HDFS, which is based on Java. When you query data you use Map/Reduce.

Map-Reduce

In the functional programming world the concept of “map/reduce” is used often when processing data. It is simply a two-step operation:

- Fetch the data you want to use and shape it as you need. This is the “mapping” stage.
- Run calculations on the data, rolling up numbers that are meaningful. This is the reduction stage.

The easiest way to think about this is by example. Let's do a map/reduce run on some sample data. This table represents the yearly sales figures for each country and region in our company:

country	region	total
usa	west	1000
usa	east	2200
usa	north	1600
usa	south	3200
can	north	999
can	east	745
can	south	554
can	west	233
uk	--	552
irl	--	1000

We need to calculate a monthly average, so the first step is to map and transform data:

```
var total = sales.map(function(sale){  
    return sale.total/12;  
});
```

The result is an array of floats:

```
[ 83.33333333333333,
  183.33333333333334,
  133.33333333333334,
  266.6666666666667,
  83.25,
  62.08333333333336,
  46.16666666666664,
  19.41666666666668,
  46,
  83.33333333333333 ]
```

The second step is to reduce this transformation:

```
var total = sales.map(function(sale){
  return sale.total/12;
}).reduce(function(previous, current){
  return previous + current;
},0);

console.log(total); //1006.91
```

This is the map/reduce pattern in its basic form. If you need to do something in more detail (like grouping) – it gets a bit more difficult and you need to write code to support it. Future versions of JavaScript will support grouping with **Array.prototype.map**, for now you need to roll your own.

This is just a pattern, however. We're more interested in [MapReduce, the algorithm](#):

MapReduce is a programming model and an associated implementation for processing and generating large data sets with a parallel, distributed algorithm on a cluster.

The functional nature of the map/reduce pattern (no mutations, no side effects) makes it a great fit for massive data crunching systems like Hadoop.

When you use MapReduce (the algorithm) on massive data sets, each node in your cluster can run your mapping, grouping, and reduction processes concurrently. This means that if you have 1000 servers at your disposal, you can reduce the processing time of a petabyte of data (effectively) down to terabyte processing times.

Many modern database systems will run MapReduce for you by default, across your cluster. Meaning you don't have to setup anything special – just execute a query. [RethinkDB is one of these systems](#):

Map-reduce is a way to summarize and run aggregation functions on large data sets, potentially stored across many servers, in an efficient fashion. It works by processing the data on each server in parallel and then combining those results into one set. It was originally designed by Google and later implemented in database systems such as Apache Hadoop and MongoDB.

...

Some other map-reduce implementations, like Hadoop's, use the mapping step to perform grouping as well; RethinkDB's implementation explicitly separates them. This is sometimes referred to as "group-map-reduce," or GMR. RethinkDB distributes GMR queries over tables

and shards efficiently.

Disadvantages of Map/Reduce

Processing gigantic data dumps is obviously the main selling point of using MapReduce, however it comes at a price: it's not very intelligent. When processing large business applications, you often need to get into trends, clustering and predictive analysis – these kinds of things are not very well suited to simplistic mapping and reduction functions.

It's possible to churn through massive data dumps and then feed them into a data mart that you can then process with a high-end analytical system. Business Analytics is a gigantic ecosystem and, as you can imagine, it's a good business to be in if you like statistics and management consulting.

I've run clustering algorithms before and tried my hand at predictive analysis with some sales data and an OLAP cube – and it was quite an interesting adventure. It takes a while to get your data prepped correctly, but when you do, you have visualizations you might never expect.

If you want to see this in action, head over to Amazon.com. If you shop there on a routine basis (as I do), you'll see some interesting offers right on the home page under New for you.

As I write this, I'm looking at Amazon's home page (I'm logged in, currently) and I see offers for wireless Xbox 360 ear buds, a new video game, and a really groovy coffee grinder next to a very, very interesting cold brew coffee maker.

As I browse their site, they present me with some interesting teasers. Things like Customers Who Bought This Item Also Bought and Frequently Bought Together. They compile this information using clustering algorithms, and they put that data in front of me using behavioral analysis.

It's really fascinating stuff when you get into it, and it's very, very easy to get wrong.

In The Real World

I worked at a business analytics company for about 4 years, working with companies who wanted us to sift through their warranty claims information, looking for patterns. This involved natural language processing (NLP), where we split claim information into sentence structures and then ran various algorithms over it.

It was fun, but 90% or more of my work was trying to figure out which data were good. Emails, phone calls, sifting through and correcting millions upon millions of records ... it's a lot of work.

Even then, I wouldn't call that Big Data. That was just basic analysis over a large set of data. As Sean Parker's character said in The Social Network:

A million dollars isn't cool, you know what's cool? ... A billion dollars.

These days a billion records of anything doesn't even mean much. Terabytes and ... yeah you're getting there. You know what's cool? A petabyte is cool.

That's when you break out Hadoop.

CHAPTER FOURTEEN: TEST-DRIVEN DEVELOPMENT

I am, strictly speaking, not a practitioner of Test-driven Design, or TDD. I know what it is, and I know that people like to argue about what they think it is and what they think it is not.

Just show the essence of the idea. I think we can all agree on that, can't we? For this chapter I approached a number of friends and asked them about what they considered the essence of TDD to be, and how they think of using it. By the way – each of them hedged their opinions with a variation of “this isn’t strictly TDD... but...”.

TDD requires discipline and you’re not alone if you sort of do it. As long as you’re testing your code!

As with many of the chapters in this book, the code for everything you’re about to read is [at Github](#). Clone or download if you want to play along.



Some Opinion About Testing In General

Before we begin, let’s have a think about testing your code in general, aside from TDD. No matter what, **test your code**. There really is no excuse to not test what you create, to make sure it’s correct.

I will say this, directly: if you’re not testing the code that people are paying you money to write *you deserve to be fired*. I won’t justify that remark - if you don’t believe me you should move on. Better yet: close this book and rethink your career. The fact that we should test is a given. It’s how we test that is the subject of this chapter.

In that spirit, let’s engage with Test-driven Design (TDD) and we’ll focus on the fun parts. Hopefully you’ll see the essence of what it is and how it can help, without getting hung up in the details.

The Nuts and Bolts of TDD

Unless you’ve been living under a rather large rock over the last 10 years, you’ve heard of TDD. Maybe in good ways, possibly in bad ones. At its core it’s a simple practice:

- You think about what you need to create
- You write a simple test to get yourself started
- You run that test and watch it fail
- You write some code to make the test pass
- You write another test for the next step, and repeat the process

As you go along, you’re constantly refactoring what you’ve written – and this is the somewhat goofy part: you write the bare minimum to make a test pass.

As Close To a Real Example As I Can Get

A few years back I recorded a video with my friend Brad Wilson and the idea was to capture him doing “real” TDD. No to-do list, no fake blog demo example ... *real stuff*. Brad is the creator of XUnit (along with Jim Newkirk)

and is an every day practitioner of TDD. I couldn't imagine a better person for that video.

The video was for my former company, Tekpub, and it was entitled "Full Throttle: TDD With Brad Wilson". Brad didn't know what I was going to ask him to do – so we sat together, I recorded his desktop, and then asked him to create a subscription billing system for me.

What Brad did next changed the way I thought about TDD. Unfortunately the video is no longer available (it belongs to Pluralsight who has retired it) – but I will recount the highlights for you here.

Just Start

Brad used Visual Studio 2012 and C# 4.5, creating a library project (BillingSystem) and a test project (BillingSystem.Tests) in a matter of seconds. He added XUnit to his test project and then paused to think about a few things:

```
1 using System;
2 using System.Collections.Generic;
3 using System.Linq;
4 using System.Text;
5 using Xunit;
6
7 namespace BillingSystem.Tests
8 {
9     public class BillingDoohickeyTests
10    {
11        //Monthly billing
12        //Grace period for missed payments
13        //Not all customers are subscribers
14        //Idle customers are unsubscribed
15    }
16 }
```

Two things about this code struck me. The first is Brad's use of comments (lines 11 through 14) to get the concepts in his head out onto the screen and, more importantly, he's not getting in his own way thinking about names and structure.

Brad had no idea what to call his test suite just yet, so he called it **BillingDoohickeyTests** in part for fun, but also to remind himself to rename it once things started rolling.

What comes next? What classes should we create right off the bat? This plagues many developers who can't even get past this point.

Patience, Discipline

Here's the thing with TDD that causes anxiety almost immediately: it takes rigor and it feels pretty silly, if I'm honest. So far nothing we've done is overly goofy (apart from the naming thing) – but in a second you'll see what I mean.

You can do a little thinking upfront, but TDD tries to discourage over-engineering by pushing you to let your tests tell you what to write. That's where we're going to start.

The Customer

We're building this system so we can charge customers, so why not start there? This is exactly what Brad does:

```

1 using System;
2 using System.Collections.Generic;
3 using System.Linq;
4 using System.Text;
5 using Xunit;
6
7 namespace BillingSystem.Tests2
8 {
9     public class BillingDoohickeyTests
10    {
11        //Monthly billing
12        //Grace period for missed payments
13        //Not all customers are subscribers
14        //Idle customers are unsubscribed
15    }
16
17    public class Customer{}
18 }

```

He just put the class to test right there, next to his test code. Why not? TDD is a rigorous process, but it doesn't need to be slow.

OK, so we have a Customer, now we need to charge the customer on a monthly basis. At this point: stop thinking. Let's put this idea in motion with a test:

```

1 using Moq;
2 using Xunit;
3 namespace BillingSystem.Tests3
4 {
5     public interface ICustomerRepository { }
6     public interface ICreditCardCharger { }
7     public class BillingDoohickeyTests3
8    {
9        [Fact]
10       public void Monkey_()
11        {
12            var repo = new Mock<ICustomerRepository> ();
13            var charger = new Mock<ICreditCardCharger> ();
14            var thing =
15                new BillingDoohickey (
16                    repo.Object,
17                    charger.Object
18                );
19            thing.ProcessMonth (2016, 8);
20        }
21    }
22    public class BillingDoohickey
23    {
24        public BillingDoohickey (
25            ICustomerRepository repo,
26            ICreditCardCharger charger){
27        }
28        public int ProcessMonth (int year, int month) {return 0;}
29    }
30    public class Customer{}
31 }

```

A lot just happened here. Let's step through it.

You'll notice that Brad isn't concerning himself, again, with names. We have **Monkey** and **thing**, which might be making you cringe – but for Brad, he's removing obstacles to his design process.

Which is what TDD is supposed to be: **a design process**.

Next he's using mocks as you can see on lines 12 and 13, provided by the Moq project (fake classes for testing) upfront so he doesn't need to think about implementation just yet – he's leaving that for later.

The Happy Path

At this point we have a little machinery to play with, but we still don't know what we're doing completely. When I write tests, the very first thing I do is to create what some people call the happy path: one or more tests that pass when everything works as we expect it to work.

In other words, if we're building a registration system then our happy path would be something like **User_is_registered_with_a_valid_login_and_password**. This test should always pass.

Once our happy path is set, we go about trying to break it. We'll do that later on. Right now let's create a "happy path" for ourselves to get us off the ground, renaming Monkey to focus ourselves:

```
1 using Moq;
2 using Xunit;
3 namespace BillingSystem.Tests4 {
4     public interface ICustomerRepository { }
5     public interface ICreditCardCharger { }
6     public class MonthlyChargeTests {
7
8         [Fact]
9         public void Customers_With_Subscriptions_Due_Are_Charged () {
10            var repo = new Mock<ICustomerRepository> ();
11            var charger = new Mock<ICreditCardCharger> ();
12            var thing =
13                new BillingDoohickey (
14                    repo.Object,
15                    charger.Object
16                );
17        }
18    }
19    public class BillingDoohickey{
20        public BillingDoohickey (
21            ICustomerRepository repo,
22            ICreditCardCharger charger){
23        }
24        public int ProcessMonth (int year, int month) {
25            return 0;
26        }
27    }
28    public class Customer{}
29 }
```

In this code I've renamed a few things because I now have an idea what I'm doing. The test suite is called **MonthlyChargeTests** and my test name makes it clear what it's going to test for. This name is far too broad, but it will change later on.

The simple renaming, however, has forced me to consider a few more bits of functionality:

- What is a **Subscription**?
- What does it mean for a **Subscription** to be Due?
- A **Customer**, apparently, needs to have a **Subscriptions** property

I stop right here – thinking about YAGNI (You Aint Gonna Need It). It's tempting to plow ahead and add a **Subscription** class and a **Subscriptions** property to my **Customer**... but do I need to just yet? It does seem

obvious, but this is the rigor part.

Let's focus on the test, add an assertion, and move on from there:

```
1 using Moq;
2 using Xunit;
3 namespace BillingSystem.Tests5 {
4     public interface ICustomerRepository { }
5     public interface ICreditCardCharger { }
6     public class MonthlyChargeTests {
7
8         [Fact]
9         public void Customers_With_Subscriptions_Due_Are_Charged () {
10            var repo = new Mock<ICustomerRepository> ();
11            var charger = new Mock<ICreditCardCharger> ();
12            var thing =
13                new BillingDoohickey (
14                    repo.Object,
15                    charger.Object
16                );
17            var processed = thing.ProcessMonth(2016,8);
18            Assert.True(processed > 0);
19        }
20    }
21    public class BillingDoohickey {
22        public BillingDoohickey (
23            ICustomerRepository repo,
24            ICreditCardCharger charger){
25        }
26        public int ProcessMonth (int year, int month){
27            return 0;
28        }
29    }
30    public class Customer{}
31 }
```

This is where we venture into silly land. Our first goal is to make sure this code can compile – so I've added the interfaces and classes that we need. I also added a **ProcessMonth** method to **BillingDoohickey** and, for now, I'm returning 0 because I don't know what else to return.

If we run this test, it will fail. We'll also probably feel a bit badly about ourselves because we might not have any customers with subscriptions due ... so what then? I'll get to that in a minute – for now we can compile our code and run this test: **watching it fail**.

That's a critical aspect here – our initial test needs to fail because we don't want to accidentally write a test that passes! Which does happen.

Now, let's get our test to pass:

```
//...
public class BillingDoohickey
{
    public BillingDoohickey (
        ICustomerRepository repo,
        ICreditCardCharger charger){
    }
    public int ProcessMonth(int year, int month){
        return 1; // just return 1
    }
}
```

Ugh. Our tests pass and, as dumb as it seems: this is TDD. We will fix this code and, in fact, the dumber it feels the better it is because it forces you to write more tests just to get this kind of thing out of your code!

For now, our happy path is set. Let's blow it up.

The Sad Path

The sad path is all about trying to blow up the happy path. It's "what happens when I do this!" The obvious first thing is to write a test that is in complete opposition to our happy path:

```
1 using Moq;
2 using Xunit;
3 //...
4 public class MonthlyChargeTests
5 {
6
7     ICustomerRepository repo;
8     ICreditCardCharger charger;
9     BillingDoohickey thing;
10
11    public MonthlyChargeTests ()
12    {
13        repo = new Mock<ICustomerRepository>().Object;
14        charger = new Mock<ICreditCardCharger>().Object;
15        thing = new BillingDoohickey (repo, charger);
16    }
17
18    [Fact]
19    public void Customers_With_Subscriptions_Due_Are_Charged ()
20    {
21        var processed = thing.ProcessMonth (2016, 8);
22        Assert.True (processed > 0);
23    }
24
25    [Fact] //this test fails
26    public void Customers_With_No_Subscriptions_Due_Are_Not_Charged ()
27    {
28        var processed = thing.ProcessMonth (2016, 8);
29        Assert.True (processed == 0);
30    }
31
32 }
33 //...
```

I moved some declarations around in my test because I want to keep things DRY (Don't Repeat Yourself) – a messy test suite is one you'll want to stay away from – so I moved everything up top and into the constructor.

Next, I created the exact opposite test, asserting that customers without subscriptions would not be charged – which fails because I've hard-coded the result into the **BillingDoohickey**.

Now I need to think about a few things. Specifically: *what does it mean to have no subscriptions?* **Who cares!** For now let's get this test to pass.

I'll start by setting the mock for **ICustomerRepository** available for orchestration using the **repoMock** variable and then I'll add a method called **Customers** to the **ICustomerRepository** interface because my test told me I needed to:

```

1 using Moq;
2 using Xunit;
3
4 namespace BillingSystem.Tests
5 {
6     public interface ICustomerRepository {
7         Customer[] Customers();
8     }
9     public interface ICreditCardCharger { }
10    public class MonthlyChargeTests
11    {
12
13        ICustomerRepository repo;
14        ICreditCardCharger charger;
15        BillingDoohickey thing;
16        Mock<ICustomerRepository> repoMock;
17
18        public MonthlyChargeTests()
19        {
20            repoMock = new Mock<ICustomerRepository>();
21            repo = repoMock.Object;
22            charger = new Mock<ICreditCardCharger>().Object;
23            thing = new BillingDoohickey(repo, charger);
24        }
25    //...

```

Now that I've done this, I can refactor the **BillingDoohickey.ProcessMonth** method to return a count of the records in the repo. This is a prime example of "just doing enough to get the tests to pass":

```

1 //...
2 public class BillingDoohickey
3 {
4     ICustomerRepository _repo;
5     public BillingDoohickey(
6         ICustomerRepository repo,
7         ICreditCardCharger charger)
8     {
9         _repo = repo;
10    }
11    public int ProcessMonth(int year, int month)
12    {
13        return _repo.Customers().Length;
14    }
15 }
16 //...

```

This kind of thing is really fun when pair coding. I remember pairing with a friend once and laughing so hard at just how creative she was at writing the dumbest code possible to get my tests to pass:

You keep writing tests like that, I'll keep writing code like this.

It's kind of fun to split your personality when doing this kind of thing. See if you can outsmart yourself with a more interesting test – something to break the happy path for once and for all! Then stave off the attack with some goofy way around it, like we just did here, returning the count of users.

It's tempting to scrap this test and try to write something more concise, which you're welcome to do. I typically just write another test specifically so I can get this crappy code out!

How about this:

```

1 //...
2 [Fact]
3 public void A_Customer_With_Two_Subscriptions_Due_Is_Charged_Twice ()
4 {
5     var customer = new Customer ();
6     customer.Subscriptions.Add (new Subscription ());
7     customer.Subscriptions.Add (new Subscription ());
8
9     repoMock.Setup (r => r.Customers ())
10    .Returns (new Customer [] { customer });
11
12 var processed = thing.ProcessMonth (2016, 8);
13 Assert.True (processed == 2);
14 }
15 //...

```

What do you think about this? You can play my pair who's writing code to get the tests to pass ... think you can make me sad?

Here's a way to do it:

```

1 //...
2 public class BillingDoohickey
3 {
4     ICustomerRepository _repo;
5     public BillingDoohickey (ICustomerRepository repo, ICreditCardCharger charge)
6     {
7         _repo = repo;
8     }
9     public int ProcessMonth (int year, int month)
10    {
11        var customer = _repo.Customers ().FirstOrDefault ();
12        if (customer == null) {
13            return 0;
14        } else {
15            return customer.Subscriptions.Count ();
16        }
17    }
18 }
19 public class Customer
20 {
21     public IList<Subscription> Subscriptions { get; set; }
22
23     public Customer ()
24     {
25         this.Subscriptions = new List<Subscription> ();
26     }
27 }
28 public class Subscription { }
29 //...

```

Ha! With that test I was able to push the code so that two new concepts could be added: a **Subscription** and a property on **Customer** called **Subscriptions**. You were still able to write silly code to get the tests to pass – so you win too!

Notice the ongoing *failure, pass, refactor ... failure, pass, refactor* that we're doing here. That's exactly what TDD is all about and despite the way many people make it seem – it can be quite fun. Especially if you have a pair to code with you – real or imaginary.

A final thing to notice: *I did one step at a time*. I didn't write out a set of tests upfront, which would defeat the

idea of challenging myself as I go with YAGNI.

In The Real World

Like I said above: *I get lazy sometimes.* OK a lot of times. I'll get into this in more detail in the chapter on Behavior-driven Design (BDD - coming up next) but I tend to use tests as more of a check list. Things I expect to work, etc.

I often get carried away and I just keep coding – which I know is bad. I pay for this choice often when I'm deleting code that took a while to get to work, but that I ultimately don't need.

So, I take a deep breath, maybe I'll go for a walk or get some tea. When I come back I clean up my tests and refocus myself. This is TDD to me – more of a battle with myself than anything.

CHAPTER FIFTEEN: BEHAVIOR-DRIVEN DEVELOPMENT

Behavior-driven Development (BDD) is basically the same process as TDD, but you have a specific focus: *behavior of the application*. It's a subtle shift, but an important one.



The Code

You can find the code for the examples below [at Github](#). The easiest thing to do is to clone the repo if you want to follow along:

```
git clone https://github.com/imposters-handbook/sample-code imposters-handbook
```

Getting Started With BDD

In the last chapter on TDD we began to build out a BillingSystem using TDD:

```
///...
namespace BillingSystem.Tests {
    public class MonthlyChargeTests{
        //...
        [Fact]
        public void Customers_With_Subscriptions_Due_Are_Charged(){
            repo.Setup(r => r.Customers)
                .Returns(new Customers[]{new Customer()});
            var processed = thing.ProcessMonth(2016,8);
            Assert.True(processed > 0);
        }
        //...
    }
}
```

This works, but it's a bit mechanical. In other words: Widget X will return Y when I pass in Z. This defines what we expect to happen as developers, not what we expect as humans.

Let's shift this to focus on a story instead, with some scenarios. We'll start with what our application will do when a payment is received for a monthly subscription:

```

1 /**
2 namespace BillingSystem.Specs {
3
4 [Trait("Monthly Payment Is Due", "Payment Is Received")]
5 public class PaymentReceived{
6     //... the constructor prepares the test data
7     [Fact]
8     public void An_Invoice_Is_Created(){
9         //...
10    }
11    [Fact]
12    public void Subscription_Status_Is_Updated(){
13        //...
14    }
15    [Fact]
16    public void Next_Billing_Is_Set_1_Month_From_Now(){
17        //...
18    }
19    [Fact]
20    public void A_Notification_Is_Sent_To_Subscriber(){
21        //...
22    }
23    //...
24 }
25 [Trait("Monthly Billing", "Payment Fails")]
26 public class SubscriptionPaymentIsDue{
27     //... the constructor prepares the test data
28     [Fact]
29     public void An_Invoice_Is_Not_Created(){
30         //...
31    }
32    [Fact]
33    public void Next_Billing_Is_1_Day_From_Now(){
34        //...
35    }
36    [Fact]
37    public void A_Notification_Is_Sent_To_Subscriber(){
38        //...
39    }
40    //...
41 }
42 }
```

This code is doing the same basic thing as TDD: *testing our code*. This time, however, we're doing it in the form of “how does our application respond when this thing happens”. In other words: **behavior**.

Philosophy

This approach is different from strict unit testing, which tends to be more clinical. In other words, you might put a certain class under test, vary the input data to see where it fails and then refactor until it succeeds. This is fine, but has some disadvantages, which are:

- The tests are **bound to the design of your class** by definition. That is the point of TDD. If you change your design, you have to change your tests and your code. This can be quite frustrating.
- The **focus is on engineering**, not application experience. When you're focused on code, the code wins. When you're focused on behavior, however, you're focused on the user's experience and the business wins.
- **Testing proliferates**. The tendency with TDD and unit testing is to have as much code coverage as possible. When you use BDD, you typically write few tests which are more targeted to application experience.

With BDD you tend to write your tests detailing what the application will do under certain circumstances. Given

this, BDD fans will call their tests “specifications”, as they tend to read as if dictated directly by the client. In the example above, I’m using XUnit’s **Trait** attribute to decorate my scenarios so they’re a little more readable in the test runner.

I’ve also made sure that my test names rely completely on the test class itself (called the scenario). When you run this test, you see this (or something like it):

```
[Monthly Billing]: Payment Received
- An_Invoice_Is_Created
- Subscription_Status_Is_Updated
- Next_Billing_Is_Set_1_Month_From_Now
- A_Notification_Is_Sent_To_Subscriber

[Monthly Billing]: Payment Fails
- An_Invoice_Is_Not_Created
- Next_Billing_Is_1_Day_From_Now
- A_Notification_Is_Sent_To_Subscriber
```

This is the XUnit runner output, which you can jigger in Visual Studio if you want. If you’re using a framework in another language (like Mocha for Node or RSpec for Ruby) you can have a more readable output.

Features, Scenarios, Expectations

I try to focus on the notion of *Feature*, *Scenario*, *Expectations*. I know this might seem like a syntax dance to you, but it’s incredibly easy to lapse back into unit testing mode – not that there’s a problem with that! Unit tests are indeed needed in some cases (testing utility code, parsers, etc).

Let’s revisit our example code:

```
...
namespace BillingSystem.Specs {
    [Trait("Monthly Billing", "Payment Is Received")]
    public class PaymentReceived{
        //... the constructor prepares the test data
        [Fact]
        public void An_Invoice_Is_Created(){
            //...
        }
    }
}
```

In this code you can see the features directly by examining the **Trait** attribute on the test class. The first element is the feature (“Monthly Billing”), the second is the scenario (“Payment Is Received”). Your testing library might have a different way for specifying these things.

So What?

You might be wondering, at this point, why all of this even matters? BDD does have a number of advantages:

- **Readability.** It sounds idealistic, but being able to print out a test run and read, in common language, what’s going on is quite powerful.
- **Ubiquity.** I hate that word, but it’s applicable here: you know what these tests describe and your client/boss will know as well. In this, you’re speaking the same language.
- **Focus.** If you’re focused on how your application behaves, you’re aligning yourself with the business goals. This is important for programmers! You can watch your application evolve into something exciting and understand why it’s doing what it’s doing. You might even have some questions about this, which means you can contribute your genius to the application’s design.

There's obviously some jargon to muddy up what is, otherwise, an elegant development practice. Let's talk about that now.

Should

When you write assertions in a unit test, you typically need some kind of assertion library. Often (as with XUnit), the test runner you're using will provide you with one ... which is usually called Assert.

Reading tests with assertions like this:

```
Assert.True(processed == 0);
```

... is kind of clunky. We're driving for clarity so it makes sense that our assertions should align with our specifications ... right?

```
processed.ShouldBeEqual(0);
```

This is the [should library](#) from the lovely and talented Eric Hexter. Here's the first sentence of the project description:

The Should Assertion Library provides a set of extension methods for test assertions for AAA and BDD style tests.

AAA stands for “Arrange, Act, Assert”, which is a style of writing unit tests.

When [Dan North introduced BDD back in 2006](#), it was an interesting moment. If you haven't read the linked post, please do. I'll summarize it later on – for now let's skip ahead to his thinking about assertions (emphasis mine):

Then I came across the convention of starting test method names with the word “should”. This sentence template – The class should do something – means you can only define a test for the current class. This keeps you focused. If you find yourself writing a test whose name doesn't fit this template, it suggests the behaviour may belong elsewhere.

Dan is talking about the names of your tests here as opposed to your assertions. The idea of using should in your test **names** is that it keeps you focused on describing the test class (the scenario). If you drift from that, create a new class!

Assertions using should are more of a mind game than anything else, and ultimately have little to do with the practice. That said, if it helps you – go for it. I think should sounds a bit waffley – something either does or does not. It is or is not. This is what you see in my tests above.

Feature, Context, Spec, and Scenario

BDD tests are typically called specifications or “specs”. In the theoretical world you could sit with your client, create a list of specifications for various aspects of your application, and then translate that directly into your test suite.

In the real world I've found that my clients have never cared about my tests. Maybe it's just my clients – not sure – but even when I worked at Microsoft and tried to show the progress I was making using my test suite I was laughed out of the room.

Clients like to see results, not test runs. It is good, however, to use the same language in your tests as you do with them in email or on the phone. Trying to align your thinking is really important.

To that end, you could have this conversation and you could, using BDD, translate it directly into some tests:

So let's recap the conversation: given a successful monthly billing – the billing system should generate an invoice, set the next billing date to exactly one month from now, set the user's subscription to active and then send an email to the user. Correct?

Say this out loud to yourself, as if in conversation. Feel free to use your own words.

Do you hear any problems? Your client probably will. Here's a reply I received once, when I said almost exactly that sentence to a startup client back in 1999:

Rob - yes for the most part this is correct but accounting handles the invoices so I think just notifying them will work. I don't think we should be creating our own invoices. As far as email goes, I'd want to loop in our marketing team as they own client communications...

Do you see what just happened there? Not only did I save myself some work, but I also opened up a really good conversation about the role of our application within the new company.

The feature we were discussing is the monthly billing run. The scenarios we created were payment received and payment failure; these are also called contexts. We describe the behavior of the application in response to a scenario with specifications:

```
Monthly Billing Run //feature
- Payment is successful //scenario, or context
- an invoice is created //specification
```

Is this jargon important? *Yes, and no.* It is important in the sense that you should be thinking in these ways when doing BDD. It's also important to know that just because you call a test class **MonthlyBillingFeature** and a method on that test class **SuccessfulPaymentScenario** does not mean you're doing BDD. BDD is a process of discovery for both you and your client.

That said, you can call a feature a **pancake**, a scenario a **lovely butterfly** and each specification **larry**. The naming doesn't matter – as long as your team understands what it is you're doing and why.

Given, When, Then

[Cucumber](#) is a popular Ruby test tool that helps you focus on BDD. It popularized a certain syntax, called Gherkin:

```
Feature: monthly billing run
  Scenario: payment received
    Given a charge of 20 USD
    And today is the 1st of the month
    When the charge is applied to Subscription x
    Then that subscription is considered active
    And an invoice is created
    And the customer gets an email
```

Gherkin is a specialized syntax you can use to get your head into the Given, When, Then syntax in a formal way. These rules might sound a bit goofy but when you just start getting into BDD it can help get your mind in the right place.

It's quite effective. It takes a bit of time to break things down in this way. Think about an application you're writing right now ... how would you detail the behavior of it? Give it a go!

Missing The Mark, Just a Little

As I mention above, it's easy to get trapped in focusing on syntax vs practice. Let's see some examples that do this.

This is an old example from the RSpec home page:

```
# bowling_spec.rb
require 'bowling'

describe Bowling, "#score" do
  it "returns 0 for all gutter game" do
    bowling = Bowling.new
    20.times { bowling.hit(0) }
    bowling.score.should eq(0)
  end
end
```

You can think of describe here as defining both the feature (**Bowling**, which is the class) and the scenario **#score**.

The it method is the specification. If you're scratching your head: good for you because this is a unit test, it's not a behavioral specification.

How could we write this to focus on behavior? Here's a try:

```
#bowling is not a feature of our application, scoring is
#and even then, scoring is different during and after - so let's be specific
describe "Final Scoring" do

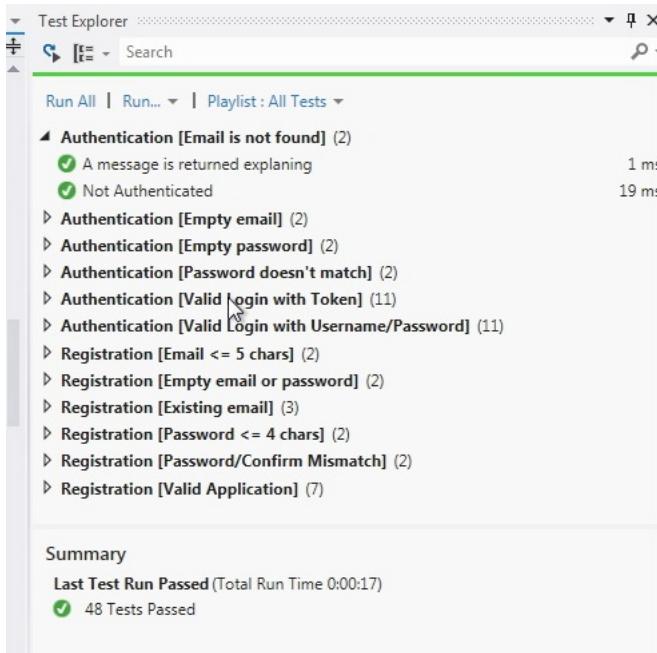
  #now we come up with a scenario
  describe "No pins knocked down for all 10 frames" do

    #what happened?
    it "returns a 0" do
      #
    end
  end
end
```

In The Real World

You might be wondering what this really looks like once a project has been up and running for a few months. Is it possible to scale this idea? To keep it readable and focused?

This is a project I did five years ago, using .NET, Visual Studio and XUnit - no additional tooling:



The feature under test is the first item in the list (Authentication, for instance) and the scenario is the second. Underneath each feature/scenario is a set of specifications. This approach worked really well for me.

If you want to check out the code, [I put it up at Github](#). The code is a few years old now but it should still work.

If you're feeling a little vague on the idea still, I don't blame you. It took me a while to get my tests (ahem, sorry: specs) to flow with behavior vs simple unit tests. The best thing I can offer you is to look at the difference between the billing system code at the start of this chapter compared with the code from the last chapter on TDD.

This is important: *both are valid and both are lovely*. Some people favor TDD, others BDD. Experiment!

APPENDIX A: ESSENTIAL UNIX

Why Do You Have An Addendum On Unix and Shell Scripts?

The simple answer to that is that there are many, many, many developers who stick to the GUI. They prefer apps and tools to commands. They click “File” and “Edit”, hunting for “Copy” and “Paste”.

You know these people. *You were one of these people.* This isn’t a judgement of any kind; I stick to the GUI myself far more than I’d care to admit. There’s a better way, *a faster more efficient way* to work with a computer, and you’ll be a better programmer all around if you learn some basic shell skills.

Unix and Unix-like systems (Linux, BSD, Solaris, RedHat, etc) have been around forever. You simply can’t expect to grow much in your career if you don’t have a basic competency with Unix and its commands. If you don’t believe me, skip right over this chapter. **It’ll be here when you come back, after you’ve realized just how true this is.**

This is an exciting thing! Crawling under the hood of your computer can increase your efficiency dramatically. Shell scripts, Makefiles, server setup routines, quick little commands to update your system, configuring your web/database server remotely over SSH … these are skills you must know.

So let’s wander through the shell. I won’t go into Unix history as I’m just not qualified to do so. I’ll also sidestep the basics of the Unix commands – that’ll be up to you.

Instead, let’s get right to the thing that will help you the most in your job: basic shell scripting skills.

Introduction

It’s a Unix world. You should have a functional knowledge of how to get around a Unix machine using the command line, as well as how to complete basic tasks using shell scripts and Make files.

Shell scripts can help make tedious programming tasks routine, like renaming/resizing images. Make is an old, reliable build tool that can replace (easily) Grunt, Gulp, Rake, Jake and any other language-dependent build tool.

As a developer, having a set of scripts, aliases and binaries in your “dotfiles” is a great way to organize the tools that help you on a daily basis.



This chapter has caused me some problems. It was the very first one I wrote for the book, but has also received the most feedback.

Quite a few people insisted that the book was better without it, and I could see their point. If you’ve based your career on Unix-based machines then you probably wouldn’t consider Unix skills to be an “essential skill for a self-taught developer”.

There are quite a few others who would, however. They let me know it too! I received 20 emails in the span of 1.5 hours right after I pushed version 0.0.4 of this book – that’s the one without the Linux section. They wanted it back.

So I’m putting it back, and I quickly want to share with you my reason for doing so.

Resources

You can find the code used in parts of this section up at [my github repo](#). In addition, you [can buy 17 video walkthroughs](#) of the things you see in this chapter and others [from here](#).

The videos cover some new ground as well, and use languages other than C#.

CHAPTER SIXTEEN: HELLO SHELL SCRIPTS

Shell scripts are little programs that your shell (typically Bash) will execute for you. When you write a shell script, you're writing little macros that can make it feel like you're programming your machine.

You can use shell scripts on Windows with Powershell - an amazing shell with a good programming language. I won't be talking about Powershell in this section – but if you're a windows user, know that you can do anything you see here with a few simple commands saved to a script file.

We're going to create shell scripts for Unix systems using Bash. It's been around for a very long time and it's easy to understand once you get passed some of the more ... arcane commands.

If you're completely new to all of this, we'll go over the basics in just a second. If you understand basic Unix "stuff" then you can probably skip ahead.



What Is a Shell?

A computer needs a way to receive data, and we're going to do that through the command line using a thing called a shell. The first computer ever conceived used punch cards to receive data, when I was in high school I used a combination of a keyboard as well as a cassette tape player to boot my computer!

Today we have visual interfaces that look quite juicy and convey information in a friendly way. We use mice to issue commands (most of the time) and, occasionally, our fingers or a stylus.

During the 1960s through 1980s, computer users entered their commands as text from a keyboard. This practice has continued today and is what you're about to do, using the command line interface.

All of these things are shells. A shell is simply a generalized way in which you give commands to a computer and receive the output. A visual shell uses a graphical interface, or a GUI, and is what I'm using right now to type this sentence on my Mac, using a visual editor.

A text-based shell has no visuals except for things you can do with ASCII symbols. To work with a text-based shell (like Bash, for instance), you use a command line interface, or CLI.

There are a number of shells that you can work with, so far we've discussed two: Powershell and Bash. You can install other ones, if you like, including:

- [Z shell](#) (or zsh). I like this one a lot and it's what I use every day together with [Oh-My-Zsh](#) from Robby Russell
- [Fish](#). They win for the best tag line: "Finally, a command line shell for the 90s"
- [Tcsh](#) (or "tc shell"). This is a common one you see on many Unix machines

Why The Name "shell"?

At this point you might be wondering why these things are called "shells". It has to do with the way Unix is constructed. There is a kernel that does all the processing, which is protected by a number of "protection rings". Each ring provides certain services, with the most sensitive being closer to the kernel and the least being on the very edge, or "shell" of the system. I won't go into Unix design at this point (mostly because I'm not qualified to); but I find that an interesting way to think about Unix.

If you look around you'll find quite a number of shells that look interesting. Bash works well for most things, but if you're looking for something a bit more friendly then I might recommend having a look at Z Shell. I've been using it for years and love it. One main reason is that it has helpful completions, spelling corrections, and you can program the prompt to be colorful and pretty.

The biggest reason, however, is the Oh-My-Zsh project, mentioned above. You get a sane way to organize scripts, aliases and other things. Here's their project description:

A delightful community-driven (with 1,000+ contributors) framework for managing your zsh configuration. Includes 200+ optional plugins (rails, git, OSX, hub, capistrano, brew, ant, php, python, etc), over 140 themes to spice up your morning, and an auto-update tool so that makes it easy to keep up with the latest updates from the community. <http://ohmyz.sh/>

It's been very useful for me.

Keeping Shell Stuff Organized

This is kind of a big deal. As you learn to work with the shell more and more, you find good organization becomes really important. Oh-My-Zsh can help with most things, but not everything.

For example: settings. If you ever work with Vim, Git, Atom, AWS, etc – you know they have various startup settings in an rc file somewhere. These files (typically ending in “rc”) need to live somewhere. This is where strong organizational skills will help a lot.

Why RC?

The more you work with command line tools, the more you'll come across “rc files”, which we briefly discuss above. You don't have to name startup scripts with an “rc” ending, but if you do, people will know what it's supposed to do by convention.

Like so many things in Unix land, the origin of the term “rc” is a bit cloudy. Google it if you like, but the actual meaning of it doesn't matter. Just know that .thingrc will be the startup script for the thingcommand/binary/whatever. Some people like to think it means “runtime configuration” – that sounds like a good explanation to me.

Dot Files

Most Unix-adept developers treat their scripts and settings with the same respect as any other code: organizing it carefully and versioning it with git. For example, one of my very favorite people is Gary Bernhardt and he keeps [his dot files on Github](#).

You'll hear the term “dot files” a lot – and you'll see them a lot. It's a convention to begin a file name with a period (or “dot”) to hide it from the finder and from the standard listing command, ls. These files are shell scripts that are read in by various programs and they contain settings of all kinds.

One that a lot of people obsess on (including yours truly) is .vimrc. In this file you will find settings for Vim. It's a shell script that's executed every time Vim starts.

Think about developers you follow on Twitter and see if they have a dot files repository on Github. I already mentioned Gary Bernhardt - here's another one of my favorite people: [Ryan Bates](#). Have a look at how they organize these files and what's in there. You could lose hours on this!

In case you haven't figured it yet, “rc files” are simply shell scripts that are executed by a program when it starts. They're simply text files that are, typically, well-commented and allow you to change how a program behaves.

OK, so now we understand how programs use shell scripts to configure themselves. How can you use them to help with what you do every day?

Why Script a Shell?

Think about the project you’re working on right now and the tasks you need to perform on a routine basis. Here are some that I do when building web sites with Node, Ruby, or Elixir:

- Navigate to a project
- Open up the project in an editor of some kind
- Work with a source control system, something like Git perhaps
- Work with a database, something like PostgreSQL, MongoDB, etc
- Write a blog post, perhaps
- Lint/concatenate/minify/compile your code files (CSS, JavaScript, whatever)

Sure there are plenty of tools that can do these tasks graphically for you. Code editors have dialogs for opening certain directories, database GUI tools can help you write queries, and there are plenty of tools out there for graphically working with Git. These tools look nice, but they’re horribly slow when compared to their command line (CLI) counterparts.

Every task you and I do on a daily basis can be done faster with a CLI tool. You can type much faster than you can visualize/click. We could argue that point, I suppose. I have some good friends who work in Visual Studio (Microsoft’s .NET IDE) and with the purchase of a few plugins they can write code rather quickly and have gained a decent level of efficiency.

It still doesn’t come close to what you can do using shell scripts.

Let’s say your boss comes in and asks you to make sure you have a database backup setup on a nightly basis that zips and loads the file to your Amazon S3 bucket. How would you do this with your favorite database GUI tools? This is a 20-line shell script.

You need to lint, concatenate and minify your JS files in a very particular way, using the rules your development lead has set for the team. In addition you need to create a warning when the build size exceeds 100K. This is a 15-line Makefile.

You have a new marketing manager who has decreed that your company site needs to load faster – your current YSlow Grade is a D. It’s decided that the 1200 JPEG files your site serves need to be optimized and reduced in size to no more than 600 pixels wide. This can be done in a 30-line shell script.

I know what you’re thinking: *I can do all of this from the shell using Ruby/Python/JavaScript – why do I care about your shell scripts?* It’s a good question, and a fair point. My response to that would be ... how many packages and supporting files are you going to need? How long will it take, as they say, to “shave that Yak”?

This is the great thing about shell scripts: once you know them, you know them. It’s one of those skills that you can use to do ... just about anything.

One Last Thing

I’ve made some strong claims here, and I’m sure you’re thinking one of two things:

- Shut up and show me the code, or
- Preach it, brother

In the following sections we’re going to build a few of these things. Before we do, you might want to brush up on some basic Unix commands. There are so many tutorials out there on how to work with the command line, it’s almost impossible to make any kind of recommendation. One that I’ve read in the past (and really like) is [Learn Unix In 10 Minutes](#). It’s dense, easy to follow, and will get you up to speed quickly.

Or, if you're a "learn as you go" type of person, I invite you to just start writing scripts with me. I'll discuss the commands in depth along the way and you should be able to pick up a flavor of what we're doing.

If you don't have access to a Unix-based machine, please take a second and go get one. You can set up an account with DigitalOcean, AWS or Azure and kick up a small Ubuntu VM to run the exercises we're about to do. If you're on a Mac, you can just play right along.

All you're going to need before we get going is a text editor. If you're on a Mac, you probably have Vim installed. I'll be using that to edit these things; but if you can't stand Vim than I might suggest the free [Atom editor](#) or Sublime Text.

If you've never done this kind of thing before, hang in there. It looks weird and cryptic, but we'll start slow and, hopefully, you'll understand increasingly as we move along.

Off we go.

CHAPTER SEVENTEEN: SHELL SCRIPT BASICS

Your company website has quite a few images; some of them rather large. Much larger than they should be. A new marketing director was just hired and found out the site is ridiculously slow to load, and has decided that these images are to blame. In short: *you have an image problem*.

Your boss has tasked you with auditing the images and then resizing them. What fun! Isn't this why you became a programmer? The very first thing she's asked for is a list of all the images in our site's directory. That will be our first task.

In the downloads for this section you'll find a directory called "images". You can use that directory to work on.



A Simple First Step

Let's crack open our terminal. On a Mac, this is (most likely) going to be Terminal.app, which you can find in /Applications/Utility. Or you can get your keyboard skills on by typing CMD-Space to bring up Spotlight, then type "Terminal".

It will open in your home directory, or **\$HOME** in Unix land. To navigate around you can use **cd** to change directories – just use the name of the directory you want to go to. If you want to go back one, you can use **cd ..**; if you want go all the way back to **\$HOME** you can just type **cd** followed by <Enter>.

Let's assume you downloaded the image files to your Desktop. For simplicity, let's create a directory in our **\$HOME** called "imposter", and then another inside that one called "demos". In your terminal, type:

```
mkdir -p imposter/demos
```

This command will create a directory set in your **\$HOME**. The **-p** flag tells **mkdir** to create the entire structure if its not already there.

Nice work, now let's move our demo files in there, and then change into that directory:

```
mv ~/Desktop/task-images ~/imposter/demoscd imposter/demos/task-images
```

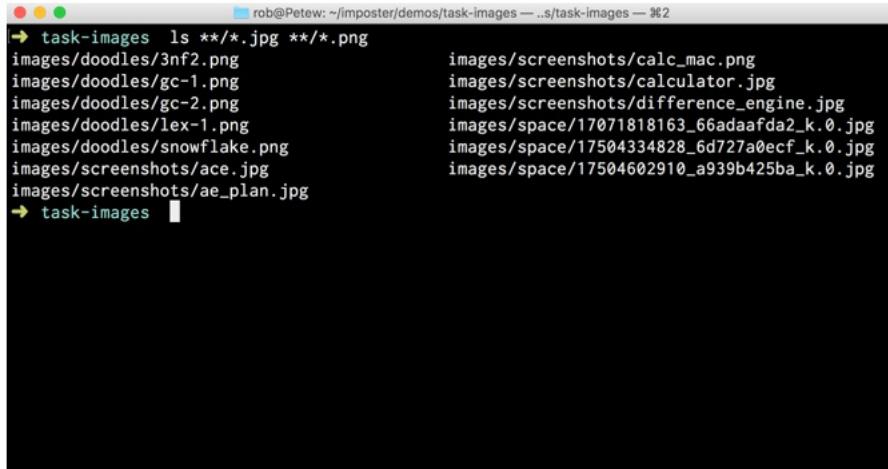
The command **mv** will move files and directories around on your machine and **cd** will change directory, which I'm sure you were able to reason out.

Is all this typing getting you down? Bash and many other shells support command completion using **TAB**. Try it! It really helps when navigating around your machine.

Now that we're here, let's list out the images. You can list files with the **ls** command, but you can also restrict it with what's known as a glob. You can think of this as a series of wildcards:

```
ls **/*.jpg **/*.png
```

This line right here says “list out the jpg and png files, any name, any directory”. Your output should look something like this:



```
rob@Petew: ~/impostor/demos/task-images — ..s/task-images — 362
→ task-images ls **/*.jpg **/*.png
images/doodles/3nf2.png           images/screenshots/calc_mac.png
images/doodles/gc-1.png          images/screenshots/calculator.jpg
images/doodles/gc-2.png          images/screenshots/difference_engine.jpg
images/doodles/lex-1.png          images/space/17071818163_66adaafda2_k.0.jpg
images/doodles/snowflake.png      images/space/17504334828_6d727a0ecf_k.0.jpg
images/screenshots/ace.jpg        images/space/17504602910_a939b425ba_k.0.jpg
images/screenshots/ae_plan.jpg
→ task-images
```

My terminal will probably look different from yours. I've styled it up a bit, and I'm also using Z shell.

If this isn't what you're seeing, make sure you're in the correct directory. Also, be sure you entered the glob correctly as well.

OK, we're almost done. What we need to do now is to create a list that we can show our boss. To do that, we'll redirect the output of the command into a text file:

```
ls **/*.jpg **/*.png > images.txt
```

And we're done! If you want to see this file, you can use the command `open images.txt`, and you'll see them in your default text editor.

That wasn't so bad, was it? That one line saved us quite a bit of work, don't you think? How would you have done this using visual tools?

I just threw a lot at you, but I'm sure it wasn't that difficult. There are two things I want to highlight, however.

Environmental Variables

I was using the term **\$HOME** a lot. This is a special place on a Unix machine – it's where you get to do whatever you want. Visually speaking, you can think of **\$HOME** as the place the Finder opens up to when you first open it. It's usually a place like `/Users/rob` (in my case). You don't ever work on the root of the machine – that's only for special users which we'll discuss later.

Take a look at your **\$HOME**. You can do this using the command **echo \$HOME**. The **echo** command simply outputs a value to the screen, in this case it will be whatever the **\$HOME** variable is set to. That's right – **\$HOME** is a variable, and a special one at that. It's called an “environmental variable” and there are many them. You can tell you're working with a variable in Unix because they have a **\$** prepended to them (this is actually parameter expansion, which I'll get into below). Other variables include **\$PATH** and **\$USER**.

We'll be working with variables of our own making later on.

The next thing I mentioned (but kind of glossed over) was that I redirected the output to a text file. I did this using the `>` operator. This is a crucial thing to understand when working with the shell: there is a standard input and standard output. The standard input is the keyboard, the standard output is the terminal.

In the same way you can refer to `$HOME`, you can refer to standard output as **STDOUT** and standard input as **STDIN**. This might seem a bit academic at this point, but if you think about working with a computer, in general, you give it information and it gives you something back. It does this with **STDOUT** and **STDIN**.

You wouldn't want to have to specify where you want the output sent every time you executed a command, would you? This is where **STDOUT** comes in. If you did want the output of a program to go somewhere, it's easy to specify. Which is what we did using `>`.

Creating An Executable Script

When you're working in a shell you're working in a REPL (Read, Eval, Print Loop). If you don't know what that is – it's a way of working directly with a language. If you have Node installed on your machine you can type in "node" and you'll be in the Node REPL. From here you can enter all kinds of JavaScript code.

If you have Ruby installed on your machine you can enter "irb" in the terminal and you'll be in the Ruby REPL. The Unix command line is the same thing. Bash (or Z shell or whatever shell you're using) will expand and then execute the commands you give it, executing them directly. We'll get into command expansion more later on; for now let's keep rolling.

Create a new file called "resizer.sh" and open it up in Vim. If you're not a Vim fan, use whatever editor you like – but I would challenge you to just give it a try, at least for this walkthrough.

You'll want to be sure you're in the same directory as before – the one with the "images" subdirectory. Then, open up Vim, passing it the file name you want:

```
vim resizer.sh
```

This will create a buffer in Vim, not an actual file. That won't be created until we save the file.

The first thing we need to enter is a *shebang*, which is a great word don't you think? It's also called a *hashbang* by some. To do this in Vim, enter "i" to go into "insert mode" (so you can type some text) and then enter this at the top:

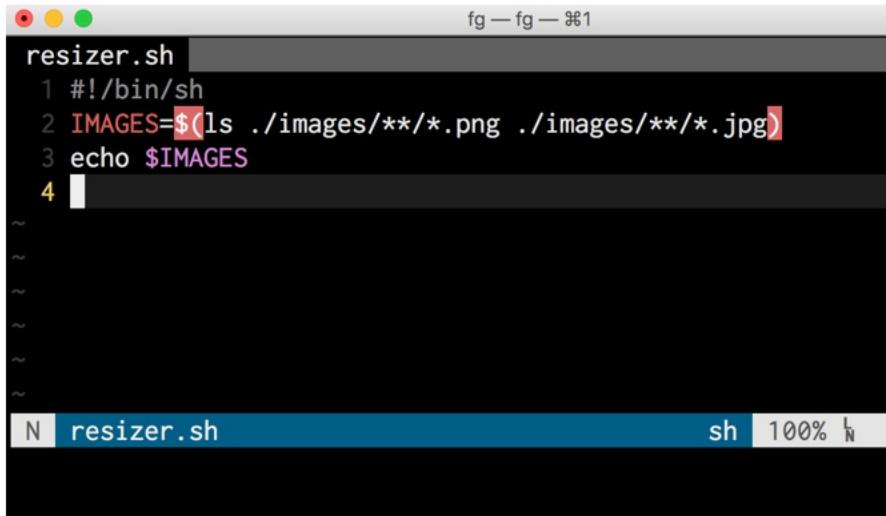
```
vim resizer.sh — vim — %1
+ resizer.sh
1 #!/bin/sh
2
~
```

When you're done typing, hit the ESCAPE key (**<ESC>**). This will return you to "normal mode".

Line #1 above is our shebang, it tells the shell what interpreter we want to use to run this script in the form of an absolute path. In this case, the interpreter is the shell itself, which uses the "sh" program to read in commands from the keyboard, a file, or **STDIN**.

The next thing we'll do is assign our image files to a variable. To make sure we've done this right, I'll output the result to the screen in the same way I might use **console.log** in JavaScript or **puts** with Ruby.

Your cursor should still be on line 1. If it is, enter "c" to create a new line and enter insert mode. You should be on line #2, where you can enter the following:



A screenshot of a terminal window titled "resizer.sh". The window shows the following code:

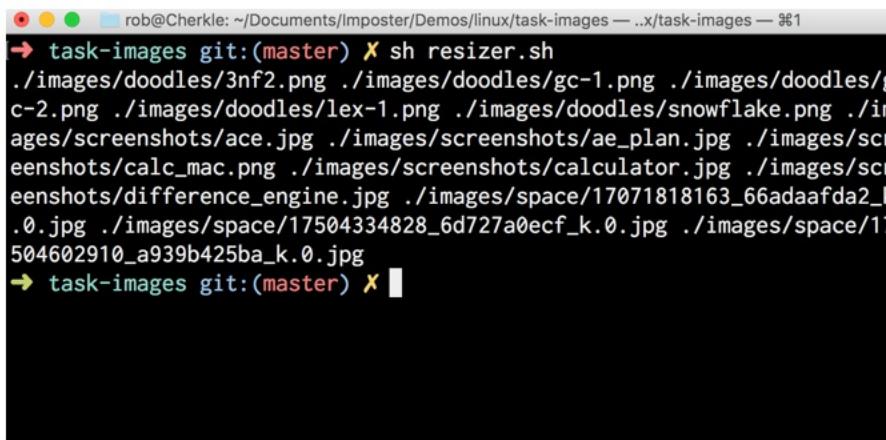
```
1 #!/bin/sh
2 IMAGES=$(ls ./images/**/*.{png,jpg})
3 echo $IMAGES
4
```

The terminal window has a blue status bar at the bottom with the text "N resizer.sh" and "sh 100%".

As with last time, make sure you hit **<ESC>** to go back to normal mode. The next thing we need to do is save the file, and you can do that by entering ":w", which means "write this buffer to disk".

Let's run it to be sure everything works, and then we'll get to the explanation. Enter CTRL-Z to suspend the Vim session, flipping back over to the terminal. If you can't get this to work for some reason, just open up a second terminal window and navigate to the same directory you've been working in – make sure you're in the shell, not Vim. Our goal is to have the shell execute what we've just written.

Enter "sh resizer.sh" into the terminal and you should see an amazing splash of text:



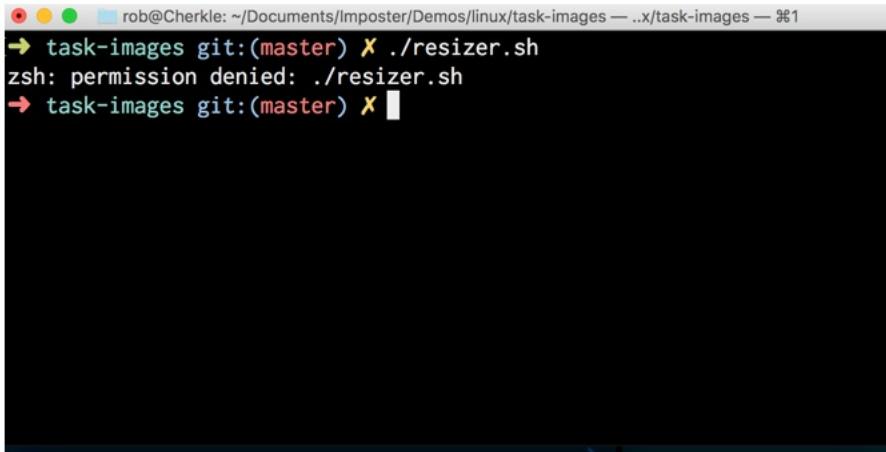
A screenshot of a terminal window showing the output of the command "sh resizer.sh". The output lists numerous image files from the directory structure:

```
rob@Cherkle: ~/Documents/Imposter/Demos/linux/task-images — ..x/task-images — %1
→ task-images git:(master) X sh resizer.sh
./images/doodles/3nf2.png ./images/doodles/gc-1.png ./images/doodles/gc-2.png ./images/doodles/lex-1.png ./images/doodles/snowflake.png ./images/screenshots/ace.jpg ./images/screenshots/ae_plan.jpg ./images/screenshots/calc_mac.png ./images/screenshots/calculator.jpg ./images/screenshots/difference_engine.jpg ./images/space/17071818163_66adaafda2_k.0.jpg ./images/space/17504334828_6d727a0ecf_k.0.jpg ./images/space/17504602910_a939b425ba_k.0.jpg
→ task-images git:(master) X
```

It worked! Or did it? We used the **sh** command, which feeds a file to the shell, the contents of which are

expanded and executed. That's kind of like executing it, but not really. It's a bit like using **eval** in Ruby or JavaScript to run a string of code, rather than executing it directly.

To properly execute a shell script, you just invoke it. Try entering “./resizer.sh” directly. This command simply gives the exact location of the shell script file, with the leading “./” indicating “this directory”. Doing this should lead to some problems:



```
rob@Cherkle: ~/Documents/Imposter/Demos/linux/task-images — ..x/task-images — %1
→ task-images git:(master) ✘ ./resizer.sh
zsh: permission denied: ./resizer.sh
→ task-images git:(master) ✘
```

A screenshot of a macOS terminal window. The title bar says "rob@Cherkle: ~/Documents/Imposter/Demos/linux/task-images — ..x/task-images — %1". The command entered was "./resizer.sh". The output shows "zsh: permission denied: ./resizer.sh", indicating that the user does not have execute permissions on the file. The cursor is at the end of the command line.

By default, you can't execute a file directly in the shell unless you specifically say it's OK to do so. This is a security feature of Unix, as you might imagine.

To grant execution permissions you need to tell the operating system, and you do that by using the **chmod** command (change file mode). By the way, if you ever want to know more about any of the commands you see here, you can use **man** in the terminal itself. This shows the “manual” for each command. Try it now – enter “man chmod” and it will tell you all about it.

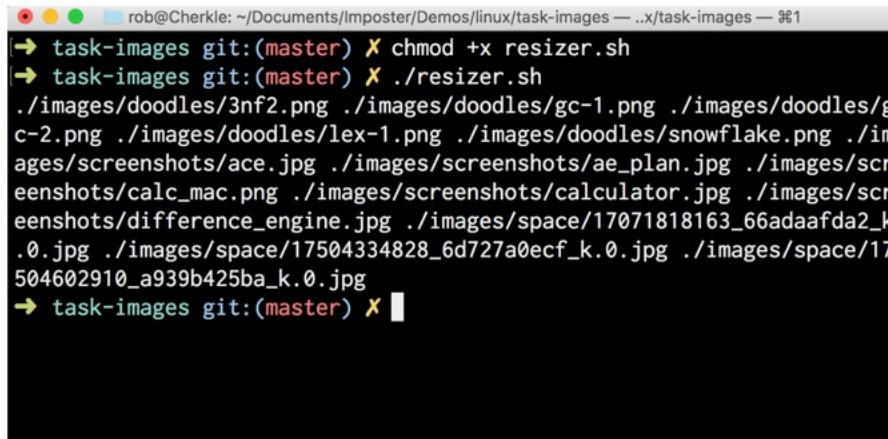
For our needs, I need to chmod +x our resizer, which means “add execute privileges to this file:



```
rob@Cherkle: ~/Documents/Imposter/Demos/linux/task-images — ..x/task-images — %1
→ task-images git:(master) ✘ chmod +x resizer.sh
→ task-images git:(master) ✘
```

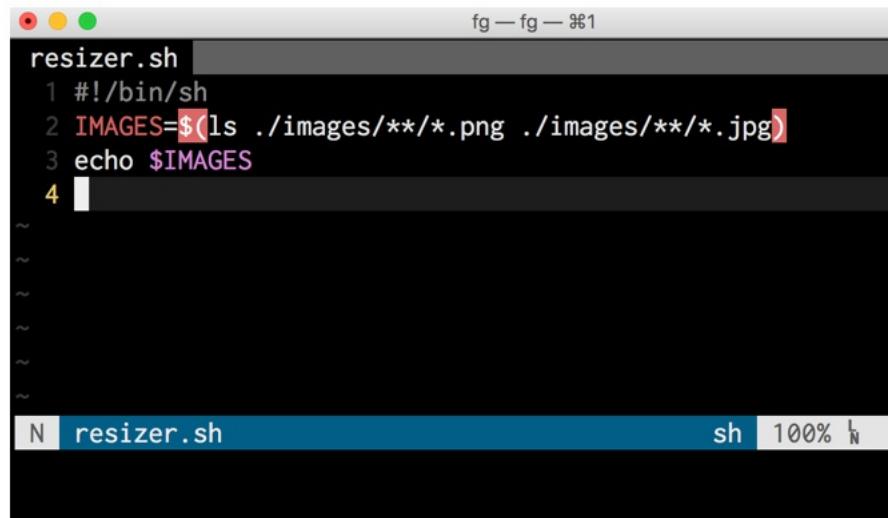
A screenshot of a macOS terminal window. The title bar says "rob@Cherkle: ~/Documents/Imposter/Demos/linux/task-images — ..x/task-images — %1". The command entered was "chmod +x resizer.sh". The output shows the command being run and then the prompt returning, indicating success. The cursor is at the end of the command line.

One thing to get used to with Unix: commands will typically not return any kind of result if they are successful. Silence, in this case, means all went well. Now we should be able to execute our script:



```
rob@Cherkle: ~/Documents/Imposter/Demos/linux/task-images — ..x/task-images — %1
→ task-images git:(master) ✘ chmod +x resizer.sh
→ task-images git:(master) ✘ ./resizer.sh
./images/doodles/3nf2.png ./images/doodles/gc-1.png ./images/doodles/gc-2.png ./images/doodles/lex-1.png ./images/doodles/snowflake.png ./images/screenshots/ace.jpg ./images/screenshots/ae_plan.jpg ./images/screenshots/calc_mac.png ./images/screenshots/calculator.jpg ./images/screenshots/difference_engine.jpg ./images/space/17071818163_66adaafda2_k.0.jpg ./images/space/17504334828_6d727a0ecf_k.0.jpg ./images/space/17504602910_a939b425ba_k.0.jpg
→ task-images git:(master) ✘
```

Nice work! Now let's dive into the code some. To get back over to Vim, enter “**fg**” in the terminal to bring up the suspended app (**fg** means “foreground”). If you get a warning about the file being modified, just enter “**1**” to load it anyway. You should see this now:



```
fg — fg — %1
resizer.sh
1 #!/bin/sh
2 IMAGES=$(ls ./images/**/*.{png,jpg})
3 echo $IMAGES
4
```

We understand the first line, which is our shebang, but the second line looks a tad cryptic. Here, we're setting a variable called **IMAGES** to the result of our image listing ... but what's that syntax?

Command and Parameter Expansion

When you surround a command with a **\$()** it's called a *subshell*. As you can probably reason, I need to set the **IMAGES** variable to the result of the list command, which means I need to invoke it in place. I can do that by wrapping it in a subshell. This subshell will be expanded, and the results returned to the **IMAGES** variable.

We'll do more with subshells in a later section; for now you can think of it as invoking a command in place and using its results directly.

The next line uses the echo command to output the value of the **IMAGES** variable to the screen. To use a variable (which should always be upper cased), you have to expand it as well using the **\$**. This is called *parameter expansion* and might seem a little weird until you have a play with it.

Open up a new terminal window and type in **THING=1**. This will set the variable **THING** to the value 1. Now let's

use **echo** one more time to have a look at this value. Try entering **echo THING**.

What happened? The **echo** command doesn't know if you're giving a literal value or a variable – it's up to you to expand the value before **echo** gets ahold of it. You do this in the same way you expand a subshell: using **\$**.

There are different ways to run subshells and expansions, and we'll get into that in a later section.

For Loops

We have our list of images and, at some point, we'll need to operate on those images individually. This means we'll need some kind of loop – a for loop specifically.

We can do this with our shell. You should be in normal mode with Vim; if you're not, just hit **<ESC>** until you are. You can move around the screen using the h, j, k and l keys. Give it a try, see what happens.

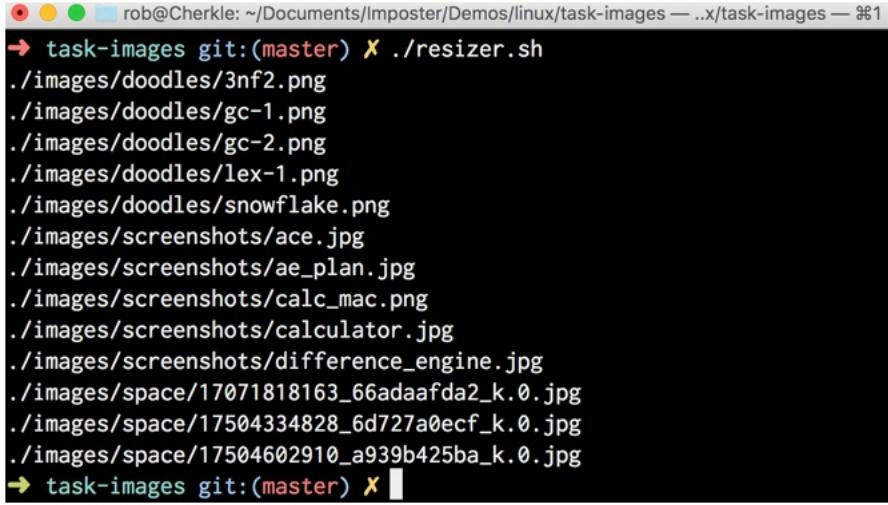
It helps me to think of the "j" as an anchor, pulling down and the "k" as a rock climber, clinging to the face of a vertical wall. Once you're on line #3, enter "dd", which will delete the line.

Now, enter "i" to get back to insert mode. We need to type some more code:

```
+ resizer.sh
1 #!/bin/sh
2 IMAGES=$(ls ./images/**/*.{png,jpg})
3
4 for IMG in $IMAGES
5 do
6   echo $IMG
7 done
8
```

This is our for loop. We declare a variable inline (**IMG**) and then create a **do** block, ending it with **done**. Inside this block we'll report the value of the **IMG** variable. Hit **<ESC>** when you're done entering this code, then type "**:w**" to save it. Now **CTRL-Z** to suspend and flip back to the terminal.

Execute again, invoking the file directly (or hit the up arrow until you see it):



```

rob@Cherkle: ~/Documents/Imposter/Demos/linux/task-images — ..x/task-images — #1
→ task-images git:(master) ✘ ./resizer.sh
./images/doodles/3nf2.png
./images/doodles/gc-1.png
./images/doodles/gc-2.png
./images/doodles/lex-1.png
./images/doodles/snowflake.png
./images/screenshots/ace.jpg
./images/screenshots/ae_plan.jpg
./images/screenshots/calc_mac.png
./images/screenshots/calculator.jpg
./images/screenshots/difference_engine.jpg
./images/space/17071818163_66adaafda2_k.0.jpg
./images/space/17504334828_6d727a0ecf_k.0.jpg
./images/space/17504602910_a939b425ba_k.0.jpg
→ task-images git:(master) ✘

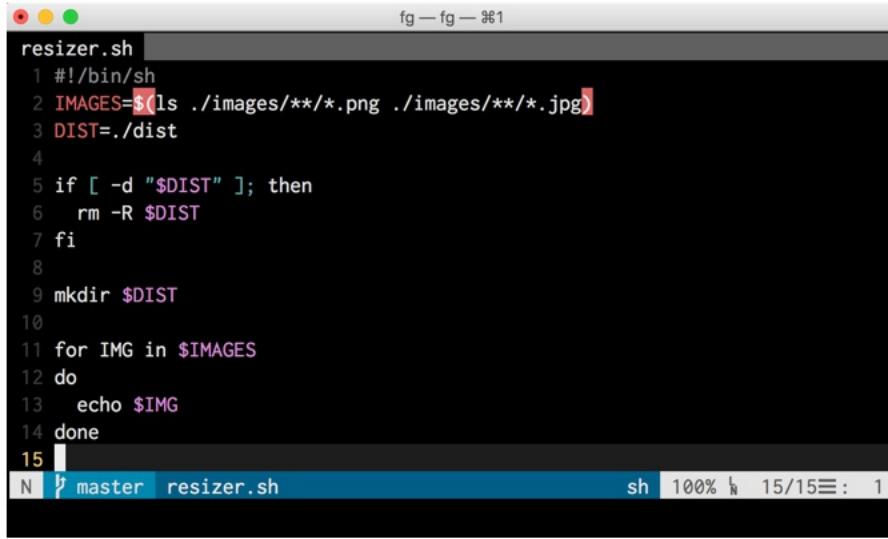
```

Nice! OK, let's keep rolling and pick up some speed.

If Statements

We don't want to blow away our original files, so we'll create a destination directory, where all the modified files will be placed. We need to check if this directory exists before we do anything. If it does, we'll delete it, otherwise we'll just create it.

Use “**fg**” again to get back into Vim and then navigate to line #2, where the **IMAGES** variable is declared. Now enter “o” to open a new line below it. Enter the following:



```

resizer.sh | fg — fg — #1
1#!/bin/sh
2IMAGES=$(ls ./images/**/*.{png,jpg})
3DIST=./dist
4
5if [ -d "$DIST" ]; then
6    rm -R $DIST
7fi
8
9mkdir $DIST
10
11for IMG in $IMAGES
12do
13    echo $IMG
14done
15
N master resizer.sh sh 100% 15/15≡: 1

```

The first thing is the most important: use variables for everything. It's bad form to hardcode values in a shell script! Here, I'm simply specifying where the output directory is going to be.

On line #5 I'm testing to see if this directory exists. Notice the spacing here? It's important! Make sure there's a single space between the brackets and the conditional statement. Also: notice the semicolon? That's optional – it signifies a code line termination. If I put them on a new line the semicolon isn't needed. I use it here because my eyes are used to reading **if** statements in this way.

Next: notice that I wrapped **\$DIST** in quotes? This is a subtle point and not one I expect you to remember entirely. If you write many scripts, however, you'll have a very interesting problem to overcome: *how do I control this expansion thing?*

Quoting

We're using the command line, which is text-based, so Bash will take you literally when you type in anything. We need a way to work with text in this environment.

For instance, let's say I create a file with a silly file name:

```
touch super-* .txt
ls super-*
ls: super-*.: No such file or directory
```

Now this file name and, indeed, this entire example is ridiculous ... sort of. I've seen some amazingly weird file names! Anyway: this command set will cause an error. We're confusing our shell because it can't tell the difference between our literal text **super-*** and our wildcard placeholder **.***. How do we get around this problem?

To "unconfuse" the shell we can use quoting:

```
ls 'super-*' .
super-* .txt
```

I'm using a single quote here, which is referred to as "strong quoting", which means absolutely nothing within the string will have any special meaning to Bash. For our needs this is actually a bit too much (or I should say too literal). We want to have some expansion in there!

For instance, we might want to set the **\$DIST** directory to be something like **\$HOME/fixed** or something. We want **\$HOME** (and other variables) to be expanded in this context.

We can do this by using "weak quoting", or double-quotes, which will allow us to access variables using **\$** (we can do other things too, like use a **\(\sim \)** for our home directory and **** to escape things).

It's good practice to use quoting when referencing variables as we're doing here. If any of our directories or files contain characters that will confuse our shell, their expansion will be ignored.

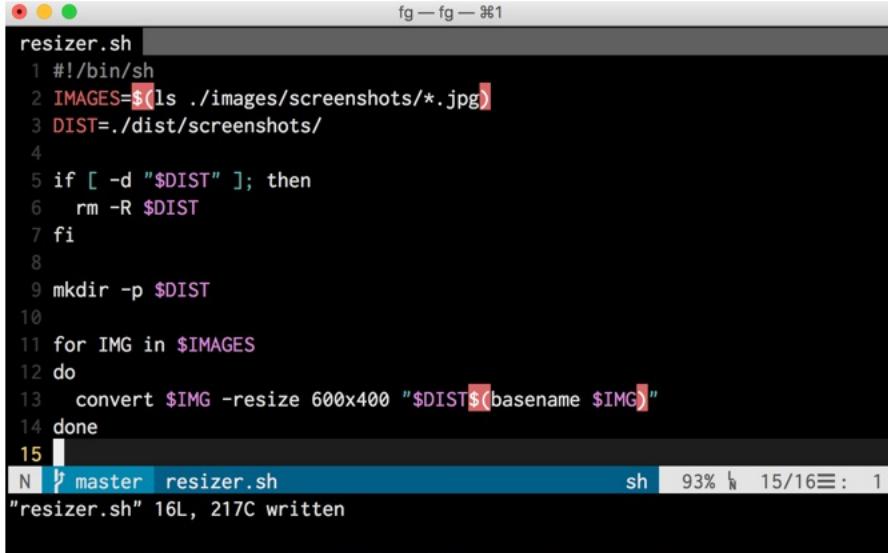
OK, we're almost done – let's rock this out!

Using ImageMagick

To run the actual image conversion I'll use ImageMagick, a popular open source image manipulation tool. You can install it on your Mac using Homebrew ("brew install imagemagick") or with MacPorts.

Once ImageMagick is installed, I simply need to use the `resize` command with some dimensions and an output.

We only need to resize our screenshots, so I'll reset the **\$DIST** variable to only include them (I'll change this later). Finally, I just call `convert` within our loop, and off we go:



```
fg — fg — #61
resizer.sh |
1#!/bin/sh
2IMAGES=$(ls ./images/screenshots/*.jpg)
3DIST=./dist/screenshots/
4
5if [ -d "$DIST" ]; then
6    rm -R $DIST
7fi
8
9mkdir -p $DIST
10
11for IMG in $IMAGES
12do
13    convert $IMG -resize 600x400 "$DIST$(basename $IMG)"
14done
15
N h master resizer.sh sh 93% 15/16≡: 1
"resizer.sh" 16L, 217C written
```

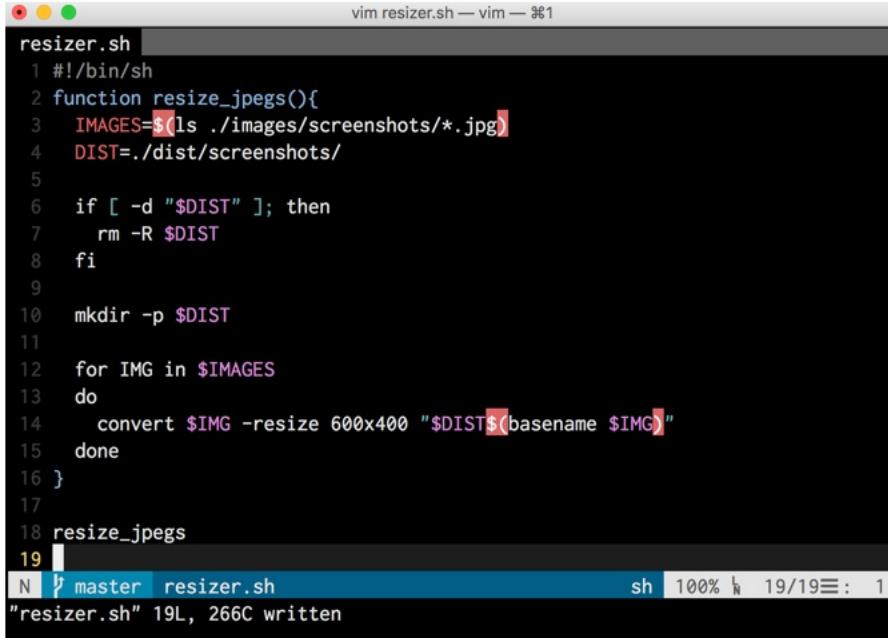
If we execute this script as before, we'll have a lovely new set of resized images waiting for us in the ./dist/screenshots directory. Not bad for 15 lines of code!

But we can do better.

Using Functions

Being the good coder that you are, you're probably wondering why I hardcoded everything for the "screenshots" directory? Good for you!

What we have here is a very workable shell script, but it could be better. Just like any code that you write, think about modularity and reuse. In our case we have a lovely bit of functionality that I'll probably want to use later. The good news is that I can do this by turning it into a function:

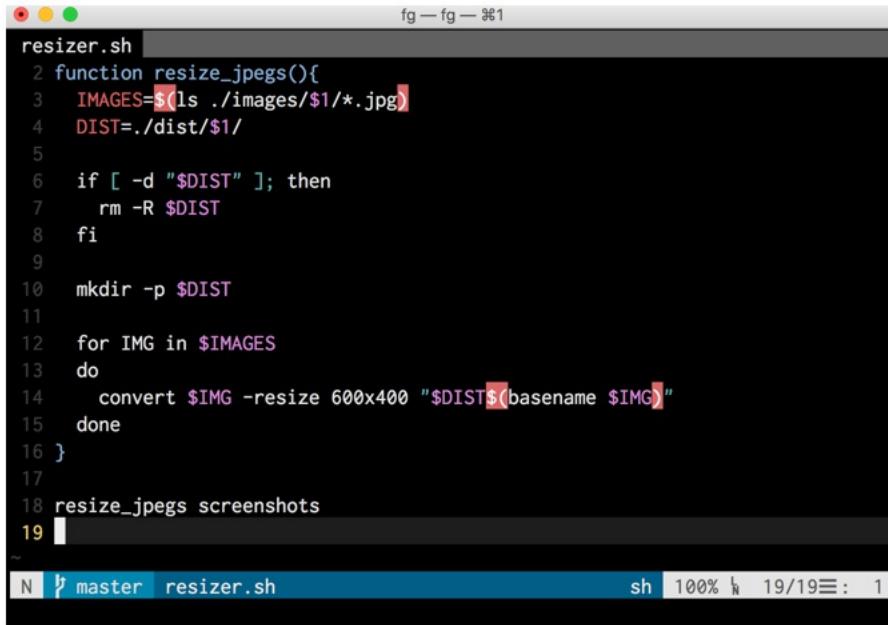


```
vim resizer.sh — vim — #61
resizer.sh |
1#!/bin/sh
2function resize_jpegs(){
3    IMAGES=$(ls ./images/screenshots/*.jpg)
4    DIST=./dist/screenshots/
5
6    if [ -d "$DIST" ]; then
7        rm -R $DIST
8    fi
9
10   mkdir -p $DIST
11
12   for IMG in $IMAGES
13   do
14       convert $IMG -resize 600x400 "$DIST$(basename $IMG)"
15   done
16 }
17
18 resize_jpegs
19
N h master resizer.sh sh 100% 19/19≡: 1
"resizer.sh" 19L, 266C written
```

On line #2 I just declare the function, which looks a bit like JavaScript doesn't it? On line #16 I close it off with a brace and then I can call it directly on line #18.

This is neat, but it's not modular! For that I'll need to send in some arguments. You work with arguments positionally in a shell script, using parameter expansion like we did before with variables.

In this case I'll change "screenshots" to reference the first parameter passed to our function (\$1):



```
fg — fg — %1
resizer.sh
2 function resize_jpegs(){
3     IMAGES=$(ls ./images/$1/*.jpg)
4     DIST=./dist/$1/
5
6     if [ -d "$DIST" ]; then
7         rm -R $DIST
8     fi
9
10    mkdir -p $DIST
11
12    for IMG in $IMAGES
13    do
14        convert $IMG -resize 600x400 "$DIST${basename $IMG}"
15    done
16 }
17
18 resize_jpegs screenshots
19
~
```

N master resizer.sh sh 100% 19/19 ≡ 1

Nice! To pass a value into the function you just tack it on to the function call, which you can see on line #18. I've replaced the hardcoded value with **\$1**, and we're good to go.

Execution vs. Loading

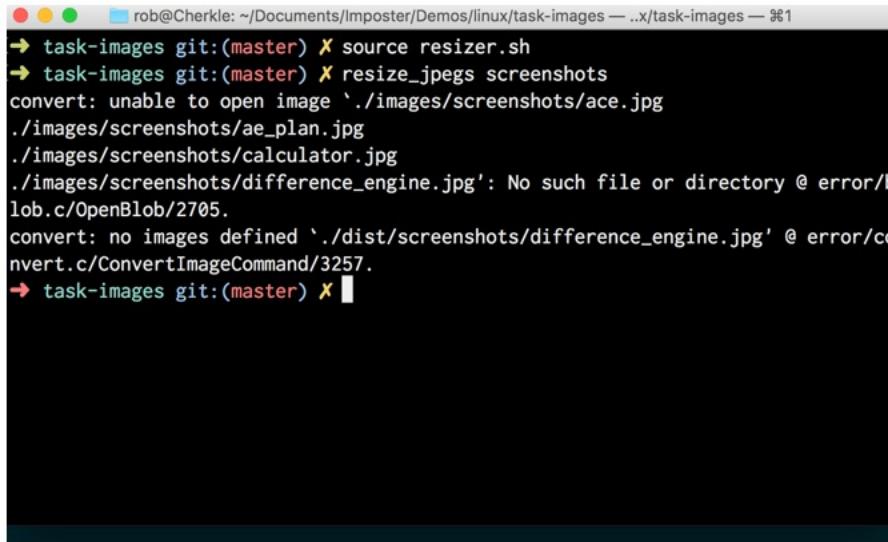
Right now our script executes every time we call it, which is fine, but we might want this command available to us from anywhere. We can do this if we just tweak a few things.

First, let's get rid of line #18. This will prevent our function from executing each time. Then, instead of invoking our script file, we can just source it:

```
source resizer.sh
```

When you source a file, as we've done here, you load its contents into the current shell. This has the same effect as execution, essentially, but it does a little bit more.

For us, we can now call our `resize_jpegs` function from anywhere:



```
rob@Cherkle: ~/Documents/Imposter/Demos/linux/task-images — ..x/task-images — %1
→ task-images git:(master) ✘ source resizer.sh
→ task-images git:(master) ✘ resize_jpegs screenshots
convert: unable to open image `./images/screenshots/ace.jpg'
./images/screenshots/ae_plan.jpg
./images/screenshots/calculator.jpg
./images/screenshots/difference_engine.jpg': No such file or directory @ error/blob.c/OpenBlob/2705.
convert: no images defined `./dist/screenshots/difference_engine.jpg' @ error/convert.c/ConvertImageCommand/3257.
→ task-images git:(master) ✘
```

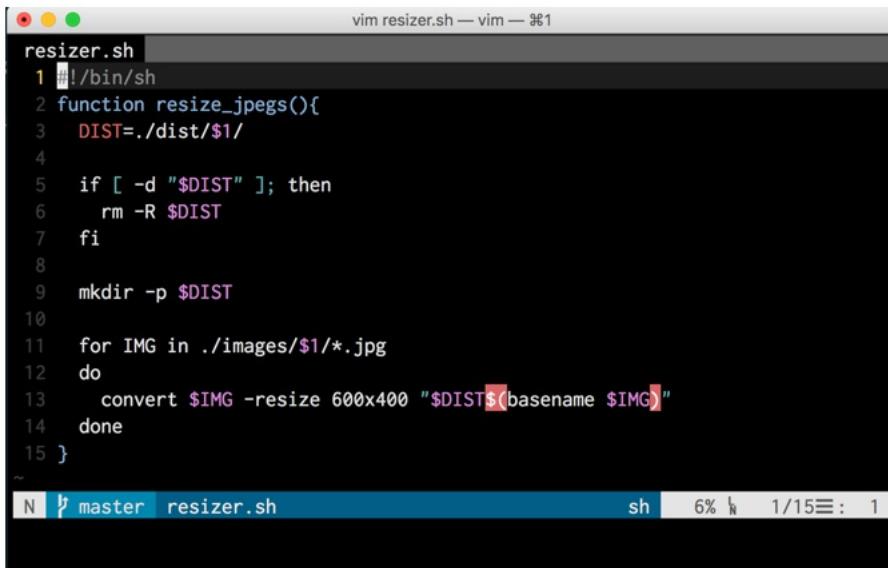
Uh oh. Looks like I was wrong about that. What happened?

Remember above when I said that executing a shell script is sort of like loading it directly into the shell? That sort of just came back to bite me.

When you execute a shell script, as we did above, it's given its own process. Our **\$IMAGES** variable executes properly in a subshell and returns a list of files, as it should. We can then loop over that and all is well.

When we source a script file, we load it into the current terminal session and the code we write in our shell script acts just as if we typed it in there ourselves. What this means is that the expansion we're using to load the **IMAGES** variable is viewed as string value. If you read the error output above, you'll see that ImageMagick is trying to convert a file identified by a huge string value!

We don't want that. The good news is that we can fix this easily by just looping over the glob directly:



```
vim resizer.sh — vim — %1
resizer.sh |
1#!/bin/sh
2 function resize_jpegs(){
3     DIST=./dist/$1/
4
5     if [ -d "$DIST" ]; then
6         rm -R $DIST
7     fi
8
9     mkdir -p $DIST
10
11    for IMG in ./images/$1/*.jpg
12    do
13        convert $IMG -resize 600x400 "$DIST$(basename $IMG)"
14    done
15 }
```

Confusing, isn't it? You'll run into these problems, and the only way to solve them is to think about which process is executing them: your current shell or a spawned shell?

If you're still confused, hang tight – we're going to do more in the next section. For now you have a usable script! Can you see improvements to be made? What about the sizing? Or the input directory?

Have at it. Playing with shell scripts can be fun, and if you get stuck or need to try something new, there are plentiful resources online.

CHAPTER EIGHTEEN: THE JEKYLL POST CREATOR

I use the static site generator Jekyll to write my blog. I store the site at Github, who then translates and hosts it all for me for free. Jekyll is simple to use and I like it a lot. There's only one problem: it's a bit manual.

Every post that you create has to have a very specific filename in the form of year-month-day-post-slug.md. Not that big of a problem, just a mere annoyance. In addition, every post you have must have a blob of text at the top called the front matter. This is the meta information for your post that the rendering engine needs to process your post, which looks something like this:

```
---
```

```
layout: post-minimal
title: 'Some Title'
image: 'featured-image.jpg'
comments: false
categories: category1, category2
summary: "lorem ipsum"
---
```

Again: not a huge problem, just kind of time consuming. The perfect candidate for a little help from a shell script.



Start With a Function

I want to have this function available to me anywhere on my machine at any given time (rather than invoking a script file), so I'll use a function:

```
function new_post() {
  #script goes here
}
```

The first thing I need to do is to navigate my console from wherever I am currently to my Jekyll _posts directory. We know how to do this already using the cd command:

```
function new_post() {
  cd ~/sites/conery-io/_posts
}
```

Creating a Slug

The next thing I need to do is to create the file name that Jekyll expects to see. As mentioned above, this is in the form of year-month-day-post-slug and looks something like this:

2016-12-22-some-title

This file name has two parts: the date bits and the slug, which is a URL-safe, somewhat-descriptive transformation of the title. I had no idea how to do this with a shell script, so I Googled *bash slugify* and found this Gist which I later modified a bit, which now looks like this:

```
"$(echo -n "$1" | sed -e 's/[^[:alnum:]]/-/g' | tr -s '-' | tr A-Z a-z)"
```

The first thing you'll probably notice is that we're surrounding this command with quotes, which we discussed in the last section. You'll also notice a few more things, some familiar, some not at all. We'll step through each one in turn – the important thing is not to get overwhelmed. Just like a complicated math equation, you have to break this command down into parts, and then it becomes a bit simpler.

The first thing you should notice is that we're using a subshell, which we discussed a bit in the last section. For the sake of review, if you surround an expression with `$()` in bash, it will create a subshell, which you can think of as opening another terminal window. You use a subshell when you want an immediate result.

The `$` operator expands the results of the subshell command. So this:

```
echo -n "A Little Brown Pig" | sed -e 's/[^[:alnum:]]/-/g' | tr -s '-' | tr A-Z a-  
a-little-brown-pig
```

... will give you the exact same result as this:

```
(echo -n "A Little Brown Pig" | sed -e 's/[^[:alnum:]]/-/g' | tr -s '-' | tr A-Z a-  
a-little-brown-pig
```

We just hard-coded the title here (the string “A Little Brown Pig”), so it worked OK. If we want to use variables, then we need to change things:

```
TITLE="A Little Brown Pig"  
(echo -n "$TITLE" | sed -e 's/[^[:alnum:]]/-/g' | tr -s '-' | tr A-Z a-z)  
a-little-brown-pig
```

Again, note the convention of uppercase variable names. This is rather important as shell commands can be difficult to read, and you'll be reminded of this if no uncertain terms should you share your scripts with someone online.

Also, it's always a good idea to surround variables with double quotes in bash (called quoting, which we also discussed in the last section), because the expanded command might hold values that you don't want further expanded (like `*`) and others that you do (like `$`):

```
AUTHOR="Rob Conery"  
TITLE="A Little Brown Pig, by $AUTHOR"  
(echo -n "$TITLE" | sed -e 's/[^[:alnum:]]/-/g' | tr -s '-' | tr A-Z a-z)  
a-little-brown-pig-by-rob-conery
```

Now, for some fun, let's add a `$` to the front of our subshell:

```
$(echo -n $TITLE | sed -e 's/[^[:alnum:]]/-/g' | tr -s '-' | tr A-Z a-z)
zsh: command not found: a-little-brown-pig-by-rob-conery
```

We get an error. The shell thinks that we're trying to execute a literal string, which is the result of our subshell command!

When you surround the subshell with a `$()`, it's as if you typed that value yourself into the terminal. This is called a "command substitution" – in other words "substitute this command with the literal value of its output".

OK, now that we understand the structure, let's figure out what those commands mean.

The sed Command

So, we're sending our title, which is the first variable in our function (represented by `$1`) through into `sed`. So, next step is to `man sed`.

What the man sed

When using Linux, `man` is going to be a good friend as you learn things. It means "manual" and it's built right into the shell. You can find help on just about any command, and you can see all of its options.

The downside is that, in some cases, the help system can be a bit impenetrable. Here's what you get when you `man sed`:

The sed utility reads the specified files, or the standard input if no files are specified, modifying the input as specified by a list of commands. The input is then written to the standard output.

If you're new to Linux this might be complete gibberish. As you get increasingly comfortable, sentences like this start making sense. I promise.

The `sed` command is the "stream editor", which means you can pass in some file contents or a raw string and then tweak it, which is precisely what we need to do. The `-e` option is used when you want to append commands to a `sed` operation, which we do because we're using the pipe here (the raw string is the first argument, which we then want to modify further).

The `sed` command is rather powerful and there's a number of things you can do with it – like editing a file in place or transforming `git log` messages. Here, we're using the substitute command (the initial `s`) followed by an "address" – which is the text you want to substitute.

We're using a regular expression `[[[:alnum:]]]`, which means "anything not alpha-numeric" (`[:alnum:]` is short for alpha-numeric if you haven't guessed that). The full expression says search globally for anything not alpha-numeric and replace it with a `-`: `s/[[[:alnum:]]]/-/g`.

Great. The result is piped into `tr`, so our next step is `man tr`.

The tr Command

The `tr` command will translate characters from the standard input (`STDIN`) to standard output (`STDOUT`). Since we're using a piped function, `STDIN` will be the result of our `sed` command.

Our goal here is to fix up what `sed` has transformed so far. Our title could contain many interesting characters, and in our drive to up our blog's SEO we could have a title like "Make Mad \$\$ With My New Framework". The `sed` command above will change this to "Make-Mad—With-My-New-Framework". This won't do – there are too many little dashes and we have capitalized letters.

We can fix the multiple dashes using `tr -s '-'`, which means “squeeze” multiple occurrences of a character into a single occurrence.

Next, we can then transform upper-cased characters to lower case by piping to another `tr` command:

```
tr A-Z a-z
```

We’re using two regular expressions here, so anything matching upper case will be transformed to lower-cased counterpart.

There’s Always a Better Way

I mentioned this in a previous chapter, but I’ll mention it again here: *there will always be a better way*. It doesn’t matter what type of code or process we’re talking about, someone will likely have a smarter idea or a terser way of doing things. I typically favor whatever is clearer – but I suspend that notion for shell scripts.

What matters to me, most, is that it works.

Adding The Date Bits To The URL

We set the result of our slugify command to a variable called `SLUGIFIED`:

```
SLUGIFIED=$(echo -n "$1" | sed -e 's/[[:alnum:]]/-/g' | tr -s '-' | tr A-Z a-z)
```

We can reuse this variable anywhere by referencing it with dollar sign notation: `$SLUGIFIED`, which you should think of, by now, as parameter expansion. Let’s do the same thing with the next part of our slug, which uses the current date to generate the year-month-day part:

```
SLUG=$(date +"%Y-%m-%d"-$SLUGIFIED.md)
```

Again we’re using a subshell with command substitution (hey, isn’t it neat you know that term!) and we’re working with the `date` command. I’ll guess you know what `date` does already.

All we want from the `date` is the current year, month and day so we can add some formatters to `date` using a `+` operator. This will give us `2016-12-22`, for example.

We can then just drop `SLUGIFIED` in there along with the `.md` extension (for markdown) and we’re done with our slug! Yay!

Adding Front matter With a Heredoc

Every Jekyll post has “front matter”, which is some meta data that the parser uses when generating your site. Here is where you put the title of your page, categories, tags, and whatever else you want.

It’s a multiple-line directive and we could, if we wanted, use escaped quoting here but that would be really hard to read and to get right – so I’m using a heredoc instead:

```
<<front_matter
---
layout: post-minimal
title: '$1'
image: ''
comments: false
categories:
summary: ''
---

### Be Brilliant
front_matter
```

You start your heredoc with `<<` and then append the termination term, which in this case is `front_matter`. It's common to see `eol`, but I like to think of it as a surrounding descriptor. Anyway, you can put whatever you want in a heredoc and it will work as if you're using double quoting.

The final steps are to redirect `STDOUT` and move to the root of my blog. We've seen how to do this already, but it looks a bit weird with a heredoc. To output the heredoc to a file, you can use `cat`, which will output the contents of the heredoc to `STDOUT`. We can redirect that using `>` and then use our `$SLUG` variable, which is now properly formatted.

We top it all off using `cd ..` and `subl .` to change into our blog directory and then open up our favorite editor (Sublime Text). Up to now, we've been working in the `_posts` directory, creating our new post file. We should be in the root when opening Sublime Text, so we navigate up a directory with `cd ..` and then call Sublime:

```
function new_post(){
  cd ~/sites/conery-io/_posts
  SLUGIFIED=$(echo -n "$1" | sed -e 's/[[:alnum:]]/-/g' | tr -s '-' | tr A-Z a-z)
  SLUG=$(date +"%Y-%m-%d"-$SLUGIFIED.md)

  cat <<front_matter > $SLUG
---
layout: post-minimal
title: '$1'
image: ''
comments: false
categories:
summary: ''
---

### Be Brilliant
front_matter
  cd ../
  subl .
}
```

Outstanding! With this simple command, you can think of a post while working anywhere on your machine and it will load up happily, pre-formatted with custom front matter.

CHAPTER NINETEEN: MAKE

Make is a build utility that works with a file called a “Makefile” and basic shell scripts. It can be used to orchestrate the output of any project that requires a build phase. It’s part of Linux and it’s easy to use.

It’s important to understand that the utility, make, is not a compiler as many people believe. It’s a build tool just like MSBuild or Ant.

Make is an extraordinarily simple tool which, combined with the power of the shell, can greatly reduce the complexity of your application’s build needs. Even if you’re a Gulp/Grunt/Whatever fan, you should understand the power of make, as well as its shortcomings.



The Basics

Make will turn one (or more) files into another file. That’s the whole purpose of the tool. If you run make and your source hasn’t changed, make won’t build your output.

Make runs on shell commands, orchestrated using the concept of “targets”. Let’s have a look.

First, create a directory where we can work and then create a Makefile:

```
mkdir make_demo && cd make_demo  
touch Makefile
```

Like many build utilities, a single file with a particular name drives the process. With Grunt, it’s a Gruntfile; with Rake it’s a Rakefile. Ever wonder where that convention came from? Yep: *it’s Make*.

Anatomy of a Makefile

All of our build commands will go into the Makefile. Let’s create the basic skeleton:

```
all:  
  
clean:  
  
install:
```

What you see here are targets: the things that Make will try to build for you. By convention they are named based on the file they are building – in some cases, like ours here, they are directives and don’t build anything.

Every Makefile should have these directives:

- **all**: builds everything
- **clean**: cleans up all the existing build artifacts; usually deleting the files built

It’s typical, as well, to have these directives as well:

- **install**: installs whatever built files are generated
- **.PHONY**: lists out the directives that don't create a file

Have you ever downloaded software from source onto a Linux box and had to run `make && make install?` This is why. The downloaded source (usually C or C++) is compiled and then installed wherever binaries go in the system, which can be /opt or somewhere else such as /usr/local or /usr/bin.

With our example we won't be installing anything so we don't need the install target. In its place we'll add another directive that tells Make which of the targets don't create a file. For this we'll use `.PHONY`, for "phony" targets:

```
all:
clean:
.PHONY: all clean
```

Understanding Targets

Make builds output files from input files. It was originally designed for C programs, which utilize both code and header files which are built into object files. These object files are then compiled to binary. This is a multistep build that requires some orchestration. That's what Make is all about.

Sometimes, however, you'll want a build step that might transform some input – it might not create a file. Letting Make know about these special (phony) targets will increase performance greatly. We'll get more into this in the next chapter.

The Simplest Operation

This is a silly demo, but it's critical you see the way things work before we dive into some useful stuff in the next chapter. Let's create a file in the root of our project directory that we want to transform and distribute. We'll leave the file empty for now – let's also create a Makefile while we're at it:

```
echo "//some code" > app.js && touch Makefile
```

The simplest first step would be to copy our app.js code file to a /dist directory. So let's do it:

```
all:
mkdir -p dist
cp app.js dist/app.js

clean:
rm -rf dist

.PHONY: all clean
```

The indents in a Makefile must use tabs. If you don't use tabs you'll get an error about an invalid separator. If you copy/paste the code here, be sure you indent with a tab that doesn't get translated to spaces.

OK, save the Makefile and run make:

Nice! Make outputs each command as it's executed – so here it's reporting that it ran mkdir and cp successfully. This can be good ... it can also be annoying.

Let's fix the chatter and provide a clean target while we're at it:

```
all:  
@ mkdir -p dist  
@ cp app.js dist/app.js  
  
clean:  
@ rm -rf dist  
  
.PHONY: all clean
```

By prepending an @ sigil to a command line in our Makefile, we're silencing the output. Also - our clean target will delete the entire build directory. We can test that by running `make clean`.

Try it out!

Target Execution

If you want to execute a single target you just specify the name when calling Make as we've done here. If you don't specify a target then make will pick the first in the Makefile; this is called the default goal.

With this Makefile, that would be the all target. It's a good idea to have this, always, as your first target. It helps avoid confusion.

Orchestrating The Build

Our code files are a bit messy - just stored right in the project root. So let's clean things up with some organization:

```
mkdir src mv app.js src
```

Great. Let's assume (for now) that all of our project code goes into the /src directory.

Our current Makefile is not really doing any build orchestrations of any kind – it's just copying a single JavaScript file to the /dist directory. Let's change that by adding a build timestamp to the output, so we know when the last build was run.

Let's clean things up a bit so we have a build target that produces a file and another that produces the

destination:

```
all: dist app.js

dist:
    @ mkdir -p dist

app.js:
    @ cp src/app.js dist/app.js

clean:
    @ rm -rf dist

.PHONY: all clean
```

A whole lot just went on there – and with it we get to learn some more jargon.

First, I created two new targets called app.js and dist. Targets in a Makefile are just the labels; what comes after the target is the recipe. So, for the dist target, **mkdir -p dist** is the recipe.

The **all** target has prerequisites that appear on the first line, but it has no recipe. Prerequisites are targets that must be built before creating the current target. So, for **all** to work, **dist** and **app.js** need to have run first.

If you put a target, recipe and prerequisites together, you have a *rule*. Which is precisely what a Makefile is: *a set of rules for building your software*.

This is the power of Make. Orchestrating what happens when, and in what order. You could say that this is all that Make does, which isn't surprising if you understand the Linux philosophy of “do one thing well”.

Using Variables

We're dealing with shell scripts here, and just like any programming effort, repeating ourselves is typically frowned upon. Let's take the hard-coded stuff out first so we can change as we need to.

By convention, variables should be declared at the top of your Makefile:

```
JS_FILES=src/app.js
DIST=dist

all: dist app.js

dist:
    @ mkdir -p $(DIST)

app.js:
    @ cp $(JS_FILES) $(DIST)/app.js

clean:
    @ rm -rf $(DIST)

.PHONY: all clean
```

Now that's starting to look like a proper Makefile! Variables in any shell script should typically be upper-cased.

Notice how I have to use **\$(DIST)** to reference the variable? Do you remember what that is? It's a *command-replacement subshell*. Make runs all the commands in a subshell, outside of its own process.

At first glance it might not look like we've done much here – but if we change our source files, which we will, it's a simple change to the **JS_FILES** variable.

OK, one last thing. We're still repeating ourselves in a couple of places – specifically with the target and the destination file names of **app.js**. If we're leaning on convention, we shouldn't have to specify things twice.

To get around this, we can use the `$@` shorthand – which means “this target name”:

```
JS_FILES=src/app.js
DIST=dist

all: dist app.js

dist:
    @ mkdir -p $@

app.js:
    @ cp $(JS_FILES) $(DIST)/app.js

clean:
    @ rm -rf $(DIST)

.PHONY: all clean
```

We're getting there. Now, let's create our timestamp. For this I'll use a variable straight away, and put it right at the top:

```
JS_FILES=src/app.js
DIST=dist
TODAY=$(shell date +%Y-%B-%d)
TIMESTAMP="//Created at $(TODAY) \n\n"
#...
```

Here I'm using the `shell` command, which, if you recall from previous sections, executes shell commands for you. We need to get the literal value of the date and the date's formatting instruction, and store it in the **TODAY** variable.

We then create our timestamp. Using variables is a great way to keep your code clean and understandable, and that's critical in Makefiles as the syntax can quickly become overwhelming.

The next task is to read the source file and timestamp it:

```
#...
app.js:
    @ echo $(TIMESTAMP) > $(DIST)/$@
    @ cat $(JS_FILES) >> $(DIST)/$@
#...
```

This is a bit roundabout, but it works. As I keep mentioning, there's almost always a better way when it comes to shell scripting – and we should most definitely extend that idea to Makefiles. This is nice and clear to me, however, and hopefully it makes sense to you as well.

I'm redirecting the output of `echo` (which is the **TIMESTAMP** variable) into **dist/app.js**, which is our output file. On the next line I'm appending **STDOUT** using `>>` to the same file – this time with the contents of the JavaScript files in our `src` directory.

Here's another go at this:

```

JS_FILES=src/app.js
DIST=dist
TODAY=$(shell date +%Y-%B-%d)
TIMESTAMP="//Created at $(TODAY) \n\n"

all: dist app.js

dist:
@ mkdir -p $@

app.js:
@ echo $(TIMESTAMP) > $(DIST)/$@
@ cat $(JS_FILES) >> $(DIST)/$@

clean:
@ rm -rf $(DIST)

.PHONY: all clean

```

This is a solid Makefile, but it's not quite there yet. Notice that I have a **dist** target and a **DIST** variable? This is redundant! You can use variables as target names, so let's do that and arrive at our final Makefile:

```

JS_FILES=src/app.js
DIST=dist
TODAY=$(shell date +%Y-%B-%d)
TIMESTAMP="//Created at $(TODAY) \n\n"

all: $(DIST) app.js

$(DIST):
@ mkdir -p $@

app.js:
@ echo $(TIMESTAMP) > $(DIST)/$@
@ cat $(JS_FILES) >> $(DIST)/$@

clean:
@ rm -rf $(DIST)

.PHONY: all clean

```

Let's give it a run using make clean && make. You should see a dist/app.js file with this in it:

```

//Created at 2016-October-19
//some code

```

Nice work! Once again, this is a bit of goofy demo, but your brain should be racing a little bit right now... thinking about all the ways you might put Make to use with your project.

In the next section we'll do even more work with JavaScript files, kicking Grunt right out of our project.

CHAPTER TWENTY: USING MAKE TO BUILD YOUR WEB PROJECT

Every language/platform seems to have its own build and automation toolset. Microsoft has MSBuild, Java has Ant, Maven, Gradle, and perhaps a few others. JavaScript has Grunt, Gulp, Jake and a few more. Ruby has Rake. Why is this?

The short answer is that Make can be a little opaque and, predictably, language users like to build things using their own language and believe it will impart some type of benefit over the decades-old Makefile.

This subject, like others in this book, is a tad volatile and I want to recognize that right up front. An easy way to start an argument between two developers is to bring up the subject of build automation! A good friend of mine, Rob Ashton, thought it would be fun to submit a pull request to the MSBuild Github Repo, [suggesting that it be replaced with Make](#). If you've ever had to wrestle with MSBuild before (which I have, numerous times), then this might make you laugh. I thought it was funny, but there are many who did not.

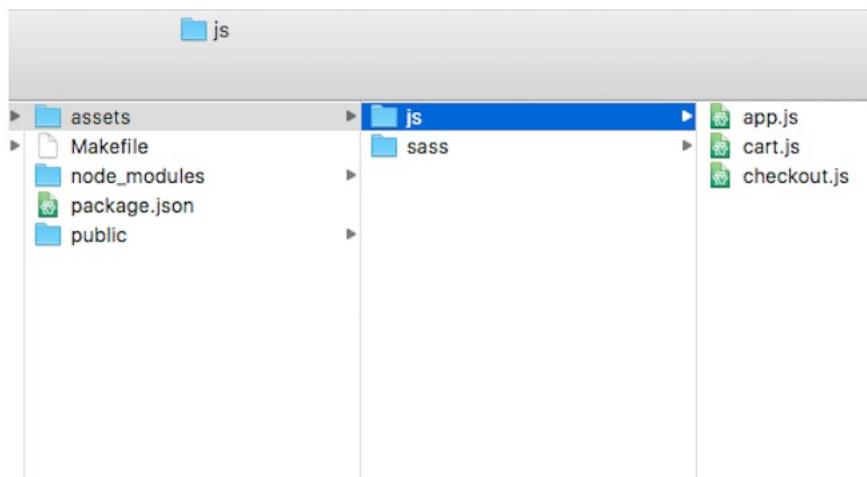
No matter which build tool you use and love, knowing Make will give you *a perspective that you need*. It's built right into Unix-based systems and you can use it to automate just about anything. Whenever you crack open that Gruntfile, Gulpfile, or Rakefile, ask yourself whether the layers of complication you're adding to your project by using these tools are really warranted. The answer, typically, is no.

The problem, typically, is that most people don't know how Make works. Let's change that now, using a fairly typical use case: compiling assets for a web project.



The Web Project

I have a typical web application that uses SASS and a bit of JavaScript:



As you can see, it's a simple app. I have a set of JavaScript files and a single SASS file that I need to build and then output to a directory somewhere.

This is my goal for this task: *concatenate, compress and build my SASS files and then concatenate and uglify my JavaScript files.*

Working With SASS

Some people dig SASS, others LESS. Other people (like myself) tend to knuckle under and just write CSS directly. Hopefully what you're about to do will translate to how you develop things.

Also: when working with SASS and LESS it's common to import all of your files into a main.scss file, which you then build with the interpreter. That's what I'll be doing today.

Step 1: The Node Modules We Need

I need to install a few Node modules to help me out:

- **node-sass**: the module I need to compile and build my SASS file
- **uglifyjs**: removes whitespace, concatenates and then minifies my JavaScript files

That's it. To install these we use NPM:

```
npm install uglifyjs node-sass --save-dev
```

Step 2: Our Makefile Skeleton

The next step is to create the Makefile in the project root and then give it a default structure:

```
touch Makefile
```

Open the Makefile and then add the default structure – the targets we know we're going to be working with:

```
all:  
clean:  
.PHONY: all
```

Hey we're getting good at this!

Step 3: Define The Variables

We don't want any hardcoded values in our rules, so let's define as much as we can at the top:

```

SASS=node_modules/.bin/node-sass
SASS_FILES=assets/sass/main.scss
JS_FILES=assets/js/*.js
UGLIFY=node_modules/.bin/uglifyjs
DIST_CSS=public/css
DIST_JS=public/js

all:

clean:

.PHONY: all

```

I'm defining the binaries, the source files, and the destinations.

Binaries

I've installed the **node-sass** and **uglifyjs** binaries locally as development dependencies. This will increase the time it takes to run npm install, but it also guarantees that the binaries will be present. It's really annoying to have to install global dependencies (my opinion).

In case this is all new to you or if you've never used Node: you can install packages from NPM and have them run locally or globally. You've probably seen a node_modules directory at some point in your career – this is the local package installation location. If I was to install node-sass locally I could run it by using node ./node_modules/.bin/node-sass. Or I could make life easy on myself and install it globally in the global NPM cache, which would allow me to run the binary anywhere on my machine.

Step 4: Define The Targets

Given that I want to smash all the JavaScript files into a single app.js file, I can create that target. Same with app.css. The destination for these files will be the public directory, which should exist already in my project. However we can't guarantee that so let's make sure that we have a **dist** target as well:

```

SASS=node_modules/.bin/node-sass
SASS_FILES=assets/sass/main.scss
JS_FILES=assets/js/*.js assets/js/**/*.*js
UGLIFY=node_modules/.bin/uglifyjs
DIST_CSS=public/css
DIST_JS=public/js

all: dist app.css app.js

dist:
  mkdir $(DIST_CSS) $(DIST_JS)

app.css:

app.js:

clean:
  @rm -rf $(DIST_CSS) $(DIST_JS)

.PHONY: all

```

Simple enough. Creating the directories we need with the **dist** rule, removing them with **clean**. Now all we need to do is to fill in the commands to create the files.

File Strategy

Loading up your source files like this is really simple when using a glob pattern, but you might have different needs. For instance: you might have an app.js file in assets/js that you want loaded first, and then all the other files loaded after that.

You can do this easily by specifying **assets/app.js** as the first file in the list, and then the glob **assets/js/**/*.js** after that.

There are other strategies which I'll discuss down below.

Step 5: Building The Files

The first step is simple: we only need to compile the main.scss file using node-sass, which comes with a binary in the node_modules/.bin directory.

This is what the command would look like normally:

```
node_modules/.bin/node-sass --output-style compressed assets/sass/main.scss > pub
```

We just need to replace this command with our variables and we're good to go:

```
app.css:  
@ $(SASS) --output-style compressed $(SASS_FILES) > $(DIST_CSS)/$@
```

Bam. That couldn't be easier!

Discovering The Proper Commands

OK, I lied: *it could be easier*. You'll find that most of your time is spent wrangling the commands together – reading the docs, trying to understand how to use the binaries properly.

When I first put this file together for a project I was working on, I had to read through the source code of some of the binary files to understand all the options and how they worked. I then ran those commands from the command line to make sure things happened the way I wanted.

The JavaScript files are a bit of a different story. I need to concatenate them together, and then uglify/compress them. Concatenating is simple – we can use cat directly to pull the code from each file – which we can then pipe directly to **uglifyjs**:

```
app.js:  
@ cat $(JS_FILES) | $(UGLIFY) > $(DIST_JS)/$@
```

We haven't discussed the `|` notation you see here. That's called the pipe operator and is a Unix operator that redirects the output of one command to the input of the following command. In the same way I redirected **STDOUT** to a text file using `>` in a previous section, the `|` operator redirects **STDOUT** and **STDIN** for two given functions.

In this example, I'm using cat to concatenate all the JavaScript files specified by my **JS_FILES** variable. I'm then piping that text into the **uglifyjs** binary, and then redirecting the output of the **uglifyjs** call to my destination.

I wish I could tell you it was harder than this ... but it's not. This is the power of Make. Run make and take a look at the built output in the public directories. If you've already run it, be sure to run **make clean && make** before you run make again to avoid errors.

Easy Not Easy

If this seems a bit cryptic or if your reaction to this code is something like "yeah great for you – you know shell

scripts!” you might be interested to know that I’m a complete n00b to this stuff. I learned it better over the last year, but creating Makefiles like this was absolutely something that was beyond me just 8 months ago.

It just takes a little time, some Googling, and some practice and you will get it too.

Ordering Of Files

Ideally your JavaScript files are not dependent on load order, but in the Real World this is typically not the case. There are a number of ways to get around the problem and they involve knowledge of some basic commands.

Personally, if I need files loaded in a particular order (as with my SQL files – tables need to be built before views and so on) – I just name them in a particular way:

```
01-init.sql  
02-tables.sql  
03-views.sql  
04-funtions.sql
```

This works fine for SQL files, but it might not work for your JavaScript build. If you need files loaded in a particular order you can pipe the result to **sed**, as we did with the Jekyll task, and then transform it from there.

You can also separate the build steps, so you build a directory first, then another, outputting interim files to assets/tmp directory, which you then concatenate later on.

Summary

In this section, we’ve touched on some very fundamental ways we can use Make to automate and simplify our lives. But we have only scratched the surface. Make has a long, long legacy, and is the spiritual great-granddaddy of newer tools like Grunt, Gulp, and other automation tools we take for granted. I can’t overstate the value that even a basic fluency with this tool will bring to your understanding of how source files are assembled into working software.

CHAPTER TWENTY-ONE: CRON

Linux has a wonderfully annoying feature referred to simply as “cron”. It’s a simple text file with an ASCII table that will execute a command on schedule.

Cron jobs are incredibly useful. They can help with various maintenance tasks such as log file cleanup or database backup.

For instance, we might want to backup a database on a nightly basis. We can do this with a simple dump command that our database tool gives us on a nightly basis, managed by cron. The dump command can be executed directly in the cron job, or (a better approach) in a shell script.



The Basics

Let’s backup a PostgreSQL database using cron. If you’re not using PostgreSQL that’s OK – it’s likely that whatever platform you’re using supports some type of binary tool that will do a backup. Translate as required.

The first thing to do is to open up our cron configuration file using crontab -e where the -e option means “edit”. This will open (can you guess?) a simple text file with some prompts to help you out.

To this file we can add our command:

```
0 0 * * * pg_dumpall -f ~/backups/db.sql
```

This is a little cryptic, but not too difficult to understand once you’re told what the 0’s and *’s mean. The entry specifies the time in a positional way, from left to right. You can read it as:

- minute (0 to 59)
- hour of the day (0 to 23)
- day of the month (1 to 31)
- month (1 to 12)
- day of week (0 to 6 with 0 being Sunday)

Then finally comes the command you want to execute, in our case `pg_dumpall -f \(\sim\)/backups/db.sql`. Again: I can execute a command directly, as I’m doing here, or I can execute a more in-depth shell script, which I strongly recommend. It’s very likely your backup command will grow, or you’ll want to do other things with it – so do yourself a favor and use an executable shell script from within cron.

I used `0 0 * * *` as the schedule, which reads “the first minute of the first hour of every day of the month, every month, every day of the week” (the stars are wildcard characters).

Execution Troubles

Cron configuration files are created for each user and, subsequently, will run as that user when executed. If your cron task doesn’t run, it’s likely due to a permissions problem. For instance, my user account might not have permission to access PostgreSQL, so this backup command will fail.

Be sure your account (or whatever account you're setting up to run Cron) can execute the command you're scheduling first.

A Real World Scenario

My last company, Tekpub.com, ran on PostgreSQL and I needed to have a backup of the database go off on a nightly basis. I started out as we have here, but soon realized that backing up a database requires a little more automation.

Every DBA has their favorite backup script and will readily share it online. The one I'm about to show you is one that I adapted long, long ago from my friend Rob Sullivan. It's a straightforward script, which I'll explain below:

```
BACKUP_DIRECTORY="/home/rob/backups"

# Date stamp (formated YYYYMMDD)
# just used in file name
CURRENT_DATE=$(date "+%Y%m%d")

# !!! Important pg_dump command does not export users/groups tables
# still need to maintain a pg_dumpall for full disaster recovery !!!

# this checks to see if the first command line argument is null
if [ -z "$1" ]
then
# No database specified, do a full backup using pg_dumpall
pg_dumpall | gzip - > $BACKUP_DIRECTORY/everything_$CURRENT_DATE.bak

else
# Database named (command line argument) use pg_dump for targeted backup
pg_dump -d $1 -Fc -f $BACKUP_DIRECTORY/$1_$CURRENT_DATE.bak

fi
```

To use this script, you invoke it directly: `\(\sim \)/scripts/backup.sh`. When it runs it will check to see if you specified a database. If you did not, the entire database will be backed up; if you did then only the select database will be backed up.

Here's a cron entry for backing up my "bigmachine" database. A quick quiz for you: do you think this work?

```
0 0 * * * ~/scripts/backup.sh bigmachine
```

The answer is: *maybe*. There are some things I need to do before this cron job will run. They are:

- Make sure the backup.sh file is executable
- Make sure the job I want to execute (in this case `pg_dumpall` and `pg_dump`) is something my account has privileges for
- Make sure that the file reference is accurate for the account executing the cron job

This cron job will run as if I'm logged in and executing the command directly, so yes it will work as long as I have permissions to run backups on my database, which I do. It is a better idea, however, to use an absolute path when invoking a script as this will break if it's ever run by another user.

So, to keep in good form, we should update the command thus:

Execution As Another User

You might run into a scenario where you want a cron job to run as a different user. In that case, you can elevate your privileges using **sudo** and edit their crontab as needed.

One word of caution, however: *please don't ever setup a cron job for the root user!* It might be tempting, especially if you're having problems getting the job to run. This can have unintended consequences – it's like trying to swat an annoying fly with a machine gun.

The root user can do anything to a Unix machine, and that one small problem in your script where you accidentally deleted all files on the server could be avoided by taking a little extra time to figure out why your script won't execute properly.

Going Further

As I mention, database backup scripts are all over the web and you can find one easily that you can plug into cron. If you run into trouble, it's almost certain to be a permissions issue. Go over the three steps above (executable script, you can perform the tasks, the file location is correct) and you'll likely solve any problem.

Have a look around, explore! Maybe you don't want to store backups locally and, instead, want to push them to your S3 bucket where they'll be safe and warm in the loving arms of the AWS cloud. You can do this with a shell script!

Maybe you want to send an email confirming that the backup went off successfully: again, doable with cron and a shell script. I used to have a script that would do just that: backup my MongoDB database every night, send it to S3, run a sales query and pass the results to me so I could read them in the morning.

Shell scripts are tons of fun.

WRAPPING UP

This book began as a series of tasks in Wunderlist: *things I need to know someday*. I didn't intend to write a book. I thought maybe I would write some blog posts, perhaps do a video or two on a few of the topics I found.

The problem I had then (and still have now) is this: **I didn't know what I didn't know**. It's getting worse, too. The more I learn – the more I feel like an imposter.

There is comfort in not-knowing something. More than that: there is magic. Even more than that I think there is also a **bit of dumb courage**. You just don't know if it will hurt – so you try it out.

I'm reminded of being 15 and jumping off the roof of my house with an umbrella. Along with wondering if the umbrella would slow me down, *I also wondered how badly it would hurt*. Turns out the umbrella did nothing and jumping off the roof of my mom's house did, in fact, hurt. I didn't break anything or seriously injure myself, which is a good thing ... but I only tried the umbrella jump once. The next time I did it without the umbrella ...

I picked the cover of this book because that's what it felt like when I learned to program on my own: *jumping off of something rather high, hoping the landing wouldn't break me somehow*. I had a job and a budding career as a geologist – but it was boring. A good friend of mine was a trainer – the very first one certified by Microsoft to teach Active Server Pages. So he taught me (thanks Dave!) and I **jumped**.

I've hit the ground hard a few times, but the joy of jumping into new things is my drug. Which is kind of a problem because new things don't stay new for very long, and soon you start to feel the magic of it all start to fade.

It's the curse of learning: remove the mystery, remove the magic. Move on to the next mystery. Pretty soon you wonder what the point is.

However.

The process of writing this book has been *transformative*. I'm going to push the whole jumping metaphor a little more – and I'm also going to reflect on the cover image a bit more too. You may not have noticed, but on the back of the "jumper" on the cover where a parachute should be is ... this book. If you look closer, you can see the same cover. I felt it kind of appropriate as recursion seems to ... recur throughout this book.



More than that, though, is the idea of a parachute. This book has become my parachute, in a sense – the things I've learned over the last year are allowing me to jump off higher rooftops – but this time I have a bit of knowledge to soften the fall.

Climb higher, jump farther. The mystery/magic/dumb courage seems to be increasing as I learn these things.

One Final Thought

I've gone deeper into these topics than I ever thought I would. I feel like I have a solid grasp on complexity theory, Big-O, graphing algorithms and Bernoulli's Weak Law of Large Numbers and ... just writing this sentence is giving me a rush! This shit is BREATHTAKING!



Sorry. I promised myself I wouldn't swear in this book – but if any of the paragraphs written deserve some caps and a four-letter word ... well it's that one.

There. Is. So. Much. Magic.

Go find it. Don't stop here. It's so easy to fall into snark and uncertainty; to feel like *you just don't know what other people are talking about*. It can really be isolating and lead one to look upon others with scorn for trying new things. As I write this sentence, Facebook just pushed a new package manager for JavaScript. As expected, quite a few people ridiculed it and laughed without even trying it!

Let's not be those people. We're the explorers who keep trying (and hopefully failing). We push the edges to find out what's possible. Enjoy this!

I urge you to explore and dig deeper, ask more questions, take a friend out and bore them to tears with all you know! Draw a picture of Traveling Salesman for your kids and ask how they would solve it. Make a Markov Chain out of popcorn for the holidays!

With that: *I leave you.* Thank you for going on this journey with me. Now let's go make a difference.

